

# Amazon — Senior Supply Chain Manager, NACF Supply Chain (L6)

**You.** Rajesh Rishiyur Chandrasekaran

**Role.** Senior Supply Chain Manager, NACF Supply Chain — Global Capacity Planning team. Job ID 3195330. Bellevue, WA.

**Team mission, in their words.** Plan and deliver North America Customer Fulfillment capacity to meet customer demand at the desired speed and quality while improving the long-term cost structure of the network.

**Their customers.** S&OP, PPT, Finance, FC Operations, WW and Operations Engineering, Real Estate, WW Design, Launch, Safety/EHS. Partners in SCOT, Product, Engineering and Construction.

**Format.** One hour. Expect two to four Leadership Principle stories with hard follow-ups, and one functional block on capacity planning. If this is the hiring manager, the functional block gets longer.

**Your position.** You have designed and consolidated DC capacity, allocated demand across a 3-DC / 100-store network through peak, and turned measured lane data into planning parameters. You have not built an Amazon FC capacity plan. Say so once, early, then show the mechanism.

Read sections 2, 3, and 4 out loud once — section 3 has "tell me about yourself" in the Amazon shape and "why Amazon / why this team / why now." Section 5 is the story bank — know which four you are leading with. Section 7 is the functional block; walk Case A on paper once before Friday.

## 1. What they are actually scoring

### The L6 bar, in one table

| L6 GUIDELINE SAYS                                                                                                                  | WHAT THE INTERVIEWER LISTENS FOR                                                     | WHERE YOU PROVE IT                                                                                                                         |
|------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Ambiguity.</b> Strategy may not be defined. You pick the stakeholders, set the goals, design the long-term solution.            | Did you start from a vague ask and define the problem yourself?                      | Elementum: the ask was a report; you built lane-level lead-time planning. Flexport: inherited an 18-month plan and re-costed the sequence. |
| <b>Scope and influence.</b> A team, sometimes an org or country. Influences 1–2 levels up.                                         | Did you move a decision above your level with evidence, not seniority?               | Roadmap flip ratified in Bellevue after private math with your manager. CIO/COO weekly pack across 60+ stakeholders.                       |
| <b>Execution.</b> Tactical and strategic. Decides where to simplify or extend. Mitigates long-term risk. Trade-offs short vs long. | Can you name the trade-off and the number that decided it?                           | Partner-first vs operator-first (CTS ~50% vs ~25%). Robots at 30% vs the launch date. Mock tolerance vs go-live.                           |
| <b>Complex.</b> Visible risks, roadblocks, constraints. You design an appropriately simple solution.                               | Was the solution simpler than the problem?<br>Did a gate exist before the crisis?    | 0.5% tolerance published before mock one. 98% pick-accuracy gate published before the robots arrived.                                      |
| <b>Designing or modeling the supply chain, not working in it.</b>                                                                  | Are your stories about the network's shape and parameters, or about running a shift? | Allocation across 3 DCs + 100 stores. Two buildings into one omni DC. 200 lanes measured. Safety-stock parameters re-set from data.        |

| L6 GUIDELINE SAYS                                                           | WHAT THE INTERVIEWER LISTENS FOR                | WHERE YOU PROVE IT                                                                                                |
|-----------------------------------------------------------------------------|-------------------------------------------------|-------------------------------------------------------------------------------------------------------------------|
| <b>Process improvement.</b><br>Drives operational-excellence best practice. | Did you leave a mechanism that ran without you? | WBR grammar at Flexport; QA-authored release gate that survived your leaving; carrier escalation and drift flags. |

## The JD, translated into what you say

- **"Take end-to-end ownership of building and manifesting our capacity plan."** Say: a plan is a set of numbers with named owners and triggers, not a deck. You have run that as a WBR.
- **"Balance capacity and demand."** Say: that is allocation at Manhattan and consolidation at Sephora UDC. Demand did not fit the building; you changed where units went, then changed the building.
- **"Define metrics, develop program management mechanisms, invent/simplify processes."** Say: input lever vs output metric, pre-wired reds, COE that changes a mechanism. Flexport Snowflake story.
- **"Craft concise written narratives to influence executive decisions."** Say: the two-column cost-to-serve table that flipped an 18-month roadmap. Offer to write it.
- **"Adherence to lean philosophy at scale."** Say: gates that can stop the line. Andon is not a poster.

## Vocabulary you should use correctly, and what to ask about

| TERM                             | USE IT AS                                                                                                                                                                                                                                                                   |
|----------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| NACF                             | North America Customer Fulfillment — the FC network that picks, packs and ships customer orders.                                                                                                                                                                            |
| Capacity plan                    | Units of demand the network must absorb, by week, by building type, against what each building can store and push out. Two constraints, not one: <b>storage</b> (cube) and <b>throughput</b> (units per hour through inbound, stow, pick, pack, ship). Say both every time. |
| S&OP                             | The weekly plan that reconciles demand forecast, inventory and capacity; the basis of labor planning and steering decisions. Your Elementum forecast-alert story lives here.                                                                                                |
| OP1 / OP2                        | Amazon's annual operating-plan cycles (fall and spring). Long-range capacity decisions — new builds, expansions, retrofits — are argued in these. You can say "the OP cycle" without pretending you have sat in one.                                                        |
| SCOT                             | Supply Chain Optimization Technologies — the science and product org that owns forecasting, buying, placement and planning systems. The interface you would manage: they own the model, you own whether the plan is honest.                                                 |
| PPT                              | Listed as a customer of the team. Do not guess the acronym. Ask: "what does PPT own relative to Global Capacity Planning?"                                                                                                                                                  |
| Sortable / non-sortable / AR     | FC types by item size and automation. Capacity is planned per type; a sortable AR building and a non-sort building are not interchangeable. Say "by building type" rather than "by FC."                                                                                     |
| Launch / Real Estate / WW Design | The teams that put new capacity on the ground. Your UDC and dual-DC stories are launch stories; speak to them in launch gates and freeze windows.                                                                                                                           |
| Speed                            | Promised delivery speed to the customer. Capacity in the wrong place is a speed miss, not just a cost miss. Your promise-equation work maps straight onto this.                                                                                                             |

---

## 2. Say this about the gap. Once.

### UNPROMPTED, INSIDE THE FIRST FIVE MINUTES

"I have not built an Amazon FC capacity plan. I have done the three things underneath one: I allocated demand across three mega-DCs and a hundred stores through peak when the network did not fit the plan; I consolidated two warehouses into one omni facility with robotics and ran the launch gates; and I replaced assumed planning parameters with measured ones across two hundred lanes and released the inventory the bad assumptions were holding. I know the scale here is different. The mechanism is not."

Then stop. Do not list what else you have not done. If they probe the gap later, the answer is in section 8 ("Why you over an internal transfer").

---

## 3. "Tell me about yourself" — Amazon format

Amazon interviewers want this in under two minutes, in three moves: **present** (what you do now, one sentence), **past** (two or three chapters, each with a number and an implied Leadership Principle), **future** (why this seat, and why you are the right risk). No adjectives about yourself. Let the numbers carry the principles. It should sound like a PR/FAQ headline, not a resume read aloud.

### PRESENT → PAST → FUTURE, ABOUT 100 SECONDS

**Present.** "Today I'm the technical program manager for retail fulfillment and inventory at Sephora. I own how inventory is defined, placed and promised across stores, the retail DC and dotcom on S/4HANA, and I ran the consolidation of two warehouses into one omnichannel DC with robotics — up fifteen percent on processing efficiency, zero customer incidents.

**Past.** Before that, three chapters that are really one job. At Flexport I scaled a shipment-visibility platform from seven pilot customers to sixty-eight behind published graduation gates — cost-to-serve from sixty-six cents to three cents a container — and I flipped an eighteen-month roadmap with a two-column cost table: sixty-six percent less engineering effort, five million avoided. At Elementum I replaced carrier-brochure lead times with measured transit distributions across two hundred lanes and planned at the tail — ninety-nine percent lead-time accuracy, 580,000 units and a million and a half of safety stock released. And at Manhattan Associates I owned order allocation and promising across three mega-DCs and a hundred stores, including the 4 AM Cyber Monday when the network didn't fit the plan — ninety-nine percent fill through peak, about eight million protected.

**Future.** The pattern across all of it: I plan capacity and promise from measured data, deliver it through gates that can stop the line, and keep it honest with a weekly review of inputs and outputs. That is how this team's charter reads to me. I have not built an Amazon FC capacity plan, and I won't pretend the scale is the same. The mechanism is. That is what I want to bring here."

**The LP signals buried in it, so you know they are there:** Customer Obsession (zero incidents, promise at the tail), Insist on the Highest Standards (graduation gates), Have Backbone (the roadmap flip), Dive Deep (200 lanes measured), Bias for Action (4 AM), Deliver Results (every number). Do not name the principles out loud. They will.

If the interviewer is the hiring manager and wants the longer chronological version, use the one below. Otherwise the version above is the whole answer.

## Why Amazon. Why this team. Why now.

You will be asked at least one of these. Amazon wants a reason that is specific to the team and honest about you, not praise of Amazon. Keep each under 45 seconds.

### WHY AMAZON

"Three reasons, all specific. First, this is the one network where the thing I have spent fifteen years working on — capacity and promise, planned from data — is the operating system, not a project. Every mechanism I built elsewhere, I built in borrowed grammar: input metrics and output metrics, a weekly review where reds arrive pre-decomposed, a correction-of-error that changes a mechanism rather than asking people to try harder. Here that grammar is native. Second, scale. I have planned capacity across three DCs and a hundred stores and consolidated two buildings into one. I want to do it where a placement decision moves millions of units and the forecast is a model, not a spreadsheet. Third, I have worked next to Amazon's standards from the outside — I studied how Seller Fulfilled Prime turned an adoption problem into an insurance contract with teeth — and I would rather run that kind of mechanism than admire it."

### WHY THIS TEAM — GLOBAL CAPACITY PLANNING

"Because the team's charter has both halves of the trade I care about: meet customer demand at the desired speed, and improve the long-term cost structure. Most roles own one half and lobby against the other. This one owns the dial. And the customers listed — S&OP, Finance, FC Ops, Engineering, Real Estate, Launch — are exactly the rooms I have sat in: I have handed a forecast alert to an S&OP that trusted it, taken a two-column cost table to Finance and let them attack it, and run launch gates that held a building through peak. I know how each of those rooms argues. I want the seat that has to make them agree on a number."

### WHY NOW — AND WHY LEAVE SEPHORA

"I've delivered what I went to Sephora to do. The WMS migration is live, the omni DC is live and fifteen percent more efficient, and the inventory operating model is the standing reference. I'm not running from anything. I'm running toward the version of this problem where the network is the product. If I wait another two years, I will have run the same play a fourth time at retail scale. I would rather run it once at this one."

**Do not say:** "Amazon is the most customer-centric company" (they wrote it; they do not need it quoted back), "I love the Leadership Principles," or anything about compensation or Bellevue in this answer. **If they ask "why not stay in retail,"** the answer is the last two sentences of "Why now."

## The longer chronological opener

Only if the hiring manager asks for the full walk. Chronological, one thread, ends on this role. Do not open with SAP governance; this room is network and capacity.

#### SAY THIS, THEN STOP

"I design and run the capacity and promise side of supply chains I don't fully control, and I run the weekly mechanism that keeps the plan honest. Five chapters, one thread.

Manhattan Associates — distribution order management. Order allocation and promising across three mega-DCs and a hundred stores. The lead time we could put in front of a customer came out of which node we sourced from. Held a 99% fill rate through peak, protected about eight million dollars one Cyber Monday, structural stockouts down twenty percent.

Elementum — supply-chain delivery lead. Replaced carrier-brochure lead times with measured transit distributions across two hundred lanes, promised at the tail. Ninety-nine percent lead-time accuracy, 580,000 units and a million and a half of safety stock released. Onboarded a thousand supplier and carrier routes on EDI, completeness sixty to ninety-five percent.

Flexport — technical program manager. Scaled a shipment-visibility platform from seven pilot customers to sixty-eight behind published graduation gates, cost-to-serve from sixty-six cents to three cents a container. Consolidated four systems into Snowflake. Re-sequenced an eighteen-month roadmap on cost-to-serve math — sixty-six percent less engineering effort, five million avoided.

Sephora — technical supply-chain program manager. Two DCs from legacy ERP to JDA WMS in eight months, on time and on budget. Two warehouses into one omnichannel DC on Blue Yonder plus Locus robotics, up fifteen percent on processing efficiency. And the S/4HANA inventory operating model — one signed definition of available-to-sell across stores, DC and dotcom.

The thread: capacity and promise, planned from data, delivered through gates, kept honest by a weekly review. That is what this team's charter reads like to me. I have not built an Amazon FC capacity plan, and I am not going to pretend the scale is the same. The mechanism is."

If they cut you off at chapter three, let them. The last two sentences are the only ones that must land.

## 4. The sixty-minute clock

| MINUTE | WHAT IS HAPPENING                                             | YOUR JOB                                                                                                                                                |
|--------|---------------------------------------------------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------|
| 0–3    | Intros. They describe the team.                               | Listen for the word they use for their pain: "peak," "utilization," "launch," "forecast," "cost per unit." That word picks your lead story.             |
| 3–7    | "Tell me about yourself." Often followed by "why Amazon."     | Section 3, Amazon format. Land the gap sentence. Keep "why Amazon" under 45 seconds.                                                                    |
| 6–40   | Two to four LP questions. Each gets four or five follow-ups.  | One story per question, 90 seconds, then stop and let them drill. Answer follow-ups in one or two sentences. Say "I" for decisions. Name the data pull. |
| 40–52  | Functional block: build a plan, handle a gap, define metrics. | Section 7. Clarify out loud for 60 seconds before you solve. Say both constraints. Do a math checkpoint. End with owners and triggers.                  |
| 52–58  | Your questions.                                               | Two from section 9. Not four.                                                                                                                           |
| 58–60  | Close.                                                        | Section 10.                                                                                                                                             |

Amazon interviewers write notes the whole time and score after. Pauses are fine. A tight 90 seconds followed by "happy to go deeper on any part" beats a four-minute story every time.

---

## 5. The story bank

Sixteen stories. Each is written to be spoken in about 90 seconds, followed by the drill the interviewer will run: *what did YOU do, why not X, how did you measure, what would you do differently, walk me through the data pull.* Every number is from your resume or your Walmart war book; nothing here is new. The "maps here" line is for your head, not the room — unless they ask.

**Lead with these four** unless their pain word says otherwise: S1 (UDC consolidation), S2 (Cyber Monday allocation), S3 (200 lanes), S7 (roadmap flip). If their pain word is "inventory," "placement," "returns," or "cutover," S10 (S/4HANA) moves into the first four.

### **S1 — Two warehouses into one omni DC, robots throttled at 30%** (Sephora UDC)

**LPs:** Customer Obsession · Insist on the Highest Standards · Bias for Action · Deliver Results. **Maps here:** this is a capacity consolidation and a launch. It is your closest analog to Launch / WW Design / FC Ops work.

"Sephora combined two warehouses — store replenishment and dotcom — into one omnichannel DC on Blue Yonder 2023.1 with Locus robotics. I ran it as a program: fifteen-plus named tracks, sixty-plus stakeholders, about a dozen systems touched — WMS, ESB and Kafka, OMS, finance, the data platform. Dated public test gates — November 22, December 20, January 30 — and a hard October-to-December change freeze because the building had to survive peak. One issues log, 836 rows over fourteen months, every row with an owner and a date. Week two after robots went live, pick accuracy printed 94% against a 98% gate I had published before the bots arrived. I throttled robots to about 30% of volume, kept the launch green on manual process, sat the zone-logic debug, and restored full volume in week three when they cleared the bar. Processing efficiency ended plus fifteen percent. Zero customer incidents. The gate existed so that decision needed no debate."

**What did YOU do?** Published the 98% gate, made the throttle call the day 94% printed, owned the issues log and the freeze.

**Why not push through to hit the automation date?** A gate that cannot stop the program is decoration. I will not ship a customer-facing miss to protect a robotics narrative.

**How did you measure?** Pick-accurate throughput versus the pre-published 98%, hourly by zone, not a building average. Customer incidents: zero. Processing efficiency +15% against the pre-consolidation baseline.

**Differently?** Shadow the bots on one zone for a week before they touched 100% of volume.

**Data pull?** WMS pick-confirm versus gold-standard location, by zone, by hour.

## S2 — Cyber Monday, 4 AM: allocation across three DCs and a hundred stores (Manhattan Associates)

**LPs:** Bias for Action · Are Right, A Lot · Dive Deep · Customer Obsession. **Maps here:** balancing capacity and demand at peak in a closed network — the core of this job's bad weeks.

"As a distribution-OMS consultant I owned order allocation and promising across three mega-DCs and about a hundred stores. The lead time we could show a customer was an output of which node we sourced from — choose the node, and you have chosen the promise. Cyber Monday, 4 AM, a VP call: backorder lines piling up, millions at risk, full diagnostic an hour away. The noise said capacity. I triaged on what we had: ranked orders by promise-at-risk, rebalanced allocation across the three DCs, one bridge, one log, every decision timestamped. The trace later showed the allocation scheduler had died quietly at the daylight-saving change — inventory was fine. Fill rate held at 99% through peak, roughly eight million dollars in sales protected, and structural stockouts were down twenty percent over the engagement. Since then I treat every peak as a different regime with its own plan."

**What did YOU do?** Took the bridge, set the promise-at-risk ranking, made the rebalance calls, kept the timestamped log.

**Why act at 70% certainty?** Waiting an hour for the full diagnostic was a customer miss with a dollar sign. The rebalance was a two-way door; the diagnostic ran in parallel.

**How did you measure?** Fill rate by hour against the 99% line; backorder lines cleared; sales protected estimated from orders that would have missed.

**Differently?** A scheduler heartbeat and a DST checklist before peak. Action without a mechanism is heroics; I do not want a second 4 AM.

**Data pull?** Allocation queue depth by node by hour; scheduler last-run timestamp; promise-at-risk by order.

## S3 — Two hundred lanes: promise at the tail, release the inventory (Elementum)

**LPs:** Customer Obsession · Ownership · Dive Deep · Think Big · Frugality. **Maps here:** turning measured data into planning parameters; the kind of input SCOT models consume and S&OP depends on.

"A planner asked me for a lead-time-variability report. The room had normalized carrier-brochure lead times — one optimistic number per lane, never measured — and the inventory plan ran on them. I treated the brochure as a lie the business was paying for in safety stock. I pulled shipment history across two hundred lanes and built mean, spread and tail per lane, then moved every planning lead time to the tail — the p99 — not the average. Ninety-nine percent lead-time accuracy held. Because the measured spread was tighter than assumed on most lanes, 580,000 units and about a million and a half of safety stock came free. Then a drift flag per lane, so a widening spread triggered a review instead of a silent re-plan. The ask was a report. The product was lane-level planning as a system."

**What did YOU do?** Pulled the history myself, built the distributions, pitched leadership with mockups, productized it, designed the drift flag.

**Why the tail and not the mean?** A lead time set at the mean is a coin flip — in a right-skewed lane, roughly two in five shipments arrive after it. The p99 is the smallest number of days by which 99 of 100 have already arrived. I set it off the empirical quantile per lane, not a normal approximation, because the normal understates the tail.

**How did you measure?** Actual arrival inside the planned lead time, by lane, by week — 99%. Safety stock released only on lanes where measured spread was tighter than assumed.

**Differently?** Start every new domain with a measurement audit in week one instead of discovering the brochure problem mid-project.

**Data pull?** Shipment-level: lane key, ship date, arrival date, carrier, source of the transit number (brochure vs scan).

## S4 — Forecast alerts that S&OP actually used (Elementum)

**LPs:** Invent and Simplify · Deliver Results · Earn Trust. **Maps here:** S&OP is a named customer of this team. This is your S&OP mechanism story.

"Planners learned about shortages when the line stopped. I did not ship a threshold email. I led the SAP ERP ingestion — inventory, production, POs, ASNs, forecasts — into one supply-and-demand view with a daily MD04-to-MB52 reconciliation held at zero drift, so the numbers were trusted before any alert fired. Forecast accuracy went from 70% to 85% at SKU-location-week. The alerts were coverage math with a thirty-second reason line, and false alarms were burned down in public. Live in three months. About two million dollars of safety stock came out — only where both forecast error and lead-time error had improved. The alert entered the weekly S&OP because planners trusted it, not because anyone mandated it."

**What did YOU do?** Owned the ingestion and the recon, wrote the alert logic and the reason-line requirement, ran the false-alarm burn-down.

**Why two signals before releasing stock?** Improving one number is not allowed to justify the dollar. Forecast accuracy alone could be mix; lead-time accuracy alone could be a quiet quarter.

**How did you measure?** Forecast accuracy at the planning grain; recon drift daily; false-alarm rate weekly; dollars released by SKU-location.

**Data pull?** MD04 vs MB52 by plant by day; forecast vs actual by SKU-location-week; alert fired vs stockout realized.

## S5 — Austria: the noise said carriers, the data said measurement (Elementum)

**LPs:** Dive Deep · Are Right, A Lot · Earn Trust. **Maps here:** capacity data is full of averages at the wrong grain. This is your "I changed my mind because of data" story.

"I walked into Austria convinced we had a carrier performance problem — the noise said carriers were late everywhere. I pulled three years of shipment history myself. The data said something more embarrassing: we could not measure lateness. True data quality was about 55% at the planning grain against the 40% the reports claimed — missing rates, lanes keyed at the wrong grain, carrier names in five spellings, transit times copied from brochures. The reports were wrong in both directions because they were not measuring the same thing. So the program flipped from 'fix carriers' to 'fix measurement.' The routing engine waited for the data. Quality went to 99% by defect class before scale. The savings method was signed with finance before a dollar was claimed. About eight million dollars in freight savings across 25 locations, every dollar auditable to a shipment."

**What did YOU do?** Pulled the history myself; did not send an analyst. Changed the program's first milestone from engine to measurement.

**Why not trust the reports?** They were at the wrong grain. A routing recommendation on 55% data is a coin flip with a confident interface.

**How did you measure?** Completeness by defect class; savings method signed before claiming; audit trail to shipment.

**Data pull?** Shipment-level: rate present, lane key, carrier spelling, transit source.

## S6 — Seven to sixty-eight customers behind gates that could say no (Flexport)

**LPs:** Insist on the Highest Standards · Deliver Results · Frugality · Success and Scale Bring Broad Responsibility. **Maps here:** scaling a program without dumping cost on the next team — the "lean at scale" line in the JD.

"I scaled a global 4PL shipment-visibility platform from seven pilot customers to sixty-eight, wave by wave, with six published graduation criteria a wave had to clear before it boarded: data quality at or above 90% against source of truth, milestone and ETA accuracy at or above 80%, verification SLA, engagement floors, stability, CSAT. I held wave three in front of everyone on an ETA-accuracy miss from one bad source while Sales wanted the pool boarded. We landed sixty-eight customers two months early. Cost-to-serve went from sixty-six cents to three cents a container — a 95% cut — while the base grew about ten times, and ticket response went from 72% to 90% at five times the volume. Unit cost falling while the book grows only happens if the gates keep bad data out."

**What did YOU do?** Wrote and published the six criteria; held the wave; named the source.

**Why a public hold?** A public hold teaches the mechanism. A private hold teaches politics.

**How did you measure?** Criteria table per wave; cost-to-serve per container; ticket response; 68 against target and date.

**Differently?** Put the bad source's kill switch into the platform earlier.

**Data pull?** Per-customer ETA accuracy against the 80% line, by data source.

## S7 — Flipping an eighteen-month roadmap with a two-column table (Flexport)

**LPs:** Have Backbone; Disagree and Commit · Think Big · Frugality · Earn Trust. **Maps here:** the JD's "concise written narrative to influence executive decisions." This is the story for it.

"I inherited an eighteen-month, operator-first launch plan that was already socialized with leadership. I costed both sequences on cost-to-serve. Partner-first cut total cost about 50% versus about 25%, and shipped value in six months instead of eighteen. I did not flip it in a meeting. I took the two-column table to my manager privately, he pressure-tested it with engineering for two days, and the Bellevue onsite ratified it. Result: 66% less engineering effort, five million dollars of cost avoidance, and value in six months. I did not win the argument with seniority. I won it with a table."

**What did YOU do?** Built the CTS table; took it private first; let engineering attack it.

**Why not both sequences in parallel?** We did not have the engineers. Sequencing was the bet.

**Why not go around your manager?** Backbone is the private table. Theater is the public ambush. Commit after the room is the principle.

**How did you measure?** Effort -66%; \$5M avoided; time to value 6 vs 18 months.

**Differently?** Price the customer-side cost (speed, breach) explicitly, not only cost-to-serve.

## S8 — Four systems into one pipeline: the WBR is only as good as its data (Flexport)

**LPs:** Invent and Simplify · Dive Deep · Frugality. **Maps here:** "define metrics, develop mechanisms." Capacity reviews die on manual pulls.

"Flexport's Monday review ran on four systems stitched together in Excel by me — two hours a week, 60% system accuracy, a single point of failure. I wrote the BRD, froze seven metric definitions exactly as the team already used them, priced doing nothing at over a hundred thousand a year in the infrastructure team's own rates, and gated cutover on a parallel run — manual versus pipeline within 1% for consecutive weeks. Four systems became one Snowflake pipeline. Accuracy went from 60% to 90%. The first automated catch was a 4% ticket-sync drop the manual era had never seen. The weekly cadence only works if the data under it is boring."

**What did YOU do?** BRD, frozen definitions, priced the status quo, parallel-run gate.

**Why not a better spreadsheet?** Excel was the single point of failure. The invention was the frozen contract plus the pipeline, not the tool.

**What went wrong?** I did not plan the data-access security review; ticket data carries customer content. It cost two weeks. I owned it upward immediately and re-sequenced — the parallel run kept Monday alive on the manual path. Every data program since has "whose approval does this data need, and what is its lead time" on the kickoff checklist.

**Data pull?** Seven frozen definitions; week-by-week parallel delta.

## S9 — Dual-DC WMS migration in eight months, tolerance published before mock one (Sephora JDA)

**LPs:** Learn and Be Curious · Deliver Results · Ownership · Earn Trust. **Maps here:** launch discipline. One-way doors get rehearsals against a published number.

"Sephora handed me a dual-DC migration from legacy ERP to JDA WMS — I had never run one — on a hard eight-month clock with a twenty-million-dollar ROI pipeline of four programs gated behind it and nothing documented. I did not fake experience. I sat the sponsor for objectives, assumptions and success criteria, held discovery's written exit until every flow had a current state, then entrance and exit criteria for every phase. Three mock cutovers against a 0.5% inventory reconciliation tolerance published before mock one. Misses went 7%, then 1.8%, then 0.3%. The go/no-go read the scorecard, not the mood. On time, about 3% under the \$1.8 million budget, four-week hypercare with numeric exit criteria. Mock one at 7% went up unvarnished with my name on it."

**What did YOU do?** Wrote the tolerance with finance before mock one; read 7% as a scorecard, not a verdict; made the go/no-go.

**Why not hire a WMS specialist and step back?** I did use specialists. I owned the gate.

**What slipped?** Training sign-off — a go-live exit criterion — slipped 1.5 weeks. I had scored the RF-gun conversion as low effort; my own pre-training test caught the upgrade script failing on newer hardware, and COVID night shifts killed the recovery window. I told leadership it was my miss. Program still landed on time. The rule since: a single-point dependency in front of an exit criterion gets a what-if plan regardless of size.

**Data pull?** Unit-level reconciliation, expected vs actual by location, after each mock; RAID ageing.

## **S10 — S/4HANA: one definition of "available," returns as first-class stock, and a cutover with a numeric rollback** (Sephora S/4HANA ↔ ISM)

**LPs:** Ownership · Invent and Simplify · Earn Trust · Dive Deep · Deliver Results · Are Right, A Lot (being wrong).

**Maps here:** this is the JD's "optimize and manage the flow of inventory within the network," literally. Returns and quality-hold are reverse flow; they consume storage and throughput and are usually invisible to the plan until they are promised to a customer. It is also your cleanest "I was wrong" story and your cleanest one-way-door cutover story.

"Sephora was standing up S/4HANA as the inventory system of record for stores, the retail DC and dotcom, with a warehouse execution system beside it. The problem was not the software. 'Available' meant something different in every system, because each definition protected a different number — fill rate for stores, labor for the DC, the customer promise for dotcom, the credit memo for finance. Customer returns, return-to-vendor stock, quality-hold and lost-and-found were being promised to customers because somewhere they counted as available.

I redesigned the inventory operating model. Every stock state — available-to-sell, reserved, in-transit, quality-inspection, blocked, lost-and-found — got one definition, one owner, and one rule about who may promise against it, signed by the business before anyone configured anything. Returns got a plant of their own. Return-to-vendor stock-transport orders were modeled as pick-reserve so credit could never run ahead of the physical unit: receive, grade, then restock, scrap or RTV. Nothing in that path was promiseable until disposition.

I ran it as a program, not a configuration: a RAID log, a weekly CIO/COO pack across sixty-plus stakeholders on frozen metrics so a red stayed red, and a 343-task weekend cutover runbook with rollback written as a number before the weekend — interface sync under 99% for four hours, or p95 latency over twice budget for one hour. It never fired. The legacy inventory system, MCS, came out only after the new definition held in production. Zero customer-facing incidents on the staged go-lives.

One rule I wrote was wrong. I forbade every channel from promising against units on a truck. It felt safe. It starved dotcom of sellable units while physical inventory in transit rose. The data showed the miss, and the fix went through the same signed policy — a defined in-transit promise window per lane, from real lane data. Conservatism has a price, paid in the same currency as risk."

**What did YOU do?** Wrote the stock-state definitions and took them to steering with the failure mode attached. Designed the returns plant and the pick-reserve RTV model. Owned the RAID, the weekly pack, the runbook, and wrote the rollback triggers. Made the MCS retirement call.

**Why a signed policy before configuration?** Most oversell "bugs" are vocabulary bugs. A config change nobody signed becomes next quarter's argument. Once the rule was signed, the design arguments got small — the plant and the STOs were just the implementation of a sentence everyone had already agreed to.

**Why rollback as a number?** A rollback that is a paragraph gets debated at 2 AM by tired people. A number gets called by whoever is on the bridge, no meeting. The triggers existed so the weekend could start.

**How did you measure?** Binary, in production: could a unit in quality-hold still be promised? It could not. Sync and latency against the published triggers through the weekend. Customer-facing incidents: zero. MCS off.

**How did you know the in-transit rule was wrong?** Sellable units by channel fell while units in transit by lane rose. Two lines diverging is the signal. Fix: a per-lane in-transit window from measured lane transit — the same idea as the 200-lane work, applied to promising instead of planning.

**Differently?** Set the in-transit window from lane data on day one instead of starting from "never." And sample the OMS-to-WMS handshake daily in hypercare from day one, not after the first discrepancy.

**Data pull?** Stock by state by plant by day; promise log joined to stock state at promise time; sellable units by channel against in-transit units by lane; interface sync rate and p95 latency by hour through cutover.

**If they ask how it maps to capacity:** "Returns are inbound volume the plan did not forecast. Quality-hold is storage you cannot sell. If the capacity plan counts them as available, it is overstating usable capacity and the promise is lying. Giving each state a definition and an owner is the same move as giving each building type its own ceiling."

## **S11 — Earn trust: the \$2.5M discrepancy I refused to fix quietly (Manhattan)**

**LPs:** Earn Trust · Customer Obsession.

"On an OMS-to-WMS handshake we found a discrepancy across a couple hundred orders, roughly two and a half million dollars. The manager's instinct was to correct it quietly. I argued for full transparency: show the customer, let them prioritize, bulk-patch the class. It was uncomfortable. They protected the orders that mattered to their business, and we later won a larger initiative with them. Silent corrections are loans against trust. Same instinct at Flexport: a truncated week printed 99.4% automation and I discarded it, because a VP steering on three data points is a decision-quality problem."

## **S12 — Giving the skeptic the gate (Sephora agile conversion)**

**LPs:** Hire and Develop the Best · Strive to be Earth's Best Employer.

"A VP sponsored converting a twenty-person product, engineering and QA group from a six-week to a four-week release. The QA lead was the most skeptical person in the room. I gave her the load-bearing artifact — the definition of done with entry criteria — and coached her on how to write a gate that can say no. Within two cycles she was defending it to engineering. Critical defects dropped by half. Neighboring teams copied the model, and the ceremony survived my leaving. Development that depends on me staying is not development."

**Honest shape of your people leadership, if asked directly:** direct reports in the classic sense at MetricStream — onshore/offshore delivery teams, SOWs, the go-live chair. Everywhere else, matrix: the 20-person conversion, UDC with 15+ tracks and 60+ stakeholders, about 100 people across nine teams sequenced at Flexport. Do not dress matrix leadership up as a headcount you did not hold.

### S13 — Two go-lives, one engineering team, no 24-hour days (Sephora)

**LPs:** Strive to be Earth's Best Employer · Deliver Results.

"Two programs compressed onto one engineering team, roughly two hundred thousand dollars a month per DC at stake. One RAID log, one 45-minute daily forum, and a two-pass war room so issues moved around the clock without anyone working around the clock. Both landed; about a million dollars per DC protected. Heroics are a process miss."

### S14 — Teaching myself a stack: the LLM triage agent (independent build)

**LPs:** Learn and Be Curious · Frugality · Insist on the Highest Standards.

"I built an LLM triage agent — Python, FastAPI, tool use — with the evaluation set written before the agent. Every prompt or model change was gated on the known past still passing. Cost per incident went from about \$25 to \$0.40, a 98% cut, with accuracy held and 100% citation coverage. Same operating system as everything else I run: instrument first, then let the scorecard teach."

Use this only if asked about learning something new or about AI. It is a side build, not a program you owned at a company. Say that.

### S15 — A scope cut that was priced, published, and scored (Flexport FF App export launch)

**LPs:** Deliver Results · Bias for Action · Customer Obsession.

"The export launch was minimum-scope by design: partner path first, Air and LCL waited a year, operators kept their old screens two more quarters. Real cuts, real complaints. What made it defensible: the success metric was set before scope — monthly operator interactions eliminated, not features shipped — so every cut was scored against that number, and a burn-up tripwire kept the door shut afterward: new work estimated and swapped, never absorbed. MVP landed April 2022, eleven partners, ten countries, and 48 of 53 import partners migrated within a month."

### S16 — A thousand routes, one gate (Elementum EDI)

**LPs:** Insist on the Highest Standards · Deliver Results.

"I program-managed about a thousand supplier and carrier routes onto X12 EDI with SAP PO and ASN integration. Cohorts graduated only at 95% or better valid ASNs. Shipment completeness went from 60% to 95%. Routing-guide skeptics were often right about specific lanes, so overrides carried reason codes — data, not sabotage."

## 6. Leadership Principles — which story, which hook

They will not say the LP's name. They will say "tell me about a time." Hear the LP, pick the row, speak the hook, tell the story, stop.

| LP                 | LEAD /<br>BACKUP | THE HOOK SENTENCE                                                                                                                                                                                                                                 |
|--------------------|------------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Customer Obsession | S3 / S1          | "I treated the brochure as a lie the customer was paying for in inventory." / "I will not ship a customer-facing miss to protect a robotics narrative."                                                                                           |
| Ownership          | S3 / S10         | "Nobody owned the seam between brochure lead times and the plan. I took the seam, not the ticket." / "Nobody owned what 'available' meant. I wrote the definitions, took them to steering with the failure mode attached, and owned the cutover." |

| LP                                                 | LEAD /<br>BACKUP | THE HOOK SENTENCE                                                                                                                                                                                                                                                  |
|----------------------------------------------------|------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Invent and Simplify                                | S8 / S10         | "New metrics without killing old ones is complexity wearing a badge. I froze seven definitions and killed the spreadsheet." / "Most oversell bugs are vocabulary bugs. One definition, one owner, one rule per stock state — then the design arguments got small." |
| Are Right, A Lot                                   | S2 / S10         | "70% is enough when waiting is a customer miss." / "I forbade promising against in-transit units. I was wrong, the data said so, and the fix went through the same signed policy."                                                                                 |
| Learn and Be Curious                               | S9 / S14         | "I had never run a WMS migration. I did not fake it. I built the gate and used the specialists."                                                                                                                                                                   |
| Hire and Develop the Best                          | S12              | "I gave the skeptic the gate. Development that depends on me staying is not development."                                                                                                                                                                          |
| Insist on the Highest Standards                    | S6 / S1          | "I held wave three in front of everyone. Standards without a stop-the-line are posters."                                                                                                                                                                           |
| Think Big                                          | S3 / S7          | "The ask was a report. The product was lane-level planning as a system." / "Think Big without sequencing is a slogan."                                                                                                                                             |
| Bias for Action                                    | S2 / S1          | "Timestamped decision, then a post-mortem that makes silence an alarm." / "Throttling the robots was a two-way door we had built. I turned it the same day." <i>One-way doors (S/4HANA cutover, JDA cutover) get rehearsals and a numeric rollback instead.</i>    |
| Frugality                                          | S6 / S7          | "Cost-to-serve down 95% while the book grew ten times. Headcount would have scaled the waste."                                                                                                                                                                     |
| Earn Trust                                         | S11 / S10        | "Silent corrections are loans against trust." / "Sixty stakeholders, one weekly pack, metrics frozen so nobody could re-litigate a red by changing the chart."                                                                                                     |
| Dive Deep                                          | S5 / S8          | "I pull until the unit is a shipment, an order, a lane or a building — not a national average."                                                                                                                                                                    |
| Have Backbone; Disagree and Commit                 | S7               | "Backbone is the private table. Theater is the public ambush."                                                                                                                                                                                                     |
| Deliver Results                                    | S9 / S10         | "Results with a rollback are not results. The gates are why it counted." / "343-task runbook, rollback as a number, never fired, legacy system off, zero customer incidents."                                                                                      |
| Strive to be Earth's Best Employer                 | S13 / S12        | "Heroics are a process miss."                                                                                                                                                                                                                                      |
| Success and Scale<br>Bring Broad<br>Responsibility | S6               | "We scaled ten times and ticket response improved. The cost of scale was paid in gates, not in someone else's backlog."                                                                                                                                            |

### The follow-up discipline

- **"What did you do, specifically?"** One sentence that starts with "I." Name the artifact: the gate, the table, the log, the SQL.
- **"Why not [the obvious alternative]?"** Name the cost of the alternative in the same currency as the result.
- **"How did you measure?"** Metric, grain, baseline, target. Never "we saw improvement."
- **"What would you do differently?"** A real change you have since made. Not "nothing" and not "communicate more."

- **"Walk me through the data pull."** Tables joined, grain, filters. If you did not pull it yourself, say who did and what you checked.

## 7. The functional block — capacity planning

You will get at least one of these. Format they expect: clarify out loud (60–90 seconds) → structure (60 seconds) → a math checkpoint → recommendation with owners, triggers and risks. Say "both constraints — storage and throughput" early in every answer. Label every number you make up as illustrative, and say what you would replace it with on day one.

### Primer: the shape of a capacity plan

|                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        |                                                                                                                                                                                                                                |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Demand side                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            | units by week × building type × speed tier, with a peak factor (forecast owner: S&OP / SCOT; your job is to interrogate its grain and mix)                                                                                     |
| Required capacity needed                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                               | storage: units on hand ≈ daily outbound units × days of cover → cube<br>throughput: units/week each process step must push (inbound, stow, pick, pack, ship)<br>both, by building type; the binding one is the one you plan to |
| Available capacity before 100%)                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        | buildings × design rate × a utilization ceiling (congestion starts minus planned downtime, retrofits, launches still ramping)                                                                                                  |
| Gap by week                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                            | required – available, by region, by type; where and when it binds                                                                                                                                                              |
| Levers, cheapest and fastest first <ol style="list-style-type: none"> <li>1. flow: move units – placement / allocation to buildings with slack (days to weeks)</li> <li>2. hours: shifts, overtime, labor plan (weeks)</li> <li>3. inventory policy: pull inbound forward or push it back; trim slow movers (weeks)</li> <li>4. overflow: 3P / temporary space (weeks to months)</li> <li>5. retrofit / automation: more throughput in the same cube (quarters)</li> <li>6. expansion / new build: more cube (OP cycle; years)</li> <li>7. demand shaping: speed tier, promise, selection (the lever nobody wants to admit)</li> </ol> |                                                                                                                                                                                                                                |
| Plan                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   | named owner per lever, trigger that fires it, review cadence, and what you will not do                                                                                                                                         |

### Case A — "Build next year's capacity plan for a region."

**Clarify (say these out loud).** Which building types are in scope? Is the demand forecast handed to me or do I own it? What is the peak factor and what week does it peak? What are the policy ceilings on utilization? Which launches and retrofits are already committed? What is the planning grain — week, building, type?

**Structure.** The primer, top to bottom. Spend the most time on the gap-by-week and the lever ranking.

**Math checkpoint (illustrative).** "If a building ships 1.0M units a week at average and the peak factor is 1.6, peak week is 1.6M. If its throughput design rate is 1.4M and the ceiling I am allowed to plan to is 90%, usable is 1.26M. The gap is 340K units in peak week. Flow first: how much of that can neighboring buildings absorb inside their own ceilings? Then hours. Overflow only for what is left. I would replace every one of those numbers with the S&OP demand file and the engineering design rates on day one."

**Recommendation.** A one-page plan: gap by week by region, the lever stack with owners, the triggers that promote a lever (e.g., "if the week-8 forecast is still 5% above plan, overflow contracts sign"), and the two things you will not do — plan to 100% of design rate, or pad every building's forecast to hide a placement problem.

**Risk:** forecast mix shifts between types; you plan by type, not by total.

### Case B — "Peak forecast just came in 10% above capacity. Go."

1. **Verify the number before you spend on it.** Grain and mix: is the 10% real demand, or a type-mix shift that makes one building type bind while another has slack? Same cohort weights as last forecast?
2. **Find where it binds.** Storage or throughput? Inbound dock, stow, pick, pack, or outbound? One region or network? "A network 10% is usually three buildings at 25% and the rest fine."
3. **Cheapest lever first, with a trigger.** Placement and allocation shifts this week. Labor plan next. Pull slow-moving inbound back; pull fast-moving inbound forward. Overflow for the residual. Speed tier only if the customer promise is the last protection.
4. **Publish the plan with a date on every lever** and a freeze on non-essential change into peak — the UDC October-to-December freeze.
5. **What you will not do.** Spread a national pad across every building. That hides the cell that is actually breaking and spends cost everywhere.

**Story to attach:** S2 — rank by promise-at-risk, rebalance across nodes, timestamped log, then a mechanism so you never discover the regime in November again.

### Case C — "Utilization looks fine but speed is missing."

Utilization is a weekly average. Speed misses live at the hour, the dock door, the sort type, the cutoff. Change the grain: by building by hour by process step. Check mix: did fast-promise demand land in buildings far from the customer? Check cutoff cliffs: units arriving an hour late lose a calendar day. **Story:** S5 — the reports were wrong in both directions because they measured the wrong grain. **Line to say:** "If I am still talking about network utilization after two follow-ups, I am failing the question."

### Case D — "Finance says the new building is not justified."

Do not argue. Price both columns. Cost of the build, annualized. Cost of not having it: overflow contracts, transfers and split shipments, speed downgrades and their conversion cost, labor premiums, the risk-adjusted cost of a peak miss. Show the sensitivity: at what forecast does the build break even? Offer the middle path — retrofit or flow redesign — priced in the same table. **Story:** S7 — the two-column table that flipped an eighteen-month plan; Finance attacked it and it survived. **Line:** "I would rather lose this argument to a better number than win it with a louder voice."

### Case E — "Define the metrics for a network capacity program."

Outputs (what the business feels): capacity delivered vs plan by week; utilization vs ceiling by type; cost per unit; customer speed attainment. Inputs (what a named owner can move this week): launch and retrofit milestones on dated gates; placement compliance; forecast accuracy at the planning grain; labor plan attainment; inbound ASN accuracy. Mechanism: a weekly review where reds arrive pre-decomposed — which input moved, whose lever, what correction, verified by the same number next week — and a correction-of-error write-up for any miss above a published threshold that ends in a mechanism change, not "try harder." **Line:** "Inputs without outputs are a hobby. Outputs without inputs are weather."

### Fermi, if they throw one

- State assumptions. Compute out loud. Label it illustrative. Say what you would replace it with.
- *"How many units does a large sortable FC ship in a week?"* Structure: shifts × hours × headcount × units per hour, cross-checked against outbound dock capacity. Do not quote an Amazon number you have not seen. "I

would take the design rate from Ops Engineering on day one."

- *"How much storage for a week of demand?"* Daily outbound × days of cover ÷ units per cubic foot ÷ bin utilization. Say the structure, pick round numbers, flag the two you are least sure of.
- *"Cost of a 4-point speed miss at peak?"* Volume × delta × cost per late unit, plus the unpriced trust cost. Name the grain. Never say "huge."

### Lines that read as judgment, not worship

- "I plan to a ceiling, not to design rate. Congestion starts before the building is full."
- "A network average hides the three buildings that are actually breaking."
- "Flow is the cheapest capacity in the building. Steel is the most expensive. I spend them in that order."
- "The forecast is an input I interrogate, not a fact I plan to."
- "Every lever gets a trigger and an owner before peak, or it is a wish."

---

## 8. The hard questions

### "You have never built an Amazon FC capacity plan. Why you?"

"Correct, and I will not pretend otherwise. An internal transfer brings Amazon scar tissue. I bring the three things a capacity plan is made of, each proven: allocating demand across a network when the plan did not fit — 99% fill through peak; consolidating buildings and running a launch behind gates — plus fifteen percent, zero incidents; and turning measured data into planning parameters — two hundred lanes, a million and a half released. I also bring the review grammar this team already runs — input levers, output metrics, pre-wired reds. Days one to thirty I learn your numbers, your building types, and the five decisions last peak that nobody wants to talk about. I will not arrive trying to rebuild it. I will arrive with the mechanism."

### "This role is in Bellevue."

Decide before Friday and say it in one sentence. Do not improvise a relocation answer in the room. [Your line here.]

### "Tell me about a failure."

Pick one, own it in the first sentence, name the mechanism it built, stop. Three ready:

- **UAT that crashed under load (Elementum).** Five users, ten scenarios; I sent screenshots instead of digging. Concurrent load crashed it; sign-off missed; sponsor escalation. The implicit requirement — performance at the customer's shape — was not in the entrance gate. Every critical UAT since has numeric load criteria; Flexport's launch proof was 500 containers in 20 minutes at 99.6% connectivity before a customer touched it. "I was the mechanism that failed; I wrote the mechanism that replaced me."
- **RF devices, 1.5-week training slip (Sephora JDA).** In S9. Single-point dependency in front of an exit criterion now gets a what-if plan regardless of size.
- **Security review, two weeks (Flexport Snowflake).** In S8. "Whose approval does this data need, and what is its lead time" is now a kickoff question.

### "Tell me about a time you had to influence without authority, two levels up."

S7. Then one sentence: "I set the instrument, not the person — a named line, a numeric target, a published grain, a date. Next Monday's number is the appeal."

## "How do you set goals for people who don't report to you?"

"I set the instrument, not the person. A named input line, a numeric target, a published grain, and a date the number gets checked. On UDC the issues log made it unemotional: an owner's row either moved or it did not. I have never won a goal argument with a title; I have won them with a line on the page everyone reads."

## "What would you do in your first 90 days?"

"Thirty days: learn the plan as it actually exists — the demand file, the design rates, the ceilings, which lever fired last peak and why. Sixty: take over the weekly review's output block and pre-wire the reds. Ninety: one output metric moved by a mechanism the review can audit — a trigger that fires a lever before the week it is needed, or a placement change that bought capacity without steel. Not a new dashboard."

## "Are you an L6 or an L7?"

If asked, do not oversell. "L6 is running the plan and the review well, independently, on complex problems with real trade-offs. That is what I have done. I have also flipped a socialized roadmap with a table, published tolerances that stopped go-lives, and throttled my own launch. I have not held a large direct-report org, and I will not pretend that is my differentiator."

---

## 9. Your questions — pick two

1. "When the capacity plan was wrong last peak, how long did it take to know which input was wrong — the forecast, the design rate, or the placement?" (*Shows you think in decomposition. Whatever they name, you have the matching story.*)
2. "Where does the storage-versus-throughput call actually get made today — in the model, in a planning table, or in a negotiation between S&OP and Ops?" (*The answer is almost always "it's complicated," which is your opening.*)
3. "What does PPT own relative to this team, and where do the two of you most often disagree?"
4. "Which lever is the team least willing to pull — flow, hours, overflow, or the promise — and why?"
5. "What does a good first ninety days look like to you, in one output metric?"

---

## 10. Close — fifteen seconds

"I want this work. I have not built an Amazon capacity plan, and I have said so. I have allocated a network through peak, consolidated buildings behind gates, and turned measured data into the parameters a plan runs on. In ninety days I will move one output with a mechanism your review can audit. If that is the job, I am here for it."

---

## 11. Numbers worth memorizing

All on your resume or in your war book. If you cannot say what sits under one, leave it out.

| NUMBER                                                   | WHAT IT IS                                                         |
|----------------------------------------------------------|--------------------------------------------------------------------|
| 3 mega-DCs, ~100 stores, 99% fill, ~\$8M, ~20% stockouts | Manhattan allocation and promising through peak; Cyber Monday 4 AM |
| 200 lanes, p99, 99% LT accuracy, 580K units, ~\$1.5M     | Elementum lead-time distributions replacing brochure numbers       |

| NUMBER                                                                           | WHAT IT IS                                                                                                |
|----------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------|
| 70% → 85%, ~\$2M, zero recon drift                                               | Elementum forecast accuracy and safety stock; MD04/MB52 daily recon                                       |
| 40% / 55% / 99%, ~\$8M, 25 locations                                             | Austria: reported vs true data quality; freight savings                                                   |
| ~1,000 routes, ≥95% valid ASN, 60% → 95%                                         | EDI onboarding gate; shipment completeness                                                                |
| 7 → 68, \$0.66 → \$0.03, 72% → 90% at 5×, 2 months early                         | Flexport visibility platform; six graduation criteria                                                     |
| ~50% vs ~25%, 66%, \$5M, 6 vs 18 months                                          | Roadmap flip on cost-to-serve                                                                             |
| 4 → 1, 60% → 90%, >\$100K/yr, 1% parallel run, 4% first catch                    | Snowflake WBR pipeline                                                                                    |
| 8 months, 7% → 1.8% → 0.3% vs 0.5%, ~3% under \$1.8M, \$20M                      | Dual-DC JDA migration; mocks; pipeline unblocked                                                          |
| 15+ tracks, 60+ stakeholders, 836 rows, Nov 22 / Dec 20 / Jan 30                 | UDC omni DC program; test gates; Oct–Dec freeze                                                           |
| 343 tasks; sync < 99% × 4 h or p95 > 2× × 1 h; never fired; MCS off; 0 incidents | S/4HANA ↔ ISM cutover runbook and rollback triggers; legacy inventory system retired                      |
| Returns plant; RTV STOs as pick-reserve; receive → grade → restock / scrap / RTV | S/4HANA stock-state model: nothing promiseable until disposition; credit never ahead of the physical unit |
| 94% vs 98%, ~30%, week 3, +15%, zero incidents                                   | Locus robots throttled; processing efficiency                                                             |
| 6-week → 4-week in 6 weeks, defects –50%                                         | Agile conversion; QA owned the gate                                                                       |
| ~\$2.5M, a couple hundred orders                                                 | OMS–WMS discrepancy disclosed, not hidden                                                                 |
| ~\$200K/month per DC, ~\$1M per DC                                               | Two compressed go-lives on one team                                                                       |
| \$25 → \$0.40, 98%, 100% citations                                               | LLM triage agent (side build)                                                                             |
| 11 partners, 10 countries, 48 of 53 in a month                                   | FF App export MVP, April 2022                                                                             |

Leave "\$27M+ cost savings" off the table unless you can show the addition live.

## 12. Traps

**Do not quote an Amazon internal number.** You do not know their design rates, ceilings or peak factors. Label every number illustrative and say what replaces it on day one.

**Do not say "we" for a decision.** "I published the gate. I made the throttle call. I wrote the tolerance."

**Do not stack stories.** One question, one story, one number. Offer depth; do not volunteer it.

**Do not let Situation eat the clock.** Fifteen seconds of context. Seventy percent of the story is what you did and what it measured.

**Do not open with SAP governance.** This room is network and capacity. S/4HANA (S10) is a strong answer to inventory-flow, returns, cutover, and being-wrong questions, and a diluted answer to "tell me about yourself." Use

it when they ask about inventory; do not lead with it when they ask about capacity.

**Do not invent a headcount.** Matrix leadership, said plainly, is fine at L6. A fake org chart is not.

**Do not solve the case before clarifying.** Sixty seconds of questions out loud is scored, not tolerated.

**"What would you do differently" needs a real answer.** You have five: shadow one zone first; scheduler heartbeat; measurement audit in week one; dual-path devices; security lead time on the kickoff checklist.

**Decide the Bellevue answer before Friday.**

---

### **13. The morning-of card**

1. Read "tell me about yourself" and "why Amazon" (section 3) out loud once. End on the gap sentence.
2. Tell S1 and S2 to a wall, with numbers, in under two minutes each.
3. Walk Case A on paper: demand → two constraints → gap by week → lever stack → owners and triggers.
4. Say the five judgment lines in section 7 once.
5. Pick your two questions. Pick your Bellevue sentence.