FRIDAY OCT 30 · APPLE SOURCE DELTA · SAY-ALOUD ADD-ON

Amazon NACF — Apple Workbook Delta

Rajesh Rishiyur Chandrasekaran · Rocky Goforth 10:00 · James Cabreros 11:00 · Venkata Josyula 12:00 PDT

What this is. A subtraction, not another master packet: Apple PM III workbook v11 compared with the Amazon master and the panel files present during build. It preserves the Apple-only operating framework, failure detail, probing language, and three compact stories that are not standalone stories in the Amazon master.

Three opening sentences.
1. “For fifteen years I have run programs the same way: baseline first, one scorecard, numeric gates, first-principles triage, then corrections in the same cadence.”
2. “I will be specific about what I owned, what engineering or operations owned, and whether a number is a record, a document figure, or a model.”
3. “The part of every story I care most about is what I would do differently, so please ask me that on any example.”

Freshness and allocation

Interviewer	Use first	Do not lead with
Rocky — infrastructure + Bar Raiser	G1 concurrent-load UAT failure; the five-part program framework; “test at the customer’s shape.”	S/4/ISM, Flexport estimation, safety stock, lane lead-time, or Shipment Visibility as a primary.
James — placement architecture + process	B3 two-program collision; one ruler, one forum, written handoff; process governance.	Already-used S/4/ISM, estimation, safety stock, lane lead-time; use them only as a sentence of proof.
Venkata — FBA Principal PM + behavioral	I1 triage product; customer-safe evidence; G2/G3 if asked for another failure. Shipment Visibility is allowed for the product lens.	Do not turn I1 into an employer deployment: it is an independent build.

Basis labels. record Apple v11 and/or resume; record-from-docs primary program artifact; user anchor earlier Rajesh packet only; modeled arithmetic or assumptions, not banked value. No bare dollar: always attach its noun.

Five numbers to keep exact. Customers **7 → 68 = 9.7×**; CTS **$0.66 → $0.03 per container, about 95%**; FF App **about $5M of avoided cost**, said once; RF fleet **“hundreds of devices”**; UAT **5 users × 10 scenarios, crash 2 days before sign-off**.

Primary sources: Apple_PMIII_Interview_Stories_v11_2026-08-26_7241.pdf/text; Rajesh_ChandrasekaranApple resume; Amazon master HTML; NUMBERS-RECON CANON v2; panel files present during build. All Apple page cites are PDF pages.

1 · Apple workbook map — every part and chapter

"In master" means the content is materially absorbed into that Amazon story, even if Apple's chapter boundary disappeared. "Panel" names a current in-flight or existing interviewer packet where the angle is already developed. **APPLE-ONLY** means no standalone Amazon-master story.

Part / ID	Apple chapter, one-line job	Amazon master	Panel / status
Part 0	Opening, five-part "How I run a program," TMAY, Why Apple, AIQB coverage.	No standalone equivalent; mechanisms distributed across stories.	**APPLE-ONLY framework**; use with Rocky and James.
A1	Shipment Visibility wave machine: calibration, six graduation criteria, 7 → 68.	S11	Josyula existing; master-primary family.
A2	WBR math: grain/date discipline, automation and cost-to-serve.	S11	Josyula / Hoerger family; exact integrity phrases remain delta.
A3	Four systems to Snowflake; BRD, frozen metrics, parallel run.	S16	Josyula S16; G3 failure remains useful.
A4	Legacy decommission: normalize unlike comparisons, then improve the residual.	S11 depth	Josyula Crux bench; not fresh as a full story.
A5	Hundred-container launch gate with bug IDs and two go/no-go checkpoints.	S11 depth	Josyula recut / master family.
B1	FF App complete arc: scope, CTS sequence, burn-up, UAT, launch.	S12	Already used today; bench only tomorrow.
B2	Cost-to-serve decision that flipped operator-first to partner-first.	S12	Already used today; one-line backbone bridge only.
B3	Two programs, one shared team, one ruler; both dates held.	**APPLE-ONLY**	Goforth bank R6 now has a full telling; allocate tomorrow's telling to James only.
C1	UDC: live building from two systemic warehouses to one omni facility.	S2	James / UDC panels; not delta as a full story.
C2	UDC counted operating system: 836-row log, build plan, cutover.	S2 depth	James process family.
C3	Locus robots earn volume; throttle at 94% against a 98% bar.	S2	Goforth PGM / UDC family.
C4	Architecture council: pre-wire every system row; ask for three decisions.	S5	James profile match; use exact phrase only.
D1	Two DCs to JDA in eight months using gates instead of prior WMS experience.	S18	Josyula recut / James process family.
D2	Three mock cutovers; tolerance published before mock one.	S18	Josyula recut / James process family.
D3	Twenty-person team adopts four-week cadence; QA owns the gate.	S15	Josyula S15 / James process family.
E1	ISM inventory operating model: policy before configuration.	S1/S3/S4	Already told today; bridge only.
F1	Thousand-lane program phased by value and a data-quality toll.	S9	Josyula S9; count conflict—use safe wording.
F2	Routing-guide engine: measurement first, Finance-auditable savings.	S10	Goforth bank R1 / Josyula S10.
F3	Reason-coded shortage alert and forecast-accuracy program.	S6	Already told today; bridge only.
F4	Measured lane distributions and tail promise; drift-flag regret.	S8	Already told today; James uses as lived bridge.

M1	Cyber Monday 4 AM: promise-at-risk triage across three DCs.	S13	Josyula S13.
I1	Independent LLM triage agent: eval-first, citations, cost routing.	**APPLE-ONLY**	Goforth bank R7 now has a full telling; allocate tomorrow's telling to Venkata only.
G1	Failure: UAT scenarios passed; concurrency crashed two days before sign-off.	**APPLE-ONLY**	Goforth / James / Josyula failure arcs overlap; allocate full telling to Rocky only.
G2	Failure: RF script broke the training gate; effort is not impact.	S18 failure beat	James training angle / Josyula failure arc.
G3	Failure: security review found mid-build; plan access lead time.	S16 failure beat	Josyula S16 / failure arc.
Part 7	Gap scripts, rapid-fire bank, risk flags, closing integrity notes.	No standalone equivalent.	**APPLE-ONLY phrasing bank**; selected deltas below.

2 · Delta list — what Apple adds tomorrow

How to read this. Each item is an Apple-specific framework, failure detail, probing answer, or exact phrase not needed as a repeated master-story telling. “Drop-in” is the sentence to say; do not recite the analysis around it.

Rocky Goforth — infrastructure program + Bar Raiser (10 items)

RΔ1 · Part 0 — the named operating system. Baseline → one scorecard → numeric gates → root-cause triage → correction verified in cadence. Apple pp4–5.
Drop-in: “My program method is baseline first, one scorecard, numeric gates, first-principles triage, and correction verified by the same number going green.”

RΔ2 · Part 0 — exact gate test. The workbook’s sharper test is not whether a gate exists but whether it can stop work. Apple p4.
Drop-in: “A gate that cannot stop the program is decoration, so I publish the stop line before the first test.”

RΔ3 · Part 0 — red-item contract. Every red arrives with a cause, owner, and date in the same review that surfaced it. Apple p4.
Drop-in: “A red without a cause, an owner, and a date is not status; it is deferred discovery.”

RΔ4 · TMAY ownership probe. Rajesh explicitly separates program machinery from code: he wrote scorecards, BRDs, gates and cadence; engineers built product and automation. Apple p5.
Drop-in: “When I say I, I mean I built the operating system and made the cutover call; engineering built the code.”

RΔ5 · D1 — “gates instead of experience.” Rajesh had never run a WMS migration; the absence of a “trust me” shortcut made written entrance criteria load-bearing. Apple pp28–29.
Drop-in: “I had no experience shortcut to sell, so every phase had to earn entry and exit on written evidence.”

RΔ6 · D2 — rehearsal philosophy. The mock produces the runbook; the live cutover is the next rehearsal with the audience watching. Apple pp30–31.
Drop-in: “A failed mock is cheap; a failed cutover is not—the cutover should feel like the next rehearsal with a larger audience.”

RΔ7 · M1 — instrument silence. The scheduler stopped without emitting a failure, so the permanent control made absence of a heartbeat the alert. Apple pp40–41.
Drop-in: “The process was not wrong; it was silent. We made silence itself the alarm.”

RΔ8 · M1 — the reliability outcome. The next peak being uneventful is framed as an engineered result, not an anticlimax. Apple p41.
Drop-in: “The next peak ran boring, and boring is the reliability deliverable.”

RΔ9 · A3 — why the BRD was the work. The workbook answers the “you only wrote a document” probe directly: complaints become funded, buildable programs through exact specifications. Apple p11.
Drop-in: “Programs die as complaints and ship as documents; my BRD made the operating pain fundable and testable.”

RΔ10 · Part 7 — tool humility. Rajesh’s strongest instrument is the hand-built model and the shared SQL layer, not decorative BI. Apple pp46–47.
Drop-in: “I build in the tool the team reads; I am not a dashboard stylist, but I can interrogate the data underneath it.”

James Cabreros — placement architecture + process improvement (10 items)

JΔ1 · E1 — policy versus configuration. Rajesh owns the business-readable state and channel rules; ERP specialists own transaction-level configuration. Apple pp32–34.
Drop-in: “The policy is mine and the transaction codes are theirs; when they disagree, the signed policy wins and configuration changes.”

JΔ2 · F1 — exception governance. Planner overrides were preserved only with reason codes; silent workarounds were rejected because they cannot improve a rule. Apple pp34–35.
Drop-in: “Overrides with reasons are data; silent workarounds are not.”

JΔ3 · F3 — trustworthy self-service. Every alert carried the input that moved, making verification a thirty-second operator action. Apple pp37–38.
Drop-in: “A self-service decision needs a reason line an operator can check in thirty seconds.”

JΔ4 · F3 — trust as a process metric. False alarms were counted and burned down publicly instead of asking planners to “adopt” the alert. Apple p38.
Drop-in: “Trust is the residue of a low false-alarm rate; I measure it instead of asking for it.”

JΔ5 · C1 — interim architecture has an operating cost. Four months with two systems of record required every report to declare its source; Rajesh would start a reporting workstream on day one. Apple pp21–23.
Drop-in: “If migration creates two temporary truths, reporting strategy is a workstream from day one, not month three.”

JΔ6 · A4 — retire the old truth respectfully. A team accepts one canonical source more readily when its old metric is credited for exposing the problem. Apple pp12–13.
Drop-in: "I gave the old metric a dignified retirement—it became evidence for the backlog, not the loser in a dashboard fight."

JΔ7 · D2 — publish the line before the data. The quality threshold is governance only if it predates the result and can force a delay. Apple pp30–31.
Drop-in: "A tolerance negotiated after results arrive is not a tolerance; publish the line before the data."

JΔ8 · D3 — change design. "Their ceremonies, my evidence": the team authored the cadence mechanics while Rajesh published whether work actually reached QA and demo. Apple pp31–32.
Drop-in: "I traded control for adoption: they authored the ceremonies; I authored the evidence."

JΔ9 · D3 — resistance as a requirement. QA's fear of a shorter testing window became the definition-of-done they owned. Apple p32.
Drop-in: "Resistance is often a design requirement wearing a complaint; I handed QA the gate instead of arguing with the fear."

JΔ10 · Part 7 — executive decision format. One page: metric, trend, driver, ask; the answer belongs in the room and the derivation in the appendix. Apple pp47–48.
Drop-in: "The room gets the answer; the appendix gets the audit—metric, trend, driver, ask."

Venkata Josyula — Principal PM / product + behavioral (10 items)

VΔ1 · A5 — launch proof versus scale proof. The 500-container performance check answered launch readiness, not the 68-customer steady state. Apple pp14–15.
Drop-in: "A launch gate that pretends to be a scale proof is lying; ours answered launch week, and waves carried the scale proof."

VΔ2 · A4 — sequence the decision. Normalize unlike data to earn time, then improve the real residual to deserve decommission. Apple pp12–13.
Drop-in: "I normalized first to earn the time, then improved the residual to deserve it."

VΔ3 · B1 — product outcome over shipping output. Monthly operator interactions eliminated was the focus metric; features shipped was explicitly rejected as the customer outcome. Apple pp15–17.
Drop-in: "A screen shipped is a build output; an operator touch that no longer happens is a product outcome."

VΔ4 · B1 — customer translation. The internal interaction count connected to five customer-facing levers: documents on platform, reports in-app, fewer impersonations, no-touch onboarding and self-booking. Apple pp16–17.
Drop-in: "I forced the internal metric to map to what the customer felt; otherwise it was vanity."

VΔ5 · F1 — adoption is behavior change. The lane was the unit of work, but the real product change was planners trusting a guide over memory, with explainable overrides. Apple pp34–35.
Drop-in: "It looked like a data program; it was behavior change wearing a data costume."

VΔ6 · A4 — analysis that cuts both ways. Normalization removed a false gap but exposed a real variance gap; that made it analysis rather than spin. Apple pp12–13.
Drop-in: "The normalization was credible because it also exposed where we still lost; analysis that only flatters is spin."

VΔ7 · A5 — shipping a measured gap. The dry run processed 100% but tracked 86%; the gap shipped named, owned and visible rather than hidden. Apple pp14–15.
Drop-in: "I will ship a measured, owned gap before I ship a surprise; 86% tracked had a named workstream and stayed on the dashboard."

VΔ8 · F4 — product maintenance regret. The distributions launched before the drift flag; a planner found a degraded lane manually. Apple pp38–40.
Drop-in: "I shipped the insight before the self-monitor; now the drift flag ships with the first distribution."

VΔ9 · M1 — incident leadership regret. The decision log was right, but Rajesh carried the first stretch alone and should have paged the on-call engineer earlier. Apple pp40–41.
Drop-in: "The triage was right; my escalation was late. I should have paged the on-call engineer before carrying the first stretch alone."

VΔ10 · B2 — disagree without ambush. The manager saw the contrary math first and pressure-tested it with engineering for two days before the room decided. Apple pp17–19.
Drop-in: "I disagreed privately, let engineering attack the math for two days, then asked the room to ratify what survived."

3 · Apple-only story A01 — I1: the LLM triage agent

Allocation note. The current Goforth bank also contains full versions of A01, A02 and A03. They remain here because this task requires the Apple-only stories as a self-contained add-on. Do not tell them twice: A02 to Rocky, A03 to James, A01 to Venkata.

Assigned: Venkata — product judgment, evals, customer trust, unit economics. **Company:** independent build, not an employer deployment. **LPs:** Invent and Simplify · Frugality · Dive Deep · Ownership.

LP proof. Invent: routed easy and hard classes differently. Frugality: unit cost was a launch KPI. Dive Deep: graded historical outcomes before model changes. Ownership: built the stack, eval harness and operating rule himself.

60-second pitch. "I built an independent LLM triage agent and treated it as an operations product, not a demo. The naive version cost $25 per incident, and an uncited recommendation made the human repeat the work. I first graded historical incidents into an eval set, then set two release bars: the past must keep passing, and every proposal must cite the ticket and runbook. I routed easy classes to smaller models and reserved expensive calls for the hard tail. Cost fell to $0.40 per incident, a 98% reduction, with accuracy held on the eval set and 100% citation coverage. My regret is that the first eval set grew while early prompts were being judged; I would freeze and version the baseline before the first change."

S I independently built a Python/FastAPI triage agent for operational incident queues. The naive path produced useful proposals but cost $25 per incident and could answer without a checkable chain of evidence. record

T I needed an operations-grade product: quality held on known cases, every proposal checkable, and unit cost below one dollar—without presenting it as an employer production deployment.

A I graded historical incidents with known-correct triage outcomes before changing prompts or models. I made the eval set the release gate; every prompt, model or routing change had to keep the past passing. I imposed "no citation, no proposal," binding each action to ticket evidence and a runbook line. I then measured cost by class and routed routine classifications to small models, reserving expensive calls for the hard tail. I reviewed eval accuracy, citation coverage and unit cost together, because optimizing any one alone would create a demo, not a product.

R Cost moved from $25 to $0.40 per triaged incident—a 98% reduction—while accuracy held on the eval set; citation coverage was 100%. record The result was an independently built product with a quality contract and unit economics, not a claim of enterprise adoption.

TRADE-OFF I MADE. I rejected one expensive model for every incident and chose tiered routing. That added routing complexity and eval maintenance, but protected quality while making the unit economics viable.

Failure / learning. The early eval set was still growing while prompts were judged. I would freeze and version the first baseline before the first experiment.

SAFE-SAY. "This was my independent build, not an employer production program: $25 to $0.40 per triaged incident, with quality held on my graded eval set and every proposal cited."
If asked for source: "Those are my build records as captured in Apple workbook I1, pages 41–42."

Q1 — How did cost fall without quality falling? Most incidents were easy classes. I proved small models held on those classes against the eval set, then reserved the expensive path for the hard tail.

Q2 — What would break at FBA scale? Access review, data boundaries, eval-set maintenance and drift become programs of their own. I would keep a human approval gate until production evidence beats the human baseline.

20-second version. "I built an independent incident-triage agent with the eval set first and one rule: no citation, no proposal. Routing easy cases to small models moved cost from $25 to $0.40 per incident, a 98% reduction, while eval quality held. I would freeze the first eval version earlier."

Question → Apple story router (16 LPs, compact)

LP	Real AIQB prompt	Apple chapter	Open / number
Customer Obsession	“Anticipate a customer need…” (AIQB: Customer Obsession Q3)	F4	Feature request became measured distributions; 200 lanes.
Ownership	“Something significant outside your area…” (AIQB: Ownership Q1)	I1 / D1	Independent build; or first WMS migration.
Invent and Simplify	“Complex problem… simple solution” (AIQB: Invent and Simplify Q1)	I1 / A3	No citation, no proposal; four systems to one pipeline.
Are Right, A Lot	“Made an error in judgment…” (AIQB: Are Right, A Lot Q5)	G1	Scenarios passed; concurrency was absent.
Learn and Be Curious	“Work outside your comfort area…” (AIQB: Learn and Be Curious Q2)	D1	Never ran WMS migration; delivered in eight months.
Hire and Develop	“Provided coaching…” (AIQB: Hire and Develop the Best Q4)	D3	QA authored the gate; do not claim direct-report coaching.
Highest Standards	“Decision between standards and delivery…” (AIQB: Insist on the Highest Standards Q4)	C3	94% vs 98%; throttled 30% to carts.
Think Big	“Opportunity… bigger than initial focus” (AIQB: Think Big Q1)	F4	Feature request became inventory program.
Bias for Action	“Calculated risk where speed was critical” (AIQB: Bias for Action Q1)	B3 / M1	One ruler under compression; or 4 AM triage.
Frugality	“Save costs or eliminate waste” (AIQB: Frugality Q1)	I1	$25 → $0.40 per incident.
Earn Trust	“Not able to meet a commitment” (AIQB: Earn Trust Q4)	G1 / G2	Lost sign-off; or training slipped 1.5 weeks.
Dive Deep	“Dig deep to get to root cause” (AIQB: Dive Deep Q2)	G1 / F2	Concurrent load; or measurement before carriers.
Backbone	“Strongly disagreed with manager…” (AIQB: Have Backbone Q1)	B2	Private check; two-day pressure test; sequence flipped.
Deliver Results	“Important project under tight deadline” (AIQB: Deliver Results Q1)	B3	Compressed program inside four weeks; other date held.
Earth’s Best Employer	“Change… improving coworkers’ work experience” (AIQB: Earth’s Best Employer Q3)	D3	Team authored ceremonies; no mandate.
Success & Scale	Use broad-consequence probe from the LP indicators.	A4 / A5	Retire parallel truth; ship a measured, owned gap.

3 · Apple-only story A02 — G1: the UAT that crashed

Assigned: Rocky — strongest real infrastructure failure and Bar Raiser answer. **Company:** Elementum. **LPs:** Are Right, A Lot · Earn Trust · Insist on the Highest Standards · Ownership.

LP proof. Are Right: names his wrong reflex and decision. Earn Trust: took the criteria miss to the sponsor as his. Standards: rebuilt UAT around numeric load criteria. Ownership: treats an implicit performance requirement as his gap, not the customer's fault.

60-second pitch. "Two days before UAT sign-off, five key users had passed all ten end-to-end scenarios on an Elementum track-and-trace app. Then the sponsor opened access more broadly and concurrent logins crashed the system. My first move made it worse: I reran my scenarios, they passed, and I sent screenshots. I was defending my test plan instead of reproducing the customer's shape. I owned the miss to the sponsor, simulated concurrency with engineering, and the team fixed it, but the sign-off date passed and confidence had to be rebuilt. I permanently changed the UAT contract: every critical scenario gets numeric entrance and exit criteria for uptime, page load and concurrent load at the customer's user count, volume and network. The later proof that the lesson stuck was a Flexport gate of 500 containers in 20 minutes at 99.6% connectivity before customer testing."

S On an Elementum track-and-trace application, UAT sign-off was two days away. Five key users had completed ten end-to-end scenarios each, and my test evidence was green. record

T I owned readiness. I needed to earn customer sign-off and protect launch confidence, including requirements the written scenario list had not made explicit.

A When the sponsor said users could not complete a flow, I first reran the same scenarios and sent screenshots showing they passed. That was my central failure: I tried to prove innocence. I then reproduced the customer's operating shape and found concurrent logins crashed a system my criteria had never load-tested. I told the sponsor the gap was mine—the customer had exposed an implicit performance requirement I failed to surface. Engineering simulated the concurrent load and worked the fix while I rewrote the UAT contract: numeric uptime, page-load and concurrent-load criteria at the customer's user count, data volume and network, with entrance and exit bars written before testing.

R The sign-off date was missed and trust had to be rebuilt. record The structural correction persisted: a later Flexport launch gate required 500 containers in 20 minutes at 99.6% connectivity before customer testing. record-from-docs I no longer accept "works for us" as readiness evidence.

TRADE-OFF I MADE. After finding the cause, I chose round-the-clock engineering effort and a transparent missed sign-off over relaxing the performance requirement. That spent schedule margin and team capacity, but did not ask the customer to accept an unproven system.

Failure / learning. The actual failure was defensive diagnosis: I mistook my coverage for the truth. The first hour of an escalation now goes to reproducing the customer's shape, not proving our innocence.

SAFE-SAY. "Five key users passed ten scenarios, but concurrent access crashed the system two days before sign-off; my criteria omitted load, and the deadline was missed."
If asked for source: "The failure details are Apple workbook G1, pages 43–44; the later performance proof is in the Flexport TPM status source."

Q1 — Wasn't the sponsor at fault for opening access broadly? No. Users exercised the system the way their work demanded. Their behavior exposed an implicit requirement my readiness criteria should have surfaced.

Q2 — Why did you send screenshots? Because the scenarios genuinely passed for me, and I confused test coverage with customer truth. That defensive reflex is the behavior I would take back.

20-second version. "Five users passed ten UAT scenarios, but broader concurrent access crashed the system two days before sign-off. I first defended with screenshots—my mistake—then owned that load was absent from my criteria. The date was missed; every gate I run now tests at the customer's concurrency and volume."

3 · Apple-only story A03 — B3: two programs, one ruler

Assigned: James — sequencing, shared engineering capacity, process governance and floor adoption. **Company:** Flexport. **LPs:** Bias for Action · Deliver Results · Earn Trust · Have Backbone; Disagree and Commit.

LP proof. Bias: made capacity decisions inside a four-week compression. Deliver: compressed program landed and the original date held. Trust: made the collision visible instead of hiding load. Backbone: challenged the resource pull with priced consequences, then committed.

60-second pitch. “Leadership pulled a strategic go-live forward one month and moved engineers off my program to support it. Rather than let two programs run separate escalation systems, I made them one operating problem: one RAID log, one 45-minute daily ruling forum, and one issue grammar—blocker, non-blocker, or requirements gap. True blockers moved through onsite and offshore passes with a written handoff, so issues moved around the clock without asking people to. I escalated only priced decisions: option A costs this; option B costs that. Both programs landed—the compressed one inside four weeks, mine on its original date. About $200K per month per DC was the program-doc value rate. The honest cost was the team’s calendar, and my miss was shipping the ten-step floor guide only a week before cutover. Now operator documentation is an entrance criterion, not a final-week task.”

S At Flexport, leadership advanced a strategic go-live by a month and shared my engineers across that work and my existing launch. Two programs had one team, an immovable date, and about $200K per month per DC at stake in the program record. record

T I had to make the resource collision explicit, protect floor adoption, and land the accelerated program without moving my original date.

A I did not fight the priority after the value was visible; I challenged how it would be executed. I put both programs in one RAID log and one hard 45-minute daily forum. I imposed one grammar—true blocker, non-blocker, or requirements gap—so teams could not win attention by labeling everything urgent. Most alleged blockers were reclassified within a week. True blockers entered a two-pass war room: onsite worked the first pass, offshore continued from a written handoff, and nobody was asked to work continuously. Go/no-go read the true-blocker burn-down to an agreed floor. Escalations went upward the same day as priced choices, not complaints. I also produced a ten-step quick-reference guide for floor adoption.

R The accelerated program landed inside four weeks and my program held its original date. record The defendable economic statement is the program-doc rate—about $200K per month per DC—not a larger cumulative claim. The operating mechanism made shared capacity visible and gave one forum authority over both programs.

TRADE-OFF I MADE. I accepted bounded compression on the team’s calendar and rejected two independent priority systems. The written handoff and time-boxed forum cost coordination effort but prevented an unmanaged slip and continuous heroics.

Failure / learning. The operator guide shipped only one week before cutover, after floor questions had already arrived. Under compression, documentation must be an entrance criterion.

SAFE-SAY. “About $200K per month per DC was at stake—the rate in the program record; the compressed program landed inside four weeks and my original date held.”
If asked for source: “That rate and the operating mechanism are recorded in Apple workbook B3, pages 19–21; I do not inflate the monthly rate into a banked annual claim.”

Q1 — Did you oppose leadership taking your engineers? I made the collision and cost visible, but once the pull-forward was priced it was probably the right business decision. I committed and designed one governance system for both.

Q2 — What actually slipped? No public date slipped. The cost was the team’s calendar and a late operator guide. Naming that cost matters; a “free” compression merely hides where the cost landed.

20-second version. “When two launches shared one team, I ran one RAID log, one 45-minute forum and one blocker grammar. True blockers followed written onsite-to-offshore handoffs. The compressed launch landed inside four weeks, mine held its date, and I learned to make the operator guide an entrance criterion.”

Part 6 in full · G2 — the RF-device script

Best interviewer: Rocky backup; James if asked about training/change control. **Company:** Sephora. **LPs:** Are Right, A Lot · Earn Trust · Deliver Results.

Opener. “I missed a training gate because I scored a dependency by effort instead of impact. The script was a small task; its position on the critical path was not.”

S During Sephora’s RP-to-JDA upgrade, training sign-off depended on hundreds of RF guns, printers and workstations being converted. The device team estimated three to four days, so I treated it as low effort and gave it low attention. record

T I owned the readiness chain for an eight-month, dual-DC migration. I had to prove device readiness before training without moving the final go-live.

A I scheduled a device test days before training. The conversion script failed because refreshed RF hardware did not run the old code. My first-principles read was not “the script is buggy”; it was that my risk model was wrong. A small task in front of a go-live exit criterion was a high-impact dependency and had no what-if plan. With key resources on COVID night shifts, the manual fallback could not finish inside the training window. I told leadership the 1.5-week training slip was mine. I then added a what-if and mitigation to every cross-functional dependency regardless of effort and moved proven device readiness into training’s entrance criteria.

R Training slipped 1.5 weeks, but the go-live date held; the broader migration landed on time and about 3% under its roughly $1.8M budget. record The later mock-cutover discipline—7% to 1.8% to 0.3% against a 0.5% line—shows the what-if habit institutionalized.

TRADE-OFF I MADE. I accepted a visible training slip rather than declare unproven devices ready. That consumed scarce SME windows but protected the live date.

Failure / learning. Catching it with my own test saved the go-live; it did not erase the planning failure. Impact, not effort, now decides risk attention.

SAFE-SAY. “Hundreds of devices; a three-to-four-day estimate; device test before training; training moved 1.5 weeks; go-live held.”
If asked for source: “Apple workbook G2, pages 44–45, reconciles two earlier device-count tellings and explicitly tells me to say ‘hundreds.’”

Q1 — Why is the miss yours if the device team estimated it? I accepted an effort estimate as a risk assessment. The impact and lack of fallback were my planning responsibility.

Q2 — Is a 1.5-week training move a big enough failure? G1 is the larger customer-facing miss. G2 matters because small unplanned dependencies are how mature programs fail.

20-second version. “A three-to-four-day device script failed before training because refreshed RF hardware was incompatible. I had scored risk by effort, not critical-path impact. Training slipped 1.5 weeks, go-live held, and proven device readiness is now an entrance criterion.”

Part 6 in full · G3 — the security review I discovered

Best interviewer: Rocky infrastructure; Venkata behavioral backup. **Company:** Flexport. **LPs:** Are Right, A Lot · Earn Trust · Dive Deep · Ownership.

Opener. “I planned the extracts, loads, metrics and parallel run, but I discovered the access review mid-build instead of planning it. A small two-week miss exposed a structural gap in how I scoped data programs.”

S At Flexport I was replacing a manual four-system weekly review with a Snowflake pipeline. The source included ticket data carrying customer content. I scoped the pipeline and correctness gate, but not the data-access security review. record

T I owned the BRD and cutover. I had to complete the access review without skipping the Monday operating review or weakening the one-percent parity gate.

A When the review surfaced mid-build, I owned the two-week miss upward as a scoping failure, not “slow security.” I had treated approval as a formality instead of a lead-time workstream. I re-sequenced rather than bypass it: the manual path stayed live during the parallel run, so the review continued every Monday. I documented the cross-reference table, named the data classes and made read access explicit. I then changed the kickoff checklist permanently: security has an owner and lead time in the BRD, and every data program starts by asking whose approval the data needs and how long it takes.

R The access review cost two weeks of schedule margin, but no weekly review was skipped. The pipeline cut over only after manual and automated metrics agreed within one percent for consecutive weeks, and system accuracy closed at 90% from a 60% baseline. record The miss cost margin, not delivery.

TRADE-OFF I MADE. I kept the labor-intensive manual path alive and rejected a premature cutover. That spent two more weeks of my time but preserved delivery and access integrity.

Failure / learning. The parallel run saved me partly by luck—it was designed for correctness, not schedule resilience. Now I deliberately keep the old path until the new path clears both access and correctness.

SAFE-SAY. “I under-planned the customer-data access review by two weeks; the manual path kept every Monday review alive, and cutover still waited for one-percent parity.”
If asked for source: “Apple workbook G3, pages 45–46, and A3; the two-week miss and parity gate are record, not modeled.”

Q1 — Why call a two-week slip a failure? Because the cause was structural. Every customer-data program needs approval; absent a checklist change, I would repeat the miss.

Q2 — Why not ask for an exception? Customer-content access was a gate. Bypassing it would convert a schedule miss into a trust and compliance risk.

20-second version. “I discovered a customer-data security review mid-build and lost two weeks because it was absent from my BRD. The manual path kept the weekly review alive; cutover still waited for one-percent parity. Security lead time is now planned at kickoff, never discovered.”

Part 0 in full · “How I run a program”

Say this in 35 seconds. “I use the same five-part system on every program. First, baseline the KPI before touching the process. Second, publish one scorecard on one cadence with a grain and date on every number. Third, write numeric entrance and exit criteria that can actually stop the program. Fourth, when a number goes red, trace it to the root before spending money on the symptom. Fifth, put the cause, owner and date into the same cadence, then use the same number to prove the fix went green.” Apple pp4–5.

Move	Entry	Delivered artifact	Exit / proof	Tomorrow’s translation
1 · Baseline	Business problem and decision owner named.	KPI definition, grain, date, owner, starting value.	Owner signs the starting number; no “before” invented later.	Rocky: deployment baseline. James: change-volume/error baseline. Venkata: customer and unit-economics baseline.
2 · One scorecard	Definitions reconciled; duplicate dashboards exposed.	WBR, burn-down or issues log with outputs over inputs.	The decision forum reads one instrument on the same day each week.	“Two dashboards is zero dashboards.”
3 · Numeric gates	Entrance evidence and exit threshold published before testing.	Go/no-go line, rollback trigger, fallback and authorized caller.	The gate has stopped or slowed work at least when evidence missed.	“A gate that cannot stop the program is decoration.”
4 · First-principles triage	A red metric, not a theory.	Defect class / causal trace; symptom separated from mechanism.	Root cause reproduced at the customer or operator’s shape.	G1 concurrency, C3 zone logic, M1 scheduler heartbeat.
5 · Correct in cadence	Cause accepted; owner empowered.	Red item with cause, owner, due date and verification method.	The same KPI that found the miss proves the correction.	No side status, no changed denominator, no victory by calendar.

Framework trade-off. This system willingly spends speed before a gate—calibration, parallel run, rehearsal—to buy faster and safer scaling after the gate. It rejects both a big-bang launch and endless analysis.

Framework failure line. “The three scars behind the system are G1: I omitted load, G2: I scored risk by effort, and G3: I discovered access approval. The framework is those misses turned into controls.”

4 · Apple resume → proof story map

Every experience bullet in the Apple-tailored resume is mapped. The aggregate summary claim is included as a guardrail, but it should not be volunteered because the underlying addition is not reconciled in the Amazon packet.

Resume bullet	Proof story	Number to land	Basis / safe line
S/4HANA ↔ ISM reverse/forward inventory; one “available”; MCS retired; zero incidents.	S1 + Apple E1; S4 for the rule failure.	Zero customer-facing incidents.	record “I owned policy and integration governance; ERP engineers owned configuration.”
Portfolio execution system: RAID, 836-item log, 343-task runbook, rollback, 60+ stakeholders.	S2/C2 for 836; S1/S2 cutover family.	836 rows; 343 tasks / ~815 task-hours; 60+ stakeholders.	record Call it the umbrella program’s log, not an ISM-only log.
UDC omni NPI on BlueYonder + robotics; staged gates; +15% efficiency.	S2; Apple C1/C3.	+15% processing efficiency; zero customer incidents.	record DC Ops’ measurement.
Re-sequenced 18-month roadmap; 66% effort reduction; about $5M avoided in six months.	S12; Apple B1/B2.	18 months → value in 6; 66%; about $5M avoided.	record State the avoided-cost noun once.
Scaled 4PL visibility 7 → 68; CTS; Snowflake accuracy.	S11 and S16; Apple A1–A3.	7 → 68 (9.7×); $0.66 → $0.03; system accuracy 60% → 90%.	record Keep S11 and S16 as separate stories.
Dual-DC legacy ERP → JDA in eight months; three mocks; $20M pipeline.	S18; Apple D1/D2/G2.	7% → 1.8% → 0.3% vs 0.5%; eight months; $20M pipeline of four programs.	record Pipeline value, not banked savings.
VP-sponsored 20-person org conversion; four-week cadence; defects down.	S15; Apple D3.	Six-week → four-week cadence in six weeks; critical defects −50%.	record No direct-report claim; QA authored the gate.
SAP supply/demand view; daily MD04/MB52; forecast accuracy; safety stock.	S6; Apple F3.	70% → 85%; about $2M safety stock released.	record Planner collaboration increased 50%; do not call it effort reduction.
1,000 supplier/carrier routes; ≥95% valid ASN; completeness.	S9; Apple F1.	A thousand-lane program across 25 European locations; completeness 60% → 95%.	record If discussing phasing: 597 first is a user anchor, not the settled delivered count.
MetricStream dual-hat PM/solution lead: SOWs, estimates, budgets, onshore/offshore UAT/cutover/hypercare.	Resume-only role proof; B3 written handoff is a transferable operating example.	No safe project-outcome number in the source set.	record “I owned SOWs and delivery; I will not invent a client or value.”
Manhattan OMS across 3 DCs/100 stores; stockouts, fill and protected sales.	S13; Apple M1.	3 DCs + 100 stores; 99% fill vs 95%; stockouts −20%; about $8M of the day’s sales protected.	record Say about $6M was on the line when paged; keep “protected” as the $8M meaning.
Summary: “$27M+ cost savings.”	Aggregate across multiple programs; no reconciled ledger in packet.	Do not land this aggregate.	Safe: “I’ll use the program-level figures individually rather than add unlike savings, avoidance and protected revenue.”

90-second resume walk — Rocky: infrastructure / program lens

S “My career has been at the seam between operational rules and the infrastructure that must carry them.”

T “The recurring job has been to move a live operation without losing the customer promise.”

A “At Sephora I currently lead inventory-integration programs across S/4HANA, WMS, OMS and warehouse automation. My operating system is one scorecard, entrance and exit gates, rehearsed cutover and numeric rollback. On the dual-DC WMS move, I had never run that migration, so I used gates instead of experience: three mocks moved reconciliation from 7% to 1.8% to 0.3% against a 0.5% line published before mock one. At Flexport I ran a global visibility rollout and specified a Snowflake pipeline that replaced a four-system manual review; the infrastructure team built it and I owned the metrics and one-percent parallel-run gate. Earlier, at Elementum and Manhattan, I learned the same lesson on customer systems: test at the customer’s shape and instrument silence as a failure.”

R “That is the infrastructure program leader I bring: I do not claim to build the pipe; I make deployment evidence, ownership and rollback unambiguous.”

90-second resume walk — James: supply-chain process lens

S “I have spent fifteen years converting supply-chain policy into a process operators can run.”

T “The throughline is one governed decision path across systems, sites and teams.”

A “At Sephora, that meant one signed definition of available stock and a live Utah facility moving from two systemic warehouses to one omni pool. I pre-wired an architecture council with one row per impacted system, then staged purchasing and inbound before the inventory flip. The program used an 836-row work-triage log and a 343-task cutover runbook; zero customer-facing incidents and DC Ops measured processing efficiency up 15%. At Flexport, I re-sequenced an 18-month roadmap using cost-to-serve and later handled two launches competing for one engineering team with one RAID log, one 45-minute daily forum and one blocker grammar. At Elementum, I phased a thousand-lane program by value and gated each cohort on data quality.”

R “My process-improvement result is not a slide: it is a rule, a self-service path with validation, and a post-change number the business owns.”

90-second resume walk — Venkata: product lens

S “I am a supply-chain program leader who treats operating mechanisms as products.”

T “I define the customer problem, the metric and the graduation gate, then sequence engineering around evidence.”

A “At Flexport, I took Shipment Visibility from seven pilots to 68 customers behind calibration and six graduation criteria; the 68-customer and 170-user targets landed two months early while cost-to-serve moved from $0.66 to $0.03 per container. On the forwarding app, I changed the roadmap from operator-first to partner-first because that sequence moved about twice the cost-to-serve reduction; value arrived in six months instead of eighteen with 66% less effort. At Sephora, I made the inventory policy business-readable before configuration and let QA author the gate on a 20-person cadence change. Independently, I built an incident-triage agent with the eval set first, citations on every proposal and unit cost from $25 to $0.40.”

R “My product test is simple: does a customer or operator make a better decision, can we prove it, and can the system scale without adding people linearly?”

5 · Numbers check — Apple v11 versus CANON v2

This table includes every Apple-involved conflict or wording drift called out by NUMBERS-RECON, plus the resume aggregate. Precedence: primary source document → Apple v11 → Amazon master. When evidence remains inconsistent, the say-aloud value is the conservative intersection.

Story / metric	Apple v11	Other canon / conflict	Safe aloud
A1/S11 customer scale	7 → 68 customers; 170 MAU landed; 32,000 shipments public target.	WBR source contains incompatible promotional multipliers; master once implied all public targets landed.	“7 to 68 customers—9.7 times—and 170 users landed two months early. 32,000 shipments was a target.”
A1/A2 CTS start	About $1.00 in hand-carried pilot weeks, then $0.66, then $0.03.	Canon uses the stable manual baseline, not the hand-carried setup phase.	“$0.66 to $0.03 per container, about 95%.” Mention “about a dollar in hand-carried weeks” only if asked.
A1 first customer	Shadow accounts include Marc Fisher / Michelin.	Pilot schedule names Ta Chen / FAM Brands as wave one; status doc names Marc Fisher onboarding.	“The first two hand-carried customers.” No names.
A2 manual effort	About four hours a week at worst.	BRD says 1.5–2 hours per snapshot and separately prices 30 hours/week across people and two WBRs.	For S11: “about four hours a week at worst.” For S16 BRD pricing, state the 30 → 6 hours/week assumptions separately.
A2 dry run	100% processed, 86% tracked, job stories 98.4%.	Master rounded tracked performance differently.	“100% processed; 86% tracked; job stories 98.4%.”
A3/S16 priced baseline	More than $100K from PM time plus decision risk.	BRD arithmetic is $103,680 from 30 → 6 hours/week at assumed $90/hour; master’s hours/risk split is modeled.	“The BRD priced the baseline at $103,680 a year on its assumptions—priced, not claimed as saving.”
B1 launch identity/date	Calls April 2022 the MVP with 11 partners / 10 countries; same chapter’s box puts export MVP in Q1 2023.	Solution doc supports the later export launch.	“The import-side app went live April 2022 with 11 partners in 10 countries; export scope followed later.”
B2 value label	Uses both annual reduction plus overhead and avoided cost.	Resume and canon support avoided cost; overhead has no primary support.	“About $5M of avoided cost,” once. Drop the overhead claim.
C1/C2 duration	C1 says 15 months; C2 says 14 months.	Dated program window depends on whether hypercare is included.	“About 15 months, discovery through hypercare exit.”
C1/C2 tracks	C1 says 15+ tracks; C2 says about 30 build tracks.	Different grains are not clearly labeled.	“About 30 build tracks” for the build plan; avoid the smaller track count.
C1/C4 systems	About a dozen impacted systems.	New council deck supports seven families / about a dozen domains / about two dozen named systems.	“About a dozen system domains; the council impact table named about two dozen systems.”
D1 budget / value	About $1.8M, 3% under; $20M pipeline of four programs.	An older loop packet carried incompatible budget and time-value wording.	Use Apple: “About $1.8M, roughly 3% under; a $20M pipeline of four programs was unblocked.”
G2 devices	Two source tellings differ on device count and estimate.	Canon explicitly resolves to nonnumeric fleet wording and the longer estimate.	“Hundreds of devices; three-to-four-day estimate; training moved 1.5 weeks; go-live held.”
F1/S9 delivered lanes	A thousand lanes onboarded across 25 European locations.	Master says a thousand asked, 597 first for 80% of spend, another wave, remainder deferred.	“A thousand-lane program across 25 European locations; I phased by value. The 597 first-wave figure is my earlier user anchor.”
F2/S10 timing	99% inside a two-month bar.	Master says about three months.	“99% within a quarter.”

F3/S6 planner metric	Planner collaboration up 50%.	Master previously described planner effort down 50%.	“Planner collaboration increased 50%.”
F4/S8 value noun	About $1.5M of carrying cost released.	Master described one-time inventory release and modeled carrying rates inconsistently.	“About $1.5M of inventory carrying cost released; I do not annualize it.”
M1/S13 $8M meaning	Uses $8M as both at risk and protected; also says about $6M on the line at page time.	Resume says about $8M protected.	“About $6M was on the line when the page fired; about $8M of the day’s sales was protected.”
Resume aggregate	$27M+ cost savings.	No reconciled addition; underlying figures mix savings, avoidance, released inventory and protected sales.	Do not use the aggregate. State program figures individually.

6 · What is usable from earlier packets tomorrow

Earlier packet / story	Interviewer	Why it is usable	Integrity condition
Walmart war book — governed delivery-promise buffer / pre-purchase WBR / seller settings.	James, only if he gives a placement/configuration case.	Strong architecture vocabulary: one governed dial, validation before change, post-change WBR, escape routing and capacity signals.	**DESIGN CASE.** Say “I designed” and “in the model.” Pair with lived F4 in one sentence. Do not claim Walmart delivery results.
Walmart war book — promise decomposition and Andon.	Venkata, if asked how to define product metrics.	Separates inputs from customer outputs and makes a response mechanism part of the product.	Illustrative case inputs only; no delivered dollar or engagement claim.
Scaia manager packet — honest capability boundaries.	James.	Useful language: BlueYonder/JDA execution experience is deep; no claim of hands-on BlueYonder planning-suite configuration. “I have not run a factory” is an honest boundary.	Use the boundary language, not that packet’s older SV automation or inventory nouns; those drift from current canon.
Scaia manager packet — newest-to-oldest career throughline.	All three.	Good structure: Sephora policy/cutover → Flexport visibility/product → Elementum planning/transport → Manhattan peak.	Use the three reconciled 90-second versions in this add-on, not the older packet’s numbers.
Scaia manager packet — MetricStream delivery bridge.	Rocky.	Proves SOW, estimates, onshore/offshore delivery, UAT, cutover and hypercare when asked about vendor governance.	No client, team-size, dollar or outcome is in the source set; keep it qualitative.
Adobe packet — named technical + functional owner on every boundary.	Rocky and James.	Directly useful for infrastructure interfaces and James’s centralized policy-platform migration.	Phrase as the governance Rajesh would install; do not imply ownership of Adobe’s ECC conversion.
Adobe packet — “I do not run a methodology religion.”	Rocky.	Clear trade: iterative design/build; phase-gated cutover where the date and task list are fixed.	Prove with D3 four-week cadence and D1/D2 rehearsed cutover.
Adobe packet — separate business rule from requirement.	James.	Placement-policy fit: business signs the rule; systems implement plants, interfaces and contracts that make it true.	Use S1/E1 as a one-line bridge only; it was already told today.
Adobe packet — adoption means the old path closes.	Venkata.	Product depth: training and usage are weak evidence if the old process remains available.	Use MCS retirement or the A3 pipeline cutover as proof; no new number.

7 · VERIFY / resolve before saying an exact disputed fact

V1 · S9 delivered count. Apple/resume and master differ on what went live. **Safe:** “It was a thousand-lane program across 25 European locations; I phased by value, with 597 high-value lanes first in my earlier record.”

V2 · S10 exact duration. Apple and master differ. **Safe:** “Data quality reached 99% within a quarter.”

V3 · B1 April launch identity. Apple’s prose and its own math box differ. **Safe:** “The import-side app went live in April 2022 with 11 partners in 10 countries; the export scope followed later.”

V4 · A1 first-customer names. Apple and pilot documents name different populations. **Safe:** “The first two hand-carried customers.”

V5 · G2 device count. Source tellings differ. **Safe:** “Hundreds of RF devices.”

V6 · MetricStream outcomes. Resume supports responsibilities, not client results. **Safe:** “I owned SOWs, estimates, budgets and onshore/offshore delivery through UAT, cutover and hypercare; I will not attach a client or dollar that is not in my records.”

V7 · Resume aggregate. The addition is not reconciled. **Safe:** “I will use each program’s number with its noun rather than blend savings, avoided cost and protected revenue.”

Source trace — every quantitative claim in §§3–4

Number	Story / resume proof	Source file	Page / row	Match
$25 → $0.40; 98%; citations 100%	A01 / I1	Apple_PMIII_Interview_Stories_v11	pp41–42	Y
5 users × 10 scenarios; 2 days before sign-off	A02 / G1	Apple_PMIII_Interview_Stories_v11	pp43–44	Y
500 containers / 20 min / 99.6%	A02 later proof	Apple v11 A5; Technology Program Management Status 11/19/21	Apple pp14–15; source p10	Y
One month pull-forward; 45-minute forum; inside 4 weeks	A03 / B3	Apple_PMIII_Interview_Stories_v11	pp19–21	Y
About $200K/month/DC	A03 / B3	Apple_PMIII_Interview_Stories_v11	pp19–21	Y; use rate only
Hundreds; 3–4 days; 1.5 weeks	G2	Apple_PMIII_Interview_Stories_v11	pp44–45	Y after conflict resolution
~$1.8M; ~3% under; 8 months	G2 / D1 resume proof	Apple v11 D1; Apple-tailored resume	pp28–29; resume p1	Y
7% → 1.8% → 0.3% vs 0.5%	G2 / D2 resume proof	Apple v11 D2; Apple-tailored resume	pp30–31; resume p1	Y
2-week access-review miss; parity within 1%	G3 / A3	Apple_PMIII_Interview_Stories_v11	pp10–11, 45–46	Y
System accuracy 60% → 90%	G3 / A3 / resume	Apple v11 A3; Apple-tailored resume	pp10–11; resume p1	Y
Zero customer-facing incidents	S1/E1 resume bullet	Apple-tailored resume; Apple v11 C1	resume p1; Apple p21	Y
836 rows	S2/C2 resume bullet	Apple v11 C2; UDC Consolidated Issues Log	pp23–24; 836 data rows	Y
343 tasks / 48,885 min ≈815 hr	S2/C2 resume bullet	Apple v11 C2; UDC dual cutover plan	pp24–25; task rows	Y-with-rounding
60+ stakeholders	C1/C4 resume bullet	Apple_PMIII_Interview_Stories_v11	pp21–22, 26–28	Y
+15% processing efficiency	S2/C1 resume bullet	Apple v11 C1; Apple-tailored resume	pp21–22; resume p1	Y
18 months → value in 6; 66%; ~$5M avoided	S12/B2 resume bullet	Apple v11 B2; Apple-tailored resume	pp17–19; resume p1	Y after label resolution
7 → 68 = 9.7×; two months early	S11/A1 resume bullet	Apple v11 A1; Apple-tailored resume	pp6–7; resume p1	Y / arithmetic
$0.66 → $0.03, about 95%	S11/A1 resume bullet	Apple v11 A1/A2; WBR Metric Culture	pp7–9; source p2	Y-with-rounding
$20M pipeline / 4 programs	S18/D1 resume bullet	Apple v11 D1; Apple-tailored resume	pp28–29; resume p1	Y
20 people; 6-week → 4-week in 6 weeks; defects −50%	S15/D3 resume bullet	Apple v11 D3; Apple-tailored resume	pp31–32; resume p1	Y (defect result via resume)
70% → 85%; about $2M	S6/F3 resume bullet	Apple v11 F3; Apple-tailored resume	pp37–38; resume p1	Y
Planner collaboration +50%	S6/F3 map	Apple_PMIII_Interview_Stories_v11	p37	Y; resolves conflict
1,000-lane program; 25 locations; 60% → 95%	S9/F1 resume bullet	Apple v11 F1; Apple-tailored resume	pp34–35; resume p1	CONFLICT on delivered count; safe wording used
3 DCs + 100 stores; 99% vs 95%; stockouts −20%	S13/M1 resume bullet	Apple v11 M1; Apple-tailored resume	pp40–41; resume p1	Y
About $6M on line; about $8M sales protected	S13/M1 resume bullet	Apple v11 M1; Apple-tailored resume	pp40–41; resume p1	CONFLICT resolved by noun

Final say-aloud guardrails

- Every answer is first person about Rajesh’s mechanism and decision; team-built code and operations work stay credited to the team.
- Walmart material is case design, not delivered client work.
- The Apple-only agent is an independent build, not a production employer deployment.
- Protected revenue, avoided cost, released inventory and priced baselines are different nouns and are never added.
- The UDC 836-row log is the umbrella/UDC program artifact, not an ISM-only artifact.
- Finished-goods planning grain only when discussing manufacturing or inventory planning.

> **Close.** “The common thread is not a tool. I baseline the number, put one instrument in the room, publish a gate that can stop me, and own the miss when the evidence says I was wrong.”

Sources used: Apple workbook v11 (49 pp); Apple-tailored resume; Amazon L6 SCM master; NUMBERS-RECON CANON v2; current panel files; Walmart war book; Scaia manager packet; Adobe 30-minute packet; primary sources named in the trace. No new client, headcount, value or outcome was inferred.