PART A0 · PHASE 1 · CROSS-LOOP ALLOCATION

Amazon NACF
Master Story Inventory & Loop Allocation

Rajesh Rishiyur Chandrasekaran · Friday, October 30, 2026 (PDT)

10:00 · Rocky Goforth — Amazon Infrastructure + Bar Raiser: deployment, cadence, stakeholder leadership, cost versus speed.

11:00 · James Cabreros — placement/configuration architecture: process redesign, policy governance, sequencing, self-service.

12:00 · Venkata Josyula — Principal PM, FBA + Bar Raiser: product judgment, customer mechanisms, LP depth.

Debrief rule: one primary owner per story. A second interviewer may receive only a materially different facet, with the shared facts unchanged.

Already spent Thursday: S/4HANA/ISM; FF App estimation; Shipment Visibility; safety stock; 200-lane analysis. Do not lead with these for Rocky or James. Shipment Visibility remains available to Venkata as an explicitly planned product story.

Phase-1 checkpoint. The six worker packets were still arriving. Rocky BR ended after R2 and James after J3 when this edition was frozen. “Unused” therefore means no standalone full STAR in the six named in-flight packets at this checkpoint. Phase 2 must re-read all final files.

Corpora read: Amazon master S1–S21/E1–E12; Google Fiber master v8; Apple PM III workbook v11 including failures; earlier Amazon loop/NACF packets. Basis vocabulary: record · record-from-docs · user anchor · modeled · illustrative.

Read this first — allocation in sixty seconds

INTERVIEWER	OPEN	THEN	FAILURE	DO NOT SPEND
Rocky	S2 UDC dual inventory: about a dozen systems, ~30 build tracks, 15 months	S18 RP→JDA; S5 council; Locus; A5 launch gate	G2 RF-device dependency	S11, S12, S10, S15, S13 as primaries
James	S17 carrier-bid process or S10 Austria if Venkata has not spent it	S15 governance; A4 one truth; E8/E9 explicitly as design cases	G3 security review or G1 load criteria	Walmart designs as delivered; Thursday's S1/S3 family
Venkata	S11 Shipment Visibility product launch: 7→68 customers, 9.7×	S14 trust; S16 ownership; S9 scale; S13 crisis; I1 AI product	G1 UAT concurrency miss	S18 full telling; S15 if James has already used it

Five facts to hold: S11 is 7→68 customers (9.7×), and 32,000 shipments was a target. S9 is a thousand-lane program with 597 carrying 80% of spend first. S18 uses "hundreds of RF devices"; training moved 1.5 weeks and go-live held. S14's $2.5M is order value and its $10M is initiative won. S16's six-figure annual value was priced, not claimed as realized savings.

A · Master story inventory

Location codes: RPM Rocky PM planned · RBR Rocky Bar Raiser checkpoint · UDC dual-inventory/Canada planned · J James checkpoint/planned · VM Venkata main/current · VX Venkata extras planned. “Today” means Thursday’s completed interviews. Apple and earlier-loop aliases are shown because the request was chapter-complete; aliases are not new achievements.

ID	ONE-LINE TITLE	COMPANY / PERIOD	TYPE	LP TAGS	SCALE + HEADLINE NUMBER (BASIS)	LENS	TOLD TODAY?	CURRENT / PLANNED PACKET	FIT R / J / V	VERDICT
S1	One signed “available”	Sephora 2024–25	REAL	Trust, Own, Backbone	Recurring oversell→0 through two go-lives; dollars are user-anchor/modeled	process	YES	Bridge only; older James	L/L/L — spent	SKIP tomorrow
S2	UDC comingle + robotics	Sephora 2024–25	REAL	Own, Big, Results	~12 systems; ~30 tracks; 15 months; +15% efficiency (record)	program	No	RPM, UDC, RBR	H/M/M — infrastructure scale	PRIMARY Rocky
S3	One Financial On-Hand	Sephora 2024–25	REAL	Trust, Simplify	Two truths→one; ~$0.6M modeled	process	YES family	older James	L/M/L	SKIP; spent family
S4	In-transit promise-rule failure	Sephora 2025	REAL	ARL, Trust, Own	Failure direction fixed; no unsupported dollar	failure	No	RPM/RBR bench	M/H/M	BACKUP James
S5	Architecture council sign-off	Sephora 2024	REAL	Trust, Big, Backbone	~12 systems; sign-off before build (record)	program	No	RPM, UDC	H/H/M	PRIMARY Rocky
S6	Supply/demand view	Elementum	REAL	Customer, Simplify	Planner collaboration +50%; finished-goods grain (record)	product	YES safety-stock family	VM bridge	L/M/L	SKIP; spent
S7	CM data/revision story	Elementum	REAL / weak	Deep, Standards	Support conflicts with finished-goods-only framing	process	No	None	L/L/L	SKIP integrity risk
S8	Measured lane distributions	Elementum	REAL	Deep, Own	~$1.5M carrying cost released (record)	process	YES	RBR inventory; VM theme	L/M/L	SKIP; spent
S9	Thousand-lane program	Elementum	REAL	Big, Backbone, Results	1,000 across 25 locations; 597/80% first (record + user anchor)	program	No	VM	M/H/H	PRIMARY Venkata
S10	Austria routing guide	Elementum	REAL	Deep, Right, Frugal	~55%→99% within a quarter; ~$8M across 25 locations (record)	process	No	RBR R1; VM	M/H/M	PRIMARY James; resolve VM clash
S11	Shipment Visibility launch	Flexport 2021–22	REAL	Customer, Standards, Results	7→68 customers = 9.7×; CTS $0.66→$0.03 (~95%) (record)	product	YES	VM planned; RBR bench	L/L/H	PRIMARY Venkata exception
S12	FF App roadmap flip	Flexport 2022	REAL	Backbone, Frugal, Results	66% effort reduction; about $5M avoided (record)	product	YES	VM recut/RBR inventory	L/L/M	BACKUP Venkata only
S13	Cyber Monday at 4 AM	Manhattan	REAL	Bias, Own, Results	99% fill; ~$8M of day’s sales protected (record)	program/crisis	No	RPM bench; J planned; VM	M/M/H	PRIMARY Venkata
S14	OMS↔WMS discrepancy	Manhattan	REAL	Trust, Customer, Backbone	~200 orders / $2.5M order value; $10M initiative won (user anchor)	behavioral	No	RPM bench; VM	M/M/H	PRIMARY Venkata
S15	Agile conversion	Sephora	REAL	Trust, Develop, Simplify	20 people; 6 weeks; four-week cadence (record)	process/people	No	RPM bench; J planned; VM	M/H/M	PRIMARY James
S16	Snowflake pipeline + late security review	Flexport	REAL	Own, Simplify, Trust	Four systems→one; 60%→90%; annual value priced (record)	product/failure	No	VM	M/H/H	PRIMARY Venkata
S17	Carrier-bid reconciliation	Elementum	REAL / anchors	Simplify, Deep, Backbone	~6 weeks→~3; no award dollar (user anchor)	process	No	RBR R2	M/H/M	PRIMARY James
S18	RP→JDA migration	Sephora 2022	REAL	Standards, Learn, Results	2 DCs; 8 months; ~$1.8M budget, ~3% under (record)	program	No	RPM, RBR; VM recut	H/M/M	PRIMARY Rocky

ID	ONE-LINE TITLE	COMPANY / PERIOD	TYPE	LP TAGS	SCALE + HEADLINE NUMBER (BASIS)	LENS	TOLD TODAY?	CURRENT / PLANNED PACKET	FIT R / J / V	VERDICT
S19	Promise buffer as a dial	Walmart prep	DESIGN CASE	Customer, Deep	All outcomes illustrative and “in the model”	design	No	J1; VM bench	L/H/M	James design probe only
S20	Promise WBR + Andon	Walmart prep	DESIGN CASE	Own, Deep	Modeled metrics only	design	No	J2; VM bench	L/H/M	James design probe only
S21	Managed-shipping shadow gate	Walmart prep	DESIGN CASE	Bias, Results	Modeled cohort gates only	design	No	J3	L/H/M	James design probe only
E1	Contractor/truck stock	GFiber prep	DESIGN CASE	Own, Broad Resp.	No delivered result; controls proposed	design	No	J planned	M/H/M	BACKUP James design
E2	Serialized chain	GFiber prep	DESIGN CASE	Standards, Deep	No delivered result	design	No	J planned	M/H/M	BACKUP James design
E3	Vendor EDI + VMI	GFiber prep	DESIGN CASE	Customer, Big	95% ASN gate is proposed, not achieved	design	No	J planned	M/H/M	BACKUP James design
E4	CRM→ERP fulfillment	GFiber prep	DESIGN CASE	Simplify, Deep	No delivered result	design	No	J planned	M/H/M	BACKUP James design
E5	3PL reconciliation	GFiber prep	DESIGN CASE	Standards, Trust	No delivered result	design	No	J planned	M/H/M	BACKUP James design
E6	PLM change→ERP	GFiber prep	DESIGN CASE	Deep, Standards	Finished-goods framing only	design	No	J planned	M/M/M	SKIP unless asked
E7	New-market supply/demand	GFiber prep	DESIGN CASE	Big, Customer	Illustrative planning method	design	No	J planned	M/H/M	BACKUP James design
E8	Monthly release/regression	GFiber prep	DESIGN CASE	Standards, Results	Illustrative release controls	process design	No	J planned	M/H/M	James design; pair S15
E9	ERP cutover	GFiber prep	DESIGN CASE	Bias, Results	Illustrative cutover controls	program design	No	J planned	M/H/M	James design; pair S18
E10	Returns / RMA	GFiber prep	DESIGN CASE	Customer, Own	No delivered result	design	No	J planned	L/H/M	BACKUP James design
E11	Direct procurement	GFiber prep	DESIGN CASE	Frugal, Own	No delivered result	design	No	J planned	L/M/M	SKIP unless asked
E12	Smart vending/consignment	GFiber prep	DESIGN CASE	Simplify, Big	No delivered result	design	No	J planned	L/M/M	SKIP unless asked
A1	SV wave launch (alias S11)	Flexport	REAL alias	Results, Standards	7→68 customers; 170 MAU target	product	YES	VM	L/L/H	PRIMARY Venkata as S11
A2	SV WBR measurement	Flexport	REAL	Deep, Own	CTS and automation measurement	process/product	YES family	VM/RBR references	L/L/M	Venkata probe only
A3	SV WBR automation (alias S16)	Flexport	REAL alias	Frugal, Simplify	Four systems→one; 60%→90%	product	No	VM	M/H/H	Venkata as S16
A4	Apples-to-apples decommission	Flexport 2021–22	REAL	Right, Deep, Own	79 customers; $185,892 ARR; 294,581 containers (docs)	process/product	No	RBR inventory; VM Crux bench	M/H/H	ADD-ON James
A5	Hundred-container launch gate	Flexport 2021	REAL	Standards, Results	500 containers/20 min/99.6%; 86% tracked (docs)	program	No	References, no standalone	H/M/H	ADD-ON Rocky

ID	ONE-LINE TITLE	COMPANY / PERIOD	TYPE	LP TAGS	SCALE + HEADLINE NUMBER (BASIS)	LENS	TOLD TODAY?	CURRENT / PLANNED PACKET	FIT R / J / V	VERDICT
B1	FF App complete launch (alias S12)	Flexport 2022	REAL alias	Customer, Results	April import-side app; export later	product	YES	bridge only	L/L/M	SKIP; spent family
B2	FF App CTS options	Flexport 2022	REAL alias	Backbone, Frugal	66% effort reduction; about $5M avoided	product	YES	bridge only	L/L/M	SKIP; spent family
B3	Compressed dual-program go-live	Flexport 2022	REAL	Bias, Results, Trust	One month pulled forward; both landed; compressed one <4 weeks	program	No	RBR inventory only	H/M/M	ADD-ON Rocky
C1	UDC rollout (alias S2)	Sephora	REAL alias	Own, Big	15 months; zero customer-facing incidents	program	No	RPM/UDC	H/M/M	Rocky as S2
C2	UDC counted issues log	Sephora	REAL	Deep, Own, Standards	836 rows; 85.4% closed; bugs 115 (record)	program/process	No	UDC planned; RBR refs	H/H/M	BACKUP Rocky
C3	Locus robotics	Sephora	REAL	Bias, Standards	98% gate; week-two 94%; full automation week 3 (record/anchor)	program	No	RPM/UDC	H/M/M	PRIMARY Rocky
C4	UDC council (alias S5)	Sephora	REAL alias	Trust, Backbone	Sign-off before build	program	No	RPM/UDC	H/H/M	Rocky as S5
D1	RP→JDA rollout (alias S18)	Sephora	REAL alias	Learn, Results	2 DCs; 8 months	program	No	RPM/RBR	H/M/M	Rocky as S18
D2	Mock-cutover rigor	Sephora	REAL	Deep, Standards	7%→1.8%→0.3% vs 0.5% tolerance (record)	program	No	RPM/RBR	H/H/M	Rocky depth
D3	Twenty people, no mandate (S15)	Sephora	REAL alias	Trust, Develop	Four-week cadence in six weeks	people/process	No	J/VM	M/H/M	James as S15
F1	Lane onboarding (alias S9)	Elementum	REAL alias	Big, Backbone	1,000 asked; 597 first	program	No	VM	M/H/H	Venkata as S9
F2	Routing-guide engine (S10)	Elementum	REAL alias	Deep, Right	~$8M validated across 25 locations	process	No	RBR/VM	M/H/M	James as S10
F3	Shortage alert	Elementum	REAL	Customer, Deep, Simplify	70%→85% forecast accuracy; 3 months; +50% collaboration	product	YES family	RBR inventory	L/M/M	SKIP; spent theme
F4	Measured lead-time distributions (S8)	Elementum	REAL alias	Deep, Own	~$1.5M carrying cost released	process	YES	bridge only	L/L/L	SKIP
M1	Cyber Monday (alias S13)	Manhattan	REAL alias	Bias, Results	99% fill; ~$8M protected	crisis	No	RPM/J/VM	M/M/H	Venkata as S13
I1	LLM triage agent	Independent build	REAL build	Simplify, Frugal, Deep	$25→$0.40/incident; 98% cut; 100% citations (record)	product	No	RBR inventory; J planned	M/H/H	ADD-ON Venkata
G1	UAT concurrency miss	Elementum	REAL FAILURE	Right, Trust, Standards	5 users×10 scenarios passed; all-user load failed 2 days before sign-off	failure	No	RBR, VM	H/H/H	PRIMARY Venkata failure
G2	RF-device script miss	Sephora	REAL FAILURE	Right, Trust, Results	Hundreds of devices; training +1.5 weeks; go-live held	failure	No	RPM/RBR/VM	H/M/M	PRIMARY Rocky failure
G3	Late security review	Flexport	REAL FAILURE	Own, Trust, Right	Review discovered late; no invented delay/dollar	failure	No	J planned; VM	M/H/H	PRIMARY James failure
L19	Trace-ID bug localization	Flexport loop cut	REAL / source cut	Deep, Simplify	Per-hop trace; no separate outcome claimed	product/process	No	RBR inventory/Tripathy	M/H/H	BACKUP James
L20	Monitoring before customer-ready	Flexport loop cut	REAL / source cut	Standards, Own	Processed-vs-tracked counters; no separate result	product	No	References	M/H/H	Venkata probe
L21	Production test + source switches	Flexport loop cut	REAL / source cut	Bias, Standards	One-week hold at 76% vs 80% gate (record/packet)	program	No	References	H/M/H	Rocky probe

ID	ONE-LINE TITLE	COMPANY / PERIOD	TYPE	LP TAGS	SCALE + HEADLINE NUMBER (BASIS)	LENS	TOLD TODAY?	CURRENT / PLANNED PACKET	FIT R / J / V	VERDICT
L22	Customs vs FFA ownership	Flexport 2022	REAL DOC	Own, Trust, Simplify	15 work items: 12 export + 3 import; final boundary signed (docs)	process/product	No	RBR inventory only	M/H/H	ADD-ON James
L23	Hundred-container test (alias A5)	Flexport 2021	REAL alias	Standards, Results	500/20 min/99.6%; 86% tracked	program	No	References	H/M/H	Rocky as A5
L24	UDC launch/log (alias C2)	Sephora	REAL alias	Own, Deep	836-row work-triage system	program	No	UDC	H/H/M	Rocky depth
L25	Interface study: refuse, then reverse	Sephora loop cut	REAL / secondary	Right, Frugal	17 interfaces→2 critical uses; $42K funded, $25K contingent (secondary)	behavioral	No	older James only	M/H/M	ADD-ON James, verify provenanc
L26	Input/output lever map	Flexport loop cut	REAL method	Deep, Simplify	No separate result; WBR method	process	No	J/VM references	L/H/H	Probe, not standalone
UDC-B1	UDC council-to-hypercare full arc (S2/S5/C2/C3)	Sephora 2024–25	REAL composite	Trust, Own, Standards, Results	44 requirement rows / 23 named systems; 2 go-lives; 0 customer incidents (docs/record)	program	No	UDC final; RPM planned	H/H/M	PRIMARY Rocky as S2/S5
CAN-B2	US-managed bulk ship to Canada	Sephora 2023	REAL PLAN; outcome VERIFY	Bias, Frugal, Customer	~13% volume; 99 brands; 20/40/39 waves; ~2.5M-unit front-load (docs)	program/process	No	UDC final	H/H/M	BACKUP Rocky; plan-only result
UDC-B3	Enable a new node type twice	Sephora 2023–25	REAL composite	Deep, Simplify, Standards	UDC 44 rows/23 systems; Canada 6 changes, 5 small (docs)	process	No	UDC final	M/H/M	BACKUP James; distinct config facet

C · Ranked unused / underdeveloped additions

RANK	STORY	WHY ADD IT NOW	BEST OWNER	QUESTION TYPES
1	A4 apples-to-apples decommission	Real product retirement and one-truth governance; not a standalone STAR in the six worker packets.	James	Are Right Q3; Dive Deep Q2; Ownership Q4
2	A5 hundred-container launch gate	Concrete readiness proof with honest imperfection; stronger than another architecture answer.	Rocky	Highest Standards Q4; Deliver Results Q2
3	L22 Customs/FFA ownership	Direct match to James's centralized-policy and cross-team boundary work; primary document exists.	James	Ownership Q3; Earn Trust Q3; Simplify Q1
4	B3 compressed dual-program go-live	Resource collision and speed/cost trade-off for an infrastructure leader.	Rocky	Bias Q1; Deliver Results Q1; Earn Trust Q5
5	I1 LLM triage agent	Hands-on product/AI example with an eval contract and explicit unit economics.	Venkata	Invent Q1; Frugality Q1; Learn Q2
6	L25 interface-study reversal	A rare clean "I changed my mind" story; useful only with conservative secondary-source wording.	James	Are Right Q4; Frugality Q4; Dive Deep Q3

ADD-ON 1 · A4 — Apples to apples before decommission REAL

Best owner: James · **LPs:** Are Right, A Lot · Dive Deep · Ownership · **Headline:** 79 legacy customers, $185,892 ARR and 294,581 annual containers received a segment path (record-from-docs).

20-SECOND VERSION

"Leadership was comparing our 25% reliability with a competitor's 63% on unlike source mixes. I normalized both to the same basis; the competitor re-scored to about 27%, exposing measurement—not product—as the immediate gap. I then assigned all 79 legacy customers a destination and retired the duplicate truth while improving reliability to 90% against an 80% north star."

S — Situation. At Flexport, the new visibility product was being judged at 25% schedule reliability against a competitor's 63%, while a legacy source still billed 79 customers and $185,892 in annual recurring revenue. The source mixes were unlike, yet the comparison could have written off the new product and preserved two competing truths.

T — Task. I owned the retirement risk: establish an honest comparison, expose any residual product gap, and give every legacy customer an explicit path rather than letting decommission drift across another planning cycle.

A — Action. I put decommission at the top of the risk register with an owner, mitigation and date. I rebuilt the benchmark on identical sources; on that basis the competitor's 63% re-scored to about 27% versus our 25%. I showed the old dashboard's defenders that their figure was valid on its own basis, then changed the basis rather than attacking the team. I did not stop at normalization: it also exposed variance of sixteen days versus eight as the real gap, so I secured six months for source additions. In parallel, I segmented all 79 customers into migrate, white-label or cancel, with one canonical reliability definition and one executive forum signing both the improvement and retirement plans.

R — Result. The 79-customer legacy base—294,581 annual containers—received a disposition, the duplicate truth retired, and reliability reached 90% against an 80% north star after source work. Leadership could now separate measurement correction from product improvement. The durable mechanism was "reconcile before compare."

TRADE-OFF I MADE: I retired $185,892 of legacy ARR rather than preserve a small revenue stream that forced two products and two scoreboards. Normalization bought time; it did not substitute for the six months of product work.

FAILURE / LEARNING: I should have established the canonical-source rule before the first executive comparison. Once dashboards have constituencies, a cheap definition becomes an expensive negotiation.

Are Right, A Lot: I changed the decision by making both products use the same evidence.

Dive Deep: I exposed the remaining variance gap instead of using normalization only to flatter us.

Ownership: every customer received a destination and date.

SAFE-SAY: "The transition document records 79 customers, $185,892 of annual recurring revenue and 294,581 annual containers. The 90% result is the workbook record; the normalized comparison is about 27% versus 25%."
If asked for source: "The Crux transition RAPID and my 2022 product-management record, reconciled in Apple workbook v11 A4."

Was normalization spin?

No. It cut both ways: it removed the false 63-versus-25 alarm and exposed the unfavorable sixteen-versus-eight-day variance. That residual drove the six-month backlog.

What was yours versus the team's?

I rebuilt the comparison, wrote the one-source rule and set the sequence. Engineering added sources; transition owners built the customer paths with me.

B · LP × interviewer allocation

Each cell is **PRIMARY / backup — landing number**. The question column uses the AIQB's wording (lightly shortened only where marked with an ellipsis). No story is primary for two interviewers in this matrix.

LP + AIQB QUESTION	ROCKY	JAMES	VENKATA
Customer Obsession Q3: "anticipate a customer need…they didn't know they needed"	S2 / A5 — +15% efficiency	A4 / S17 — 79 customers	S11 / S14 — 7→68
Ownership Q3: initiative benefiting company/customers that no group owned	S18 / S5 — 2 DCs, 8 months	L22 / S17 — 15 work items	S16 / S14 — 4 systems→1
Invent and Simplify Q1: "complex problem…simple solution"	S5 / C2 — one council gate	S17 / L22 — ~6→~3 weeks	I1 / S16 — $25→$0.40
Are Right, A Lot Q4: bad decision; impact; learning; application	G2 / S4 — training +1.5 weeks	A4 / L25 — ~27% vs 25%	G1 / S16 — 2 days before sign-off
Learn and Be Curious Q2: work outside comfort; build expertise; outcome	S18 / C3 — 3 mocks	S10 / E9 — ~55%→99%	I1 / S9 — 100% citations
Hire and Develop the Best Q4: coaching; what started it; outcome	S2 / S18 — ~30 tracks	S15 / C2 — 20 people	S9 / G1 — 25 locations
Insist on Highest Standards Q4: decision between standards and delivery	A5 / G2 — 500 in 20 min	S10 / E8 — 99% data gate	S11 / G1 — 86% tracked, named gap
Think Big Q1: opportunity to do something bigger than initial focus	S2 / S5 — one omni facility	S10 / S17 — 25 locations	S9 / S11 — 1,000 lanes
Bias for Action Q1: calculated risk where speed was critical	B3 / C3 — <4 weeks	S17 / E9 — one live parallel run	S13 / S11 — 4 AM bridge
Frugality Q4: trade-offs between quality and cost	S18 / B3 — ~3% under budget	L25 / S17 — $42K study	I1 / S16 — 98% unit-cost cut
Earn Trust Q3: influence a peer with a differing opinion	S5 / G2 — sign-off before build	S15 / A4 — cadence adopted without mandate	S14 / S16 — ~$2.5M order value
Dive Deep Q2: dig deep to root cause	C2 / A5 — 836 rows	S10 / A4 — 3 years of history	S16 / G1 — 60%→90%
Have Backbone; Disagree and Commit Q1: strongly disagreed with manager or peer	S5 / B3 — one architecture decision	L22 / S15 — 12 export + 3 import	S14 / S9 — transparent correction
Deliver Results Q1: important project under a tight deadline	S18 / B3 — 8 months	S10 / S17 — within a quarter	S13 / S11 — 99% fill
Earth's Best Employer Q5: compassion in a high-pressure situation	G2 / B3 — protected scarce SME windows	S15 / G3 — skeptics authored the gate	G1 / S16 — owned miss, no blame
Success and Scale Bring Broad Responsibility Q2: changed a decision due to unintended customer consequences	S2 / C3 — throttled automation by zone	A4 / S4 — retired two truths deliberately	S11 / S9 — held a wave one week

Clash detector — resolve before 10:00

COLLISION	WHERE IT APPEARS	RESOLUTION
S10 Austria	RBR R1 + current VM primary	Give James the full process-improvement story. Venkata uses it only if James did not; otherwise use S9.
S13 Cyber Monday	RPM bench + James Manhattan plan + VM primary	Venkata owns the full crisis STAR. Rocky may cite only the heartbeat mechanism; James skips it.
S15 agile governance	RPM bench + James plan + VM primary	James owns people/change governance. Venkata uses I1 or S9 for the relevant LP.
S18 migration	RPM primary + VM recut	Rocky owns the full program. Venkata may mention only “published tolerance before mock one” if probed.
S11 Shipment Visibility	Thursday telling + Rocky references + VM planned	Venkata gets the product/launch facet by explicit exception. Rocky and James do not retell it.
S14 discrepancy	RPM/James potential Manhattan + VM primary	Venkata owns trust/transparency. Others use S5 or A4.
G1 failure	RBR and VM full failure arcs; James may use load criteria	Venkata owns full failure. James uses G3; Rocky uses G2.
Walmart S19–S21	J1–J3 and VM bench	James alone uses them, always bannered DESIGN CASE and paired with a lived analog.
UDC-B1 / UDC-B3	New UDC/Canada pack tells the same UDC base to Rocky and James	Rocky owns the end-to-end UDC program. James may use only Canada node configuration or the 44-row governance mechanism, explicitly saying it is the same UDC program.

ADD-ON 2 · A5 — A hundred containers before a customer REAL

Best owner: Rocky · **LPs:** Insist on Highest Standards · Deliver Results · Dive Deep · **Headline:** 500 containers in 20 minutes at 99.6% connectivity; dry run 100% processed and 86% tracked (record-from-docs).

20-SECOND VERSION

"I replaced a launch date with two evidence gates. A hundred real containers carried seven checks and every miss had a defect ID. The performance proof processed 500 containers in 20 minutes at 99.6% connectivity; the end-to-end run exposed an 86% tracked gap, which I shipped only as a named, owned workstream. Both checkpoints passed and the first pilot launched on schedule."

S — Situation. Before Flexport's first paying 4PL visibility customer could launch, leadership wanted a date, but the product's promise—trustworthy tracking—had not been proven through a representative production path. If the date became the test, the customer would absorb our defects.

T — Task. I owned the readiness verdict. I had to define business and technical gates before decision week, create evidence that joined defects to real containers, and balance an imperfect tracking gap against the customer and schedule cost of indefinite delay.

A — Action. I published two go/no-go gates, both required to pass. I made a hundred real containers the UAT spine, with checks for creation, subscription, returned data, route, milestones, exceptions and webhooks. Every mismatch carried a Jira ID, so clusters—not anecdotes—burned down. I set an 80% minor-or-no-issues bar before customer-facing UAT and used named shadow accounts first. The weekly data-quality bar rose through the final four weeks, and a thin sample could not print green. I also required load evidence: 500 containers in 20 minutes at 99.6% connectivity. The end-to-end run processed everything but tracked 86%; I refused to hide that fourteen-point gap, assigned its causes and owner, and put it on the launch dashboard.

R — Result. Both checkpoints passed, defect clusters reached zero before customer exposure, and the first pilot launched on schedule. The unresolved tracking gap shipped measured and owned, not as a surprise. Decision week became a readout of pre-published evidence rather than a negotiation over the date.

TRADE-OFF I MADE: I chose evidence with a visible imperfection over either perfection-driven delay or an unmeasured launch. The cost was carrying a named follow-on workstream and defending 86% tracked in public.

FAILURE / LEARNING: I named the processed-versus-tracked gap too late. I now put the gap on the dashboard from the first test cycle.

Highest Standards: a date could not substitute for two published gates.

Deliver Results: the pilot launched on schedule after both gates passed.

Dive Deep: each container-level failure mapped to a defect cluster and owner.

SAFE-SAY: "The source records 500 containers in 20 minutes at 99.6% connectivity and an end-to-end run at 100% processed, 86% tracked. That was launch evidence, not a steady-state scale claim."
If asked for source: "The launch-status guide, product-testing approach, hundred-container CSV and November 2021 program status."

Why launch with one in seven not tracked?

The causes were known, bounded and owned. Delay would not remove weak-lane coverage; putting the gap on the dashboard created an accountable improvement path.

What was yours?

I designed the gates, UAT spine and decision week. Engineering ran the harness and fixed defects; QA co-owned test execution.

ADD-ON 3 · L22 — Customs and FFA: one boundary, fifteen work items REAL · DOC-BACKED

Best owner: James · **LPs:** Ownership · Earn Trust · Invent and Simplify · **Headline:** 15 planned 2023 work items—12 export and 3 import—mapped to one ownership agreement (record-from-docs).

20-SECOND VERSION

"Customs Platform and Freight Forwarding Applications overlapped on fifteen work items. I turned the argument into a row-by-row ownership proposal: domain decision, operational execution, data owner, API and exception. The final meeting assigned FFA all work items except ISF filing, with explicit APIs and shared cost-to-serve goals where the seam remained."

S — Situation. In 2022, Flexport's Customs Platform and Freight Forwarding Applications teams overlapped on customs-declaration work, especially AMS. The next-year roadmap contained fifteen work items—twelve export and three import—and ambiguous ownership would create duplicate builds, dependency queues and inconsistent status for operators.

T — Task. I needed to produce a durable boundary that let each organization stay in its domain while preserving an end-to-end operator flow. Neither team reported to me, and a slogan such as "Customs owns customs" was too vague to configure.

A — Action. I decomposed the roadmap work item by work item and wrote the operating decision beside each: who plans, who executes, who owns the backend and data model, which actions remain offline, what API is required, and where status returns. I separated customs-domain planning and filing from FFA's operational execution of moving cargo. For disputed seams—EAD, port filing, ISF and AMS—I documented current capability, future capability and interim ownership instead of pretending the target architecture already existed. I used the same sheet in the final cross-team meeting, captured unresolved decisions explicitly, and tied shared seams to shared cost-to-serve goals. A future capability never erased the named interim owner.

R — Result. The November 3 ownership agreement assigned FFA all work items except ISF filing. AMS used Customs APIs for create, amend, query and delete; ISF retained a query dependency; AES aligned on integration. The roadmap now had one accountable owner per row and explicit interfaces where ownership crossed organizational lines.

TRADE-OFF I MADE: I chose a pragmatic interim boundary with some offline work over waiting for a perfect platform split. That preserved delivery but left API dependencies to govern.

FAILURE / LEARNING: The first framing was organization-based and too abstract. The agreement emerged only after I moved to work-item and state-transition grain.

Ownership: I took the seam no organization could solve alone.

Earn Trust: capability gaps and unresolved rows stayed visible.

Invent and Simplify: one row-level matrix replaced a domain argument.

SAFE-SAY: "The proposal states fifteen 2023 work items: twelve export and three import. The final agreement says FFA owned all except ISF filing; it does not claim a quantified business outcome."
If asked for source: "Customs Platform / FFA Ownership Proposal, pages 1 and 14."

Did you decide the boundary yourself?

No. I created the decision artifact and drove convergence; the accountable leaders signed the ownership agreement.

How did you know it worked?

The documented result is decision closure and explicit interfaces. I would not invent a defect or cycle-time reduction that the source does not report.

ADD-ON 4 · B3 — Two programs, one engineering team REAL

Best owner: Rocky · **LPs:** Bias for Action · Deliver Results · Earn Trust · **Headline:** strategic go-live pulled forward one month; compressed program landed inside four weeks; Rajesh's program held its date (record).

20-SECOND VERSION

"Leadership pulled a strategic launch forward one month and shared my engineers across two programs. I stopped arguing priorities and ran one RAID log, one 45-minute forum and one blocker ruler. True blockers used written onsite-to-offshore handoffs. The compressed launch landed inside four weeks and mine held its original date."

S — Situation. At Flexport, leadership pulled a strategic go-live forward by one month and reassigned engineers from my program. Two programs now shared one team and one immovable business-adoption window. A rushed technical launch that the floor rejected would simply move the failure downstream.

T — Task. I had to land both programs without asking people to work around the clock, make the resource collision visible to leadership, and protect business adoption rather than optimize only for code complete on paper.

A — Action. I ran both programs as one operating system: one RAID log, one 45-minute daily forum and one ruler—true blocker, non-blocker or requirements gap. Within a week, that classification removed most false blockers from engineering's queue. True blockers entered a two-pass war room: onsite and offshore teams handed off in writing, creating round-the-clock issue velocity without a person working a full day. I published both burn-downs together and escalated only priced choices, never complaints. Every same-day escalation stated what option A cost and what option B cost. The go/no-go read the agreed blocker floor rather than the calendar. I also wrote a ten-step quick-reference guide for the floor so adoption did not become the hidden schedule risk.

R — Result. The pulled-forward program landed inside four weeks; my original program held its committed date; the business adopted on time. The primary record supports the monthly value-at-stake rate, but I lead with the dates and the operating mechanism rather than a derived cumulative dollar.

TRADE-OFF I MADE: I accepted a heavier, tightly bounded daily cadence to preserve both dates. I rejected invisible overtime and serial priority fights. The cost landed on a managed team calendar.

FAILURE / LEARNING: The floor guide shipped only the week before cutover. Under compression, documentation dies first; I would draft it when the collision begins.

Bias for Action: one forum and a common ruler were operating in days.

Deliver Results: one program landed in under four weeks and the other held date.

Earn Trust: both burn-downs and resource consequences stayed public.

SAFE-SAY: "The record says the launch moved one month earlier, the compressed program landed inside four weeks, and mine held date. I do not lead with a multiplied dollar."
If asked for source: "My 2022 ambiguity record and Apple workbook v11 B3."

Did you oppose leadership's decision?

I made the collision and value visible, then committed. My disagreement was with hidden cost, not with the chosen priority.

What was yours?

The shared RAID, forum, blocker taxonomy, escalation framing and floor guide. Engineers built and fixed; operators adopted.

ADD-ON 5 · I1 — LLM triage with a balance sheet and contract REAL BUILD

Best owner: Venkata · **LPs:** Invent and Simplify · Frugality · Learn and Be Curious · Dive Deep · **Headline:** cost per triaged incident $25→$0.40, a 98% cut, with 100% citation coverage and accuracy held on the eval set (record).

20-SECOND VERSION

"I built an incident-triage agent end to end, but the first version cost $25 per incident—a demo, not a product. I created a graded eval set, enforced 'no citation, no proposal,' and routed easy classes to smaller models. Cost fell to $0.40, 100% of proposals cited evidence, and accuracy held on the eval set."

S — Situation. I independently built an LLM agent to read operational incidents, classify them and propose the first action. The naïve version worked, but at $25 per incident it was an expensive demo. An uncited proposal would also force an operator to redo the analysis, destroying trust.

T — Task. I owned the stack and the product bar: preserve triage accuracy, make every proposal auditable, and bring unit cost below one dollar before treating the tool as operational.

A — Action. I built the eval set before optimizing the agent: historical incidents with known-correct outcomes, graded and rerun after every prompt, model or routing change. I established a hard contract—no citation, no proposal—binding each recommendation to ticket evidence and a runbook line. I then treated inference cost like a program budget. A cheap classifier routed common, easy incident classes to smaller models; only the hard tail reached the expensive model. I rejected the cheaper alternative of removing evidence because that would transfer verification cost to operators. I tracked one dashboard with three measures: eval-set accuracy, citation coverage and cost per incident. I built the Python, FastAPI, tool-use and SQLite path myself, so I could isolate each cost and failure mode rather than negotiate across owners.

R — Result. Cost fell from $25 to $0.40 per triaged incident—a 98% reduction—while accuracy held on the eval set. Citation coverage reached 100%, and the $0.40 outcome cleared the sub-dollar operations bar without transferring evidence-checking work back to the operator.

TRADE-OFF I MADE: I chose tiered routing over one premium model for every incident. The eval set paid the complexity cost by proving which classes were safe to route cheaply.

FAILURE / LEARNING: I did not freeze the first eval-set version early enough. A growing baseline made early accuracy claims harder to compare; I now version the contract before tuning.

Invent and Simplify: the router reserved expensive reasoning for the hard tail.

Frugality: unit cost fell 98% against a stated bar.

Learn and Be Curious: I learned by building and instrumenting the full stack.

SAFE-SAY: "This was my independent build. The recorded result is $25 to $0.40 per triaged incident, with accuracy held on my graded eval set and 100% citation coverage."
If asked for source: "My project record, carried into Apple workbook v11 I1; it is not an employer production claim."

Is this just another demo?

A demo has no regression contract or unit economics. This one gated every change on an eval set and engineered to a sub-dollar bar.

Where would it break at Amazon scale?

Data sensitivity, eval maintenance and volume economics become programs. I would require access review and production monitoring before deployment.

ADD-ON 6 · L25 — The interface study I first refused REAL · SECONDARY-SOURCE CAUTION

Best owner: James · **LPs:** Are Right, A Lot · Frugality · Dive Deep · **Headline:** 17 interfaces narrowed to 2 critical usages; $42K study funded with $25K contingent (secondary packet record; verify provenance).

20-SECOND VERSION

"A partner requested a $42K interface study. I first refused because seventeen interfaces sounded like broad documentation. Evidence then showed two usages sat on the critical path. I reversed my decision, funded the study with $25K contingent on findings, and narrowed the scope to those two usages. The lesson was to be frugal about scope, not evidence."

S — Situation. During the Sephora inventory program, a partner requested a $42K study across seventeen interfaces. My first reaction was no: the request looked like paying to document everything when only a subset could affect the launch.

T — Task. I had to protect scarce program funds without creating a blind spot on the critical path. More importantly, I had to be willing to reverse my own decision if the evidence changed the risk classification.

A — Action. I asked for a usage-level dependency map instead of a generic interface list: which business event crossed each interface, which go-live gate depended on it, whether a mock exercised it, and what failure would look like. The map showed that two usages were critical, not all seventeen. That changed the decision. I went back to the sponsor, stated that my initial refusal had treated interface count as a proxy for risk, and recommended funding the $42K study with $25K contingent on what the first pass found. I narrowed the charter to the two critical usages and required the study to produce test cases and an owner—not a shelf document. I preserved a visible stop decision if evidence did not support further work.

R — Result. The study was funded on a gated basis and the broad seventeen-interface concern became two testable critical usages with named owners and gates. The value of the story is the judgment reversal and scoped evidence; the source does not support a separate delivered-savings claim.

TRADE-OFF I MADE: I spent $42K to reduce a launch blind spot but kept $25K contingent. I rejected both a blanket refusal and an open-ended study.

FAILURE / LEARNING: My first no was too fast. I confused a large inventory with low marginal value. I now classify usage criticality before pricing assurance.

Are Right, A Lot: I reversed my own call when the dependency evidence changed.

Frugality: scope and follow-on funding were gated.

Dive Deep: business-event usage replaced interface count.

SAFE-SAY: "An earlier loop packet records seventeen interfaces, two critical usages, a $42K funded study and $25K contingent. I would verify the original artifact before using the dollars; the safe story is that I reversed a refusal after usage-level evidence."
If asked for source: "The earlier Amazon loop packet, story S25; the primary artifact is not in the current source set."

Why include a story with secondary provenance?

Only as a backup after verification. The behavior is useful, but I will not upgrade packet figures into primary-record facts.

What if Rajesh cannot verify it?

Drop the dollars and say: "I narrowed a broad interface study to two critical usages and reversed my initial refusal."

D · Swap table

IF THIS HAS ALREADY BEEN SPENT…	DO NOT REPEAT IT TO…	SWAP TO	DIFFERENT NUMBER TO LAND
S2 UDC to Rocky	James / Venkata	James S17; Venkata S9	~6→~3 weeks / 597 first
S18 migration to Rocky	Venkata	G1 or S16	2 days before sign-off / 4 systems→1
S10 Austria to James	Venkata	S9	1,000 asked; 597 first
S15 governance to James	Venkata	I1	$25→$0.40 per incident
A4 decommission to James	Venkata	S14	~200 orders / $2.5M order value
S11 product launch to Venkata	Rocky / James	Rocky A5; James L22	500/20 min/99.6% / 15 work items
S13 crisis to Venkata	Rocky / James	Rocky B3; James S17	<4 weeks / cycle about halved
G1 failure to Venkata	Rocky / James	Rocky G2; James G3	training +1.5 weeks / late review owned
I1 AI product to Venkata	James	L22 ownership boundary	12 export + 3 import
Thursday's S1/S3/S6/S8/S12	Rocky / James	Rocky S2; James S17	15 months / ~6→~3 weeks

Top 10 per interviewer — order of play

#	ROCKY	JAMES	VENKATA
1	S2 UDC — 15 months, ~30 tracks	S17 carrier bid — cycle about halved	S11 SV — 7→68, 9.7×
2	S18 migration — 2 DCs, 8 months	S10 Austria — ~55%→99%	S14 trust — ~$2.5M order value
3	S5 council — sign-off before build	S15 change — 20 people, 6 weeks	S16 pipeline — 4 systems→1
4	C3 Locus — 94% vs 98% gate	A4 decommission — 79 customers	S9 scale — 597/80% first
5	A5 launch gate — 500/20 min/99.6%	L22 ownership — 15 work items	S13 crisis — 99% fill
6	B3 collision — <4 weeks	G3 security failure — own it early	G1 failure — 2 days before sign-off
7	C2 issue mechanism — 836 rows	E8 DESIGN CASE + S15 lived analog	I1 AI — $25→$0.40
8	G2 failure — training +1.5 weeks	E9 DESIGN CASE + S18 lived analog	S15 only if James did not spend
9	D2 cutover — 7%→0.3%	L25 reversal — verify first	S12 different facet only
10	CAN-B2 — 99-brand plan; outcome VERIFY	UDC-B3 — node configuration, same UDC umbrella	A4 only if James did not spend

E · Source trace for Section C

NUMBER / CLAIM	STORY	SOURCE	PAGE / ROW	MATCH
25% vs 63%; normalized ~27%; variance 16 vs 8 days; +6 months; 80%→90%	A4	Apple workbook v11 A4; 2022 PM record	Apple pp12–13 / text 600–683	Y
79 customers; $15,491 MRR; $185,892 ARR; 294,581 containers	A4	RAPID — Crux Customer Transition; Apple A4	Apple p13 / text 640–683	Y
100-container UAT; seven checks; named defect IDs	A5	100 Container Debugging Session CSV; Apple A5	Apple pp14–15 / text 688–775	Y
500 containers / 20 min / 99.6%; 100% processed / 86% tracked	A5	Technology Program Management Status, Nov 19 2021; Apple A5	Apple p14 / text 728–765	Y
80% minor/no-issues UAT bar; sample ≥20	A5	SV Product Testing Approach; Launch Status Guidance	Apple pp14–15 / text 724–764	Y
15 work items: 12 export, 3 import	L22	Customs Platform / FFA Ownership Proposal	p1 / text rows 5–17	Y
FFA owns all except ISF; Nov 3 2022; AMS API actions	L22	Customs Platform / FFA Ownership Proposal	p14 / text rows 342–359	Y
One month earlier; two programs; 45-minute daily; <4 weeks	B3	2022 Senior TPM Ambiguity record; Apple B3	Apple pp20–21 / text 1028–1103	Y
~$200K/month/DC; cumulative protection framing	B3	Same as above	Apple pp20–21	Y; rate safer than cumulative
$25→$0.40 per incident; 98% cut; 100% citations	I1	Rajesh TPM stories; Apple I1	Apple pp42–43 / text 2159–2224	Y
17 interfaces→2 critical uses; $42K funded; $25K contingent	L25	Earlier Amazon loop packet S25 / older James packet S1-e	secondary packet; old HTML row 190	VERIFY primary provenance

VERIFY checklist and exact safe sentences

- **L25 primary artifact.** Safe: “An earlier loop packet records seventeen interfaces narrowed to two critical usages. Until I confirm the source artifact, I will describe the reversal and omit the dollar amounts.”
- **S17 cycle-time figures.** Safe: “My program anchor is about six weeks to about three; I claim no award savings.”
- **A4 ARR retirement interpretation.** Safe: “The document records the legacy revenue base; I am not claiming that the full amount was lost in one period.”
- **B3 value framing.** Safe: “I can defend the monthly per-DC rate in the program record; the strongest result is that both dates held.”
- **Earth’s Best Employer coverage.** Safe: “I did not manage a direct-report organization in these stories; my evidence is how I protected operators and SMEs under pressure.”

Phase-1 report

What the mining found. The master has enough real stories to keep designs off the behavioral path. The strongest underused records are A4's product-retirement governance, A5's launch evidence, L22's ownership boundary, B3's resource collision and I1's AI unit economics. L25 is useful only after provenance verification. The newly arrived UDC/Canada pack adds a real 99-brand cross-border plan and a 44-row node-governance mechanism; post-launch Canada outcomes remain VERIFY.

Six add-on headlines. A4: 79 customers / $185,892 ARR / 294,581 containers. A5: 500 containers in 20 minutes at 99.6%, with 86% tracked named honestly. L22: 15 work items, one signed boundary. B3: one month pulled forward, compressed launch inside four weeks, both dates held. I1: $25→$0.40 per incident, 100% citations. L25: 17 interfaces→2 critical usages, secondary-source dollars to verify.

Primary collision fixes. S10→James, S13→Venkata, S15→James, S18→Rocky, S11→Venkata, S14→Venkata; G2/G3/G1 become Rocky/James/Venkata's distinct failures.

Design discipline. S19–S21 and E1–E12 remain DESIGN CASES, never delivered results. Pair E8 with lived S15, E9 with lived S18, and Walmart promise ideas with lived F4 only if needed.