INTERVIEW-PREP PACKET 4 OF 5 · PER-INTERVIEWER · BUSINESS-PROCESS LENS

Amazon NACF Supply Chain loop — Packet 4: Zhaowen Lan

Amazon — Senior Supply Chain Manager, NACF Supply Chain (L6)

Zhaowen Lan

Senior Manager, Supply Chain — NACF Supply Chain (peer of the hiring manager)

Thursday, October 30, 2026 · 11:00 am PDT · 55-minute video interview

Strategy in one line. A peer Senior Manager tests whether you can reach one truth across organizations that don't report to you and build a process that survives its author. Eight Sephora S/4HANA primaries, each a different interview answer from one program: open with S1 (three orgs, three meanings of "available" → one signed definition per node per channel; oversells recurring → zero) and hold S3 (Finance and Operations on one Financial On-Hand, the failure direction designed before the happy path); the six variants (S1-a data contract · S3-a de-aggregation to the cent · S1-b control screens and operating model · S1-c the priced no on merchandise reserve · S1-d the counted-rows go/no-go · S1-e the seventeen-interface study) answer the other LPs. Program-lead voice; S4 and S5 are bench stories. Start at the cram sheet and the router.

LP focus — my guess, not confirmed: Earn Trust · Ownership · Invent and Simplify or Have Backbone. Amazon does not disclose the assignment; the guesses are reasoned from the interviewer's title and the JD in "Read me first" and should be held loosely. No interviewer biography is asserted anywhere in this packet — only title-based inference, labeled as such.

Canon: amazon/amazon-l6-scm-master-packet.html (§1.2 roster, §4.1 ledger). Every number below is that packet's value; every dollar carries its basis label (record · user anchor · modeled · priced).

Contents NEW

Combined table of contents. Page numbers are for the printed PDF. Every NEW pill marks content added in this v2 upgrade; UPGRADED marks an existing story deepened in place; VERIFY-WITH-RAJESH marks an unsupported specific listed on the back checklist.

Front matter

5-minute cram sheet (1 of 2): hooks, stories in order, openers, five numbers 3

Cram sheet (2 of 2): tell the S/4 story in 90 seconds, 5 minutes or 10 minutes 4

Question → story router: 16 Amazon LPs and the common non-LP prompts 5

Read me first: the interviewer's lens, LP guesses, answer strategy 7

The primaries — eight S/4HANA stories, each a different interview answer

S1 — one signed meaning of "available" (Earn Trust · Ownership · Backbone · Broad Responsibility) 9

S3 — Shipped-Not-Billed and one Financial On-Hand (Earn Trust · Invent and Simplify · Customer Obsession) 20

S1-a — the data contract: ten buckets, one owner each (Dive Deep · Highest Standards) 30

S3-a — de-aggregation to the cent (Highest Standards · Dive Deep · Learn and Be Curious) 38

S1-b — control screens and the operating model (Ownership · Earn Trust · Deliver Results) 45

S1-c — the priced no on merchandise reserve (Have Backbone · Customer Obsession · Are Right, A Lot) 52

S1-d — go/no-go from counted rows, two go-lives (Deliver Results · Bias for Action · Think Big) 59

S1-e — seventeen interfaces and the $42K I first refused (Are Right, A Lot · Frugality) 66

Peer operating judgment and NACF

Peer operating drills: scope, dependencies, what I would cut, what I would never cut 73

NACF bridge — AFS on the inbound side: what the promise can see 75

Question bank — twenty-four likely questions with a one-line plan 77

Cheat-numbers box — canon values with basis labels 79

Backup bench — only if a primary does not fit

Backup bench: S4 (the rule I got wrong) · S5 (the council) · S9 (597 lanes) · Flexport Audit Logs 81

Back of book

VERIFY BEFORE YOU SAY IT — every tagged item, with a keep / drop decision 92

NUMBERS I MUST NOT GET WRONG — with basis labels 93

5-minute cram sheet — read this, then walk in NEW

Profile hooks (inference from the title plus the inbound booklet's lens — not a biography). A peer Senior Manager of Supply Chain who shares the cadence. Listening for: how you run mechanisms other people own; the handoffs you do not want reopened; a disagreement that stays professional; what you would just fix versus bring to them; who speaks first when a shared metric goes red; and cross-functional operating judgment — scope, dependencies and what you would cut.

STORY	WHAT IT IS (ONE LINE)	LPS	THE ONE NUMBER IT LANDS
S1	One signed meaning of "available" across three node types; policy before configuration	Earn Trust · Ownership · Backbone	**oversell recurring → zero through two go-lives; ≈$2.5M/yr S1 program impact = $1.0M MCS run cost (user anchor) + ≈$1.5M modeled avoided cost**
S3	Shipped-Not-Billed: Finance and Operations on one Financial On-Hand; failure direction first	Earn Trust · Invent and Simplify · Customer	**FOH truths 2 → 1; late line = −2, never +2; ≈$0.6M modeled**
S1-a	The data contract: ten buckets, one owner each, every event into exactly one bucket	Dive Deep · Highest Standards	**five-unit order: AFS 555 held; bad events give 550 · 550 · 565**
S3-a	De-aggregation to the cent: one posting, many lines, Σ lines = posting	Highest Standards · Dive Deep	**$33.33 + $33.33 + $33.34 = $100.00 (naive: $99.99)**
S1-b	Control screens and the operating model: every AFS-changing act has a screen, a role, an audit row	Ownership · Earn Trust · Deliver Results	**Negative AFS clean at T+1; drift = 0 SKUs on consecutive days**
S1-c	The priced no on merchandise reserve; reserve-copy invariants	Have Backbone · Customer Obsession	**560 vs 710 (150 unpickable units); break-even 30 ÷ 49 = 61%**
S1-d	Go/no-go from counted rows; two go-lives; rollback as numbers	Deliver Results · Bias for Action	**UDC log: 836 rows, 441 severe (52.8%); triggers: sync < 99% for 4 h, p95 > 2× for 1 h — never fired**
S1-e	Seventeen interfaces and the $42K I first refused: refuse, study, reverse	Are Right, A Lot · Frugality	**$42K funded, $25K contingent; 17 interfaces → 2 critical usages**

Three sentences to open with.

1. "Most oversell bugs are vocabulary bugs: three organizations meant three different things by 'available,' I owned no system, and I got one definition signed on 6/5/25 before any configuration."
2. "I choose the failure direction before the happy path — wrong on the safe side, visible, with a named owner."
3. "I read go/no-go from counted rows, and I say out loud which of my dollars are measured and which are my model."

Five numbers never to get wrong.

1. ≈$2.5M a year of S1 program impact (user anchor — not in Apple v11) = $1.0M MCS run cost retired (user anchor) + ≈$1.5M of modeled avoided cost (oversell ≈$0.9M · reconciliation ≈$0.33M · change premium ≈$0.3M). Never record for the $1.5M.
2. FSD v1.1 (3/5/25) → v1.2 (5/16/25) → signed 6/5/25 by Supply Chain Strategy, Inventory Management and FSC Inventory Accounting — dates and org names are a user anchor (GFiber v8 only; Apple E1 says only "the business signed the sellable table"). If unsure, say "signed in the first week of June 2025, by three orgs."
3. AFS = Basis − PickRes − EcomPick − EcomPack − SNB, with Basis = min(S4_Avail, WMS_OnHand); worked SKU 980 − 420 = 560.
4. Financial On-Hand truths 2 → 1, ≈$0.6M/yr modeled, no pre-state drift dollar in the record.
5. UDC issues log (not an ISM log): 836 rows / 85.4% closed / Blocker 164 + Critical 277 = 52.8%; rollback triggers sync < 99% for 4 h or order-flow p95 > 2× for 1 h — armed both go-lives, never fired.

Cram sheet (2 of 2) — tell the S/4 story in 90 seconds, 5 minutes or 10 minutes

NEW

The S/4HANA program is one umbrella with eight different answers. Pick the length from the interviewer's body language and the clock, then let the probe choose the depth beat.

LENGTH	WHAT GOES IN, IN ORDER	STOP LINE	IF THEY INTERRUPT
90 seconds	S1 pitch only: situation (three meanings of "available") → task (define every bucket, sign before configuration) → one action (pre-wired every row; three signatories 6/5/25) → result (oversell recurring → zero through two go-lives; ≈$2.5M a year of program impact: $1M MCS run cost (user anchor), ≈$1.5M of modeled avoided cost).	"…the first number is my program's anchor; the other three are my model."	Answer the question, then return to the sentence you were on.
5 minutes	0:00–1:30 S1 pitch. 1:30–2:30 the one-SKU walk: 980 − 420 = 560; raw S/4 over-promises 440. 2:30–4:00 *one* depth beat chosen by the probe: Finance/controls → S3; data → S1-a; control screens → S1-b; conflict → S1-c; cutover → S1-d; testing → S1-e. 4:00–5:00 what I would do differently (the in-transit rule, S4) and the NACF bridge sentence.	"…and at NACF the same discipline is what the promise can see."	Skip the depth beat; the walk and the result are the spine.
10 minutes	0:00–1:30 S1 pitch. 1:30–3:30 the five gated phases with their exit bars (4/30 → FSD 6/5/25 → E2E → go-lives → hypercare exit). 3:30–5:00 the math: one SKU, the decrement contract (AFS 555 through five events), drift 0. 5:00–8:00 *two* depth beats: one control or finance (S3 or S1-b) and one conflict or evidence (S1-c or S1-e). 8:00–9:00 the dollar decomposition with basis words. 9:00–10:00 failure/learning (S4) and the NACF bridge.	"…which is why I'd bring the same bar to inbound."	Compress phases first, never the math or the basis words.

Which depth beat for which probe. "How did you get them to agree?" → S1 pre-wiring. "How does Finance trust it?" → S3. "How is it defined?" → S1-a. "Who owns it after you?" → S1-b. "Tell me about a no" → S1-c. "How did you decide to go?" → S1-d. "How did you scope testing?" → S1-e. "One cent?" → S3-a.

Never both ends of one program to one question. Eight answers, one umbrella — never tell two of them as if they were unrelated programs; say "same program" when you move between them. S2/S18 (UDC and WMS floor stories) are Goforth's; S14, S10, S15 and S16 are Josyula's; S11/S12 are Tripathy's; S6/S8 are Hoerger's.

Question → story router — find the story in ten seconds NEW

If the interviewer asks a behavioral question on an LP, read across: the primary story and the angle to take, the backup, the 15-second opener, and the number to land. Every one of the 16 LPs routes somewhere; two are thin and say so. Only the stories in this packet are routed — never a story allocated to another interviewer.

IF ASKED ON…	PRIMARY STORY · ANGLE	BACKUP	15-SECOND OPENER	NUMBER TO LAND
Customer Obsession	**S1-c** — priced the ask: unpickable stock sold is a cancellation	S3 — failure direction under-sell	The customer cost of selling held stock is a cancellation, so I priced it in units.	**560 vs 710 (150 units)**
Ownership	**S1-b** — every act: screen, role, audit row; owners named	S1 — took the definition nobody owned	On handoff day I wanted no AFS-changing act that only I understood.	**0 unowned buckets; drift tolerance 0**
Invent and Simplify	**S3** — one ledger, one settlement authority; the bridge has no job	S1-a — one event, one bucket	I did not improve the month-end bridge — I removed the reason for it.	**FOH 2 → 1**
Are Right, A Lot	**S1-e** — refuse, study, reverse on evidence	S4 (bench) — the in-transit rule I got wrong	My first answer was no, and then the study came back.	**$42K funded after a no**
Learn and Be Curious	**S3-a** — learned how S/4 posts aggregates before arguing	S4 — back-test before publishing	I had to learn how the book is posted before I could argue about it.	**$100.00 = 33.33 + 33.33 + 33.34**
Hire and Develop the Best	**S1-b** — develop the control owners by handing them the load-bearing artifact thin	S5 — council pre-wiring (pointer)	I measure development by whether the drift report runs without me.	**drift tolerance 0, checked daily**
Insist on the Highest Standards	**S3-a** — Σ lines = posting to the cent	S1-a (tolerance 0) · S1-d (counted rows)	A one-cent break is a failed test, not noise.	**$0.00 break in every cell**
Think Big	**S1** — book-of-record migration, ~15 functions, ~60 people	S1-d — two staged go-lives	The scale was not the system count; every promise had to keep holding while the book changed.	**2 go-lives; rollback triggers never fired**
Bias for Action	**S1-d** — reversible go: rollback as numbers	S3-a — sequenced by event source	I could move because I wrote the way back first.	**sync < 99% 4 h · p95 > 2× 1 h**
Frugality	**S1-e** — fund the study, not the fear	S9 (bench) — 597 lanes = 80% of spend	I refused to fund fear, but I did not refuse to fund evidence.	**$42K, $25K held**
Earn Trust	**S1** — three signatures; pre-wiring; objections as rows	S3 — the aging report given to Finance	Three organizations each owned a different truth; none reported to me.	**signed 6/5/25 by three orgs**
Dive Deep	**S1-a** — a five-unit order walked event by event	S3-a — the grain; S1-e count/sum/per-store	I stopped arguing about the definition and walked one order through four feeds.	**550 · 550 · 565 vs 555**
Have Backbone; Disagree and Commit	**S1-c** — the no lives in an FSD row	S3 — two refusals · S5 (fence)	Merchandising wanted held stock sellable and the room was leaning yes.	**break-even 61%**
Deliver Results	**S1-d** — dated chain held; two go-lives	S1 — MCS retired, ≈$2.5M	The chain was dated before we started, and it held through a peak freeze.	**E2E 28 then 41 days apart**
Strive to be Earth's Best Employer	**S1** — nobody learned their cost in the council room thin	S1-b — owners given real screens	Respect for people's time was the method: every row pre-wired before the room.	**zero surprise rows at the council**
Success and Scale Bring Broad Responsibility	**S1** — oversell is a customer harm, not a ledger error	S3 — under-sell on purpose · S1-c	An oversell is a customer who was told yes and then gets a cancellation.	**oversell recurring → zero**

Common non-LP prompts

IF ASKED…	ROUTE TO	NUMBER
Tell me about a time you failed / were wrong	S4 (bench): v1 rule; or S1-e: my first no	**σ 0.3 d → 0.49-d window ≈ 95% · $42K**
Disagreed with your manager	No manager dispute in the S/4 record — use S1-c (senior stakeholder) and say so no manager story	**61% break-even**

IF ASKED…	ROUTE TO	NUMBER
Influence without authority	S1: ~60 people, three signatories; S5 pointer	**signed 6/5/25**
Reduced cost	S1: $1.0M MCS run cost (user anchor); S9 bench (freight −10% ≈ $2M in six months, user anchor)	**$1.0M anchor**
Dealt with ambiguity	S1: nobody owned the definition; S1-a: ten buckets	**0 open “who owns this” rows**
Simplified a process	S3: no bridge spreadsheet; S1-a: one event, one bucket	**FOH 2 → 1**
Decision with incomplete data	S1-c: no release-reliability data, priced both sides; S1-e: threshold, not price	**p = 61% · L* = $420,000 at 10%**
Missed a deadline	No missed S/4 date in the record — do not invent; pivot to S4 (a rule I corrected) **none on record**	**28 and 41 days**
Mentored someone	Thin: S1-b handoff to control owners **no record of mentoring**	**drift tolerance 0**
Delivered under a tight timeline	S1-d: production freeze Oct–Dec, three E2E days, two go-lives	**11/22 · 12/20 · 1/30**
Changed your mind with data	S1-e: refuse → study → reverse; S4	**$42K**

Read me first

The interviewer's lens (inference from title + JD — not a biography)

Inference from the title. A second Senior Manager of Supply Chain in the same organization is the peer interviewer: someone whose team will share processes, data definitions and escalation paths with this role. Peers test whether the candidate can get to *one truth* across organizations that do not report to them, design a process that survives its author, and disagree without breaking the relationship. They also notice whether the candidate claims more than they owned — a peer has to live with the person afterwards.

What that means for the answers. Eight Sephora S/4HANA governance primaries from one program, each a different interview answer, told as process redesigns of the Plan/Deliver seam. The two originals: S1 (three organizations meant three different things by "available"; one signed definition per node per channel, policy before configuration; oversell incidents recurring → zero through two go-lives) and S3 (Finance and Operations reconciling on-hand by argument; one Financial On-Hand with a safe failure direction and a control report owned by Finance). Program-lead voice throughout — no system reported to me; SAP teams and configurators built to the FSD three orgs signed. Six variants cut from the same record answer the other LPs: S1-a the data contract and the ten buckets · S3-a de-aggregation to the cent · S1-b the control screens and operating model · S1-c the priced no on merchandise reserve · S1-d the counted-rows go/no-go across two go-lives · S1-e the seventeen-interface regression study. Their dollars are slices of S1 and S3 and are never added to them. S4 (the in-transit rule I got wrong) and S5 (the architecture council) are the same program family and appear only as backup pointers.

Three guessed LP focus areas, with rationale

> **Labeled guess.** Amazon assigns 2–3 Leadership Principles per interviewer and does not disclose them. These three are inferred from the title and the JD; if the first question lands on a different LP, the twist maps under each story carry the pivot.

GUESSED LP	WHY THIS INTERVIEWER WOULD CARRY IT	STORY TO OPEN · BACKUP
Earn Trust	The peer LP: one signed truth across Supply Chain Strategy, Inventory Management and FSC Inventory Accounting (S1); Finance given the real control report to own instead of a bridge spreadsheet (S3). Both are Earn Trust backups in the master map — the primary (S14) is the bar raiser's.	S1 (open) · S3 (Finance and Ops on one number) · S1-b (screens everyone can read)
Ownership	Taking the definition nobody owned without owning any system — the bucket inventory, the asymmetry, the signature chain. A peer wants to know the candidate will own the seam, not the org chart.	S1 (the definition nobody owned) · S3 (the failure direction designed before the happy path) · S1-b (who owns each screen)
Invent and Simplify / Have Backbone (one of the two)	De-aggregation with a safe failure direction is the simplification (S3); pricing Merchandising's "merch reserve promisable" ask and putting the no in the FSD as a row is the backbone (S1).	S3 (one FOH, failure direction) · S3-a (de-aggregation) · S1-c (the priced no)

Answer strategy — business first, systems as enablers

- **Open with the process problem and its cost, not the system.** First sentence names the broken process (Plan, Source, Make, Deliver or Return), who it hurt (customer, Finance, the floor) and the baseline number. The platform appears only as the enabler of the redesign, in one clause.

- **Use the JD's four outcome words as the results frame:** reliability (did the promise hold?), responsiveness (how much earlier did we see it?), cost efficiency ($ with its basis label), asset utilization (capital, labor-hours, capacity released).
- **Action = operating-model redesign in gated phases** — Discovery → Plan & Design → Build & Test → Cutover → Hypercare, each spoken as entry bar · what I delivered · exit bar. Say the gate that fired, not the meeting that happened.
- **Every dollar carries its basis word in the same breath** — record · user anchor · modeled · priced. A record dollar is said flat; a modeled dollar is said with its inputs. Never promote a modeled number to record.
- **Program-lead voice.** I owned definitions, rules, gates and the signature chain; engineers, SIs and vendor teams configured and built to the spec I signed. The only platform personally configured in canon is Manhattan DOM.
- **The pitch ends on the primary metric.** Then stop. Silence after the pitch is the interviewer filling their STAR worksheet; their probes are the second half of the answer.
- **For a peer specifically:** name the three organizations and what each owned (Strategy — policy; Inventory Management — execution; Inventory Accounting — the book); say what was pre-wired privately before any council; convert every objection into a document row. "A no that lives in a document survives; a no that lives in a meeting gets re-litigated weekly."
- **Say the decomposition only from this packet's model:** S1 is ≈$2.5M/yr of program impact = $1M MCS run cost retired (user anchor) + ≈$1.5M of modeled avoided cost (oversell ≈$0.9M · reconciliation labor ≈$0.33M · change premium ≈$0.3M). S3 is ≈$0.6M/yr, entirely modeled; the record has no pre-state drift dollar. If the room wants one cancel-rate number, keep S1's and drop S3's row.
- **Variants are slices, not additions.** S1-a to S1-e re-cut S1 and S3 facts for other LPs; any dollar in a variant is a labeled slice or an illustrative number, never added to the ≈$2.5M or ≈$0.6M. Use one story per question; if you have told S1 for Earn Trust, answer the next LP with a variant, not a re-run.

Loop-consistency rule. One question, one story; never the same story to two interviewers — interviewers compare notes. Held for other packets (do not spend here): S11/S12 (Tripathy) · S6/S8 (Hoerger) · S2/S18 (Goforth) · S16/S15/S10/S14 (Josyula). S4 and S5 are bench stories (backup bench at the end), not primaries; S2 shares the UDC program's dated chain with S1/S5 — if asked, say it is one umbrella program. S1-e draws on the loop packet's S25 interface-regression facts (a secondary source, not in the master §4 ledger; the loop tracker also lists S25 as a Bar Raiser backup for Josyula) — if his panel spends it, drop S1-e and use S1-a or S1-d. If pushed for "another example" on an LP, use the named backup in the question bank or the backup bench at the end of this packet, never a story allocated to another interviewer.

How to read the "L6 level evidence" block under each story. Amazon's L6 functional bar (user-supplied rubric): *Ambiguity* — strategy or network undefined; I select the stakeholders, set the goals, deliver independently. *Scope & influence* — team-level scope; influences one to two levels up. *Execution* — tactical and strategic; simplifies or extends; applies best practices; force-multiplies; mitigates long-term risk; makes the short-versus-long-term trade-off explicit. *Impact* — long-term team goals. *Process improvement* — operational-excellence best practices institutionalized. *Complex* — visible risks, roadblocks and constraints mitigated through expertise, judgment and influence to consensus. Overview themes the block is tuned to: designing and modeling the supply chain rather than working in it; developing new processes for others to own; root-cause mechanisms that resolve trending issues at scale; recommendations to leadership that outline competing metrics. The block sits right after Results so the evidence is on the page when the probing starts.

The stories — business-first, canon numbers NEW

Eight Sephora S/4HANA governance primaries from one umbrella program: S1 and S3 (upgraded with LP proof lines, 20-second versions, four-table math, a behavioral pack, three angles and trap drills) and six new variants, each a complete standalone answer with its own situation, math, conflict and result. Each: LP proofs → 60-second pitch → 20-second version + NACF bridge → Situation/Task → design facets → Action in gated phases → Results → L6 evidence → Probing → four-table math → system map → sensitivity → spoken statement → follow-ups → four behavioral answers → three angles → three trap drills.

S1 — Sephora S/4HANA ↔ ISM — one signed meaning of "available" across Store, Retail DC and Dotcom DC: policy before configuration (oversell incidents recurring → zero through two go-lives) UPGRADED

LPs this story can serve here: Earn Trust (backup — three orgs signed one definition) · Ownership (backup — took the definition nobody owned) · Think Big (third — the book-of-record migration) · Success & Scale Bring Broad Responsibility (oversell is a customer harm, not a ledger error) · Have Backbone (the priced no to Merchandising)

Process family: Plan / Deliver seam — inventory availability governance (available-to-promise policy, bucket ownership, real-time decrement contract, daily reconciliation); book-of-record migration

Honesty label. Integration program lead — no system reported to me; every configurator sat in another org. I owned the definitions, the asymmetry, the decrement and reconciliation contract, the signature chain and the council ask; SAP teams and configurators built to the FSD; planners executed the business acts. The record has no pre-state oversell dollar — the ≈$1.5M is my model and I say so. Dated gate chain, issues log and rollback triggers are program-level artifacts shared with S2/S5.

LEADERSHIP PRINCIPLES CARRIED — ONE-LINE PROOF EACH NEW

Ownership — Nobody owned the definition and I owned no system — I wrote the bucket inventory with exactly one owning system per bucket and carried “vocabulary drift” at the top of my own risk register.

Earn Trust — Three organizations signed the same FSD on 6/5/25 because each owned a different truth and I pre-wired every row of the systems-impacted table before the council.

Have Backbone — Merchandising asked for merch reserve to be promisable; I priced it (unpickable stock sold) and wrote the no into the FSD as a row, not a meeting minute.

Success & Scale Bring Broad Responsibility — An oversell is a customer harm, not a ledger error — the asymmetry accepts visible undersell to make oversell zero by construction.

60-SECOND BUSINESS-FIRST PITCH — STOP ON THE PRIMARY METRIC

Situation. Sephora was moving the inventory book of record from the in-house MCS to SAP S/4HANA, with a real-time layer beside it for the digital demand S/4 cannot see in time — a web order placed, a BOPIS order packed, a POS sale rung but not yet posted. Store, Retail DC and Dotcom DC each meant something different by "available": Merchandising wanted merchandise reserve visible, Finance wanted on-hand to tie to the book, dotcom wanted in-transit promisable. Most oversell bugs are vocabulary bugs, and oversells were recurring across all three node types — a customer harm, not a ledger error. **Task.** As integration program lead for inventory supply — no system reported to me — define every bucket, its owning system, and how available-for-sale is computed and decremented in real time, and get Planning, Finance and operations to sign it before any configuration. **Action.** I built the bucket inventory with exactly one owning system per bucket; wrote the lower-of-two basis and the available-for-sale equation, the real-time decrement contract with an idempotency key per event, and a daily reconciliation of the real-time ledger to the book; pre-wired every row of the systems-impacted table with its owner's level of effort before the council so nobody learned their cost in the room, and asked the council for three decisions rather than a review. Merchandising's ask to make merch reserve promisable was priced — unpickable stock sold — and the no went into the specification as a row, not a meeting minute. Three organizations — Supply Chain Strategy, Inventory Management and FSC Inventory Accounting — signed the FSD on 6/5/25. Then a node × channel × bucket-transition regression matrix, dated end-to-end days, a January planning go-live and a May execution go-live, with the order system reading one published availability number and never computing its own. **Result.** Primary metric: customer-facing oversell incidents — recurring across three node types → zero through two go-lives; one signed availability number per node per channel; MCS retired as book of record. About $2.5M a year of program impact: $1M of MCS run cost retired, which is my program's anchor number, plus about $1.5M of modeled avoided cost — avoided oversell, dual-book reconciliation labor, legacy change premium — and I say which is which.

20-SECOND VERSION NEW

Sephora moved the inventory book to S/4HANA and Store, Retail DC and Dotcom DC each meant something different by "available." I owned no system, but I got one signed definition per node per channel — lower of two books minus committed buckets — before any configuration. Oversell incidents went from recurring to zero through two go-lives and MCS retired: about $2.5M a year of program impact — $1M of it the MCS run-cost anchor, the rest modeled.

NACF BRIDGE NEW

The NACF inbound version is "what can the promise see." Units on an ASN or a dock are not available; only units that cleared receive and stow are. Same discipline: one signed definition of each state, one owner per state, and a daily check that the physical count and the system agree.

SITUATION / TASK — THE BUSINESS PROCESS PROBLEM, SCOPE, SCALE, BASELINES

S — Sephora moved the book of record from MCS to S/4HANA, with ISM beside it to see digital demand S/4 cannot see in real time. Three node types, three working definitions of "available." Merchandising wanted merchandise reserve visible, Finance wanted on-hand to tie to the book, dotcom wanted in-transit promisable. Nobody owned the definition; I owned no system.

The process cost: recurring vocabulary-driven oversells — a customer harm, not a ledger error; two books reconciled by hand (≈4 FTE, modeled); the MCS run cost continuing (≈$1M/yr, user anchor); a legacy change premium on an undocumented stack. Vocabulary drift — a bucket without an owner — sat at the top of my risk register.

T — Own the integration design and data contract of inventory supply: which buckets exist, which system owns each, how available-for-sale is computed per node, and how it feeds promising, replenishment and Finance — signed by Planning, Finance and operations before configuration. Sponsor memo with the VP Supply Chain: one signed AFS per node per channel, one Financial On-Hand, zero customer-impacting incidents.

Baselines on the table: oversell incidents recurring across three node types (target zero) · availability truths: two books reconciled by hand (target one signed AFS per node per channel with daily drift as the standing check) · MCS running as book of record (target retired) · ~15 functions, ~60 people to move without authority.

FACET 1 — DATA ANALYSIS — WHAT I PROFILED, AT WHAT GRAIN, WHAT IT REVEALED NEW

- **Data.** Both books — MCS and S/4 — profiled for SKU overlap and order profiles per node type (Store, Retail DC, Dotcom DC); node-by-node sessions with ~15 functions; the live feeds that would carry the real-time layer: Aptos POS sale lines, OMS (Yantriks) commit/pick/pack/cancel events, BlueYonder WMS on-hand (≤ 15 min), FRAME channel reserves (nightly + holds), S/4 MB52 daily plus MSEG movements as events.
- **Grain and joins.** SKU × node × channel × bucket. Every event had to map to exactly one bucket; MB52 ↔ ISM copy joined daily on SKU × node (tolerance 0); WMS on-hand joined to the S/4 copy to produce OOB.
- **Cuts.** By node type, by "who says available" (Strategy, Inventory Management, Inventory Accounting), and by the bucket a unit could sit in.
- **What it revealed.** The worked SKU gives three answers to one question — raw S/4 says 1,000 (440 over-promised against the 560 that is truly promisable), raw WMS says 980 (420 over), and the committed buckets (420 units) were invisible to both. The cause was not data quality; it was that no bucket had an owner, so each reader built its own equation.

FACET 2 — SOLUTION DESIGN — OPTIONS I PUT ON THE TABLE, AND WHY ONE WON NEW

Same SKU, two cases (illustrative units): Case A — S/4 1,000, WMS 980, committed 420. Case B — S/4 1,000, WMS 1,020, committed 420. "Exposure" is units promised above the lower of the two books after committed buckets.

OPTION	CASE A PROMISES	CASE A EXPOSURE	CASE B PROMISES	CASE B EXPOSURE	VERDICT
Raw S/4 Available (Definition A)	1,000	440	1,000	420	Rejected — committed buckets invisible
Raw WMS on-hand (Definition B)	980	420	1,020	440	Rejected — same blind spot
S/4 minus committed buckets	580	20	580	0	Rejected — over-promises when WMS is the lower book
Average of the two, minus committed	570	10	590	10	Rejected — the error always has a sign; the average hides it
min(S/4, WMS) minus committed	**560**	**0**	**580**	**0**	Chosen — oversell is a customer defect, undersell a planning variance

Other options rejected in writing: letting OMS compute its own AFS (two publishers — the one rule is that OMS reads the published number, never computes); making merch reserve promisable (priced in story S1-c); configuring to the legacy definitions first (the legacy definitions were the problem — policy before configuration).

ACTION — FACET 4: IMPLEMENTATION IN FIVE GATED PHASES (PHASES, OWNERS, CONTROLS, RISKS)

- **Discovery** — profiled both books (MCS and S/4) for SKU overlap and order profiles per node; node-by-node sessions with ~15 functions; the bucket inventory with one proposed owning system per bucket. *Exit bar:* every bucket owned by exactly one system; sessions closed 4/30; open questions carried with owner and date.
- **Plan & Design** — the lower-of-two basis; the AFS/FOH equations; the real-time decrement contract (event → bucket delta; idempotency key per event; a reversal must reference its original); daily ISM-vs-MB52 reconciliation with tolerance zero; the RAID log; the council pack — FSD v1.1 (3/5/25) → v1.2 (5/16/25). *Exit bar:* FSD signed 6/5/25 by Supply Chain Strategy, Inventory Management and FSC Inventory Accounting with zero open "who owns this bucket" rows.

- **Build & Test** — the node × channel × bucket-transition matrix with before/after assertions — no story enters QA without a matrix cell reference; SIT on the interface contracts before UAT; dated E2E days — Inbound 11/22, Dotcom 12/20, Retail 1/30 — business and engineering in one room. *Exit bar:* 0 open contract defects in SIT; E2E green on three dates; Blocker + Critical open = 0 at go/no-go.
- **Cutover & Go-Live** — go/no-go recommendation from counted rows; the publish-switch plan — OMS reads the published AFS, never computes; reserve-copy invariants; control-screen roles assigned; rollback triggers (sync < 99% for 4 h; order-flow p95 > 2× budget for 1 h) armed for both go-lives, never fired. *Exit bar:* January go-live (reserve rules live against demand); May go-live (execution); no SKU with two published AFS values; Negative AFS clean at T+1.
- **Hypercare / Close** — daily drift report (ISM vs MB52); Negative AFS / AFS Correction audit review; control-owner map — Negative AFS to Inventory Control, Internal Audit and Finance; AFS Correction to the ISM app team; FSD as-built note. *Exit bar:* drift = 0 SKUs on consecutive days; blockers/day under threshold; the design enters the monthly release train.

Influence and ownership. Sixty people across three orgs, none reporting to me. The method: pre-wire every row of the systems-impacted table with its owner's level of effort before the council; run working sessions to one table instead of arbitrating in meetings; convert objections into FSD rows; ask the council for three decisions rather than a review. Three signatories because each owned a different truth — Strategy the policy, Inventory Management the execution, Inventory Accounting the book; one signature would have left two orgs free to disagree later. The asymmetry held: oversell is a customer defect, undersell a planning variance Inventory Control chases through the control screens.

> **Designed the process; others own it now.** I designed the availability definition, the decrement contract and the control-owner map; Inventory Control, Internal Audit, Finance and the ISM app team now own the controls, and the drift report runs daily without me. **Trade-off put to leadership.** The recommendation to the council outlined the competing metrics — availability to sell vs oversell exposure — and resolved them as an asymmetry in the signed FSD: oversell is a customer defect, undersell a planning variance.

RESULTS — $ · # · % · TIME · CUSTOMER · QUALITY (CANON NUMBERS ONLY)

- **$** — ≈$2.5M/yr of S1 program impact (*user anchor* headline) — $1M MCS run/support cost retired (*user anchor*) + ≈$1.5M of avoided cost, *modeled* (avoided oversell ≈$0.9M · dual-book reconciliation labor ≈$0.33M · legacy change premium ≈$0.3M); range $1.7–5.2M; upside not in the headline: shrink ≈$0.3M, time-to-market ≈$0.2M. The record has no pre-state oversell dollar.
- **#** — oversell incidents recurring across three node types → zero through two go-lives (record); dual-book reconciliation ≈4 FTE by hand → ≈1 FTE reading the drift report (modeled).
- **%** — UDC issues log 85.4% of 836 rows closed (record-from-docs: the UDC BlueYonder WMS + Locus log the shared go-lives were gated on — not an ISM log); the oversell rate behind the ≈$0.9M row is 0.10% of orders — modeled, benchmark 0.05–0.25%; the record movement is incidents → zero.
- **Time** — FSD v1.1 3/5/25 → v1.2 5/16/25 → signed 6/5/25 (*user anchor* — GFiber v8 only, no Apple support); E2E Inbound 11/22 · Dotcom 12/20 · Retail 1/30 and January PO-creation / May combined go-lives (record dates, Apple C1 p22).
- **Customer** — customer-facing oversell incidents zero through two go-lives; the asymmetry treats oversell as a customer defect and undersell as a planning variance (record).
- **Quality** — one signed AFS per node per channel; no SKU with two published AFS values; daily ISM-vs-MB52 drift = 0 SKUs on consecutive days at hypercare exit; MCS retired as book of record (record).

L6 LEVEL EVIDENCE — what the interviewer pattern-matches against the Amazon L6 functional rubric; every line is this story's canon, no new numbers

L6 DIMENSION	EVIDENCE IN THIS STORY
Ambiguity	Three organizations with three meanings of "available" and no owner of the definition. I defined the lower-of-two basis, the AFS/FOH equations and the bucket inventory with one owning system per bucket, and chose the ~15 functions for node-by-node sessions.
Scope & influence	Sixty people across three orgs, none reporting to me; the council took three decisions I framed; the FSD was signed 6/5/25 by three signatories — influence up to the council and across orgs.
Execution trade-off + mechanism others now own	Oversell protection vs availability, made explicit as an asymmetry — oversell is a customer defect, undersell a planning variance. Mechanisms others own now: the daily ISM-vs-MB52 drift report and the control-owner map (Negative AFS → Inventory Control, Internal Audit and Finance; AFS Correction → the ISM app team).
Impact (long-term)	Oversell incidents recurring across three node types → zero through two go-lives; MCS retired as book of record; ≈$2.5M/yr ($1M run cost user anchor + ≈$1.5M modeled).
Process improvement (standard work institutionalized)	The node × channel × bucket-transition matrix — no story enters QA without a matrix cell reference; SIT on interface contracts before UAT; dated E2E days as entry bars.
Complexity & consensus	Every row of the systems-impacted table pre-wired with its owner's level of effort before the council; objections converted into FSD rows instead of arbitrated in meetings.
Root cause at scale	Oversell was traced to two books carrying two meanings — fixed once in a signed definition and a real-time decrement contract (idempotency key per event; a reversal references its original), not incident by incident.

PROBING — CHALLENGE Q&A TUNED TO THIS INTERVIEWER

> **Peer — You owned no system. How did you get three organizations to sign one definition?**
>
> Pre-wired every row of the systems-impacted table with its owner's level of effort before the council — nobody learned their cost in the room — ran working sessions to one table, converted objections into FSD rows, and asked the council for three decisions rather than a review. Three signatories because each owned a different truth; one signature would have left two orgs free to disagree later.

> **Peer — Why lower-of-two and not an average or the book?**
>
> Because oversell is a customer defect and undersell is a planning variance. A positive out-of-balance means the book is above the floor; promising the difference risks oversell, so the basis is the minimum, and Inventory Control chases the variance through the Negative AFS and AFS Correction screens. The asymmetry is the policy.

> **Peer — Merchandising wanted merch reserve promisable. What happened?**
>
> I priced it — unpickable stock sold — and offered visibility without promisability. The no went into the FSD as a row, not into a meeting minute. A no that lives in a document survives; one that lives in a meeting gets re-litigated weekly.

> **Peer — How do you know it was a tenth of a percent?**
>
> I don't, to the decimal. The record says incidents were recurring and went to zero; the rate is a benchmark set deliberately low — a twentieth of a typical e-commerce inventory-cancel rate. First thing I would measure: the "item unavailable after order" cancel-reason code for the twelve months before January. If the $1M MCS figure already includes enhancement labor, I drop the change-premium row and the headline is about $2.2M.

> **Peer — What did you get wrong in this program?**
>
> The first in-transit rule I wrote barred every channel from promising units in transit — safe-feeling, and it starved dotcom of fast movers a day from the building. I corrected it through the same three signatories with a per-lane promise window and a kill switch, released on the normal cadence. Lesson: both sides of a rule get a number before it ships, and back-test a ruleset against history before publishing it. (S4 — pointer; full story on request.)

> **Peer — How would this work between your team and mine here?**
> Every promise rests on one published availability number per node that downstream systems read and never recompute. Between two teams that means: one definition with one owner per bucket, the policy signed before the code, a daily drift report with tolerance zero and a named control owner — and objections turned into rows we both sign.

FACET 3 — MATH & DATA LOGIC, FOUR TABLES, EVERY PLUG-IN SHOWN NEW

Table 1 — before: the problem in numbers.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Two books for one SKU	`OOB = S4_Avail − WMS_OnHand = 1,000 − 980 = 20`	S/4 Available 1,000; WMS on-hand 980 (illustrative)	+20 units — the books disagree; book above floor
2	Units already spoken for	`Σ buckets = 200 + 120 + 60 + 40 = 420`	PickRes 200 · EcomPick 120 · EcomPack 60 · SNB 40 (illustrative)	420 committed or sold-not-posted
3	Truly promisable (the yardstick)	`min(1,000, 980) − 420 = 980 − 420 = 560`	the number the fix will compute	560 units
4	Definition A — channel reads raw S/4	`1,000 − 560 = 440`	committed buckets invisible to it (record problem)	440 units over-promised
5	Definition B — channel reads WMS on-hand	`980 − 560 = 420`	same buckets invisible (record problem)	420 units over-promised
6	Oversell outcome	`30,000,000 × 0.10% = 30,000`	orders illustrative; rate benchmark 0.05–0.25%; incidents recurring across three node types (record)	≈30,000 incidents/yr — modeled
7	Reconciliation burn	`≈4 FTE reconciling two books by hand`	record baseline	≈4 FTE

Table 2 — the fix: step by step, values substituted.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	One supply basis	`Basis = min(S4_Avail, WMS_OnHand) = min(1,000, 980) = 980`	asymmetric lower-of-two (record design)	Basis = 980 units
2	Why the minimum	`Basis − S4_Avail = 980 − 1,000 = −20`	a positive out-of-balance reduces supply (record design)	the +20 is never promised
3	Available-for-sale	`AFS = 980 − 200 − 120 − 60 − 40 = 560`	AFS = Basis − PickRes − EcomPick − EcomPack − SNB; MerchRes and Lost/Found sit outside S4_Avail — no double-subtract path (record design)	AFS = 560 units
4	Floor-above-book case	`Basis = min(1,000, 1,020) = 1,000; AFS = 1,000 − 420 = 580`	WMS 1,020 → OOB = −20 (illustrative)	AFS = 580, not 600 — the 20-unit surplus stays unpromised
5	Real-time decrement	`AFS = 560 − 3 = 557; SNB = 40 + 3 = 43`	store rings 3 units at 10:02 (illustrative); each event maps to exactly one bucket	AFS 557 within the 60-second publish budget
6	Idempotency	`AFS = 557 − 0 = 557`	replayed key at 10:04 is a no-op; a reversal must reference its original (record design)	AFS stays 557
7	Overnight posting	`SNB 43 → 0; S4_Avail = 1,000 − 43 = 957`	aggregate posting lands in S/4 for 43 units (illustrative)	books move together
8	Daily drift check	`Drift = ISM.S4_Avail − MB52 = 957 − 957 = 0`	tolerance 0 — a non-zero drift is a Sev, not a planning variance (record design)	drift = 0

Table 3 — after: the metric, before → after.

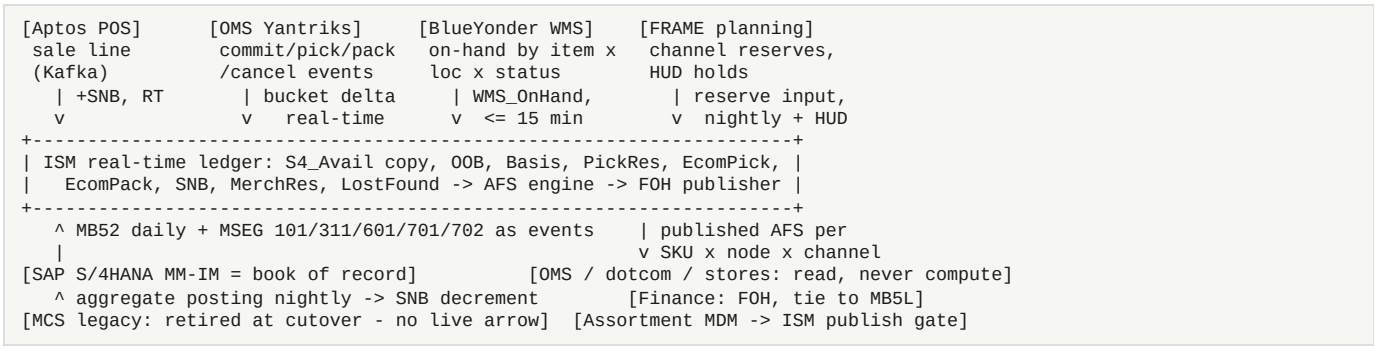
STEP	METRIC	BEFORE	AFTER	Δ / CHECK	BASIS
1	Working definitions of “available”	3	1 signed AFS	`3 − 1 = 2`	record
2	Worked SKU, over-promise under Definition A	440 units	0 units	`440 − 0 = 440`	record problem; units illustrative
3	Oversell incidents	recurring across three node types	zero through two go-lives	`recurring − 0 (count not in record)`	record
4	Dual-book reconciliation	≈4 FTE by hand	≈1 FTE reading the drift report	`4 − 1 = 3 FTE`	modeled on the record mechanism
5	Daily drift at hypercare exit	unmeasured	0 SKUs on consecutive days	`0 − 0 = 0`	record
6	MCS run / support cost	≈$1.0M/yr	retired	`$1.0M − $0 = $1.0M`	user anchor

Table 4 — the dollars, every multiplication shown.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	MCS run / support retired	`≈$1.0M/yr, taken as stated`	user anchor	$1.0M/yr
2	Incidents avoided	`30,000,000 × 0.10% = 30,000`	orders illustrative; rate benchmark 0.05–0.25%; record movement: → zero	30,000 incidents/yr
3	Cost per incident	`$6 + $4 + $19 + $1 = $30`	CS contact $6 · refund $4 · lost margin ~$19 · goodwill ~$1 (benchmark)	$30 each
4	Oversell dollars	`30,000 × $30 = $900,000`	rows 2 × 3 (modeled)	≈$0.9M/yr
5	Reconciliation labor	`4 − 1 = 3; 3 × $110,000 = $330,000`	$110K loaded (benchmark)	≈$0.33M/yr
6	Tech-debt change premium	`40 × 60 h = 2,400 h; 2,400 × $125 = $300,000`	≈40 changes/yr × +60 h each on MCS × $125/h (modeled)	≈$0.3M/yr
7	Headline	`$1,000,000 + $900,000 + $330,000 + $300,000 = $2,530,000`	rows 1, 4, 5, 6	≈$2.5M/yr
8	Not in headline	`$0.3M + $0.2M = $0.5M excluded`	shrink ≈$0.3M + time-to-market ≈$0.2M (modeled upside)	excluded

Say the decomposition only from this packet's model (the loop packet says ≈$1M/yr support cost — the $1M is the shared core).

SOLUTION FLOW — THE SYSTEM & DATA MAP NEW



SENSITIVITY / IMPACT — LOW / BASE / HIGH NEW

INPUT	LOW	BASE	HIGH
Vocabulary-driven oversell rate	0.05%	0.10%	0.25%
Cost per oversell incident	$25	$30	$40
Oversell cost avoided	$0.4M	$0.9M	$3.0M
Reconciliation FTE released × loaded	2 × $95K = $0.19M	3 × $110K = $0.33M	4 × $130K = $0.52M
Change premium (changes × hrs × rate)	24 × 40 × $100 = $0.1M	40 × 60 × $125 = $0.3M	60 × 80 × $150 = $0.72M
MCS run cost (user anchor, fixed)	$1.0M	$1.0M	$1.0M
Total modeled $/yr (headline, excl. upside)	≈$1.7M	≈$2.5M	≈$5.2M

SPOKEN TOTAL-IMPACT STATEMENT — SAY IT LIKE THIS

Conservatively, about two and a half million dollars a year of program impact. One million of that is my program's number: the MCS support cost that went away when S/4 became the book of record. About nine hundred thousand is avoided oversell — on roughly thirty million dotcom orders, a tenth of a percent going wrong at about thirty dollars each — and the record shows those incidents went to zero through both go-lives. About three hundred thousand is reconciliation labor and about three hundred thousand is the change premium we stopped paying on an undocumented stack. The first number is my program's anchor, not a ledger I can show you; the other three are my model, and I can show you every input.

FACET 5 — HARDEST FOLLOW-UPS, WITH NUMERIC ANSWERS NEW

Numeric — Walk me through the AFS on one SKU.

Basis is the lower of the two books: min(1,000, 980) = 980. Subtract what is already spoken for — pick reserve 200, e-commerce pick 120, e-commerce pack 60, shipped-not-billed 40 = 420 — so AFS is 560. A channel reading raw S/4 would have promised 1,000 and over-promised 440; reading WMS, 980 and 420. If the floor is above the book (WMS 1,020), the basis stays 1,000 and AFS is 580, not 600 — the surplus is never promised.

Numeric — What does a sale do, and what if the message is delivered twice?

A store rings 3 units: AFS 560 → 557, SNB 40 → 43, inside the 60-second publish budget. The replayed message carries the same idempotency key, so it is a no-op — AFS stays 557. A reversal must reference its original event, so a cancel can never decrement twice.

Numeric — Which of the $2.5M a year of program impact is real and what is the downside case?

Only the $1.0M MCS run cost is a program number (user anchor). The $1.5M is modeled: $0.9M oversell + $0.33M reconciliation + $0.3M change premium. The low case on the same model is about $1.7M (0.05% rate, $25 per incident, two FTE, 24 changes); if the $1M already includes enhancement labor I drop the premium row and the headline is about $2.2M.

Numeric — What did the gate chain look like, and where was your slack?

Kickoff 3/29, business sessions closed 4/30, council 5/15, LOE signed 7/2, sign-off 7/31, FSD 8/14; ISM FSD v1.1 3/5/25 → v1.2 5/16/25 → signed 6/5/25. E2E Inbound 11/22, Dotcom 12/20, Retail 1/30 — 28 days then 41 days apart — then January planning and May execution go-lives. The slack was the peak freeze: October to December nothing moved.

BEHAVIORAL PACK — FOUR QUESTIONS THIS STORY ANSWERS, EACH A DIFFERENT 90-SECOND MOMENT NEW

EARN TRUST **1. Tell me about a time you built trust with teams that don't report to you.**

Open (first two sentences): Three organizations each meant something different by "available," and I owned none of their systems. Before the council ever met, I sat with the owner of every row on the systems-impacted table and got their level of effort in writing, so nobody learned their cost in the room.

A. I ran working sessions to one table instead of arbitrating in meetings, turned every objection into a row in the FSD, and asked the council for three decisions rather than a review. I used three signatories on purpose — Strategy owned policy, Inventory Management owned execution, Inventory Accounting owned the book — because one signature would have left two orgs free to disagree later. **R.** The FSD was signed 6/5/25 by all three, and none of the three re-opened it through two go-lives.

Land the number: three signatures, 6/5/25 | **Pivot to NACF:** At NACF the same handoff is placement's recommendation versus what the FC actually receives — it gets signed at reservation, while the information is still ahead of the freight.
Do NOT say: "I aligned everyone" without naming the artifact; "I configured S/4"; S14 (the bar raiser's Earn Trust story).

OWNERSHIP **2. Tell me about something you took on that was outside your remit.**

Open (first two sentences): The riskiest line on my risk register wasn't a system — it was a bucket with no owner. I took the definition nobody owned: a bucket inventory where every bucket had exactly one owning system, and I wouldn't exit discovery until there were zero open "who owns this" rows.

A. I wrote the lower-of-two basis, the AFS equation and the real-time decrement contract before any configuration, carried "vocabulary drift" as my own top risk with a trigger and a mitigation, and left open questions each with an owner and a date at the 4/30 session close. **R.** The FSD went to signature with zero open ownership rows, and nobody had asked me to own it — the definition was the whole problem.

Land the number: zero open "who owns this bucket" rows at signature | **Pivot to NACF:** In inbound a stranded unit is the same orphan bucket — physically present, no owner, not available. I would want each stock state in the building to have exactly one owner.
Do NOT say: "It was my system," "I owned S/4," or claiming the SAP team's configuration.

THINK BIG **3. Describe the largest-scale change you have led.**

Open (first two sentences): We moved the inventory book of record from the in-house MCS to S/4HANA with a real-time layer beside it, across about fifteen functions and sixty people I had no authority over. The scale wasn't the system count — it was that every promise to every customer on three node types had to keep holding while the book underneath changed.

A. Five gated phases with dated exits — discovery to 4/30, council 5/15, LOE 7/2, sign-off 7/31 — then two go-lives so each carried one defect class: January put the reserve rules against live demand, May moved execution. Rollback was written as numbers (sync under 99% for four hours, or order-flow p95 over twice budget for an hour) and armed for both. **R.** Zero customer-facing oversell incidents through both go-lives, and MCS retired as book of record.

Land the number: two go-lives, zero customer-facing incidents; ≈$1M/yr run cost retired | **Pivot to NACF:** A new-FC launch is the same shape: stage it so each go-live carries one defect class, and write the stop condition as a number before the first unit moves.
Do NOT say: The Utah DC comingle or robotics gate — that is another interviewer's story; say "one umbrella program" if asked.

SUCCESS & SCALE BRING BROAD RESPONSIBILITY **4. Tell me about a decision where you weighed the broader impact, not just the metric.**

Open (first two sentences): Most oversell bugs are vocabulary bugs, but an oversell isn't a ledger error — it is a customer who was told yes and then gets a cancellation. So I wrote the asymmetry into the policy: oversell is a customer defect, undersell is a planning variance that Inventory Control can chase.

A. That meant lower-of-two, accepting that some sellable units stay hidden at the margin, and giving Inventory Control the Negative AFS and AFS Correction screens to chase the variance. I also priced Merchandising's ask to make merch reserve promisable — unpickable stock sold — and gave them visibility instead. **R.** Oversell incidents recurring → zero through two go-lives; the trade was visible undersell, accepted on purpose.

Land the number: oversell incidents recurring → zero through two go-lives | **Pivot to NACF:** Amazon's version is the delivery promise: a date shown honestly later is cheaper for the customer than a date broken.
Do NOT say: a measured oversell rate — 0.10% is a benchmark I set; the record is incidents → zero.

SAME STORY, 3 ANGLES — DIFFERENT OPENER, HEADLINE, LP NEW

DATA / ANALYTICAL

"Before I wrote one definition, I put one SKU through every reader in the building."

440 · 420 · 560

LP: Dive Deep

One SKU, S/4 1,000, WMS 980, 420 committed. Reading raw S/4 over-promises 440; reading WMS over-promises 420; the lower-of-two minus buckets says 560.

The finding was not bad data — it was that nobody owned the buckets, so every reader built its own equation.

That table, not my opinion, is what the council saw.

CONFLICT / INFLUENCE / BACKBONE

"Merchandising wanted merch reserve promisable, and the room was leaning toward giving it to them."

3 signatures · 6/5/25

LP: Have Backbone · Earn Trust

I priced the ask — unpickable stock sold — and offered visibility without promisability.

The no went into the FSD as a row, not a meeting minute: a no that lives in a document survives; one in a meeting is re-litigated weekly.

Their visibility need was met; I did not caricature them.

FAILURE / LEARNING

"The first ruleset I authored barred every channel from promising in-transit units — it felt safe and it was wrong."

σ 0.3 d → 0.49-d window ≈ 95%

LP: Are Right, A Lot · Learn and Be Curious

The Aging In-Transit report showed dotcom starved of fast movers one day from the building.

I took it back through the same three signatories with a per-lane promise window and a kill switch — never fired.

Lesson: both sides of a rule get a number before it ships, and back-test against history first. (Full walk: S4.)

TRAP DRILLS — THE THREE NASTIEST FOLLOW-UPS, WITH THE EXACT SENTENCE NEW

Trap — "How do you know it was your definition and not seasonality, or the other programs, that took oversells to zero?"

I can't isolate it statistically from the record: there is no pre-state oversell rate in the record, only "recurring across three node types → zero through two go-lives." What I defend is the mechanism: committed buckets were invisible to both raw reads (440 and 420 units on the worked SKU), and the AFS equation plus an idempotent decrement removes that path. The cross-check I would run is the "item unavailable after order" cancel-reason code for the twelve months before January, by node, against the months after; and any SKU with non-zero daily drift would be the first place I looked if incidents had persisted.

Say exactly: "I can't separate it from seasonality using the record, and I won't pretend to — I can show you the mechanism and the one query that would."

Trap — "What's the counterfactual for the $2.5M — and what would make it smaller?"

The counterfactual is the status quo: two books, MCS running. Only the $1.0M run cost is measured. The $1.5M is modeled — 30,000 incidents × $30 = $0.9M, 3 FTE × $110K = $0.33M, 40 × 60 h × $125 = $0.3M. On the model's low inputs the headline is about $1.7M; if the $1M already includes enhancement labor it is about $2.2M.

Say exactly: "One million is measured, one and a half is my model, and if the cancel data comes back at the low end the headline is one point seven."

Trap — "What number would have changed your mind about lower-of-two?"

The share of time the higher book is actually right. Promise one extra unit and you earn about $19 of margin if it exists and pay about $30 of incident cost if it does not: break-even is 30 ÷ (19 + 30) = 61%. If you net the lost margin out of the incident (pay only $11), it is 11 ÷ (19 + 11) = 37%. The books disagree by about 2% (20 of 1,000), so I would measure, from cycle-count outcomes over a quarter, how often the higher book wins the count; above roughly 61% (or 37% on the netted view) lower-of-two is too conservative and I would relax it.

Say exactly: "If the higher book wins the physical count more than about six times in ten, lower-of-two is too conservative and I would change it — and I would measure that from cycle counts, not opinion."

Where this lands at Amazon. Every fulfillment promise rests on one published availability number per node that downstream systems read and never recompute — at Amazon scale the ATP/promise seam between planning and the order pipeline. One definition, one owner, policy signed before code. Reliability (oversells to zero), cost efficiency ($1M of run cost retired), responsiveness (a real-time ledger beside the book).

TWIST MAP — LIKELY LP QUESTIONS FROM THIS INTERVIEWER → THE ANGLE TO TAKE WITH S1

LIKELY QUESTION (THIS INTERVIEWER)	LP	ANGLE — WHAT TO EMPHASIZE	SKIP / NEVER SAY
Tell me about a time you built trust across teams that didn't report to you.	Earn Trust	Three orgs, three truths, one signed FSD on 6/5/25; pre-wiring; objections as rows.	S14 (the OMS↔WMS discrepancy — $2.5M of order value, a different $2.5M from S1's) is the Earn Trust primary — it is the bar raiser's; do not use it here.
Tell me about a time you took ownership of something outside your remit.	Ownership	Nobody owned the definition of "available"; I owned no system; built the bucket inventory and the signature chain.	Never "I configured S/4" — the SAP teams configured to the signed spec.
Tell me about a decision where you had to say no to a senior stakeholder.	Have Backbone	Merch reserve promisable — priced, refused, written as an FSD row with visibility offered instead.	Do not caricature Merchandising; their visibility need was met.
Tell me about a large-scale change you led.	Think Big	The book-of-record migration — MCS retired, one real-time ledger, two go-lives, peak untouched.	Keep S2's floor story out — that is Goforth's; say "one umbrella program" if asked.
How do you define and govern a shared metric / definition between two teams?	Functional — process governance	Bucket inventory with one owner each; the equation written; signatures before configuration; daily drift with tolerance zero; control owners named.	Keep the oversell $ labeled modeled; the movement is incidents → zero.
Why does this matter beyond the ledger?	Broad Responsibility / Customer Obsession	Oversell is a customer harm, not a ledger error; the asymmetry accepts visible undersell to make oversell zero by construction.	No YoY revenue claim exists.

S3 — Sephora S/4HANA — Shipped-Not-Billed and de-aggregation: Finance and Operations on one Financial On-Hand number, with the failure direction designed before the happy path UPGRADED

LPs this story can serve here: Earn Trust (backup — Finance and Ops on one number) · Dive Deep (the grain problem — one posting per aggregate vs transaction lines) · Invent and Simplify (one ledger, one settlement authority, a safe failure direction) · Insist on the Highest Standards (Σ per-line COGS = posting COGS to the cent)

Process family: Deliver / Return-to-book seam — sale-to-posting reconciliation, financial on-hand governance, month-end close control

Honesty label. The flow, controls and ownership map were mine; S/4 and OMS engineers implemented; Inventory Accounting confirmed in UAT and took over the aging report. The record has no canonical dollar for the pre-state drift; the ≈$0.6M/yr is my model and I say so. Rode the ISM program's dated chain (FSD 6/5/25; May execution go-live).

LEADERSHIP PRINCIPLES CARRIED — ONE-LINE PROOF EACH NEW

Earn Trust — Finance owned the book and Operations owned the count; I showed Finance the same number arriving from the physical side first, then handed them the unmatched-posting aging report to own — signed inside the FSD on 6/5/25.

Invent and Simplify — One real-time ledger, S/4 as the settlement authority, one failure direction, one control report — instead of two books and a bridge spreadsheet at every close.

Have Backbone — Two refusals held in writing on mechanism, not preference: real-time per-line posting to S/4, and a Finance-derived SNB (two numbers when the point was one).

Customer Obsession — The failure direction was chosen before the happy path: an unmatched posting under-states availability — under-sell, never over-sell — so the customer is never told yes to a unit that is gone.

60-SECOND BUSINESS-FIRST PITCH — STOP ON THE PRIMARY METRIC

Situation. S/4 posts sales as aggregates — one posting per aggregation number with total quantity and cost — while the point-of-sale sends transaction lines. Between the sale and the posting, units were gone but the book was not decremented: a phantom-availability window in which sold units still showed as available, and Finance and Operations read different on-hand numbers and reconciled by argument, with a bridge spreadsheet at every month-end close. **Task.** Define how the real-time layer tracks that gap and reconciles it so promising, replenishment and Finance agree — under the program's zero-incident constraint, which applied to the book as much as to the customer. **Action.** I named the window and inventoried the events that move units before posting — POS flash sales, omni pending, ship-to-home confirmations — and got Finance and Operations to agree on the description of the gap before any design. The Shipped-Not-Billed bucket increments on those events and decrements only on a matched S/4 aggregate posting; de-aggregation pro-rates cost per line so the sum ties to the cent; and the failure direction was designed before the happy path — an unmatched posting keeps the bucket high, so availability under-states and Replenishment over-orders visibly instead of the customer being oversold invisibly. Finance objected at design review and wanted its own reconciliation; I showed the same number arriving from the physical side first and gave them the unmatched-posting aging report to own — a parallel spreadsheet dies when its owner owns the real report. FSC Inventory Accounting signed inside the FSD on 6/5/25; it rode the May go-live. **Result.** Primary metric: Financial On-Hand definitions — two, Finance's book versus the operations count → one shared number, availability decremented at sale or ship rather than at posting; over-sell from unmatched postings zero by construction; zero incidents. About $0.6M a year modeled — phantom-availability cancels avoided, Finance reconciliation labor, month-end close time; the record has no pre-state drift dollar and I am not claiming one.

20-SECOND VERSION NEW

S/4 posts sales as aggregates while the POS sends lines, so sold units stayed “available” until the posting and Finance and Operations reconciled on-hand by argument. I defined a Shipped-Not-Billed bucket that drops availability at the sale and clears only on a matched posting, and I picked the failure direction first: an unmatched posting makes availability low and visible, never high and invisible. Financial On-Hand went from two numbers to one; about $0.6M a year, all modeled.

NACF BRIDGE NEW

The inbound mirror is received-not-stowed. Units on an ASN, in a trailer or on a dock are the SNB of a receiving dock: not yet available. Inbound available-to-promise should count only units that cleared receive and stow, and an unmatched receipt should leave the number low and on an aging list — the same “lines must sum to the header” check I would run on ASN counts against dock counts.

SITUATION / TASK — THE BUSINESS PROCESS PROBLEM, SCOPE, SCALE, BASELINES

S — S/4 posts sales as aggregates — one posting per Aggregation# with total quantity and COGS; the Aptos POS sends transaction lines with an Order#. Between the sale and the posting, units are gone but the book is not decremented — the phantom-availability window — and Finance and Operations read different on-hand numbers and reconcile by argument.

The process cost: sold units shown as available for up to a day on the overnight batch; two books reconciled by argument; a bridge spreadsheet at every month-end close; and the gap hiding as oversell if an estimate ever decremented the book.

T — Define how ISM tracks the gap in real time and reconciles it so promising, replenishment and Finance agree, with the program's zero-incident constraint applying to the book as much as to the customer. Goal: one Financial On-Hand, AFS decremented at sale rather than posting, a safe failure direction.

Baselines on the table: Financial On-Hand truths: two (target one) · AFS decrement point: at the S/4 posting (target the sale/ship event) · over-sell from unmatched postings: a live risk (target zero by construction; aging report empty or explained at T+1).

FACET 1 — DATA ANALYSIS — WHAT I PROFILED, AT WHAT GRAIN, WHAT IT REVEALED NEW

- **Data.** S/4 aggregate postings (Aggregation#, total quantity, total COGS); Aptos POS RetailSaleTransactions lines (Order#, store, SKU, quantity); OMS omni-pending sale lines; WMS ShipmentConfirmation lines for ship-to-home.
- **Grain and joins.** One posting row covers many lines — the join runs Aggregation# ↔ POS header ↔ lines. The grain mismatch (one posting vs many lines) is the whole problem: the sum of the lines must equal the header, in quantity and in cost, or something is missing or doubled.
- **Cuts.** By event source (POS flash sale, omni pending, ship-to-home), by posting age (same day, next day, older), and by who reads which number (Finance reads the book, Operations reads the count).
- **What it revealed.** The phantom-availability window is about one day on the overnight batch: at an illustrative 12 units a day on one SKU, 12 sold units sit “available” until the posting. And the two on-hand numbers were not wrong — Finance’s was right about the book, Operations’ was right about the floor; the defect was that nobody owned the gap between them.

FACET 2 — SOLUTION DESIGN — OPTIONS I PUT ON THE TABLE, AND WHY ONE WON NEW

One SKU, 12 units sold today (illustrative). The posting arrives for 10 units; one line (q = 2) is late. “Exposure” is the AFS error versus truth, in units.

OPTION	DAY-OF-SALE EXPOSURE	LATE-LINE (Q = 2) ERROR	VERDICT
Status quo — decrement at the S/4 posting	+12 (all sold units still “available”)	+2 until the next aggregate	Rejected — this is the problem

OPTION	DAY-OF-SALE EXPOSURE	LATE-LINE (Q = 2) ERROR	VERDICT
Real-time per-line posting to S/4	0	0	Rejected — throughput and posting cadence are why S/4 posts aggregates
Estimate-clear — release SNB at batch close	0	+2, invisible (reads as oversell later)	Rejected — the error hides inside AFS
Finance-derived SNB	0	depends on whose number	Rejected — two numbers; the point was one
event-fed SNB, cleared only on a matched posting	**0**	**−2, visible: aged row at 24 h**	Chosen — under-sell is a planning variance, over-sell is a customer defect

Why the chosen option is right even though it is "wrong" by −2: the error has a sign, the sign points at under-sell, and it lands on a report with a named owner.

ACTION — FACET 4: IMPLEMENTATION IN FIVE GATED PHASES (PHASES, OWNERS, CONTROLS, RISKS)

- **Discovery** — problem statement (the phantom-availability window); event inventory (POS flash sale, omni pending, ship-to-home confirm, S/4 Aggregation# posting); who-reads-which-number map. *Exit bar:* Finance and Ops agree on the description of the gap and which events move units before posting; Internal Audit named as reviewer.
- **Plan & Design** — SNB bucket definition (increment on events; decrement only on matched posting); join spec Aggregation# ↔ POS header ↔ lines; per-line COGS pro-rata de-aggregation with the last line absorbing rounding; the failure-direction statement (unmatched ⇒ SNB stays high); control-owner map. *Exit bar:* signed inside the FSD 6/5/25; Finance accepts one FOH; cost ownership confirmed S/4-owned, ISM visibility-only.
- **Build & Test** — de-aggregation test family (n lines; missing line; duplicate header); the SNB-never-negative assertion; the unmatched-posting aging report; sequenced by event source — POS flash sales, then omni pending, then ship-to-home — each added to SIT only when its increments and matched decrements reconciled. *Exit bar:* every posting matches or ages visibly; Σ per-line COGS = posting COGS to the cent in every cell; UAT = confirmation by Inventory Accounting.
- **Cutover & Go-Live** — riding the May go-live: SNB opening balance reconciled to unposted sales at the snapshot; SNB publish back to S/4 enabled; rollback cheap because S/4 stayed settlement authority. *Exit bar:* aging report empty or explained at T+1; no negative SNB; Finance reads FOH from ISM.
- **Hypercare / Close** — daily aging review with Inventory Accounting; control-owner handoff. *Exit bar:* aging report owned by Finance without me; month-end FOH ties to MB5L through the ISM number, not a bridge spreadsheet (mechanism canonical; days-to-tie not in record).

Influence and ownership. Mine: the SNB definition, the join spec, the de-aggregation rule, the failure-direction statement, the control-owner map and the test families; S/4 and OMS engineers implemented; Inventory Accounting confirmed in UAT and took over the aging report. Two refusals: real-time per-line posting to S/4 (throughput and posting cadence are why S/4 posts aggregates — ISM is the real-time ledger, S/4 stays settlement authority) and a Finance-derived SNB (that is two numbers; the whole point was one).

> **Designed the process; others own it now.** I designed the SNB definition, the join spec and the de-aggregation rule; Inventory Accounting owns the aging report and the daily review, and S/4 remained settlement authority. **Trade-off put to leadership.** The recommendation to Finance and Operations outlined the competing metrics — availability accuracy at the sale event vs the integrity of the financial posting — and chose the failure direction before the happy path: under-sell, never over-sell.

RESULTS — $ · # · % · TIME · CUSTOMER · QUALITY (CANON NUMBERS ONLY)

- **$** — ≈$0.6M/yr *modeled* (phantom-availability cancels avoided ≈$0.45M · Finance reconciliation labor ≈$0.14M · month-end close time ≈$0.05M; range $0.25–1.2M); the record has no pre-state drift dollar and I am not claiming one.

- # — Financial On-Hand truths 2 → 1 (record, FSD signed 6/5/25); over-sell from unmatched postings = 0 by construction.
- **%** — design KPI: match rate ≥ 99.5% within 24 h; the modeled phantom-cancel rate is ~0.10 pt of orders → ~0 (modeled, labeled).
- **Time** — AFS decrement point moved from the S/4 posting to the sale/ship event (record); rode the May execution go-live; aging report empty or explained at T+1 (record exit bar).
- **Customer** — the failure direction protects the customer: an unmatched posting under-states availability — under-sell, never over-sell (record mechanism).
- **Quality** — Σ per-line COGS = posting COGS to the cent; SNB never negative; month-end FOH ties to MB5L through the ISM number, not a bridge spreadsheet (record mechanism; days-to-tie not in record).

L6 LEVEL EVIDENCE — what the interviewer pattern-matches against the Amazon L6 functional rubric; every line is this story's canon, no new numbers

L6 DIMENSION	EVIDENCE IN THIS STORY
Ambiguity	A phantom-availability window between the sale event and the S/4 aggregation posting, with Finance and Operations reading different numbers and nobody owning the gap. I wrote the problem statement, the event inventory and the who-reads-which-number map.
Scope & influence	Finance (Inventory Accounting), Operations, S/4 and OMS engineers; rode the May go-live; Inventory Accounting confirmed in UAT and took the aging report — influence across a Finance organization I did not sit in.
Execution trade-off + mechanism others now own	Real-time per-line posting to S/4 (refused — S/4 stays settlement authority) vs an SNB bucket in ISM; the failure direction chosen before the happy path — under-sell, never over-sell. Mechanism others own now: the unmatched-posting aging report, owned by Finance without me.
Impact (long-term)	Financial On-Hand truths 2 → 1 (FSD signed 6/5/25); month-end FOH ties to MB5L through the ISM number, not a bridge spreadsheet; ≈$0.6M/yr modeled.
Process improvement (standard work institutionalized)	The de-aggregation test family (n lines; missing line; duplicate header); the SNB-never-negative assertion; the match-rate design KPI (≥ 99.5% within 24 h); a daily aging review handed to Inventory Accounting.
Complexity & consensus	Two refusals held in writing; sequencing by event source (POS flash sales, then omni pending, then ship-to-home) so each source carried its own defect class.
Root cause at scale	The structural cause — the AFS decrement point sitting at the S/4 posting instead of the sale/ship event — moved once, for every channel, rather than reconciled by hand each month.

PROBING — CHALLENGE Q&A TUNED TO THIS INTERVIEWER

> **Peer — Finance wanted their own reconciliation. How did you turn that around?**
>
> I showed the same number arriving from the physical side first and gave them the real report to own — the unmatched-posting aging report, reviewed daily with Inventory Accounting. A parallel spreadsheet dies when its owner owns the real report. They signed inside the FSD on 6/5/25.

> **Peer — Why decrement only on a matched posting? Isn't that just moving cost to Replenishment?**
>
> Yes, and I priced it that way on purpose: an unmatched posting under-states availability and shows up as an aged row Finance clears the same day, so the over-order is small and visible. Decrementing on an estimate would hide the gap as oversell. The measure is the aging report's p95 in hours and the match rate inside 24 hours; above 99.5% the Replenishment side rounds to zero.

> **Peer — Your cancel row looks borrowed from S1. Double-counted?**
>
> Different window: S1's rate is the vocabulary problem across node types; this is the posting-lag window inside one node. If the room wants one number, I keep S1's and drop this row — the mechanism still stands on the record.

> **Peer — Why not post every line to S/4?**
>
> Throughput and posting cadence — S/4 posts aggregates for a reason. ISM is the real-time ledger; S/4 stays the settlement authority. The price is a matching cycle to maintain; the benefit is one book.

> **Peer — 32 hours a week of Finance time — source?**
>
> Benchmark, not record — four people at eight hours a week is what Inventory Accounting described, not what a timesheet says. I would count manual journal entries and bridge lines per month before and after; that is in the posting log and it is not opinion.

> **Peer — Lesson?**
>
> Design the failure direction before the happy path. Differently: Internal Audit in discovery, not review — the day Inventory Accounting asked me why a posting had aged before I noticed it, the control had an owner.

FACET 3 — MATH & DATA LOGIC, FOUR TABLES, EVERY PLUG-IN SHOWN NEW

Table 1 — before: the problem in numbers.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Financial On-Hand truths	`truths = 2 (Finance's number vs Operations' number)`	reconciled by argument (record)	2
2	Phantom window	`W = t_post − t_sale ≈ 1 day`	AFS decremented at posting, not at sale (record design problem)	sold units "available" up to a day
3	Exposure per SKU	`12 units/day × 1 day = 12`	illustrative sales rate (matches the posting example below)	12 phantom units per SKU-day
4	Cancels it causes	`30,000,000 × 0.10% = 30,000`	orders illustrative; rate modeled	≈30,000 phantom cancels/yr
5	Bridge labor	`32 h/wk × 52 = 1,664`	bridge spreadsheet between the two books (modeled)	1,664 h/yr
6	Month-end bridging	`6 × 8 h × 12 = 576`	+1 day of close spent bridging (modeled)	576 h/yr

Table 2 — the fix: step by step, values substituted.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	SNB ledger rule	`SNB += each sale/ship event; SNB −= matched posting only`	POS / OMS / WMS events; decrement needs the Aggregation# match (record design)	walked on 12 units below
2	Walk — sales land	`SNB = 0 + 12 = 12`	12 units sold through the day (illustrative); AFS falls at each sale	AFS honest all day
3	Window closed	`exposure = 12 units/day × 0 days = 0`	decrement moved from t_post to t_sale	phantom window = 0
4	De-aggregation	`COGS_i = $300 × q_i / 12 → $75.00 + $50.00 + $100.00 + $25.00 + $50.00 = $300.00`	posting q_P = 12, C_P = $300.00; five lines q = 3, 2, 4, 1, 2 (illustrative); last line absorbs rounding	Σ = $300.00 to the cent
5	Matched posting	`SNB = 12 − 12 = 0`	12 posted with an Aggregation# match	books tie
6	Late line	`SNB = 12 − 10 = 2`	one line (q = 2) late, 10 matched (illustrative)	aged row visible at 24 h
7	Failure direction	`AFS_reported − AFS_true = −2`	unmatched ⇒ SNB high ⇒ AFS_reported ≤ AFS_true (record design)	under-sell, never over-sell

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
8	Match-rate KPI	`matched ÷ received ≥ 99.5% within 24 h`	design target (record)	≥ 99.5%

Table 3 — after: the metric, before → after.

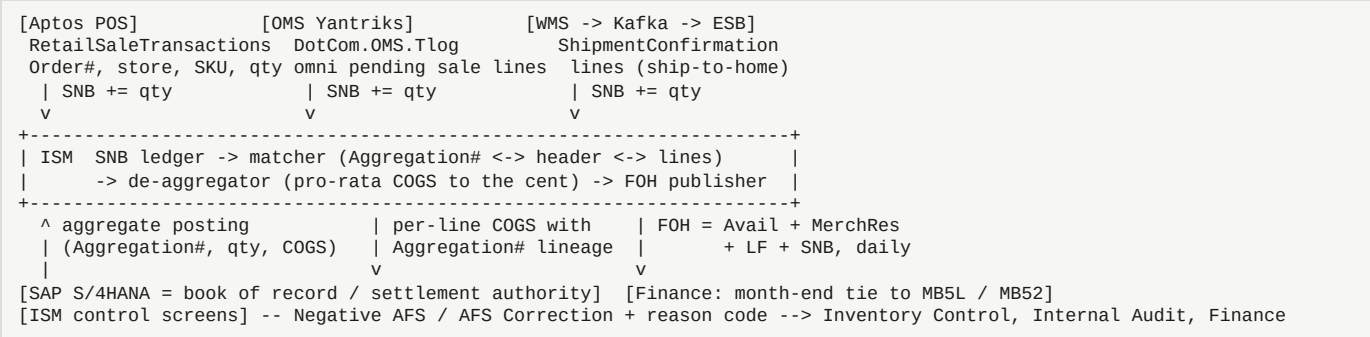
STEP	METRIC	BEFORE	AFTER	Δ / CHECK	BASIS
1	Financial On-Hand truths	2	1 — signed inside the FSD 6/5/25	`2 − 1 = 1`	record
2	AFS decrement point	at S/4 posting	at sale / ship	`moved`	record
3	Phantom-window exposure per SKU-day	12 units	0 units	`12 − 0 = 12`	record design; units illustrative
4	Phantom-availability cancels	~0.10 pt of orders	~0	`0.10 − 0 = 0.10 pt`	modeled
5	Finance bridge hours	32 h/wk	~0	`32 − 0 = 32 h/wk`	modeled
6	Month-end close	+1 day bridge	ties to MB5L	`1 − 0 = 1 day`	record design; days-to-tie not in record

Table 4 — the dollars, every multiplication shown.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Phantom cancels avoided	`30,000 × $15 = $450,000`	$15 all-in per cancel, range $10–20 (modeled); 30,000 from Table 1	≈$0.45M/yr
2	Reconciliation labor	`1,664 × $85 = $141,440`	hours from Table 1; $85/h benchmark (modeled)	≈$0.14M/yr
3	Month-end close	`576 × $85 = $48,960`	hours from Table 1 (modeled)	≈$0.05M/yr
4	Total	`$450,000 + $141,440 + $48,960 = $640,400`	all modeled; the record has no pre-state drift dollar	≈$0.6M/yr (range $0.25–1.2M)

Finance reads FOH from ISM; S/4 remains settlement authority; the aging report with a named Finance owner is the real product.

SOLUTION FLOW — THE SYSTEM & DATA MAP NEW



SENSITIVITY / IMPACT — LOW / BASE / HIGH NEW

INPUT	LOW	BASE	HIGH
Phantom-cancel rate avoided	0.05 pt	0.10 pt	0.15 pt
Cost per cancel (benchmark)	$10	$15	$20
Cancel $ avoided	$0.15M	$0.45M	$0.90M

INPUT	LOW	BASE	HIGH
Bridge hours per week × loaded rate	20 × $75 = $0.08M	32 × $85 = $0.14M	48 × $95 = $0.24M
Close days saved × rate	0.5 day → $0.02M	1 day → $0.05M	2 days → $0.11M
Total modeled $/yr	≈$0.25M	≈$0.6M	≈$1.2M

SPOKEN TOTAL-IMPACT STATEMENT — SAY IT LIKE THIS

Conservatively, about six hundred thousand a year: roughly four hundred fifty thousand of cancels that stop happening because availability drops at the sale instead of at the posting, about a hundred forty thousand of Finance hours that were going into a bridge spreadsheet, and about fifty thousand of month-end close time. What is measured is that Finance and Operations went from two on-hand numbers to one, decremented at sale or ship, with the failure direction pointing at under-sell. Every dollar is my model on an illustrative thirty-million-order scale; the record has no pre-state drift dollar and I am not claiming one.

FACET 5 — HARDEST FOLLOW-UPS, WITH NUMERIC ANSWERS NEW

Numeric — Walk me through twelve units and a late line.

Twelve units sold through the day, so SNB goes 0 → 12 and AFS falls at each sale — the window is 12 × 0 = 0. The posting arrives for 10 matched units: SNB = 12 − 10 = 2, and those two sit on the aging report at 24 hours. Because S/4 already decremented the posting, AFS under-states by 2 — under-sell, never over-sell.

Numeric — Prove the de-aggregation ties.

Posting: 12 units, $300.00. Cost per unit is $300 ÷ 12 = $25.00, so lines of 3, 2, 4, 1 and 2 units carry $75, $50, $100, $25 and $50 — and 75 + 50 + 100 + 25 + 50 = $300.00. When the cost per unit does not divide evenly the last line absorbs the cent, so the sum is exact by construction.

Numeric — Which of the $0.6M would you defend and which would you drop?

Defend the mechanism, price the rest: $0.45M cancels + $0.14M labor + $0.05M close = $0.64M, all modeled. The cancel row is three quarters of it and is the weakest; if the room wants one cancel number I keep S1's and drop this row, leaving $0.14M + $0.05M = $0.19M — and the mechanism still stands on the record.

Numeric — What is the match rate and why 99.5%?

Matched postings ÷ received within 24 hours ≥ 99.5% is a design KPI, not a measured result. At an illustrative 2,000 postings a day, 0.5% is 10 unmatched rows a day on the aging report — a list one analyst can clear the same day; at 5% it would be 100 rows and the report stops being a control.

BEHAVIORAL PACK — FOUR QUESTIONS THIS STORY ANSWERS, EACH A DIFFERENT 90-SECOND MOMENT NEW

EARN TRUST **1. Tell me about a time you got two teams with competing numbers to agree.**

Open (first two sentences): Finance and Operations read different on-hand numbers and reconciled by argument at every close. I did not ask Finance to trust mine — I showed them the same number arriving from the physical side first.

A. Finance wanted its own reconciliation. I put both numbers on one table, showed the gap was a window — sold units not yet posted — not an error, and then gave them the unmatched-posting aging report to own, with Inventory Accounting reviewing it daily. A parallel spreadsheet dies when its owner owns the real report. **R.** Financial On-Hand went from two numbers to one, signed inside the FSD on 6/5/25.

Land the number: two numbers → one, signed 6/5/25 | **Pivot to NACF:** The same move at an inbound dock: let the receiving team own the received-not-stowed aging list so the planning number and the dock number are one.
Do NOT say: "Finance signed because I was persuasive"; a claimed days-to-tie figure — the record has none.

INVENT AND SIMPLIFY **2. Tell me about a time you simplified something complex.**

Open (first two sentences): The month-end bridge was a spreadsheet between two books that disagreed for a mechanical reason. Instead of improving the spreadsheet, I removed the reason.

A. Four moving parts became one rule: SNB goes up on every sale or ship event and goes down only on a matched posting; S/4 stays settlement authority; the failure direction is under-sell; one aging report. The simplification is that the bridge no longer has a job. **R.** One Financial On-Hand number; month-end FOH ties to MB5L through the ISM number rather than a bridge; about $0.6M a year modeled.

Land the number: one rule, one report, one number | **Pivot to NACF:** For inbound: replace the three-way reconciliation of ASN, dock count and system on-hand with one ledger of received-not-stowed.
Do NOT say: "I built the matcher" — engineers implemented the design.

HAVE BACKBONE; DISAGREE AND COMMIT **3. Tell me about a time you refused something a stakeholder preferred.**

Open (first two sentences): Two proposals were on the table that would each have shipped faster: post every line to S/4 in real time, and let Finance derive the SNB itself. I refused both, on mechanism.

A. Per-line posting fights the throughput and cadence S/4 posts aggregates for — ISM is the real-time ledger and S/4 stays settlement authority. A Finance-derived SNB is two numbers, and the whole point was one. Both refusals went into the FSD as rows with the reason, and each person who proposed one got a replacement that served their need — the real-time view, and a report they owned. **R.** No rework of posting cadence; FSD signed 6/5/25.

Land the number: two refusals, both in the FSD as rows | **Pivot to NACF:** In placement, "just send everything to the nearest FC" is the same shape: fast, and it moves cost somewhere nobody is watching.
Do NOT say: Naming the people who proposed them; implying S/4 is slow.

CUSTOMER OBSESSION **4. Tell me about a time you designed for the failure case to protect the customer.**

Open (first two sentences): Before I designed the happy path I wrote down which way the system should be wrong. An unmatched posting could make availability too high — the customer is told yes to a unit that is gone — or too low — a unit sits unsold for a day.

A. I picked too low, on purpose: SNB stays high until the posting matches, so AFS under-states and Replenishment may over-order visibly. On the worked SKU the late line of 2 units is a −2, not a +2. I priced both sides: an over-stated unit costs about $30 as an incident, an under-stated one about $7.60 of expected lost margin (40% lost-sale share × $19). **R.** Over-sell from unmatched postings is zero by construction.

Land the number: −2 visible, never +2 invisible | **Pivot to NACF:** For inbound: if a receipt is uncertain, show the lower available-to-promise; the cost of a broken date is higher than a late-shown one.
Do NOT say: "Zero oversell" as a measured rate; the 40% and $19 are benchmark inputs, say so.

SAME STORY, 3 ANGLES — DIFFERENT OPENER, HEADLINE, LP NEW

DATA / ANALYTICAL

"One posting of twelve units, five lines, three hundred dollars — and they have to sum to the cent."

$300.00 = 75 + 50 + 100 + 25 + 50

LP: Dive Deep

The grain mismatch — one posting, many lines — was the whole problem.

Cost per unit $300 ÷ 12 = $25.00; each line is q × $25; the last line absorbs rounding.

The test family: n lines, a missing line, a duplicate header — each has one expected SNB.

CONFLICT / INFLUENCE / BACKBONE

"Finance wanted to keep its own reconciliation, and I understood why."

2 numbers → 1

LP: Earn Trust · Have Backbone

I showed the same number arriving from the physical side first.

Then I refused a Finance-derived SNB — two numbers — and gave them the aging report to own.

A parallel spreadsheet dies when its owner owns the real report.

FAILURE / LEARNING

"The thing I would do differently is when I brought Internal Audit in."

T+1 aging report — owner named late

LP: Learn and Be Curious

I treated Internal Audit as a reviewer at the end; the day Inventory Accounting asked why a posting had aged before I noticed, the control had an owner — but it should have had one in discovery.

No days-to-tie figure exists in the record; I do not claim one.

The cancel row is the weakest dollar; I keep it labeled modeled.

TRAP DRILLS — THE THREE NASTIEST FOLLOW-UPS, WITH THE EXACT SENTENCE NEW

Trap — "Your biggest dollar row borrows S1's rate. Strip it — what is left?"

Different window — S1 is vocabulary across node types, this is posting lag inside one node — but if the room wants one number I keep S1's and drop this row: $141,440 + $48,960 = $190,400, about $0.19M, entirely modeled, and the mechanism still stands on the record.

Say exactly: "If you want one cancel number I keep S1's; what stays here is about nineteen cents on the dollar of that — a hundred ninety thousand — and the mechanism."

Trap — "Why is under-sell the right failure direction? What number would change your mind?"

Compare unit costs: an over-stated unit that turns into an incident costs about $30 (S1's benchmark build); an under-stated unit costs expected lost margin, 40% lost-sale share × $19 = $7.60. Under-sell is right whenever the chance that an over-stated unit becomes an incident exceeds 7.60 ÷ 30 ≈ 25%. If cycle data showed fewer than one in four over-stated units ever became incidents, I would loosen the rule.

Say exactly: "Under-sell wins as long as more than about one in four over-stated units would turn into a customer incident — and I would measure that, not assume it."

Trap — "Halve the S/4 team's capacity. What do you cut, and what can you never cut?"

I cut by event source, last first: POS flash sales stay, omni pending next, ship-to-home last — each source was already sequenced to enter SIT only when its increments and matched decrements reconciled, so a cut source simply stays on the old overnight path. I never cut the matched-only clear, the SNB-never-negative assertion or the aging report with a named owner — without them the failure direction is gone.

Say exactly: "I would cut the sources, never the controls: ship-to-home waits; matched-only clear and the aging report do not."

Where this lands at Amazon. Every fulfillment event moves units before Finance posts them — the same window at a vastly larger grain. The transferable design: one real-time ledger, the financial system as settlement authority, a safe failure direction chosen before the happy path, and a control report with a named owner instead of a bridge spreadsheet. Reliability (over-sell zero by construction), cost efficiency (bridge labor and close time), one number for two organizations.

TWIST MAP — LIKELY LP QUESTIONS FROM THIS INTERVIEWER → THE ANGLE TO TAKE WITH S3

LIKELY QUESTION (THIS INTERVIEWER)	LP	ANGLE — WHAT TO EMPHASIZE	SKIP / NEVER SAY
Tell me about a time you got Finance and Operations to agree.	Earn Trust	FOH 2 → 1; the aging report given to Finance to own; signed inside the FSD on 6/5/25.	Never claim a pre-state drift dollar — the ≈$0.6M is modeled.
Tell me about a problem that needed you to understand the data at its grain.	Dive Deep	One posting per aggregate vs transaction lines; the Aggregation# ↔ header ↔ lines join; COGS pro-rated to the cent.	Skip the table names unless asked.
Tell me about simplifying something complex.	Invent and Simplify	One ledger, one settlement authority, one failure direction, one control report — instead of two books and a bridge.	Do not say "I built the matcher" — engineers implemented the design.
How do you design controls for a process?	Functional — control design	Choose the failure direction first; a never-negative assertion; a KPI with a denominator (match rate ≥ 99.5% in 24 h); an aging report with a named owner; rollback that is cheap because the authority never moved.	Keep S14's standing reconciliation report out — bar raiser's.
What did you refuse and why?	Have Backbone	Per-line posting to S/4; a Finance-derived SNB — both refused on mechanism, not preference.	—
What would you do differently?	Learn and Be Curious	Internal Audit in discovery, not review.	No days-to-tie figure exists — do not invent one.

S1-a — Sephora S/4HANA ↔ ISM — the inventory supply data contract: ten buckets, one owning system each, and every event into exactly one bucket (the double-count path closed on paper before any code) NEW

LPs this story carries (with proof lines below): Dive Deep (primary — three failure classes found at event grain) · Insist on the Highest Standards (drift tolerance zero) · Invent and Simplify (one event, one bucket) · Ownership (zero open "who owns this bucket" rows at signature)
Process family: Plan / Deliver seam — data contract design: buckets, systems of record, event-to-bucket mapping, publish budgets, daily reconciliation

Honesty label. I owned the contract — the bucket inventory, the owner column, the event-to-bucket map, the idempotency and reversal rules, the publish and lag budgets and the drift check — not the build. The ISM app team and SAP configurators implemented to the signed FSD. The unit walk below is illustrative; the equation, rules and budgets are record design. Two ownership cells are tagged VERIFY-WITH-RAJESH because the master names the systems that feed each bucket but not every bucket-to-system pairing.

LEADERSHIP PRINCIPLES CARRIED — ONE-LINE PROOF EACH NEW

Dive Deep — I stopped arguing about "available" and walked a five-unit order event by event until I found three distinct ways one bad event moves AFS by ±5 — double-bucket, replay, double reversal.

Insist on the Highest Standards — The daily drift check has tolerance zero — a non-zero ISM-vs-MB52 row is a Sev, not a planning variance — and I refused to write a "reasonable tolerance" into the FSD.

Invent and Simplify — The rule that replaced a page of exceptions: every event maps to exactly one bucket, and a move between buckets is one atomic event, never two.

Ownership — The contract has an owner column; I did not exit design until it had zero open "who owns this bucket" rows, and I carried vocabulary drift as my own top risk.

60-SECOND BUSINESS-FIRST PITCH — STOP ON THE PRIMARY METRIC

Situation. ISM sits beside S/4 as a real-time ledger, so for every SKU at every node a unit can sit in one of ten buckets — S/4 available, WMS on-hand, the out-of-balance between them, the basis, three committed buckets, shipped-not-billed, merchandise reserve and lost-and-found. If two systems can each move the same unit, or one system can move it twice, available-for-sale is wrong in a way no one can trace. **Task.** Write the data contract: which buckets exist, which system owns each, how every event moves units, and how fast each number must reach the order system — signed before configuration. **Action.** I wrote a bucket inventory with one owner per bucket; an event-to-bucket map where every event touches exactly one bucket; an idempotency key per event and a rule that a reversal must reference its original; a 60-second publish budget, a WMS feed of at most 15 minutes and consumer lag under 5 minutes; and a daily ISM-versus-MB52 comparison with tolerance zero. I proved it on a five-unit order walked event by event, because that walk is where the three failure classes showed up. **Result.** The FSD was signed 6/5/25 with zero open "who owns this bucket" rows; no SKU carries two published AFS values; drift was zero SKUs on consecutive days at hypercare exit.

20-SECOND VERSION NEW

ISM tracks ten inventory buckets beside S/4. I wrote the contract that gives every bucket one owning system and forces every event into exactly one bucket, with an idempotency key so a replay is a no-op and a reversal that must name its original. On a five-unit order the contract keeps AFS at 555 where three bad events would have moved it ±5; the FSD was signed 6/5/25 with zero unowned buckets.

NACF BRIDGE NEW

Inbound needs the same contract: a unit on an ASN, at the dock door, in receiving, in stow and on a rack is five states with one owner each, and a pallet moves between them as one atomic event. Inbound available-to-promise is then opening on-hand plus units that cleared receive and stow — and a daily count-versus-system check with a stated tolerance.

SITUATION / TASK — THE BUSINESS PROCESS PROBLEM, SCOPE, SCALE, BASELINES

S — Three readers of the same SKU — Store, Retail DC and Dotcom DC — each computed "available" from a different subset of what is really ten buckets, and nothing said which system was allowed to move a unit between them. Events came from the POS (Aptos), the order system (OMS, Yantriks), the WMS (BlueYonder) and S/4 itself (MB52 daily plus movement events).

T — Own the data contract of inventory supply: buckets, owners, event rules, budgets and the standing reconciliation, signed by Supply Chain Strategy, Inventory Management and FSC Inventory Accounting before any configuration.

Baselines on the table: bucket owners — undefined (target 1 each, zero open rows) · events able to touch more than one bucket — unknown (target 0) · replay and reversal behavior — undefined (target idempotent, reversal references its original) · publish latency — unbudgeted (target 60 s) · WMS feed ≤ 15 min · consumer lag < 5 min · daily drift — unmeasured (target 0 SKUs).

FACET 1 — DATA ANALYSIS — WHAT I PROFILED, AT WHAT GRAIN, WHAT IT REVEALED NEW

- **Data.** Every feed that moves a unit: Aptos POS sale lines, OMS commit/pick/pack/cancel events, BlueYonder on-hand by item × location × status (≤ 15 min), FRAME channel reserves and HUD holds, and S/4 (MB52 daily plus MSEG movements as events — receipts, transfers, deliveries, count differences).
- **Grain.** SKU × node × bucket × event. One order of five units was traced across all four feeds, one event at a time, with the bucket balance after each.
- **Cuts.** By feed, by bucket, and by failure class: what happens when an event touches two buckets, arrives twice, or is reversed twice.
- **What it revealed.** Three distinct ways a single five-unit order moves AFS by five units off truth: a pick recorded as a new pack without releasing the pick (−5), a commit delivered twice with no key (−5), and a cancel applied twice (+5 — an oversell in waiting). None needs bad data; each only needs an unwritten rule.

FACET 2 — SOLUTION DESIGN — OPTIONS I PUT ON THE TABLE, AND WHY ONE WON NEW

Same SKU, state AFS = 560 (Basis 980 − 200 − 120 − 60 − 40; illustrative). A 5-unit web order commits, picks, packs and ships; then a bad event arrives. "Error" is AFS minus truth (555).

OPTION	WHAT IT DOES	AFS ERROR	VERDICT
Rich events — one event updates several buckets	pick event writes EcomPack +5 and forgets EcomPick −5	−5	Rejected — a half-applied move is a double count
At-least-once delivery, dedupe downstream	commit arrives twice; the consumer sees two orders	−5	Rejected — dedupe belongs at the event, not in each reader
Reversals as free-standing events	a cancel applied twice returns 10 units for a 5-unit order	+5	Rejected — a reversal must reference its original
Each consumer computes its own AFS from raw buckets	OMS and the dotcom site each run the equation	two published values	Rejected — two publishers; OMS reads, never computes

OPTION	WHAT IT DOES	AFS ERROR	VERDICT
15-minute snapshot polling instead of events	sales appear when the snapshot lands	up to a 15-min window	Rejected — misses the real-time purpose
one event, one bucket delta; key per event; reversal names its original; one publisher	atomic move between buckets, replay is a no-op	**0**	Chosen — the error cannot be written

ACTION — FACET 4: IMPLEMENTATION IN FIVE GATED PHASES (PHASES, OWNERS, CONTROLS, RISKS)

- **Discovery** — the bucket inventory — ten rows with one proposed owning system each — and the event inventory by feed. *Exit bar:* every bucket owned by exactly one system; sessions closed 4/30; open questions carried with owner and date.
- **Plan & Design** — the equation and the contract: lower-of-two basis, AFS = Basis − PickRes − EcomPick − EcomPack − SNB; event → one bucket; idempotency key per event; reversal references original; 60-second publish budget; WMS feed ≤ 15 min; consumer lag < 5 min; drift tolerance zero; FSD v1.1 (3/5/25) → v1.2 (5/16/25). *Exit bar:* signed 6/5/25 by three orgs with zero open “who owns this bucket” rows.
- **Build & Test** — the node × channel × bucket-transition matrix with before/after assertions — the five-unit order is the template cell; test families for replay, duplicate, out-of-order and double reversal; SIT on the interface contracts before UAT. *Exit bar:* 0 open contract defects in SIT; every matrix cell has an owner.
- **Cutover & Go-Live** — the publish-switch plan — OMS reads the published AFS and never computes; no SKU with two published values; control-screen roles assigned. *Exit bar:* January and May go-lives; Negative AFS clean at T+1.
- **Hypercare / Close** — daily ISM-vs-MB52 drift to Inventory Control; consumer-lag and publish-latency monitors against budget. *Exit bar:* drift = 0 SKUs on consecutive days; blockers/day under threshold; exited into the monthly release train.

Influence and ownership. Mine: the bucket inventory, the owner column, the event map, the idempotency and reversal rules, the budgets, the drift definition and the matrix template cell. Built by: the ISM application team and SAP configurators; the POS, OMS and WMS teams exposed the events. The hard part was not the rules — it was getting three orgs to accept that a row without an owner blocks signature.

RESULTS — $ · # · % · TIME · CUSTOMER · QUALITY (BASIS-LABELED)

- **#** — Buckets with an owning system: ten of ten at signature; zero open “who owns this bucket” rows (record). Events able to touch more than one bucket: 0 by contract (record design).
- **%** — Budgets: publish ≤ 60 s; WMS feed ≤ 15 min; consumer lag < 5 min; drift tolerance 0 — all record design. No SKU with two published AFS values.
- **Time** — FSD v1.1 3/5/25 → v1.2 5/16/25 → signed 6/5/25; E2E Inbound 11/22 · Dotcom 12/20 · Retail 1/30; January and May go-lives (record dates).
- **Quality** — Daily ISM-vs-MB52 drift = 0 SKUs on consecutive days at hypercare exit (record). The three failure classes closed: double-bucket −5, replay −5, double-reversal +5 each reduce to 0 on the five-unit walk.
- **$** — No new dollars. The contract’s value is a slice of S1’s modeled oversell row — illustrative 25% × $900,000 = $225,000 VERIFY-WITH-RAJESH — and is never added to S1’s ≈$2.5M.

L6 LEVEL EVIDENCE — what the interviewer pattern-matches against the Amazon L6 functional rubric; every line is this story's record or labeled model

L6 DIMENSION	EVIDENCE IN THIS STORY
Ambiguity	Three readers, ten buckets, no owner for the definition of a move. I defined the buckets, owners and the event map from first principles by walking one order.
Scope & influence	Four feeds (POS, OMS, WMS, S/4) and three signatory orgs; none reported to me; the contract was signed 6/5/25.
Execution trade-off	Atomic one-bucket events cost more events and a stricter producer contract; I accepted that to make double-counting unwritable. Under-sell accepted over over-sell.
Impact	No SKU with two published AFS values; drift zero on consecutive days; ≈$0.225M illustrative slice of S1's oversell row, never added.
Process improvement	The five-unit order is the template cell of the transition matrix; no story enters QA without a matrix cell reference.
Complexity	Real-time events, a ≤ 15-minute floor feed and a nightly aggregate posting reconciled under one definition.
Root cause	The cause was three missing rules (one bucket per event, a key per event, reversals name originals), fixed once in the contract rather than per incident.

PROBING — FACET 3: MATH & DATA LOGIC, FOUR TABLES, EVERY PLUG-IN SHOWN

Table 1 — before: the problem in numbers.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Starting state	`AFS = 980 − 200 − 120 − 60 − 40 = 560`	S4_Avail 1,000; WMS 980; PickRes 200, EcomPick 120, EcomPack 60, SNB 40 (illustrative; Basis = min(1,000, 980))	560 units
2	Truth after a 5-unit web order commits	`AFS = 980 − 200 − 125 − 60 − 40 = 555`	EcomPick 120 + 5 = 125 (illustrative)	555 units
3	Failure class A — one event writes two places	`AFS = 980 − 200 − 125 − 65 − 40 = 550`	pick recorded as EcomPack +5 without EcomPick −5: both 125 and 65 carry the 5 units	550; error 550 − 555 = −5
4	Failure class B — commit replayed, no key	`AFS = 980 − 200 − 130 − 60 − 40 = 550`	EcomPick 125 + 5 = 130	550; error −5
5	Failure class C — cancel applied twice	`AFS = 980 − 200 − 115 − 60 − 40 = 565`	EcomPick 125 − 5 − 5 = 115 for a 5-unit order	565; error 565 − 555 = +5 (oversell in waiting)
6	Floor staleness at the 15-minute WMS feed	`2 units/h × (15 ÷ 60) h = 0.5`	floor-only changes (damage, lost, relocation) — illustrative 2 units/h; sales do not wait for the feed	≈0.5 unit worst case

Table 2 — the fix: step by step, values substituted.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Rule 1 — one event, one bucket	`commit: EcomPick 120 → 125; AFS = 980 − 200 − 125 − 60 − 40 = 555`	event maps to exactly one bucket delta (record design)	AFS 555
2	Rule 1 — a move is one atomic event	`pick: EcomPick 125 − 5 = 120 and EcomPack 60 + 5 = 65 together; AFS = 980 − 200 − 120 − 65 − 40 = 555`	both legs commit or neither does	AFS stays 555
3	Ship confirm	`EcomPack 65 − 5 = 60 and SNB 40 + 5 = 45; AFS = 980 − 200 − 120 − 60 − 45 = 555`	event → SNB; WMS and S/4 books have not moved yet	AFS stays 555

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
4	Overnight posting and floor feed	`S4_Avail 1,000 − 5 = 995; WMS 980 − 5 = 975; SNB 45 − 5 = 40; AFS = min(995, 975) − 200 − 120 − 60 − 40 = 555`	both books and SNB move together at the posting (record design)	AFS stays 555
5	Rule 2 — replay with the same key	`AFS = 555 − 0 = 555`	second delivery of the commit key is a no-op (record design)	AFS 555
6	Rule 3 — reversal names its original	`cancel of the 5-unit order: EcomPick 125 − 5 = 120; a second reversal finds the original already reversed → rejected; AFS = 980 − 200 − 120 − 60 − 40 = 560`	a reversal must reference its original (record design)	AFS 560 = starting state
7	Latency budget	`60 s + 15 min = 60 + 900 = 960 s`	publish budget plus WMS feed — worst case from a floor-only change to a published AFS	960 s = 16 min
8	Daily drift	`drift = ISM.S4_Avail − MB52 = 995 − 995 = 0 for every SKU × node row`	tolerance 0 (record design); a non-zero row is a Sev	0 SKUs

Table 3 — after: the metric, before → after.

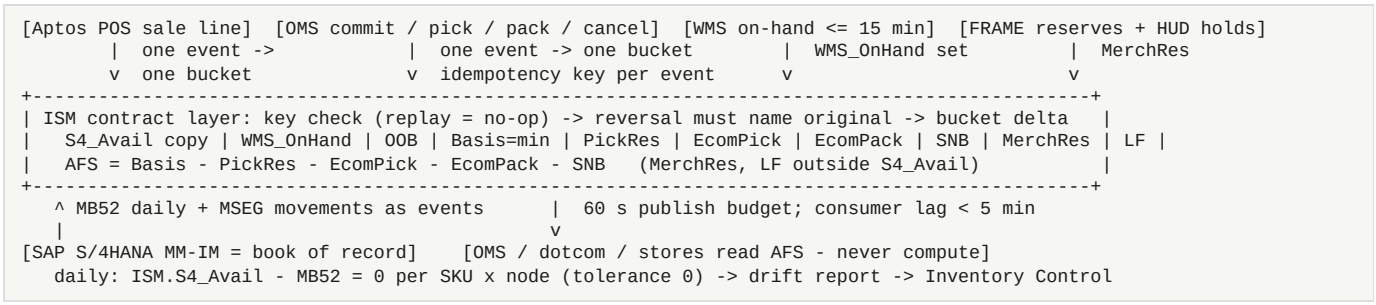
STEP	METRIC	BEFORE	AFTER	Δ / CHECK	BASIS
1	Owners per bucket	undefined	1 owning system each; 0 open rows	`10 − 10 = 0 unowned`	record
2	Events able to touch > 1 bucket	unknown	0 by contract	`0`	record design
3	Replay of the same key	AFS 550 (−5)	AFS 555 (0)	`5 units`	record design; units illustrative
4	Reversal applied twice	AFS 565 (+5)	rejected; AFS exact	`5 units`	record design; units illustrative
5	Two published AFS values per SKU	possible	0 SKUs	`0`	record
6	Daily drift	unmeasured	0 SKUs on consecutive days	`0`	record

Table 4 — the dollars, every multiplication shown.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	S1 oversell row (parent)	`30,000,000 × 0.10% = 30,000; 30,000 × $30 = $900,000`	orders illustrative; rate benchmark 0.05–0.25%; $30 = $6 + $4 + $19 + $1 (modeled)	$900,000 (S1's row)
2	Event-level share (illustrative)	`$900,000 × 25% = $225,000`	25% attributable to event-level double counting VERIFY-WITH-RAJESH; illustrative	≈$0.23M — a slice of S1, not additive
3	Price of one class-C event	`1 incident × $30 = $30`	one +5 over-state becomes at most one incident per order (benchmark)	$30 per event
4	Check	`$225,000 ÷ $30 = 7,500`	implied incidents in the slice	7,500 of 30,000 incidents

Walk values are illustrative units; the equation, the three rules, the budgets and the tolerance are record design. The dollar slice is illustrative, flagged for verification, and sits inside S1's ≈$2.5M — never on top of it.

SOLUTION FLOW — THE SYSTEM & DATA MAP



SENSITIVITY / IMPACT — LOW / BASE / HIGH

INPUT	LOW	BASE	HIGH
Share of oversell from event-level double counting	10% → $0.09M	25% → $0.23M	40% → $0.36M
Parent oversell row (S1)	$0.4M	$0.9M	$3.0M
WMS feed interval	5 min	15 min	30 min
Worst-case floor-only exposure at 2 units/h	0.17 unit	0.5 unit	1.0 unit
Publish budget + feed (seconds)	60 + 300 = 360	60 + 900 = 960	60 + 1,800 = 1,860

SPOKEN TOTAL-IMPACT STATEMENT — SAY IT LIKE THIS

Conservatively, there is no new dollar here, and I will not manufacture one: this contract is the mechanism inside S1's two and a half million, and my illustrative slice is about a quarter of the oversell row, about two hundred twenty-five thousand, which I would drop before defending. What I can defend is on the record: ten buckets, ten owners, zero open rows when three organizations signed on 6/5/25, and drift of zero SKUs on consecutive days at hypercare exit.

FACET 5 — HARDEST FOLLOW-UPS, WITH NUMERIC ANSWERS NEW

Numeric — Walk me through the five-unit order.

Start AFS 560. The order commits: EcomPick 120 → 125, AFS 555. Pick moves 5 from EcomPick to EcomPack in one event, 555. Ship moves 5 from EcomPack to SNB, 555. The overnight posting moves S/4, the WMS feed and SNB together by 5 — min(995, 975) − 200 − 120 − 60 − 40 = 555. AFS never moves except at the commit.

Numeric — What breaks if one rule is missing?

Without "one event, one bucket" the pick leaves 5 units in two buckets: AFS 550, −5. Without the key, a replayed commit does the same, −5. Without "a reversal names its original", a cancel applied twice gives back 10 for a 5-unit order: 565, +5 — that one is an oversell waiting to happen.

Numeric — How stale can the number be?

Events are real time inside a 60-second publish budget. The floor feed is up to 15 minutes, so the worst case for a change that only the floor sees — a damage or a relocation — is 60 + 900 = 960 seconds, about 16 minutes; at an illustrative 2 units an hour that is half a unit. Sales do not wait for the feed because they decrement the committed buckets immediately.

Numeric — Why tolerance zero? Isn't that unrealistic?

Because both books are event-fed, a non-zero drift means a missing or duplicated event — a defect with a cause — not noise. At an illustrative 5,000 SKUs × 10 nodes = 50,000 rows VERIFY-WITH-RAJESH, "within 1%" would tolerate 500 rows of unexplained difference; zero forces each into the report, where Inventory Control names the cause.

BEHAVIORAL PACK — FOUR QUESTIONS THIS STORY ANSWERS, EACH A DIFFERENT 90-SECOND

MOMENT NEW

DIVE DEEP **1. Tell me about a time you dug into the details and found something others missed.**

Open (first two sentences): Everyone was arguing about the definition of “available.” I stopped, took one five-unit order, and walked it event by event through four feeds, writing the bucket balance after each one.

A. The walk showed three separate ways the same order could leave AFS five units off — a pick written to two buckets, a commit replayed without a key, a cancel applied twice. None needed bad data, only a rule nobody had written. I turned each into a rule and the walk into the template cell of the test matrix. **R.** The contract was signed 6/5/25 with zero open ownership rows.

Land the number: three failure classes, ±5 units each | **Pivot to NACF:** The same walk on an inbound pallet — ASN, door, receive, stow — shows where available-to-promise could be double counted.
Do NOT say: Claiming a measured defect rate — the walk is illustrative; the rules are record design.

INSIST ON THE HIGHEST STANDARDS **2. Tell me about a standard you held when others wanted to relax it.**

Open (first two sentences): In design review someone asked for a drift tolerance — “within 1% is fine.” I held at zero.

A. Both books are fed by events, so a non-zero difference means a missing or duplicated event. I showed that 1% of an illustrative 50,000 SKU × node rows is 500 unexplained rows a day, and that zero turns each into a named row for Inventory Control. I did write a rule for the exceptions: the row is a Sev, with a reason code, not a quiet variance. **R.** Drift = 0 SKUs on consecutive days at hypercare exit.

Land the number: tolerance 0 vs 500 rows | **Pivot to NACF:** At inbound: a receive count compared with the ASN should have a stated tolerance and a named owner — not “close enough.”
Do NOT say: Claiming the 50,000 rows is real VERIFY-WITH-RAJESH.

INVENT AND SIMPLIFY **3. Tell me about a time you replaced a complicated process with a simple rule.**

Open (first two sentences): The first draft of the event handling had a page of exceptions. I replaced it with one sentence: every event touches exactly one bucket.

A. A move between buckets became one atomic event with two legs that commit together, so “half-applied” could not exist. The exceptions page collapsed into the idempotency key and the reversal-names-original rule. **R.** On the five-unit walk AFS holds at 555 through every event.

Land the number: one page of exceptions → one sentence | **Pivot to NACF:** In placement, the same simplification is “each unit has exactly one planned destination.”
Do NOT say: “I invented a new data model” — it is the SAP/OMS events under one rule.

OWNERSHIP **4. Tell me about a time you took on something no one owned.**

Open (first two sentences): The contract needed an owner column, and the awkward finding was that for some buckets nobody could say who was allowed to move them.

A. I did not exit design until the column had no blanks, and I carried “a bucket with no owner” as my own risk. Where two systems both claimed a bucket I wrote the tie-break into the FSD with the signatories. **R.** Zero open rows at signature; after handoff Inventory Control owns the drift report. Two pairings are flagged for verification VERIFY-WITH-RAJESH.

Land the number: 0 unowned buckets at signature | **Pivot to NACF:** Stranded inbound stock is the same orphan bucket.
Do NOT say: Claiming I owned any of the systems.

SAME STORY, 3 ANGLES — DIFFERENT OPENER, HEADLINE, LP NEW

DATA / ANALYTICAL

"I traced one five-unit order through four feeds and wrote the balance after every event."

560 → 555 → 555

LP: Dive Deep

Truth after the commit is 555; three bad events give 550, 550 and 565.

None needed bad data — each needed a missing rule.

The walk became the template cell for the matrix.

CONFLICT / INFLUENCE / BACKBONE

"The room asked for "within one percent" on drift and I said no."

tolerance 0

LP: Insist on the Highest Standards · Have Backbone

Event-fed books drifting means a missing or duplicate event.

I priced the alternative: an illustrative 500 unexplained rows a day.

The compromise was a reason-coded Sev row, not a tolerance.

FAILURE / LEARNING

"The contract held; the first promising rule I put on top of it — no channel may promise units in transit — was wrong."

120 hidden units → 120 promisable

LP: Are Right, A Lot

A contract that says what a bucket is does not say whether a rule on top of it is right.

The in-transit rule (S4) starved fast movers; I corrected it through the same signatories with a window and a kill switch.

Lesson: back-test any rule against history before publishing; the data contract is the right place to write the kill switch.

TRAP DRILLS — THE THREE NASTIEST FOLLOW-UPS, WITH THE EXACT SENTENCE NEW

Trap — "Your 25% slice — where does it come from?"

It does not come from the record and I say so: it is an illustrative share I would drop before defending. What I can defend is the mechanism: three failure classes at event grain. To get a real share I would take a quarter of reason-coded incident rows and count how many trace to a replay, a double-bucket or a double reversal.

Say exactly: "The twenty-five percent is a placeholder, flagged for verification; the real number is a count of coded incidents by failure class."

Trap — "What is the dependency risk in this contract?"

Four producers have to honour it — POS, OMS, WMS, S/4 — and I own none. Dependency order: the key and reversal rules need the OMS and POS teams first because they emit the events; the 15-minute feed is the WMS team's; the drift check needs S/4's MB52. The one I would cut if a team slipped is the 60-second budget on non-sale events — never the key or the reversal rule.

Say exactly: "I would relax a latency budget before I relaxed idempotency or the reversal rule — one costs minutes, the other costs trust."

Trap — "How would you know the contract is working without waiting for a customer incident?"

Three leading checks: the drift report (target 0 rows), the count of events rejected as duplicates (non-zero is healthy; a spike means a producer is retrying), and the count of reversals rejected for a missing original (should be zero; any is a producer bug).

Say exactly: "A healthy contract rejects duplicates quietly and rejects zero orphan reversals."

Where this lands at Amazon / NACF. Amazon's promise seam is the same contract at a vastly larger grain: every event moves a unit between states exactly once, and the order system reads one published number. The NACF version is inbound states from ASN to stow with one owner each.

S3-a — Sephora S/4HANA — de-aggregation to the cent: one posting, many lines, and the rule that makes the sum of the parts equal the whole (the grain problem behind Financial On-Hand) NEW

LPs this story carries (with proof lines below): Insist on the Highest Standards (primary — Σ per-line COGS = posting COGS to the cent) · Dive Deep (the grain: one posting vs transaction lines) · Learn and Be Curious (learning how S/4 posts) · Bias for Action (sequenced by event source)
Process family: Deliver / Return-to-book seam — cost allocation and matching at the aggregate-versus-line grain; Finance control design

Honesty label. The de-aggregation rule, the join spec, the test families and the aging-report spec were mine; S/4 and OMS engineers implemented; Inventory Accounting confirmed in UAT and owns the aging report. Cost ownership stayed with S/4 (WAC$ S/4-owned, ISM visibility-only). The $100.00, $175.00 and 2,000-postings-a-day figures are illustrative; the rounding mode and the line-order key are design choices I flag VERIFY-WITH-RAJESH because the record says "the last line absorbs rounding" without naming either.

LEADERSHIP PRINCIPLES CARRIED — ONE-LINE PROOF EACH NEW

Insist on the Highest Standards — Σ per-line COGS must equal the posting COGS to the cent in every test cell — a one-cent break is a failed row, not a rounding variance.

Dive Deep — The problem lived one level below the dashboards: one S/4 posting per Aggregation# against many POS lines per Order#, so the join and the allocation had to be written at line grain.

Learn and Be Curious — I learned how S/4 posts sales — aggregates with total quantity and COGS — well enough to explain to Finance why per-line posting was never the answer.

Bias for Action — I sequenced the build by event source — POS flash sales first, then omni pending, then ship-to-home — each added to SIT only when its increments and matched decrements reconciled.

60-SECOND BUSINESS-FIRST PITCH — STOP ON THE PRIMARY METRIC

Situation. S/4 posts sales as aggregates — one posting per aggregation number with a total quantity and a total cost — while the POS sends transaction lines. ISM needs cost and quantity per line to decrement the Shipped-Not-Billed bucket and to give Finance a number it can trust, so every posting has to be spread back onto its lines. **Task.** Define the de-aggregation exactly: how a posting's quantity and cost are matched to lines, how the cent is handled, and what happens when a line is missing, late or duplicated. **Action.** I wrote the join Aggregation# to POS header to lines; the pro-rata cost rule — each line gets posting COGS × its quantity ÷ the posting quantity, rounded, with the last line taking the difference so the sum is exact; a SNB-never-negative assertion; and three test families — n lines, a missing line, a duplicate header — each with one expected answer. Finance owns the aging report for postings that do not match. **Result.** Every posting either matches to the cent or ages visibly on a report with an owner; the unmatched rate is a design KPI (match ≥ 99.5% within 24 hours); Financial On-Hand moved from two truths to one on the signed FSD of 6/5/25.

20-SECOND VERSION NEW

S/4 posts one aggregate per aggregation number; the POS sends many lines. I wrote the rule that spreads a posting's cost back onto its lines — quantity-weighted, rounded, the last line absorbing the cent — so the lines always sum to the posting to the cent, and any posting that cannot match ages on a report Finance owns. It is the engine behind Financial On-Hand moving from two numbers to one.

NACF BRIDGE NEW

Inbound has the same grain problem: an ASN header with many lines versus a dock count by pallet. "Lines sum to the header" is the check that would have caught the hidden 98.0% on one class in a 99.0% overall accuracy number — an aggregate hides what the lines show. I would run the same spread-and-tie on received quantities by ASN line.

SITUATION / TASK — THE BUSINESS PROCESS PROBLEM, SCOPE, SCALE, BASELINES

S — One S/4 posting per Aggregation# carries a total quantity and a total COGS; the Aptos POS holds the lines with an Order#. Between sale and posting the book is behind; once the posting lands, ISM has to retire exactly the units and dollars that posting covers — no more, no less — against lines that may arrive late, twice or not at all.

T — Define the de-aggregation and matching rule, the exception handling and the control report so that Σ line cost = posting cost to the cent and SNB can never go negative — inside the ISM FSD signed 6/5/25.

Baselines on the table: Σ per-line COGS vs posting COGS — untested at the grain (target = $0.00 difference in every cell) · SNB negative values — a possible silent path (target never negative) · unmatched postings — invisible (target aged row at 24 h, owner named) · match rate — unmeasured (target ≥ 99.5% within 24 h).

FACET 1 — DATA ANALYSIS — WHAT I PROFILED, AT WHAT GRAIN, WHAT IT REVEALED NEW

- **Data.** S/4 aggregate postings (Aggregation#, quantity, COGS); Aptos RetailSaleTransactions headers and lines (Order#, store, SKU, qty); OMS and WMS shipment lines; S/4 cost stays S/4-owned (WAC$ read-only to ISM).
- **Grain and joins.** Posting ↔ header ↔ lines is one-to-many-to-many; the quantity must sum exactly, the cost must sum to the cent. Cut by lines per posting and by whether the unit cost divides evenly into cents.
- **What it revealed.** A break is not an exception, it is guaranteed: $100.00 over three equal units is $33.333… a unit, and no way of rounding each line independently makes the three lines sum to $100.00 — they sum to $99.99. Left alone, every such posting leaves a one-cent hole that Finance would chase by hand.

FACET 2 — SOLUTION DESIGN — OPTIONS I PUT ON THE TABLE, AND WHY ONE WON NEW

Posting of 3 units, $100.00; three lines of 1 unit (illustrative). “Break” is Σ lines − posting.

OPTION	LINES (ROUNDED TO THE CENT)	Σ LINES	BREAK	VERDICT
Round each line independently	$33.33 + $33.33 + $33.33	$99.99	−$0.01	Rejected — a break on every non-divisible posting
Bank the cents in a rounding account	$33.33 × 3 and $0.01 to a GL residual	$100.00 with a residual	$0.00 + a $0.01 residual	Rejected — creates a balance Finance must explain
Largest-remainder allocation	equal remainders, so an order key is needed anyway	$100.00	$0.00	Rejected — more logic for the same answer
last line absorbs the difference	$33.33 + $33.33 + $33.34	**$100.00**	**$0.00**	Chosen — deterministic if the line order is written down VERIFY-WITH-RAJESH

ACTION — FACET 4: IMPLEMENTATION IN FIVE GATED PHASES (PHASES, OWNERS, CONTROLS, RISKS)

- **Discovery** — the event inventory (POS flash sale, omni pending, ship-to-home confirm, S/4 Aggregation# posting), the who-reads-which-number map, the lines-per-posting profile and the rounding-class profile. *Exit bar:* Finance and Operations agree on the description of the gap; Internal Audit named as reviewer.
- **Plan & Design** — join spec Aggregation# ↔ header ↔ lines; the pro-rata cost rule with the last line absorbing; rounding mode and line order written into the spec VERIFY-WITH-RAJESH; SNB-never-negative assertion; unmatched-posting aging report spec; WAC$ S/4-owned, ISM visibility-only. *Exit bar:* signed inside the FSD 6/5/25; Finance accepts one FOH.
- **Build & Test** — three test families — n lines, a missing line, a duplicate header — with expected SNB and cost for each; sequenced by event source, each added to SIT only when its increments and matched decrements reconciled. *Exit bar:* every posting matches or ages visibly; Σ per-line COGS = posting COGS to the cent in every cell; UAT = confirmation by Inventory Accounting.
- **Cutover & Go-Live** — rode the May go-live: SNB opening balance reconciled to unposted sales at the snapshot; SNB publish back to S/4 enabled. *Exit bar:* aging report empty or explained at T+1; no negative SNB.

- **Hypercare / Close** — daily aging review with Inventory Accounting; handoff of the report. *Exit bar:* aging report owned by Finance without me; match ≥ 99.5% within 24 h as the design KPI.

Influence and ownership. Mine: the join spec, the allocation rule, the exception cases, the test families and the aging-report spec. Built by: S/4 and OMS engineers. Confirmed and owned after: Inventory Accounting. Decision I held: cost stays S/4-owned — ISM carries it for visibility, so there is never a second cost number.

RESULTS — $ · # · % · TIME · CUSTOMER · QUALITY (BASIS-LABELED)

- **Quality** — Σ per-line COGS = posting COGS to the cent in every test cell; SNB never negative; every unmatched posting visible on the aging report (record mechanism).
- **#** — Financial On-Hand truths 2 → 1 (record, FSD signed 6/5/25); three test families, each with one expected SNB and cost (record design).
- **%** — Design KPI: match ≥ 99.5% within 24 h (design target). At an illustrative 2,000 postings a day, 0.5% = 10 aged rows a day — a list one analyst clears VERIFY-WITH-RAJESH.
- **Time** — SNB opening balance reconciled at the May go-live snapshot; aging report empty or explained at T+1 (record exit bar).
- **$** — No new dollars. Illustrative slice of S3's modeled Finance bridge-labor row: 1,664 h × 30% × $85 = $42,432 VERIFY-WITH-RAJESH — inside S3's ≈$0.6M, never added to it.

L6 LEVEL EVIDENCE — what the interviewer pattern-matches against the Amazon L6 functional rubric; every line is this story's record or labeled model

L6 DIMENSION	EVIDENCE IN THIS STORY
Ambiguity	Nobody had written what "spread a posting onto its lines" means when the cents do not divide. I wrote it, including the cases.
Scope & influence	Finance, Inventory Accounting, S/4 and OMS engineers, Internal Audit as reviewer; Inventory Accounting took the report.
Execution trade-off	Last-line-absorbs over a rounding account: one line carries a cent's distortion in exchange for no residual balance anywhere.
Impact	The grain problem behind Financial On-Hand truths 2 → 1; every posting matches or ages visibly.
Process improvement	Test families as standard work: n lines, missing line, duplicate header — each with one expected answer.
Complexity	Three event sources with different latencies, one posting grain; sequenced by source.
Root cause	The cause was grain, not data quality; fixed in the allocation rule once, not in a monthly bridge.

PROBING — FACET 3: MATH & DATA LOGIC, FOUR TABLES, EVERY PLUG-IN SHOWN

Table 1 — before: the problem in numbers.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Naive rounding	`3 × $33.33 = $99.99`	posting $100.00 over three equal units; each line rounded alone (illustrative)	$99.99; break −$0.01
2	Unit cost does not divide	`$100.00 ÷ 3 = $33.3333`	posting COGS ÷ posting quantity	$33.3333 a unit
3	A second case	`$175.00 ÷ 8 = $21.875`	posting 8 units, $175.00 (illustrative)	$21.875 a unit
4	Naive rounding, second case	`$65.63 + $65.63 + $43.75 = $175.01`	lines of 3, 3, 2 units rounded independently, half-up (illustrative)	$175.01; break +$0.01
5	Missing line	`unmatched units = 8 − 6 = 2`	posting 8 units; lines present 3 + 3 = 6 (illustrative)	2 units no line explains
6	Duplicate header	`Σ q = (3 + 3 + 2) × 2 = 16 vs posting 8`	the same Order# arrives twice (illustrative)	16 ≠ 8 — doubled

Table 2 — the fix: step by step, values substituted.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Rule	`COGS_i = round(C_P × q_i ÷ q_P, 2); last line = C_P − Σ prior lines`	record design; rounding mode and line order flagged VERIFY-WITH-RAJESH	defined for every case
2	Case 1, lines 1–2	`round($100.00 × 1 ÷ 3, 2) = $33.33`	each of two lines	$33.33 + $33.33 = $66.66
3	Case 1, last line	`$100.00 − $66.66 = $33.34`	the last line absorbs the cent	$33.34
4	Case 1, tie	`$33.33 + $33.33 + $33.34 = $100.00`	Σ lines = posting	$100.00 exact
5	Case 2, lines 1–2	`round($175.00 × 3 ÷ 8, 2) = round($65.625, 2) = $65.63`	half-up (illustrative mode); two lines	$65.63 + $65.63 = $131.26
6	Case 2, last line	`$175.00 − $131.26 = $43.74`	unrounded would be $43.75; the −$0.01 is absorbed here	$43.74
7	Case 2, tie	`$65.63 + $65.63 + $43.74 = $175.00`	Σ lines = posting	$175.00 exact
8	Missing line — guard	`attempted clear 8 against SNB 6: 6 − 8 = −2 → blocked`	SNB-never-negative assertion (record design)	clear 6; 2 to the aging report
9	Duplicate header — guard	`Σ q = 16 > 8 → hold the second set`	quantity may not exceed the posting (record design)	8 matched; 8 held as duplicate
10	SNB opening at cutover	`SNB = 1,200 − 1,150 = 50`	unposted sales at the snapshot: lines received 1,200, matched postings 1,150 (illustrative)	opening SNB 50

Table 3 — after: the metric, before → after.

STEP	METRIC	BEFORE	AFTER	Δ / CHECK	BASIS
1	Case 1 break	−$0.01 ($99.99)	$0.00 ($100.00)	`$0.00 − (−$0.01) = $0.01`	record design; units illustrative
2	Case 2 break	+$0.01 ($175.01)	$0.00 ($175.00)	`$0.01`	record design; units illustrative
3	Residual GL balance from rounding	a rounding account	none	`$0.01 per break → $0`	record design
4	Missing-line unmatched units	invisible	2 units on the aging report	`2`	record design
5	Duplicate set	double count	8 held	`8`	record design
6	Match rate	unmeasured	≥ 99.5% in 24 h	`design KPI`	design target

Table 4 — the dollars, every multiplication shown.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Aged rows a day at KPI	`2,000 × 0.5% = 10`	illustrative 2,000 postings a day VERIFY-WITH-RAJESH	10 rows
2	At 5% unmatched	`2,000 × 5% = 100`	the report stops being a control at this volume	100 rows
3	S3 reconciliation labor (parent)	`32 h/wk × 52 = 1,664; 1,664 × $85 = $141,440`	modeled (S3)	$141,440
4	Slice from line-level breaks	`1,664 × 30% = 499.2 h; 499.2 × $85 = $42,432`	30% illustrative VERIFY-WITH-RAJESH; slice, not additive	≈$0.04M

Rounding is half-up here (illustrative); the point is that any mode, with the last line absorbing the difference, ties exactly. Case 2 shows a different shape: the naive sum is +$0.01, the chosen rule is exact.

SOLUTION FLOW — THE SYSTEM & DATA MAP

```
[Aptos POS lines]        [OMS omni lines]       [WMS ShipmentConfirmation]
 Order#, SKU, qty ----+---- sale / ship events ----+---- (all SNB += qty) ---->  SNB ledger
                                                                                    |
[S/4 aggregate posting: Aggregation#, q_P, C_P] --> matcher: Aggregation# <-> header <-> lines
                                                                |
     +------------------------------+-----------------+-----------------------+
     |                              |                                         |
 q matched = q_P ?             q matched < q_P                          q matched > q_P
 spread cost pro-rata          clear matched only                       hold duplicate set
 last line absorbs cent        SNB never < 0 (assert)                   aging report
 SNB -= matched                unmatched -> aging report                (owner: Inventory Accounting)
     |
 per-line COGS with Aggregation# lineage -> FOH publisher -> Finance: month-end tie to MB5L / MB52
```

SENSITIVITY / IMPACT — LOW / BASE / HIGH

INPUT	LOW	BASE	HIGH
Postings a day	1,000	2,000	5,000
Unmatched at 0.5%	5 rows	10 rows	25 rows
Unmatched at 5%	50 rows	100 rows	250 rows
Share of bridge hours from line-level breaks	15% → $0.02M	30% → $0.04M	50% → $0.07M
Parent labor row (S3)	$0.08M	$0.14M	$0.24M

SPOKEN TOTAL-IMPACT STATEMENT — SAY IT LIKE THIS

Conservatively, I am not adding a dollar to the story I already told: this is the engine inside S3's six hundred thousand, and the slice I can put on it is about forty thousand — illustrative, and I would drop it before defending it. What is on the record is the control: every posting matches to the cent or ages on a report Finance owns, SNB can never go negative, and Financial On-Hand moved from two numbers to one.

FACET 5 — HARDEST FOLLOW-UPS, WITH NUMERIC ANSWERS NEW

Numeric — Why can't each line just be rounded?

Because $100.00 over three units is $33.3333 a unit; rounding each line alone gives $33.33 three times, $99.99, a one-cent hole on every non-divisible posting. With the last line absorbing it, $33.33 + $33.33 + $33.34 = $100.00.

Numeric — Show me a case where naive rounding goes the other way.

Posting 8 units, $175.00: $21.875 a unit, lines of 3, 3 and 2 units. Naive half-up gives $65.63 + $65.63 + $43.75 = $175.01 — a plus one-cent break. The rule gives the last line $175.00 − $131.26 = $43.74, so it sums to $175.00 exactly.

Numeric — A posting says 8 units but only 6 lines are in. What happens?

SNB would go 6 − 8 = −2 if I let it clear, so the assertion blocks it: 6 units clear, 2 go to the aging report as an unmatched posting, and Finance sees them the same day. A negative SNB would quietly add supply to AFS — an invisible over-sell.

Numeric — How many exceptions a day can your team actually clear?

At an illustrative 2,000 postings a day VERIFY-WITH-RAJESH, a 99.5% match rate is 10 aged rows a day; at 95% it is 100. The KPI is chosen so the report stays a list a person can clear — a control only works while it is small enough to read.

BEHAVIORAL PACK — FOUR QUESTIONS THIS STORY ANSWERS, EACH A DIFFERENT 90-SECOND MOMENT NEW

INSIST ON THE HIGHEST STANDARDS **1. Tell me about a time you insisted on a standard that others thought was too strict.**

Open (first two sentences): Someone said a one-cent difference on a posting is rounding noise. I said it is a failed test.

A. I showed that a one-cent break is guaranteed on every non-divisible posting, so “noise” is a hole on every one of them, and that Finance would chase each by hand. I wrote the last-line-absorbs rule and made Σ line cost = posting cost to the cent the pass condition for every cell in every test family. **R.** Every cell tied; nothing accumulates in a rounding account.

Land the number: $0.00 break in every cell | **Pivot to NACF:** On inbound receipts: the standard is lines sum to the ASN header exactly, or the row ages.
Do NOT say: Quantifying the cent as a dollar claim.

DIVE DEEP **2. Tell me about a time the answer was hidden below the aggregate.**

Open (first two sentences): The Finance number and the Operations number were both right; the problem was one grain down.

A. I laid one posting next to its lines and showed that the join is one-to-many and the quantity and cost must both sum. The three failure shapes — a missing line, a duplicate header, a non-divisible cost — each have one expected result. **R.** The three test families became standard work.

Land the number: one posting, three failure shapes | **Pivot to NACF:** Same pattern as an overall inbound accuracy that hides the class that is off.
Do NOT say: Table names unless asked.

LEARN AND BE CURIOUS **3. Tell me about something you had to learn quickly to do the job.**

Open (first two sentences): I was designing a bucket that depended on how S/4 posts sales, and I had to learn it before I could argue about it.

A. I sat with the S/4 and Finance teams and learned why S/4 posts aggregates — throughput and posting cadence — and what the posting carries: an Aggregation#, a quantity, a cost. That is why I refused per-line real-time posting and designed the match on the ISM side. **R.** The design held S/4 as the settlement authority; I never claim to have configured it.

Land the number: S/4 stays the settlement authority | **Pivot to NACF:** The same: learn the receiving system’s posting cadence before promising a number off it.
Do NOT say: “I configured S/4.”

BIAS FOR ACTION **4. Tell me about a time you had to deliver in sequence under a hard date.**

Open (first two sentences): There were three event sources and one date, so I did not try to land them together.

A. I sequenced by source — POS flash sales first, then omni pending, then ship-to-home — and each entered SIT only after its increments and matched decrements reconciled. A late source slipped alone, not the date. **R.** Rode the May go-live with SNB opening balance reconciled at the snapshot.

Land the number: three sources, one date, sequenced | **Pivot to NACF:** Inbound launches the same way: one door class or one lane group at a time, each reconciled before the next.
Do NOT say: Dates beyond the May go-live.

SAME STORY, 3 ANGLES — DIFFERENT OPENER, HEADLINE, LP NEW

DATA / ANALYTICAL

"I put one posting next to its three lines and asked what the cents add up to."

$99.99 vs $100.00

LP: Dive Deep

$100.00 ÷ 3 = $33.3333; independent rounding gives $99.99.

The break is guaranteed on every non-divisible posting.

Last-line-absorbs gives $33.33 + $33.33 + $33.34 = $100.00.

CONFLICT / INFLUENCE / BACKBONE

"The easy answer to a penny is a rounding account, and I did not want to create one."

$0.00 residual

LP: Insist on the Highest Standards · Have Backbone

A rounding account creates a balance someone must explain every month.

I wanted the sum right on the line, not corrected downstream.

Who proposed the account, if anyone, is VERIFY-WITH-RAJESH; I do not claim a specific disagreement.

FAILURE / LEARNING

"What I would write on day one: which rounding mode, and which line is last."

rounding mode · line order

LP: Are Right, A Lot

The rule says "the last line absorbs"; it needs a stated mode and a stated order or two runs could differ by a cent.

The record does not carry a miss here — VERIFY-WITH-RAJESH whether one happened.

Differently: write the ordering key and mode into the spec and the test cell on day one.

TRAP DRILLS — THE THREE NASTIEST FOLLOW-UPS, WITH THE EXACT SENTENCE NEW

Trap — "Is the cent worth a rule? What would change your mind?"

The rule costs one line of spec and one test cell. What would change my mind is if breaks were rare — they are not: a cost that divides evenly is rare; $33.33-type splits happen on every posting whose unit cost is not a whole cent. If Finance showed me a rounding account they already reconcile for free, I would revisit it.

Say exactly: "It costs one line of spec to avoid a cent on every non-divisible posting; if Finance already had a free rounding account I would revisit it."

Trap — "What is the dependency, and what do you cut if S/4's team has no capacity?"

The dependency chain: the Aggregation# on the posting (S/4), the Order# join (POS), and the shipment lines (WMS). I cut ship-to-home last-in, first-out — it rides the old overnight path longer. I never cut the SNB-never-negative assertion or the aging report.

Say exactly: "I would cut a source, never the assertion or the report."

Trap — "A posting arrives twice. What happens to Financial On-Hand?"

Nothing, by two guards: the posting's Aggregation# is idempotent — a second delivery is a no-op — and quantity matched can never exceed the posting quantity. On the walk with 8 units, a doubled set is 16 > 8, so 8 clear and 8 hold as a duplicate on the aging report.

Say exactly: "A double delivery is a no-op, and a doubled set of lines holds on the aging report — it cannot move Financial On-Hand."

Where this lands at Amazon / NACF. At Amazon scale the lines-to-header tie is the same control wherever aggregates meet transactions: ASN lines to the header, dock pallets to receipts. It is how an aggregate accuracy number stops hiding the class that is off.

S1-b — Sephora S/4HANA ↔ ISM — control screens and the operating model: Negative AFS, AFS Correction and the drift report, so every change to available-for-sale has a screen, a role and an audit row before I stepped back NEW

LPs this story carries (with proof lines below): Ownership (primary — a mechanism others own after I leave) · Insist on the Highest Standards (an audit row for every AFS-changing act) · Earn Trust (Inventory Control, Internal Audit and Finance as named owners) · Deliver Results (hypercare exit on counted rows)
Process family: Plan / Deliver seam — control design and operating model: exception screens, role map, reason codes, hypercare handoff

Honesty label. I designed the control-owner map, the screens' purpose and exit bars, the reason-code requirement and the hypercare handoff; the ISM application team built the screens, and Inventory Control, Internal Audit, Finance and the ISM app team run them. The SKU example and the daily row counts are illustrative; the roles, the T+1 bar and the zero drift tolerance are record design. Whether the screens enforce a separate approver for a correction is VERIFY-WITH-RAJESH — the record says "a screen, a role and an audit row," not maker-checker.

LEADERSHIP PRINCIPLES CARRIED — ONE-LINE PROOF EACH NEW

Ownership — Every AFS-changing act got a screen, a role and an audit row before I handed off — the drift report runs daily without me and the owners are named in the FSD.

Insist on the Highest Standards — A negative AFS is never clamped to zero silently: it appears on a screen, with a reason code, and the exit bar is "clean at T+1."

Earn Trust — I gave Internal Audit and Finance seats in the control design rather than a review at the end — they own the screens that would catch my own design's errors.

Deliver Results — Hypercare exited on counted rows — drift = 0 SKUs on consecutive days and blockers per day under threshold — not on a calendar.

60-SECOND BUSINESS-FIRST PITCH — STOP ON THE PRIMARY METRIC

Situation. ISM computes available-for-sale from ten buckets in real time. That is only as good as what happens on the days the number is wrong — a negative AFS because commitments exceed the lower of the two books, a stale reservation that no longer corresponds to an order, a drift between ISM and S/4. If anyone can change an AFS by hand, or nobody owns the exception, the real-time ledger becomes the next spreadsheet. **Task.** Design the control layer and the operating model so the system stays trustworthy after the program team leaves. **Action.** I defined three controls — a Negative AFS screen, an AFS Correction screen with a reason code and an audit row, and a daily ISM-versus-S/4 drift report with tolerance zero — and a control-owner map: Negative AFS to Inventory Control, Internal Audit and Finance; AFS Correction to the ISM application team; the drift report to Inventory Control. I set the exit bars before go-live: Negative AFS clean at T+1; drift zero SKUs on consecutive days. **Result.** At hypercare exit every AFS-changing act had a screen, a role and an audit row; the drift report runs without me; and the design entered the monthly release train. The dual-book reconciliation that took about four people by hand became about one reading a report — my model of the record mechanism.

20-SECOND VERSION NEW

A real-time availability number is only trustworthy if the days it is wrong have an owner. I designed three controls — a negative-AFS screen, a reason-coded AFS correction and a daily drift report at tolerance zero — and put a named owner on each in the signed spec. When I handed off, every change to available-for-sale had a screen, a role and an audit row; reconciliation went from about four people by hand to about one reading a report (modeled).

NACF BRIDGE NEW

Inbound needs the same operating model: a negative or implausible available-to-promise (more promised than received-and-stowed) must land on a screen with an owner, and any manual change to inbound availability needs a reason code and an audit row. A daily "system versus count" drift report for each building closes the loop.

SITUATION / TASK — THE BUSINESS PROCESS PROBLEM, SCOPE, SCALE, BASELINES

S — A real-time availability ledger beside the book of record creates a new class of exception: a negative AFS, a stale reservation, a drift between the ledger and S/4 — and a new temptation: a manual fix nobody can trace. Before the program, reconciling the two books took about four people by hand (modeled).

T — Design controls and a role map so every AFS-changing act is visible and owned, with exit bars that hold up in hypercare, and hand the mechanism to the people who would run it.

Baselines on the table: AFS-changing acts with a screen, a role and an audit row — none defined (target 100%) · Negative AFS — unmonitored (target clean at T+1) · ISM-vs-MB52 drift — unmeasured (target 0 SKUs on consecutive days) · reconciliation ≈4 FTE by hand (target ≈1 FTE reading a report, modeled).

FACET 1 — DATA ANALYSIS — WHAT I PROFILED, AT WHAT GRAIN, WHAT IT REVEALED NEW

- **Data.** The bucket ledger (SKU × node × bucket), the published AFS stream, S/4 MB52 snapshots and MSEG movements, and the cause classes of exceptions seen in testing: committed above supply, stale reservation, feed lag, count adjustment.
- **Grain.** SKU × node, per publish cycle; one row on the screen per SKU × node where AFS < 0 or where drift ≠ 0.
- **What it revealed.** On a SKU with S/4 500, WMS 460 and 480 committed, AFS is −20 on the lower-of-two basis but +20 if a reader trusts the S/4 book. The negative is the screen working, not a bug — and clamping it to zero would hide 20 units of over-commitment.

FACET 2 — SOLUTION DESIGN — OPTIONS I PUT ON THE TABLE, AND WHY ONE WON NEW

The question is who may change AFS and how. Same SKU, AFS −20 (illustrative).

OPTION	WHO CAN CHANGE AFS	TRACEABILITY	VERDICT
Anyone with order-system access adjusts by hand	many	none	Rejected — the next spreadsheet
App team fixes by direct database change on request	one team, off-screen	ticket only	Rejected — no reason code or audit row at the act
Clamp negatives to zero and move on	nobody	hides the cause	Rejected — hides 20 units of over-commitment
No manual change ever	nobody	n/a	Rejected — a count error needs a legitimate path
one AFS Correction screen: reason code, audit row, a named owning role	ISM app team	**every act: who, what, why, when**	Chosen — with Negative AFS to Inventory Control, Internal Audit and Finance

ACTION — FACET 4: IMPLEMENTATION IN FIVE GATED PHASES (PHASES, OWNERS, CONTROLS, RISKS)

- **Discovery** — inventory every way AFS can change and every exception class; name the readers — Inventory Control, Internal Audit, Finance, the ISM app team. *Exit bar:* every AFS-changing act mapped to a screen or removed.
- **Plan & Design** — the control-owner map — Negative AFS → Inventory Control, Internal Audit and Finance; AFS Correction → ISM app team; drift report → Inventory Control; reason-code requirement on every correction; exit bars (Negative AFS clean at T+1; drift 0). *Exit bar:* in the FSD signed 6/5/25.
- **Build & Test** — screen specs and the matrix cells that exercise a negative, a correction and a drift row; UAT = Inventory Control confirming it can work the list. *Exit bar:* 0 open contract defects in SIT; owners trained on a rehearsal list VERIFY-WITH-RAJESH.
- **Cutover & Go-Live** — control-screen roles assigned before January and May go-lives; rollback triggers armed. *Exit bar:* Negative AFS clean at T+1 after each go-live.
- **Hypercare / Close** — daily drift report; Negative AFS / AFS Correction audit review; hypercare exit memo and FSD as-built note. *Exit bar:* drift = 0 SKUs on consecutive days; blockers/day under threshold; the design enters the monthly release train.

Influence and ownership. Mine: the map of AFS-changing acts, the owner map, the reason-code rule, the exit bars and the handoff memo. Built by: the ISM application team. Run by: Inventory Control (negative AFS and drift), Internal Audit (sampling), Finance (value review) and the ISM app team (corrections). I deliberately did not make myself an owner of any screen.

RESULTS — $ · # · % · TIME · CUSTOMER · QUALITY (BASIS-LABELED)

- **Quality** — Every AFS-changing act has a screen, a role and an audit row (record); Negative AFS clean at T+1 (record exit bar); drift = 0 SKUs on consecutive days at hypercare exit (record).
- **#** — Three controls, four owning roles; reconciliation ≈4 FTE by hand → ≈1 FTE reading the drift report (modeled on the record mechanism).
- **Time** — Controls assigned before the January (planning) and May (execution) go-lives; exit into the monthly release train after hypercare (record).
- **$** — No new dollars. The dollar behind this story is S1's reconciliation row: 3 FTE × $110,000 = $330,000 (modeled), inside S1's ≈$2.5M — never added to it.

L6 LEVEL EVIDENCE — what the interviewer pattern-matches against the Amazon L6 functional rubric; every line is this story's record or labeled model

L6 DIMENSION	EVIDENCE IN THIS STORY
Ambiguity	No one had listed the ways AFS could change. I wrote the list and removed the ones that had no owner.
Scope & influence	Four owning roles across Inventory Control, Internal Audit, Finance and the ISM team; signed 6/5/25.
Execution trade-off	A screen with a reason code slows a correction by a minute; I accepted that for an audit row on every act.
Impact	The reconciliation burden shifted from argument to a report (≈4 → ≈1 FTE, modeled).
Process improvement	Exit bars on counted rows: T+1 negative clean; drift zero on consecutive days.
Complexity	Two go-lives, each with the controls assigned before it, each with a rollback trigger armed.
Root cause	The cause of "shadow spreadsheets" was an unowned exception; fixed by naming owners, not by policing.

PROBING — FACET 3: MATH & DATA LOGIC, FOUR TABLES, EVERY PLUG-IN SHOWN

Table 1 — before: the problem in numbers.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Basis, lower of two	`Basis = min(500, 460) = 460`	S4_Avail 500; WMS_OnHand 460 (illustrative)	460
2	Out-of-balance	`OOB = 500 − 460 = 40`	book above floor	+40
3	Committed	`200 + 150 + 80 + 50 = 480`	PickRes 200, EcomPick 150, EcomPack 80, SNB 50 (illustrative)	480
4	AFS on the lower-of-two	`AFS = 460 − 480 = −20`	Basis − committed	−20 — a negative
5	AFS if a reader trusts S/4	`AFS = 500 − 480 = 20`	raw S/4 read (the old definition A)	+20 — masks −20
6	Exposure if clamped to zero and ignored	`20 units × $30 = $600`	$30 incident cost (S1 benchmark build)	$600 per SKU-day

Table 2 — the fix: step by step, values substituted.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Negative AFS row	`AFS = −20 < 0 → row on the Negative AFS screen at the next publish`	60-second publish budget (record design); reason-code classes: committed above supply · stale reservation · feed lag · count	1 row, owner Inventory Control
2	Resolve by AFS Correction	`PickRes = 200 − 25 = 175`	25 stale reservations for orders already cancelled (illustrative); reason “stale reservation”	PickRes 175
3	AFS after correction	`AFS = 460 − 175 − 150 − 80 − 50 = 5`	same Basis 460	+5 — clean
4	Audit row	`SKU, node, bucket PickRes, before 200, after 175, Δ −25, reason, user, time`	every AFS-changing act (record: screen, role, audit row)	1 audit row
5	Drift after the correction	`drift = ISM.S4_Avail − MB52 = 500 − 500 = 0`	the correction moves a committed bucket, not a book	0
6	T+1 exit check	`open negative rows at T+1 = 0`	exit bar: Negative AFS clean at T+1 (record)	0 rows
7	Workload check — negatives	`40 × 6 ÷ 60 = 4.0 h`	illustrative 40 rows a day, 6 min each VERIFY-WITH-RAJESH	4.0 h/day
8	Workload check — drift	`10 × 15 ÷ 60 = 2.5 h`	illustrative 10 rows a day, 15 min each VERIFY-WITH-RAJESH	2.5 h/day

Table 3 — after: the metric, before → after.

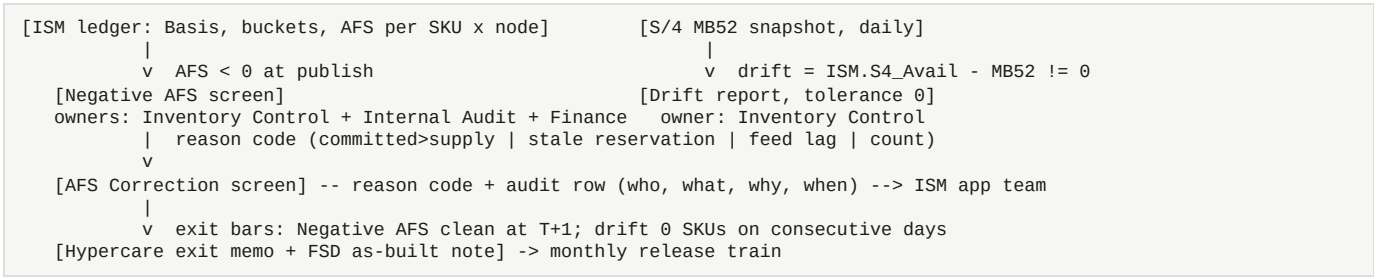
STEP	METRIC	BEFORE	AFTER	Δ / CHECK	BASIS
1	AFS-changing acts with screen, role and audit row	none defined	100%	`all`	record
2	Negative AFS at T+1	unmonitored	0 rows open	`0`	record exit bar
3	Daily drift	unmeasured	0 SKUs on consecutive days	`0`	record
4	Reconciliation people	≈4 FTE	≈1 FTE	`4 − 1 = 3 FTE`	modeled on the record mechanism
5	Daily workload on one reader	—	6.5 h of a 7.2 h day	`7.2 − 6.5 = 0.7 h slack`	illustrative

Table 4 — the dollars, every multiplication shown.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Workload sum	`4.0 + 2.5 = 6.5`	rows above (illustrative)	6.5 h/day
2	One person's productive day	`1,800 h ÷ 250 days = 7.2`	1,800 productive h a year (illustrative)	7.2 h/day
3	Fits one FTE?	`7.2 − 6.5 = 0.7`	slack > 0	yes, with 0.7 h slack
4	Reconciliation labor released	`4 − 1 = 3; 3 × $110,000 = $330,000`	S1 modeled row; $110K loaded (benchmark)	≈$0.33M/yr (inside S1)

The lower-of-two shows −20 where the S/4 book shows +20 — that gap is the whole reason the screen exists. The workload check tests whether "about one FTE" is plausible; it does not replace a time study.

SOLUTION FLOW — THE SYSTEM & DATA MAP



SENSITIVITY / IMPACT — LOW / BASE / HIGH

INPUT	LOW	BASE	HIGH
Negative rows a day × minutes each	20 × 6 = 2.0 h	40 × 6 = 4.0 h	80 × 6 = 8.0 h
Drift rows a day × minutes each	5 × 15 = 1.25 h	10 × 15 = 2.5 h	20 × 15 = 5.0 h
Total daily workload	3.25 h	6.5 h	13.0 h
FTE needed at 7.2 h a day	0.45	0.90	1.81
Reconciliation $ if 3 FTE released × loaded	2 × $95K = $0.19M	3 × $110K = $0.33M	4 × $130K = $0.52M

SPOKEN TOTAL-IMPACT STATEMENT — SAY IT LIKE THIS

Conservatively, there is no new dollar here; the dollar is S1's reconciliation row, about three hundred thirty thousand, modeled. What I can say flat is the record: Negative AFS clean at T+1, drift at zero SKUs on consecutive days, and every AFS-changing act with a screen, a role and an audit row when I stepped back. And I can show you the check that "about one person" is plausible — six and a half hours of a seven-point-two hour day, on illustrative volumes.

FACET 5 — HARDEST FOLLOW-UPS, WITH NUMERIC ANSWERS NEW

Numeric — Show me a negative AFS and how it clears.

S/4 500, WMS 460, committed 480: Basis = min(500, 460) = 460, AFS = 460 − 480 = −20. A reader on raw S/4 would see +20. The screen shows the row; Inventory Control finds 25 stale reservations, a correction releases them — PickRes 200 → 175 — and AFS = 460 − 175 − 150 − 80 − 50 = +5, with an audit row.

Numeric — Why not clamp negatives to zero?

Because clamping hides the cause: −20 means 20 units are over-committed, and at about $30 an incident that is $600 on one SKU-day that nobody would see. A negative is the lower-of-two doing its job.

Numeric — Is “about one FTE” real?

Plausibility check, not a time study: at an illustrative 40 negative rows at 6 minutes and 10 drift rows at 15 minutes, 4.0 + 2.5 = 6.5 hours against a 7.2-hour productive day. At double the volume it needs 1.8 people. I would run a two-week time study before committing the headcount.

Numeric — What proves the controls work?

Three counted rows: open negative rows at T+1 = 0, drift SKUs = 0 on consecutive days, and AFS-changing acts without an audit row = 0. Anything else is a feeling.

BEHAVIORAL PACK — FOUR QUESTIONS THIS STORY ANSWERS, EACH A DIFFERENT 90-SECOND MOMENT NEW

OWNERSHIP **1. Tell me about a time you built something that outlasted your involvement.**

Open (first two sentences): On the day I handed off, I wanted to be able to leave without a single AFS-changing act that only I understood.

A. I listed every way AFS could change, gave each a screen, a role and an audit row, and put the owners in the signed spec — Inventory Control for negatives and drift, Internal Audit and Finance as reviewers, the ISM app team for corrections. I did not make myself an owner of any of them. **R.** The drift report runs daily without me and the design entered the monthly release train.

Land the number: every act: screen, role, audit row | **Pivot to NACF:** At NACF I would design inbound controls so the building runs them: one screen for implausible available-to-promise, one for drift versus the dock count.
Do NOT say: “I ran the screens.”

EARN TRUST **2. Tell me about a time you built trust with a control function like audit or finance.**

Open (first two sentences): Internal Audit and Finance are usually reviewers at the end; I put them in the design at the start.

A. The Negative AFS screen has Inventory Control, Internal Audit and Finance as owners, so the people who would test the design own the thing that catches its errors. A correction carries a reason code and an audit row. **R.** No control owner asked for a parallel report.

Land the number: three control owners on one screen | **Pivot to NACF:** At NACF: the finance and ops partners both read the inbound exception list.
Do NOT say: Claiming audit findings or opinions — none in the record.

INSIST ON THE HIGHEST STANDARDS **3. Tell me about a time you wouldn’t accept a quick fix.**

Open (first two sentences): The quick fix to a negative AFS was to clamp it to zero. I would not accept it.

A. A negative is a lower-of-two result; clamping hides over-commitment. I required the row on a screen with a reason code and an owner, and made “clean at T+1” the exit bar. **R.** Negative AFS clean at T+1 after each go-live.

Land the number: −20 not 0 | **Pivot to NACF:** In inbound: never clamp a negative promise; show it.
Do NOT say: Counting incidents prevented — none measured.

DELIVER RESULTS **4. Tell me about a time you delivered a result on a hard exit criterion.**

Open (first two sentences): Hypercare exit was defined before go-live as numbers, not a date.

A. Exit on drift = 0 SKUs on consecutive days and blockers per day under threshold; each go-live had its controls assigned and rollback triggers armed — sync under 99% for four hours, or order-flow p95 above twice budget for one hour; neither fired. **R.** Exited into the monthly release train.

Land the number: drift 0 on consecutive days | **Pivot to NACF:** The same: an inbound launch exits on a door-to-stock time and a count tie, not a week.
Do NOT say: A hypercare duration — not in the record.

SAME STORY, 3 ANGLES — DIFFERENT OPENER, HEADLINE, LP NEW

DATA / ANALYTICAL

“Same SKU, two readers: one sees plus twenty, the other minus twenty.”

+20 vs −20

LP: Dive Deep

S/4 500, WMS 460, committed 480: lower-of-two −20; raw S/4 +20.

Clamping to zero would hide 20 units; $600 at $30 an incident.

The screen is the lower-of-two doing its job.

CONFLICT / INFLUENCE / BACKBONE

“Everyone wanted me to be the owner of the screens, and I said no.”

4 owning roles

LP: Ownership · Have Backbone

A control owned by the author does not survive the author.

I named Inventory Control, Internal Audit, Finance and the ISM app team in the spec.

Handoff on counted rows: T+1 clean, drift 0.

FAILURE / LEARNING

“The honest gap: the record says “a screen, a role and an audit row” and does not say who approves a correction.”

maker-checker?

LP: Insist on the Highest Standards

If one person can propose and approve a correction, the audit row is a record, not a control.

Whether the screens separated the two is VERIFY-WITH-RAJESH.

What I would do now: require a second role for corrections above a unit threshold.

TRAP DRILLS — THE THREE NASTIEST FOLLOW-UPS, WITH THE EXACT SENTENCE NEW

Trap — “How do you know the controls caused the result and not the other programs?”

I cannot isolate it statistically and I would not claim to. What I defend is the mechanism: every AFS-changing act traceable, negatives visible, drift measured. The counterfactual is the status quo — a ledger where the same exceptions would be fixed by hand and traced by nobody.

Say exactly: “I can’t separate it from the rest of the program, and I won’t pretend to; I can show what each control would have caught.”

Trap — “What would you cut if Internal Audit had no time?”

I would cut their seat on the Negative AFS screen to a monthly sample of corrections, not their review of the reason-code list. I would never cut the audit row on a correction, the T+1 bar, or the drift report.

Say exactly: “Audit sampling can be thinner; the audit row, T+1 and drift cannot.”

Trap — “What is the dependency risk — who must act for this to work?”

Inventory Control must work the lists daily; the ISM team must keep corrections available; Finance must read the value impact. If Inventory Control’s capacity is the constraint — 6.5 of 7.2 hours on one person at base volume — the first thing I cut is non-exception drift cases, never negatives.

Say exactly: “The binding dependency is one team’s daily capacity; I protect negatives first.”

Where this lands at Amazon / NACF. Amazon’s equivalent is the exception review on promise data: a negative or implausible available-to-promise on a screen with an owner, a reason-coded override with an audit trail, and a daily system-versus-physical drift check.

S1-c — Sephora S/4HANA ↔ ISM — the priced no: Merchandising wanted merchandise reserve promisable; I priced it, wrote the no into the FSD as a row, and made the reserve copy provably equal to the plan (reserve-copy invariants and the old-versus-new AFS reconciliation) NEW

LPs this story carries (with proof lines below): Have Backbone, Disagree and Commit (primary) · Customer Obsession (unpickable stock sold is a broken promise) · Are Right, A Lot (priced both sides; named what would change my mind) · Earn Trust (visibility offered; Merchandising not caricatured)
Process family: Plan / Deliver seam — reserve policy and cutover: merchandise reserve, channel reserves, publish switch, old-versus-new reconciliation

Honesty label. The ask, the pricing, the no-as-a-row and the reserve-copy invariants were mine; FRAME planning, the ISM app team and OMS configurators built to the signed FSD. The record says the no was priced as "unpickable stock sold" and written as an FSD row and that merch reserve got visibility; it does not carry the units, the sell-through or the number of reserve holds below, which are illustrative and flagged VERIFY-WITH-RAJESH. The four invariants are my statement of the "reserve-copy invariants" the record names; their exact wording is VERIFY-WITH-RAJESH.

LEADERSHIP PRINCIPLES CARRIED — ONE-LINE PROOF EACH NEW

Have Backbone, Disagree and Commit — I said no to Merchandising's promisable-reserve ask in the room and in writing — priced, with the alternative on the table — and I was ready to commit if the evidence changed.

Customer Obsession — Held reserve is stock nobody can pick; selling it is a cancellation waiting to happen, and the customer is the one who finds out.

Are Right, A Lot — I priced both sides — the margin earned if the hold releases in time, the incident if it does not — and wrote the break-even that would flip my answer.

Earn Trust — Merchandising's underlying need was visibility of the reserve; I met it with its own column in FOH and the screens, so the no did not leave them blind.

60-SECOND BUSINESS-FIRST PITCH — STOP ON THE PRIMARY METRIC

Situation. Merchandising holds stock back — launches, allocations, channel reserves — and that stock sits in the building but cannot be picked until a hold is released. Merchandising wanted it counted as sellable, so a product could take orders before the hold lifted. In the council the room was leaning toward yes. **Task.** Decide, and make the reserve data trustworthy enough that either answer was safe at cutover. **Action.** I priced the ask: every promisable unit that is unpickable becomes an oversell incident if the hold does not release in time. I proposed instead that merch reserve is a visible bucket — its own column in the Financial On-Hand and on the control screens — but outside available-for-sale; the no went into the specification as a row with the pricing, not a meeting minute. For the January planning go-live I wrote the reserve-copy invariants — the reserve in ISM equals the reserve in the plan at every nightly cut; reserve never sits inside the S/4 available figure; it never enters AFS as supply; there is one published AFS and the order system reads it — and a reconciliation of old versus new AFS by SKU, node and channel where every difference has a reason code. **Result.** Merchandising kept full visibility; unpickable stock was never sold; January and May go-lives with oversell incidents recurring → zero; the no survived because it lived in the FSD.

20-SECOND VERSION NEW

Merchandising wanted held-back stock to count as sellable. I priced it — held stock is unpickable, so every unit promised is an oversell incident if the hold doesn't lift — wrote the no into the FSD as a row, and gave them what they actually needed: reserve visible in its own column, outside available-for-sale. I then wrote four invariants so the reserve copy provably matched the plan at the January go-live.

NACF BRIDGE NEW

At NACF the equivalent is inbound stock that is physically present but not promisable — held for QA, cross-dock allocation or a launch. It belongs in a visible column, outside inbound available-to-promise, until the hold lifts; promising it is promising what the floor cannot yet release.

SITUATION / TASK — THE BUSINESS PROCESS PROBLEM, SCOPE, SCALE, BASELINES

S — Merchandise reserve — stock held back for launches, allocations and channel plans — is in the building but unpickable. It sits outside S/4 available and is fed from the planning system's channel reserves nightly plus intraday holds. Merchandising asked for it to be promisable; the council leaned that way.

T — Decide the policy with evidence, and make the reserve data safe to go live on: the copy of reserves in ISM had to equal the plan, old and new AFS had to reconcile, and the order system had to read one number.

Baselines on the table: promisable reserve — requested (target: visible, not promisable) · reserve copy vs plan — untested (target 0 differences at each nightly cut) · old-vs-new AFS differences without a reason code — unknown (target 0 at go/no-go) · SKUs with two published AFS values — possible (target 0).

FACET 1 — DATA ANALYSIS — WHAT I PROFILED, AT WHAT GRAIN, WHAT IT REVEALED NEW

- **Data.** The channel-reserve and hold feed from planning (nightly plus intraday holds); the reserve column in the ISM ledger; S/4 available (which excludes merch reserve); order history for the demand that would see a promisable reserve.
- **Grain.** SKU × node × channel × bucket. The reconciliation compared AFS under the legacy definition and under the new one for every SKU × node × channel row.
- **What it revealed.** Counting reserve as sellable adds exactly the held units to AFS (illustratively 560 → 710 on the worked SKU), which is the size of the over-promise if the hold is not lifted before demand arrives. The size of that error, not the principle, carried the decision.

FACET 2 — SOLUTION DESIGN — OPTIONS I PUT ON THE TABLE, AND WHY ONE WON NEW

Worked SKU: Basis 980, committed 420, AFS 560; merch reserve 150 held for a launch (illustrative). If demand takes the whole promisable amount before release, "over-promise" is the units sold that cannot be picked.

OPTION	AFS SHOWN	UNITS SOLD ABOVE PICKABLE	VERDICT
Promisable reserve (the ask)	560 + 150 = 710	up to 150	Rejected — unpickable stock sold
Promisable if the hold releases within a window	560 + 150 while the window is open	up to 150 if release slips	Rejected — the release is Merchandising's decision, not a supply fact
Invisible reserve (folded into S/4 or hidden)	560	0	Rejected — Merchandising blind; Finance cannot read FOH
Reserve inside the S/4 figure and subtracted again	560 − 150 = 410	0 but under-states by 150	Rejected — a double subtract
own bucket: visible in FOH and screens, outside AFS	**560**	**0**	Chosen — visibility without promisability

ACTION — FACET 4: IMPLEMENTATION IN FIVE GATED PHASES (PHASES, OWNERS, CONTROLS,

RISKS)

- **Discovery** — who sets reserves (Merchandising through planning), how often (nightly plus intraday holds), where they sit relative to S/4 available (outside). *Exit bar:* the reserve data path mapped to one owner per step.
- **Plan & Design** — the pricing of the ask; the policy (visible, not promisable) written as an FSD row; reserve-copy invariants; the old-versus-new AFS reconciliation by SKU × node × channel; council ask framed as three decisions. *Exit bar:* FSD signed 6/5/25 with the no as a row.
- **Build & Test** — node × channel × bucket-transition matrix cells for reserve set, hold, release and expire; the invariant checks as assertions. *Exit bar:* every reserve transition has a matrix cell; E2E green 11/22 · 12/20 · 1/30.
- **Cutover & Go-Live** — January planning go-live — the reserve rules live against demand; the publish-switch plan: OMS reads ISM AFS, never computes; go/no-go read from counted rows of the old-versus-new reconciliation. *Exit bar:* no SKU with two published AFS values; every difference has a reason code.
- **Hypercare / Close** — daily reserve-copy equality check; visibility reports to Merchandising; drift report. *Exit bar:* reserve-copy differences 0 on consecutive days.

Influence and ownership. Mine: the pricing, the policy row, the four invariants, the reconciliation design and the council framing. Built by: the planning-system and ISM teams. Decision-maker for any future change: the three FSD signatories, through the regression matrix and the release train.

RESULTS — $ · # · % · TIME · CUSTOMER · QUALITY (BASIS-LABELED)

- **Customer** — Unpickable stock was never sold; oversell incidents recurring → zero through two go-lives (record).
- # — One policy row in the FSD; four invariants; zero SKUs with two published AFS values; reserve visible as its own column (record design / record).
- **Time** — FSD signed 6/5/25; January planning go-live with reserve rules live against demand; May execution go-live (record).
- **Quality** — Old-versus-new AFS difference rows unexplained at go/no-go: 0 (record exit bar); reserve copy equals the plan at the nightly cut.
- **$** — No new dollars. Illustrative per-hold price of the no: 150 × 60% × $30 = $2,700 VERIFY-WITH-RAJESH — a worked example, not a savings claim, never added to S1's ≈$2.5M.

L6 LEVEL EVIDENCE — what the interviewer pattern-matches against the Amazon L6 functional rubric; every line is this story's record or labeled model

L6 DIMENSION	EVIDENCE IN THIS STORY
Ambiguity	Whether reserve counts as available was undefined; I defined it with a price.
Scope & influence	Merchandising, Planning, Finance, the SAP and ISM teams; the council leaned toward yes and the FSD said no.
Execution trade-off	Visibility without promisability: some potential early sales are forgone to make oversell zero by construction.
Impact	Oversell recurring → zero through two go-lives; MCS retired.
Process improvement	Reserve transitions in the matrix; invariants as assertions; old-versus-new reconciliation as the go/no-go evidence.
Complexity	Nightly feed plus intraday holds, an outside-S/4 bucket, three node types and channel reserves.
Root cause	The root cause was a definition that mixed "in the building" with "pickable"; fixed once.

PROBING — FACET 3: MATH & DATA LOGIC, FOUR TABLES, EVERY PLUG-IN SHOWN

Table 1 — before: the problem in numbers.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	AFS without reserve	`AFS = 980 − 200 − 120 − 60 − 40 = 560`	Basis min(1,000, 980) = 980; committed 420 (illustrative, S1 worked SKU)	560

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
2	AFS if reserve were promisable	`560 + 150 = 710`	merch reserve 150 held for a launch (illustrative) VERIFY-WITH-RAJESH	710
3	Over-promise if the hold slips	`710 − 560 = 150`	units shown that cannot be picked	150 units
4	Incidents if 60% sell before release	`150 × 60% = 90`	60% sell-through before release (illustrative) VERIFY-WITH-RAJESH	90 incidents
5	Cost of those incidents	`90 × $30 = $2,700`	$30 = $6 + $4 + $19 + $1 (S1 benchmark build)	$2,700 per hold
6	Double-subtract trap	`560 − 150 = 410`	reserve already outside S4_Avail, subtracted again (record design says no double-subtract path)	under-states by 150

Table 2 — the fix: step by step, values substituted.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Policy	`MerchRes sits outside S4_Avail; AFS = Basis − PickRes − EcomPick − EcomPack − SNB`	record design	reserve never enters AFS
2	Invariant 1 — copy equals plan	`Σ reserve (ISM) − Σ reserve (plan) = 0 per SKU × node at each nightly cut`	copy equality, tolerance 0 VERIFY-WITH-RAJESH	0 differences
3	Invariant 2 — never inside S/4 available	`AFS = 980 − 420 = 560, not 560 − 150`	no double-subtract	560
4	Invariant 3 — never promisable	`AFS(with reserve held) − AFS(reserve = 0) = 0`	changing MerchRes does not move AFS	0
5	Invariant 4 — one publisher	`published AFS values per SKU = 1`	OMS reads ISM AFS, never computes	1
6	Release	`reserve 150 → 0; stock moves to the sellable pool and AFS rises by the released units after the S/4 side posts`	visible in the FOH column until then	visible, then promisable
7	Old-vs-new reconciliation	`20,000 − 18,600 = 1,400 rows differ; explained 1,400; unexplained = 1,400 − 1,400 = 0`	illustrative counted rows VERIFY-WITH-RAJESH; go/no-go bar: 0 unexplained	0 unexplained → go
8	Break-even that would flip the no	`p × $19 = (1 − p) × $30 → p = 30 ÷ 49 = 0.612`	$19 margin if released in time; $30 incident if not (benchmarks)	p = 61%

Table 3 — after: the metric, before → after.

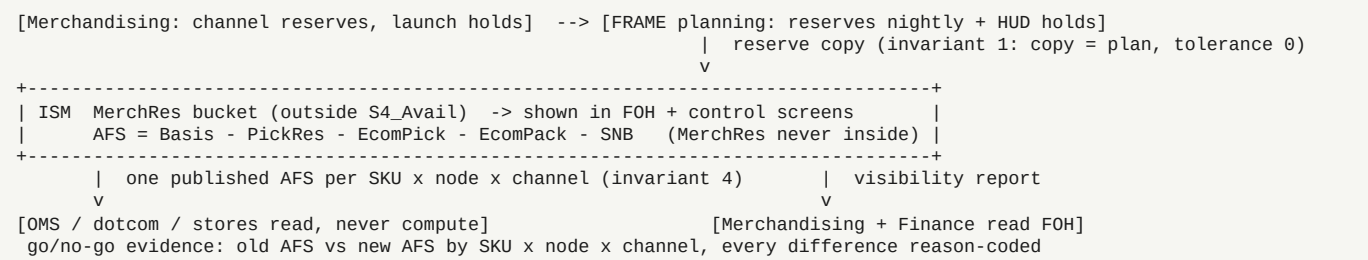
STEP	METRIC	BEFORE	AFTER	Δ / CHECK	BASIS
1	Reserve promisable	requested	visible, not promisable	`policy`	record
2	Over-promise exposure on the worked hold	150 units	0 units	`150 − 0 = 150`	illustrative
3	Unexplained old-vs-new differences at go/no-go	unknown	0	`0`	record exit bar; counts illustrative
4	SKUs with two published AFS values	possible	0	`0`	record
5	Oversell incidents	recurring	zero through two go-lives	`recurring − 0`	record

Table 4 — the dollars, every multiplication shown.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Per-hold price of the no	`150 × 60% = 90; 90 × $30 = $2,700`	illustrative units, sell-through and benchmark incident cost	$2,700 per hold
2	If 200 holds a year	`200 × $2,700 = $540,000`	200 holds a year VERIFY-WITH-RAJESH — a bound on exposure, not a savings claim	$0.54M exposure bound
3	Margin forgone by not promising (if all release in time)	`150 × $19 = $2,850`	$19 margin per unit (benchmark)	$2,850 per hold
4	Break-even reliability	`p × $19 = (1 − p) × $30 → p = 30 ÷ 49 = 0.612`	$19 margin if released in time; $30 incident if not (benchmarks); promising wins only above this reliability	p = 61% — unmeasured, so no

Merch reserve counts and sell-through are illustrative; the break-even 30 ÷ (19 + 30) = 61% reuses S1's benchmark inputs. The $0.54M is an exposure bound for a hypothetical number of holds, not a result.

SOLUTION FLOW — THE SYSTEM & DATA MAP



SENSITIVITY / IMPACT — LOW / BASE / HIGH

INPUT	LOW	BASE	HIGH
Sell-through before release	30%	60%	90%
Incidents on a 150-unit hold	45	90	135
Cost per hold at $30	$1,350	$2,700	$4,050
Break-even release reliability	30 ÷ (19 + 30) = 61%	same	same
Holds a year (exposure bound)	100 → $0.27M	200 → $0.54M	400 → $1.08M

SPOKEN TOTAL-IMPACT STATEMENT — SAY IT LIKE THIS

Conservatively, there is no new dollar in this story and I will not manufacture one. What I will say is the decision rule: promising a held unit pays about nineteen dollars if the hold releases in time and costs about thirty if it does not, so it only wins if release is reliable more than sixty-one percent of the time — and nobody measured that, so the answer was no, in writing, with visibility instead. The record is that oversell incidents went from recurring to zero through two go-lives.

FACET 5 — HARDEST FOLLOW-UPS, WITH NUMERIC ANSWERS NEW

Numeric — What did the ask do to AFS?

On the worked SKU AFS is 560. Counting a 150-unit reserve as sellable shows 710 — an over-promise of 150 if the hold slips. Subtracting it again when it is already outside S/4 would show 410 and under-state by 150; neither is right, which is why reserve is its own bucket outside AFS.

Numeric — What number would have changed your mind?

Release reliability. Promising earns about $19 a unit if the hold lifts in time and costs about $30 an incident if it does not: p × 19 = (1 − p) × 30, so p = 30 ÷ 49 = 61%. If Merchandising could show holds lifting on time more than about six times in ten, promisable would pay; nobody had that number.

Numeric — How did you know the reserve copy matched the plan?

Invariant one: Σ reserve in ISM minus Σ reserve in the plan is zero per SKU and node at every nightly cut, tolerance zero. At go/no-go I read the old-versus-new AFS reconciliation as counted rows — on illustrative numbers, 20,000 rows, 1,400 different, every one with a reason code, zero unexplained.

Numeric — Are 200 holds a year real?

No — it is a bound for the question "how big could the exposure be." At 60% sell-through and $30 an incident a 150-unit hold exposes $2,700; the record has no hold count and I do not claim one.

BEHAVIORAL PACK — FOUR QUESTIONS THIS STORY ANSWERS, EACH A DIFFERENT 90-SECOND MOMENT NEW

HAVE BACKBONE; DISAGREE AND COMMIT **1. Tell me about a time you disagreed with a senior stakeholder.**

Open (first two sentences): Merchandising wanted held-back stock to be sellable, and the room was leaning toward giving it to them.

A. I priced the ask instead of arguing it: unpickable stock sold is a cancellation. I offered visibility — reserve in its own column in the Financial On-Hand and on the control screens — and put the no into the FSD as a row with the pricing, so it could not be re-litigated every week. I told them the number that would change my mind. **R.** Merchandising's need was met, the FSD was signed 6/5/25, and oversell went from recurring to zero through two go-lives.

Land the number: the no lives in an FSD row | **Pivot to NACF:** At NACF: when planning wants held inbound stock promised, the same price-and-visibility answer.
Do NOT say: Caricaturing Merchandising; "they didn't understand."

CUSTOMER OBSESSION **2. Tell me about a time you protected the customer against an internal request.**

Open (first two sentences): The customer cost of selling held stock is a cancellation email, which is why I kept coming back to the unit nobody can pick.

A. I turned "reserve counts" into units: 560 shown with the reserve out, 710 with it in; 150 units the customer can be told yes to and cannot receive. **R.** Unpickable stock was never sold.

Land the number: 560 vs 710 | **Pivot to NACF:** The delivery date a customer is shown must be one the building can keep.
Do NOT say: Claiming a CSAT or NPS change.

ARE RIGHT, A LOT **3. Tell me about a decision you made with incomplete data.**

Open (first two sentences): I did not have the number that settles whether reserve should be promisable — release reliability — so I priced both sides.

A. $19 of margin if the hold lifts in time, $30 of incident if not: the break-even is 61%. Nobody had the reliability data, so the default went to the side with the customer cost; I wrote the break-even into the row so the decision could be reopened with evidence. **R.** The decision held through both go-lives.

Land the number: 61% break-even | **Pivot to NACF:** Make the number that would change your mind part of the decision.
Do NOT say: Presenting 61% as a measured release rate.

EARN TRUST **4. Tell me about a time you said no and kept the relationship.**

Open (first two sentences): Merchandising's need was real — they needed to see their reserve — so the no had to come with a yes.

A. The reserve got its own column in FOH and on the screens and a visibility report; Merchandising owns the holds, I owned the policy. **R.** The no did not become a standing dispute.

Land the number: visible, not promisable | **Pivot to NACF:** At NACF: the planner's hold is visible in the same view as the inbound plan.
Do NOT say: Naming individuals.

SAME STORY, 3 ANGLES — DIFFERENT OPENER, HEADLINE, LP NEW

DATA / ANALYTICAL

"Same SKU, two answers: 560 or 710."

560 vs 710

LP: Dive Deep

AFS = 980 − 420 = 560 with the reserve out; 710 with a 150-unit hold in.

The difference is the over-promise if the hold slips.

Break-even release reliability 30 ÷ (19 + 30) = 61%.

CONFLICT / INFLUENCE / BACKBONE

"The room was leaning toward yes."

1 FSD row

LP: Have Backbone

I priced the ask and offered visibility.

The no went into the FSD as a row, not a minute.

I named the number that would reopen it.

FAILURE / LEARNING

"The no I wrote had no review trigger."

p = 61%

LP: Are Right, A Lot

The record does not say the row carried a review condition — VERIFY-WITH-RAJESH.

What I would do differently: write the break-even and a review date into the row itself.

A no with no "what would change my mind" is just a veto.

TRAP DRILLS — THE THREE NASTIEST FOLLOW-UPS, WITH THE EXACT SENTENCE NEW

Trap — "You said no to your own business partner. What did Merchandising lose?"

The ability to take orders before a hold released — a potential early-sales margin of up to 150 × $19 = $2,850 on the worked hold if everything released in time. I did not hide it; I compared it with a 90-incident exposure of $2,700 at 60% sell-through and said the decision flips if reliability exceeds 61%.

Say exactly: "They lose early orders on held stock — about two thousand eight hundred fifty dollars on the worked hold if the release is on time; the decision flips above sixty-one percent reliability."

Trap — "What would you cut if cutover capacity halved?"

The invariants stay; what I cut is breadth of the old-versus-new reconciliation — start with the highest-volume nodes and channels, and defer long-tail rows to the day-after check. I would never cut invariant 1 (copy equals plan) or the one-publisher rule.

Say exactly: "I would thin the reconciliation scope, never the reserve-copy check or the single publisher."

Trap — "How do you know the reserve copy was right at go-live?"

Counted rows, not confidence: Σ reserve difference at the nightly cut is zero per SKU and node, and every old-versus-new difference row has a reason code. The go/no-go is a recommendation read from those counts, taken by the forum.

Say exactly: "Two counted rows: reserve-copy difference zero and unexplained differences zero — read, not argued."

Where this lands at Amazon / NACF. Amazon's version is the buy-box or detail-page promise on stock that is held, quarantined or pre-allocated: it should be visible to planners and excluded from the promise until released.

S1-d — Sephora S/4HANA ↔ ISM — go/no-go from counted rows: two go-lives, the UDC program's 836-row issues log read by severity, and two rollback triggers armed in numbers and never fired NEW

LPs this story carries (with proof lines below): Deliver Results (primary — dated gates held through two go-lives) · Bias for Action (reversible decisions: rollback written as numbers before go-live) · Insist on the Highest Standards (a phase graduates on counted rows, never on confidence) · Think Big (book-of-record migration across ~15 functions)

Process family: Plan / Deliver seam — release and cutover governance: go/no-go design, rollback triggers, reconciliation evidence, staged go-live

Honesty label. I wrote the go/no-go recommendation from counted rows, the readiness checklist, the rollback triggers and the publish-switch plan; the forum took the decision and the SAP and ISM teams executed the cutover. The 836-row issues log is the UDC conversion's log (source file "UDC Consolidated Issues Log for BY WMS-Locus"; Apple v11 C2) — the program both go-lives were sequenced around. It is not an ISM log: a pandas scan of the file finds no S/4HANA rows. I say "the UDC log" aloud and cite it as the gate the shared go-lives read. The rollback triggers are resume record; the dated gate chain is shared with S2 and S5 — one umbrella program, so if two of these stories go to one loop I say so. Row counts below are record (836, 85.4%, 164, 277) or derived from them; anything else is illustrative and flagged VERIFY-WITH-RAJESH.

LEADERSHIP PRINCIPLES CARRIED — ONE-LINE PROOF EACH NEW

Deliver Results — FSD v1.1 3/5/25 → v1.2 5/16/25 → signed 6/5/25; E2E Inbound 11/22 · Dotcom 12/20 · Retail 1/30; January planning and May execution go-lives — the dated chain held and oversell stayed at zero.

Bias for Action — Rollback was written as numbers — sync under 99% for four hours, or order-flow p95 above twice budget for one hour — and armed for both go-lives, so deciding to go was reversible.

Insist on the Highest Standards — A phase graduated on counted rows, never confidence: Blocker + Critical open = 0 at go/no-go, with an overall-percent-closed figure explicitly not accepted as a substitute.

Think Big — A book-of-record migration across about fifteen functions and sixty people, staged as two go-lives so each carried one defect class.

60-SECOND BUSINESS-FIRST PITCH — STOP ON THE PRIMARY METRIC

Situation. Moving the inventory book of record from MCS to S/4 with a real-time layer beside it meant every promise on three node types had to keep holding through the switch; the UDC program log those go-lives were gated on held 836 issues, and more than half — 441 — were Blocker or Critical at some point. **Task.** Own the go/no-go recommendation and the cutover design for the inventory supply integration, with no authority over the teams that would execute. **Action.** I wrote the go/no-go as a checklist read from counted rows: Blocker plus Critical open equal to zero; reconciliation clean for a stated run of days; monitors live; rollback rehearsed; control owners named; comms drafted. I split the cutover into two go-lives — January for planning, with the reserve rules live against demand, May for execution — so each carried one defect class. And I wrote rollback as two numbers before anyone asked: ISM-to-S/4 sync under 99% for four hours, or order-flow p95 above twice its budget for one hour. Both go-lives armed both triggers; neither fired. **Result.** The dated chain held — E2E Inbound 11/22, Dotcom 12/20, Retail 1/30 — oversell incidents recurring across three node types went to zero through both go-lives, and MCS was retired as book of record.

20-SECOND VERSION NEW

I turned the UDC program's 836-row issues log into a go/no-go: Blocker plus Critical open must be zero, reconciliation clean for a stated number of days, rollback rehearsed. I split cutover into two go-lives and wrote rollback as numbers — sync under 99% for four hours, or order-flow p95 over twice budget for an hour — armed both times, never fired; oversell stayed at zero.

NACF BRIDGE NEW

A new-FC or lane launch needs exactly this: a go/no-go from counted rows (door-to-stock time inside budget, dock-count-to-system tie clean N days) and a rollback written as a number before the first trailer arrives — such as redirecting inbound to the prior building if stow lag stays above twice its budget for an hour.

SITUATION / TASK — THE BUSINESS PROCESS PROBLEM, SCOPE, SCALE, BASELINES

S — The cutover touched every promise on Store, Retail DC and Dotcom DC. The UDC log the shared go-lives read held 836 issues; Blocker 164 plus Critical 277 is 441, or 52.8% of the log, so "percent closed" alone could not be trusted to describe readiness. Production changes were frozen October through December.

T — Write the go/no-go recommendation, the publish-switch plan and the rollback triggers for the inventory supply integration, and carry the program through two go-lives with zero customer-impacting incidents.

Baselines on the table: issues log 836 rows, 85.4% closed · Blocker + Critical 441 rows (52.8%) · oversell incidents recurring (target zero) · MCS as book of record (target retired) · E2E chain dated 11/22, 12/20, 1/30 (28 then 41 days apart).

FACET 1 — DATA ANALYSIS — WHAT I PROFILED, AT WHAT GRAIN, WHAT IT REVEALED NEW

- **Data.** The issues log by severity and status; E2E results by matrix cell; reconciliation results by day; monitor coverage; rollback rehearsal evidence; control-owner assignments.
- **Grain.** One row per issue, cut by severity and status; one row per matrix cell for test evidence.
- **What it revealed.** "85.4% closed" is true and uninformative: 836 × 85.4% ≈ 714 closed and 122 open — but if those 122 were all severe, 122 ÷ 441 = 27.7% of the severe rows would still be open at the same overall figure. The overall percentage hides the severity mix, the same way an aggregate accuracy hides the class that is off. So the gate had to be a count at the severity grain.

FACET 2 — SOLUTION DESIGN — OPTIONS I PUT ON THE TABLE, AND WHY ONE WON NEW

Four rejected ways to decide "go," and the one I chose, tested on the same log (836 rows, 441 severe, 122 open — illustrative worst case: all 122 severe).

OPTION	RULE	WHAT IT WOULD HAVE SAID	VERDICT
Calendar	go on the date	go — regardless of 122 open	Rejected — the date is a plan, not evidence
Executive confidence	go when leaders feel ready	depends on the room	Rejected — confidence is not a count
Overall % closed ≥ 85%	go at 85.4%	go — even if all 122 open rows were severe	Rejected — hides the severity mix
Single reconciliation snapshot	go if drift = 0 today	go — one clean day proves little	Rejected — N days running, not one
severity-grain counted rows plus the readiness checklist	Blocker + Critical open = 0; reconciliation clean N days; monitors live; rollback rehearsed; owners named; comms drafted	**no-go until each row is a count**	Chosen — a recommendation read from the checklist; the forum decides

ACTION — FACET 4: IMPLEMENTATION IN FIVE GATED PHASES (PHASES, OWNERS, CONTROLS, RISKS)

- **Discovery** — the program's dependency map and the list of defect classes by node and bucket; sessions closed 4/30. *Exit bar:* every bucket owned; open questions with owner and date.
- **Plan & Design** — FSD signed 6/5/25; the go/no-go checklist and thresholds written before testing; rollback triggers as numbers; the two-go-live split (January planning with reserve rules live, May execution). *Exit bar:* thresholds signed by the same three orgs.
- **Build & Test** — E2E Inbound 11/22, Dotcom 12/20, Retail 1/30 — 28 then 41 days apart — each green before the next; the matrix assertions. *Exit bar:* Blocker + Critical open = 0 at go/no-go; 0 open contract defects in SIT.
- **Cutover & Go-Live** — go/no-go recommendation from counted rows; the publish-switch plan — OMS reads ISM AFS and never computes; reserve-copy invariants; control-screen roles assigned; both triggers armed. *Exit bar:* January go-live and May go-live; no SKU with two published AFS values; Negative AFS clean at T+1.
- **Hypercare / Close** — daily drift; blockers per day against a threshold; triggers monitored; hypercare exit memo and FSD as-built note. *Exit bar:* drift = 0 SKUs on consecutive days; exit into the monthly release train.

Influence and ownership. Mine: the checklist, the thresholds, the recommendation, the trigger numbers and the publish-switch plan. Decision-makers: the go/no-go forum. Executors: SAP and ISM teams, WMS and OMS owners. The recommendation was read from counts, so the forum's decision could be made on evidence rather than on my confidence.

RESULTS — $ · # · % · TIME · CUSTOMER · QUALITY (BASIS-LABELED)

- **Customer** — Customer-facing oversell incidents recurring across three node types → zero through two go-lives (record).
- **Time** — FSD v1.1 3/5/25 → v1.2 5/16/25 → signed 6/5/25; E2E 11/22 · 12/20 · 1/30; January planning go-live; May execution go-live (record).
- **#** — Rollback triggers armed at both go-lives, never fired (record); Blocker + Critical open = 0 at go/no-go (record exit bar); 836-row log, 85.4% closed (record, shared with S2/S5).
- **Quality** — No SKU with two published AFS values; drift = 0 SKUs on consecutive days at hypercare exit; MCS retired as book of record (record).
- **$** — No new dollars. The cost of a one-month slip in retiring MCS is the user anchor over twelve months: $1.0M ÷ 12 = $83,333 a month — modeled, inside S1's anchor, and VERIFY-WITH-RAJESH on when MCS run cost actually stops.

L6 LEVEL EVIDENCE — what the interviewer pattern-matches against the Amazon L6 functional rubric; every line is this story's record or labeled model

L6 DIMENSION	EVIDENCE IN THIS STORY
Ambiguity	What "ready" means was not defined; I defined it as a checklist of counts.
Scope & influence	About fifteen functions and sixty people; a forum took the decision on my recommendation.
Execution trade-off	Two go-lives (more cutover work) for isolating defect classes; a higher bar than overall percent closed.
Impact	Oversell recurring → zero; MCS retired.
Process improvement	The checklist and the "graduate on counted rows" rule were standard work for later releases.
Complexity	A production freeze October–December, dated E2E days, and three node types.
Root cause	The risk was deciding on an aggregate; fixed by deciding at severity grain.

PROBING — FACET 3: MATH & DATA LOGIC, FOUR TABLES, EVERY PLUG-IN SHOWN

Table 1 — before: the problem in numbers.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Issues log	`836 rows; 836 × 0.854 = 714`	record: 836 rows, 85.4% closed	≈714 closed; 836 − 714 = 122 open

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
2	Severe rows	`164 + 277 = 441; 441 ÷ 836 = 0.527`	record: Blocker 164, Critical 277	441 rows = 52.8% of the log
3	Non-severe rows	`836 − 441 = 395`	derived	395 rows
4	What 85.4% could hide	`122 ÷ 441 = 0.277`	worst case: all 122 open rows severe (illustrative worst case)	27.7% of severe rows open at 85.4% overall
5	E2E spacing	`Inbound 11/22 → Dotcom 12/20 = 28 days; Dotcom 12/20 → Retail 1/30 = 41 days`	record dates	28 and 41 days
6	Rollback — sync trigger	`99% of 12,000 events/h = 11,880; 12,000 − 11,880 = 120`	illustrative 12,000 events an hour VERIFY-WITH-RAJESH	trigger at more than 120 unapplied per hour for 4 h

Table 2 — the fix: step by step, values substituted.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Gate 1 — severity	`Blocker + Critical open = 441 − 441 = 0`	all severe rows closed at go/no-go (record exit bar)	0 severe open
2	Gate 1 — implication	`open rows are non-severe: 395 − (714 − 441) = 122`	derived; snapshot of the final log VERIFY-WITH-RAJESH	122 open, all below Critical
3	Gate 2 — reconciliation	`drift rows = 0 on N consecutive days`	N stated before testing VERIFY-WITH-RAJESH	0 rows × N days
4	Gate 3 — checklist	`monitors live + rollback rehearsed + owners named + comms drafted = 4 of 4`	launch readiness checklist (record design)	4 of 4
5	Trigger A — sync	`trigger if sync < 99% for 4 h; 4 × 120 = 480 unapplied`	at the illustrative 12,000 events an hour	480 unapplied events
6	Trigger B — order flow	`trigger if p95 > 2 × budget for 1 h; 2 × 60 = 120 s`	budget = the 60-second publish budget (my reading) VERIFY-WITH-RAJESH	p95 > 120 s for 1 h
7	Two go-lives	`January: reserve rules live against demand; May: execution`	one defect class each	2 go-lives
8	Result check	`triggers fired = 0 of 2 per go-live`	armed both, never fired (record)	0 firings

Table 3 — after: the metric, before → after.

STEP	METRIC	BEFORE	AFTER	Δ / CHECK	BASIS
1	Go/no-go basis	confidence / % closed	counted rows by severity	`changed`	record design
2	Blocker + Critical open at go/no-go	441 rows in the log	0	`441 − 441 = 0`	record
3	Rollback triggers	none written	2 armed	`2`	record
4	Go-lives	one big switch	2 staged	`2 − 1 = 1`	record
5	Oversell incidents	recurring	zero through two go-lives	`recurring − 0`	record

Table 4 — the dollars, every multiplication shown.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Monthly MCS run cost	`$1,000,000 ÷ 12 = $83,333`	user anchor ÷ 12	≈$0.08M a month

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
2	Cost of a one-month slip	`1 × $83,333 = $83,333`	assumes run cost stops at cutover VERIFY-WITH-RAJESH	≈$0.08M
3	Cost of a three-month slip	`3 × $83,333 = $250,000`	same assumption	≈$0.25M
4	Check against the anchor	`$83,333 × 12 = $999,996`	rounding	≈$1.0M/yr

Only the log counts (836, 85.4%, 164, 277, 52.8%) and dates are record; 122 open and the worst-case 27.7% are derived arithmetic; the event volume and N days are illustrative. The slip cost is the $1M user anchor spread over twelve months — not a result.

SOLUTION FLOW — THE SYSTEM & DATA MAP

```
[Issues log 836 rows] -> by severity: B 164 + C 277 = 441 (52.8%) -> go/no-go gate 1: B + C open = 0
[E2E 11/22 Inbound -> 12/20 Dotcom -> 1/30 Retail]  -> gate 2: reconciliation clean N days
[Checklist: monitors live | rollback rehearsed | owners named | comms drafted]  -> gate 3
                                    |
                      recommendation (mine) --> forum decision
                                    |
     Jan go-live (planning, reserve rules live)  ->  May go-live (execution)
                                    |
armed during each: sync < 99% for 4 h  OR  order-flow p95 > 2x budget for 1 h  -> rollback (never fired)
                                    |
hypercare exit: drift = 0 SKUs on consecutive days; blockers/day under threshold -> release train
```

SENSITIVITY / IMPACT — LOW / BASE / HIGH

INPUT	LOW	BASE	HIGH
Open severe rows hidden at 85.4%	0 → 0%	60 → 13.6%	122 → 27.7%
Unapplied-events trigger per hour (at 12,000/h)	5,000/h → 50	12,000/h → 120	30,000/h → 300
Unapplied events across the 4-hour window	200	480	1,200
MCS slip (months) × $83,333	1 → $0.08M	3 → $0.25M	6 → $0.5M

SPOKEN TOTAL-IMPACT STATEMENT — SAY IT LIKE THIS

Conservatively, I do not claim a dollar for the cutover; the record is outcomes — two go-lives, oversell incidents going from recurring to zero, both rollback triggers armed and never fired. The only number I will put on timing is the anchor: a million dollars a year of MCS run cost is about eighty-three thousand a month, and I would say that is modeled from the anchor and depends on when the cost actually stops.

FACET 5 — HARDEST FOLLOW-UPS, WITH NUMERIC ANSWERS NEW

Numeric — 85.4% of 836 is closed. Why is that not good enough?

836 × 85.4% ≈ 714 closed and 122 open. If all 122 were severe, 122 of the 441 severe rows — 27.7% — would be open at the same overall figure. So the rule is a count at severity grain: Blocker plus Critical open equals zero. At that point the 122 open rows are all below Critical — an arithmetic fact I would verify against the final log.

Numeric — How do you write a rollback trigger?

As a threshold, a duration and an owner before go-live: sync under 99% for four hours, or order-flow p95 above twice budget for one hour. At an illustrative 12,000 events an hour, 99% is 11,880, so more than 120 unapplied an hour — 480 across four hours — pulls the trigger. Both armed at both go-lives, never fired.

Numeric — Why two go-lives, not one?

So each carries one defect class. January put the reserve rules against live demand, planning-side; May moved execution. If something broke in January it was a rules problem, not a floor problem. The E2E days were spaced 28 and 41 days apart so each cycle had time to close before the next.

Numeric — What did a delay cost?

I do not claim a cost of delay; I can bound it: the $1M MCS anchor is about $83,333 a month, so a three-month slip is about $250,000 if the run cost stops at cutover — which I would verify.

BEHAVIORAL PACK — FOUR QUESTIONS THIS STORY ANSWERS, EACH A DIFFERENT 90-SECOND MOMENT NEW

DELIVER RESULTS **1. Tell me about a time you delivered a complex result on a fixed date.**

Open (first two sentences): The chain was dated before we started: FSD signed 6/5/25, three E2E days, two go-lives — and an October to December production freeze that removed any slack.

A. I wrote the exit bar for every phase as a count and recommended go or no-go from the checklist. I split cutover so each go-live carried one defect class and kept both triggers armed. When something slipped it slipped inside a gate, not over one. **R.** Both go-lives landed; oversell recurring → zero; MCS retired.

Land the number: zero incidents through two go-lives | **Pivot to NACF:** Inbound launches need the same chain: dated test days, a freeze, a staged opening.
Do NOT say: The Utah DC floor story or the robotics gate — Goforth's.

BIAS FOR ACTION **2. Tell me about a time you moved fast on a risky decision.**

Open (first two sentences): The decision to go was risky; what made it fast was that it was reversible.

A. I wrote the rollback before go-live as two numbers — sync under 99% for four hours, or order-flow p95 above twice budget for an hour — so the decision to go did not require certainty, only a way back. **R.** Armed at both go-lives; never fired.

Land the number: two triggers, never fired | **Pivot to NACF:** The same: launch a lane with a written redirect trigger.
Do NOT say: Saying the triggers "would have worked" — they were never tested in anger.

INSIST ON THE HIGHEST STANDARDS **3. Tell me about a time you wouldn't sign off on a number others accepted.**

Open (first two sentences): 85.4% closed sounded good, and I would not accept it.

A. I showed the same overall figure with all 122 open rows severe is 27.7% of the severe rows still open, and made Blocker plus Critical open equal zero the gate. **R.** Go/no-go read from counts; the forum decided.

Land the number: 27.7% hidden | **Pivot to NACF:** On inbound: 99.0% overall accuracy hiding 98.0% on one class.
Do NOT say: Presenting the 27.7% as what the log showed — it is a worst case.

THINK BIG **4. Tell me about the largest change you have led.**

Open (first two sentences): We moved the inventory book of record while every promise on three node types had to keep holding.

A. I broke it into five gated phases and two go-lives, with each dated gate written before the work and the whole thing frozen through peak. About fifteen functions, sixty people, none reporting to me. **R.** Zero customer-facing incidents through both go-lives.

Land the number: ~15 functions · ~60 people · 2 go-lives | **Pivot to NACF:** The scale at NACF is the number of buildings; the discipline is the same.
Do NOT say: The 80 names and a dozen core systems — UDC program figures; say them only if asked about the umbrella program.

SAME STORY, 3 ANGLES — DIFFERENT OPENER, HEADLINE, LP NEW

DATA / ANALYTICAL

"The log said 85.4% closed; I asked how many of the open ones were severe."

27.7%

LP: Dive Deep

836 × 0.854 ≈ 714 closed; 122 open.

Worst case all 122 are severe: 122 ÷ 441 = 27.7%.

The gate is a count at severity grain, not a percent.

CONFLICT / INFLUENCE / BACKBONE

"The date was the most powerful number in the room."

441 → 0

LP: Have Backbone · Insist on the Highest Standards

I made the checklist a recommendation, not a decision.

The forum could still decide to go; it could not do so without reading the counts.

Whether any forum disagreed with my recommendation is VERIFY-WITH-RAJESH.

FAILURE / LEARNING

"The triggers never fired — which means I never learned whether the thresholds were right."

99% · 4 h · 2× · 1 h

LP: Are Right, A Lot

Armed and never fired is a good outcome and an unverified calibration.

If I did it again I would add a rehearsal that injects a breach to prove the trigger fires.

Whether a rehearsal was run is VERIFY-WITH-RAJESH.

TRAP DRILLS — THE THREE NASTIEST FOLLOW-UPS, WITH THE EXACT SENTENCE NEW

Trap — "You said go/no-go was your recommendation. What if the forum overruled you?"

Then the checklist rows are on the page: Blocker plus Critical open, reconciliation days, monitors, rollback, owners, comms. A forum can accept a risk on a row; it cannot do so unseen. The recommendation is read from the checklist; the decision is theirs.

Say exactly: "The forum owns the decision; I own making every row a count they can see."

Trap — "What do you cut if the date is fixed and the log is not clean?"

I cut scope before I cut the gate: defer a node type or a channel to the second go-live — January and May were already a staged split — and never relax Blocker plus Critical equals zero, rollback armed, or owners named.

Say exactly: "Move scope between go-lives; never relax the severity gate or the rollback."

Trap — "What is the dependency that could have broken this?"

The WMS feed and the order system had to be ready to the same contract; the S/4 side owned MB52 and posting. The single point I watched was the sync rate to S/4, because it was one of the two triggers: below 99% for four hours means ISM and the book are diverging, and that is a rollback.

Say exactly: "The dependency I watched most was the sync to S/4, because it was a rollback trigger."

Where this lands at Amazon / NACF. A staged opening is how a new building ramps: dated test days, a go/no-go from counted rows, and a rollback written as a number — the FC launch playbook in one page.

S1-e — Sephora S/4HANA migration — seventeen interfaces, two global dashboards and the $42K I first refused: refuse the unscoped ask, fund the study, reverse on evidence, and test with count-and-sum arithmetic NEW

LPs this story carries (with proof lines below): Are Right, A Lot (primary — changed position on evidence, in public) · Frugality (fund the study, not the fear) · Earn Trust (the public reversal bought credibility) · Dive Deep (three cheap checks, each owning one failure mode)
Process family: Plan / Deliver seam — cross-program arbitration and regression-test design: scope, budget, non-regression testing before UAT

Honesty label. I owned the arbitration — scope, budget and a testing strategy both programs would sign — not the global reporting layer. Basis. *Record-from-docs*: the ROMA "S!ONE NORAM — Testing Strategy Review" deck (Feb 2026, pp. 2–5) carries the 17 ROMA interfaces "potentially impacted" (as-is: 10 EDWH files + 7 Databricks files), the summer-2025 refusal of ROMA's non-regression budget, the fall-2025 study, the two critical usages (Blueberry for worldwide executive management; Lumena, the global brand portal), Option 1 "mandatory" 42k non-regression tests and Option 2 "optional" 25k mapping adaptations, delta files from S/4 for US and Canada, count and total-per-store KPIs, and NRT in April / UAT in May 2026. The deck prints "42k" and "25k" with no currency sign. *User anchor* (loop packet S25 only): that I owned the arbitration, reversed my no, that Option 1 was funded and that the 25k was not pre-spent. *CONFLICT*: the deck says the fall study had NORAM teams "saying no impacts"; so I say "the study turned seventeen into a named file list for two critical usages", never "the study found defects." LOOP-CONSISTENCY FLAG: the loop packet lists S25 as a Bar Raiser backup; use it here only if it is not being spent with the Bar Raiser. I never price the avoided regression.

LEADERSHIP PRINCIPLES CARRIED — ONE-LINE PROOF EACH NEW

Are Right, A Lot — "Potentially impacted" is not a scope, so I refused — then funded a study, and when it named the files behind the two critical usages I reversed my own refusal on the record (reversal: user anchor).

Frugality — I refused to fund fear — an unscoped ask — and funded exactly what the study proved: $42K for the files feeding two critical usages, with $25K held contingent and never pre-spent.

Earn Trust — The public reversal bought more credibility with the global program than the original refusal had cost; the first no shipped with a funded study so being wrong was cheap.

Dive Deep — Count, sum and per-store matching — three checks, each naming the failure it catches — instead of clicking through screens.

60-SECOND BUSINESS-FIRST PITCH — STOP ON THE PRIMARY METRIC

Situation. During the North America S/4 migration, the global reporting program — the layer behind the dashboard worldwide leadership reads and the portal the global brands use — said our migration potentially touched seventeen of its interfaces and asked for a non-regression testing budget. **Task.** I owned the arbitration: scope, budget and a testing strategy both programs would sign. **Action.** I said no: "potentially impacted" is not a scope. But I funded a joint study to turn it into a list. It came back as a named list: our side expected no impact, but the reporting team showed exactly which files fed two critical global usages — six core files for the executive dashboard, and those plus the CRM and supply-chain files for the brand portal. I reversed myself on the record and funded the $42K non-regression option, scoped to those files, and held the second option — $25K to adapt mappings — strictly contingent on a mapping actually changing. The tests were arithmetic: for each file, a delta extract from the S/4 environment checked on count, sum and per-store match against legacy, and sequenced before business UAT. **Result.** $42K spent against a risk my first answer would have left naked; the two global-critical usages crossed the migration behind named gates; migration dates held; and the contingency was not pre-spent. I do not put a dollar on the regression that did not happen.

20-SECOND VERSION NEW

The global reporting program said the S/4 migration "potentially" touched seventeen of its interfaces and wanted test budget. I refused the unscoped ask but funded a study, which named the files behind two critical usages; I reversed myself publicly, funded $42K scoped to those files, held $25K contingent, and tested with count, sum and per-store arithmetic before UAT.

NACF BRIDGE NEW

Any NACF change that touches a shared feed — a new inbound system, a re-slotted receiving flow — has downstream readers like planners' dashboards. Same move: do not fund fear; fund a study that converts "potentially impacted" to a list, then test the list with counts, sums and per-node matching so aggregate matches cannot hide a shifted mapping.

SITUATION / TASK — THE BUSINESS PROCESS PROBLEM, SCOPE, SCALE, BASELINES

S — The NORAM S/4 migration potentially affected 17 interfaces feeding the global reporting layer, with two critical usages on top — the worldwide executive dashboard and the global brand portal. The reporting program asked for regression-test budget; the migration program — mine — had initially refused.

T — Own the arbitration: the scope, the budget and a testing strategy both programs would sign, without delaying the migration's hard dates.

Baselines on the table: candidate interfaces — 17 ("potentially impacted") · critical usages — 2 · regression-test budget — unscoped ask · options on the table — $42K non-regression (Option 1), $25K mapping-change adaptation (Option 2) · test method — undefined.

FACET 1 — DATA ANALYSIS — WHAT I PROFILED, AT WHAT GRAIN, WHAT IT REVEALED NEW

- **Data.** The interface inventory (17 candidates), the mapping from interfaces to the two critical usages, the file formats each usage consumes, and delta extracts from the S/4 environment (US and Canada data) against legacy.
- **Grain.** Interface → file → usage; within a file, store or storage-location rows. Checks ran at the store level, not the total.
- **What it revealed.** The 17 were the as-is inventory (10 EDWH + 7 Databricks files). The test list was usage-driven: six core files behind Blueberry (5 EDWH + 1 Databricks), and for the brand portal those plus the CRM client file and a twelve-file supply-chain family, 19 in all, with the supply-chain files new via Lumena (record-from-docs, deck pp. 3–4). A regression almost never invents plausible data: it drops rows, doubles rows through a bad join, or shifts a mapping.

FACET 2 — SOLUTION DESIGN — OPTIONS I PUT ON THE TABLE, AND WHY ONE WON NEW

Budget options as presented to the two programs; the avoided regression is deliberately not priced.

OPTION	COST	COVERS	VERDICT
Fund the unscoped ask (all 17)	unscoped	"potentially" everything	Rejected — fear is not a line item
Fund nothing	$0	nothing	Rejected — leaves the two critical usages naked
Fund a joint study first	not priced in the deck — a study "GO" only (fall 2025)	converts "potentially" to a list	Chosen as step one
Option 1 — non-regression tests on the critical files	$42K	files behind the two critical usages	**funded on the study's evidence**
Option 2 — adapt mappings if one changes	$25K	only if a mapping actually changes	**as contingent — not pre-spent**

ACTION — FACET 4: IMPLEMENTATION IN FIVE GATED PHASES (PHASES, OWNERS, CONTROLS, RISKS)

- **Discovery** — the unscoped ask and my first answer — no; the joint study converting 17 “potentially impacted” interfaces to a named list. *Exit bar:* a file list tied to the two critical usages.
- **Plan & Design** — the reversal, on the record; Option 1 funded at $42K scoped to the MVP file list; Option 2 held at $25K contingent; test design — delta extract per interface, count-and-sum invariants per store or storage location. *Exit bar:* both programs sign the testing strategy.
- **Build & Test** — delta extract from S/4 (US and Canada) per interface; three checks per file; non-regression tests sequenced ahead of business UAT and both ahead of the migration’s hard dates. *Exit bar:* every file passes count, sum and per-store checks.
- **Cutover & Go-Live** — the two critical usages cross the migration behind named gates. *Exit bar:* gates green; contingency untouched unless a mapping changed.
- **Hypercare / Close** — the contingency stays contingent; lessons carried into the testing strategy. *Exit bar:* no regression on the two critical usages.

Influence and ownership. Mine: the refusal, the study’s framing, the reversal, the two budget options, the test design and the sequencing. Built and run by: the migration team and the global reporting program’s owners. The global program decided what its critical usages were; I decided what I would pay for.

RESULTS — $ · # · % · TIME · CUSTOMER · QUALITY (BASIS-LABELED)

- **$** — $42K spent against a risk my first answer would have left naked (*user anchor*; the 42k option itself is record-from-docs, deck p4); the $25K contingency never pre-spent (*user anchor*); the avoided regression is deliberately unpriced.
- **Quality** — The two global-critical usages crossed the migration behind named arithmetic gates instead of hope (*user anchor*; the count and total-per-store checks are record-from-docs, deck p5).
- **Time** — Migration dates held with the global layer protected (*user anchor*); tests sequenced ahead of business UAT (record-from-docs: data from US early April, NRT April, UAT May 2026).
- **Team** — The public reversal bought more credibility with the global program than the original refusal had cost (*user anchor*).
- **Whether the $25K was ever spent** — the deck says Option 2 is for mapping changes and that “we expect no data mapping change”; the loop packet says not pre-spent (*user anchor*). Say “held contingent,” nothing more.

L6 LEVEL EVIDENCE — what the interviewer pattern-matches against the Amazon L6 functional rubric; every line is this story's record or labeled model

L6 DIMENSION	EVIDENCE IN THIS STORY
Ambiguity	“Potentially impacted” — I converted it to a named list by funding a study.
Scope & influence	Two programs, one budget, a global reporting layer I did not own.
Execution trade-off	$42K and calendar spent on tests that, if all goes perfectly, find nothing — against the program’s own budget discipline.
Impact	Two critical global usages protected; dates held.
Process improvement	Count, sum, per-store checks as a reusable non-regression pattern.
Complexity	Seventeen interfaces, two dashboards, a hard migration date and two separate programs.
Root cause	The root cause of the conflict was an unscoped ask; fixed by a list.

PROBING — FACET 3: MATH & DATA LOGIC, FOUR TABLES, EVERY PLUG-IN SHOWN

Table 1 — before: the problem in numbers.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Candidate interfaces	17	record	17

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
2	Budget on the table	`$42,000 + $25,000 = $67,000`	Option 1 + Option 2 (record)	$67,000 maximum
3	Contingent share	`$25,000 ÷ $67,000 = 0.373`	Option 2 ÷ both	37%
4	Funded share at the start	`$42,000 ÷ $67,000 = 0.627`	Option 1 ÷ both	63%
5	Failure A — rows dropped	`5,000 − 4,950 = 50`	legacy 5,000 rows; S/4 delta 4,950 (illustrative)	50 rows dropped
6	Failure B — rows doubled by a bad join	`$4,000,000 ÷ $2,000,000 = 2.0`	S/4 sum $4,000,000 vs legacy $2,000,000 (illustrative)	ratio 2.0 — doubled
7	Failure C — a shift that cancels	`1,200 + 800 = 2,000 and 1,000 + 1,000 = 2,000`	stores A and B, legacy vs S/4 (illustrative)	totals match; stores do not

Table 2 — the fix: step by step, values substituted.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Check 1 — count	`legacy rows − S/4 rows = 5,000 − 4,950 = 50`	catches drops (record design); pass = 0	FAIL on A
2	Check 2 — sum	`S/4 sum − legacy sum = $4,000,000 − $2,000,000 = $2,000,000`	catches doubles; pass = 0	FAIL on B
3	Check 3 — per store (A; B)	`1,000 − 1,200 = −200; 1,000 − 800 = 200`	catches shifts that cancel in the total; pass = 0 per store	FAIL on C, though 2,000 − 2,000 = 0
4	Pass case	`5,000 − 5,000 = 0; $2,000,000 − $2,000,000 = $0; 1,200 − 1,200 = 0; 800 − 800 = 0`	all three zero	file passes
5	Sequencing	`non-regression tests before UAT; both before the migration's hard dates`	record	tests first
6	Contingency	`Option 2 spent only if a mapping changes: expected mapping changes = 0`	the strategy assumed none (record)	$0 pre-spent
7	Break-even question (not a price)	`p × L* = $42,000 → L* = $42,000 ÷ 10% = $420,000`	at p = 10% chance of a regression (illustrative); L* is the loss above which funding pays — never a claim of the regression's value	L* = $420,000 at p = 10%

Table 3 — after: the metric, before → after.

STEP	METRIC	BEFORE	AFTER	Δ / CHECK	BASIS
1	Interfaces in test scope	17 "potentially impacted"	files behind 2 critical usages (6 core + supply-chain family)	`scope cut`	record
2	Budget committed	unscoped	$42,000	`scoped`	record
3	Budget held contingent	$0 defined	$25,000 not pre-spent	`$25,000`	record
4	Checks per file	none defined	3 (count, sum, per-store)	3	record design
5	Critical usages protected	2 at risk	2 behind named gates	2	record

Table 4 — the dollars, every multiplication shown.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Committed	`$42,000 spent`	record	$42K
2	Contingent, not spent	`$67,000 − $42,000 = $25,000`	contingency held	$25K
3	Threshold at p = 5%	`$42,000 ÷ 5% = $840,000`	illustrative probability	$840,000
4	Threshold at p = 25%	`$42,000 ÷ 25% = $168,000`	illustrative probability	$168,000

Row counts and dollar sums in the failure cases are illustrative. The break-even thresholds are a way to read the decision — at what loss $42K pays at a given probability — not a valuation of the regression, which I deliberately never price.

SOLUTION FLOW — THE SYSTEM & DATA MAP

```
[Global reporting program: 17 candidate interfaces]  --"potentially impacted"-->  ask: test budget
        |  my first answer: NO (unscoped)  +  fund a joint study
        v
[Study: interface -> file -> usage]  --> 2 critical usages: executive dashboard, global brand portal
        |  list = 6 core files + supply-chain file family
        v
[Option 1: $42K non-regression tests, scoped to the list]   [Option 2: $25K mapping adaptations - contingent]
        |
        v  per file: delta extract from S/4 (US + Canada) vs legacy
   count (drops) -> sum (doubles) -> per-store match (shifts that cancel)  -> gate
        |  tests ahead of business UAT; both ahead of the migration's hard dates
```

SENSITIVITY / IMPACT — LOW / BASE / HIGH

INPUT	LOW	BASE	HIGH
Probability of a regression without tests	5%	10%	25%
Break-even loss L* = $42,000 ÷ p	$840,000	$420,000	$168,000
Contingent share of total budget	25 ÷ 67 = 37%	37%	37%
Files on the MVP test list (deck p4)	6 (Blueberry)	19 (brand portal: 5 EDWH + 14 Databricks)	19 (union; Blueberry's 6 are a subset)

SPOKEN TOTAL-IMPACT STATEMENT — SAY IT LIKE THIS

Conservatively, there is no dollar to claim here and I would not take one: forty-two thousand was spent, twenty-five thousand was held and not pre-spent, and the value of the executive dashboard not breaking is deliberately unpriced. What is on the record is the sequence — refuse the unscoped ask, fund the study, reverse on evidence — and that the two critical usages crossed behind arithmetic gates.

FACET 5 — HARDEST FOLLOW-UPS, WITH NUMERIC ANSWERS NEW

Numeric — Why $42K, and why not the whole seventeen?

Because the study showed that most of the seventeen did not feed anything critical. $42K bought count, sum and per-store tests for the files behind the two critical usages — six core files plus the supply-chain family. The $25K for mapping changes stayed contingent because the strategy assumed no mapping would change.

Numeric — Show me a regression that totals hide.

Stores A and B: legacy 1,200 and 800 visits, S/4 1,000 and 1,000. Both totals are 2,000, so count and sum pass in aggregate — but store A is off by −200 and B by +200. Per-store matching is the check that sees a mapping shift cancelling in the total.

Numeric — How do you read $42K against a risk you never priced?

As a threshold, not a price: at a 10% chance of a silent regression, $42K pays for itself against any loss above $420,000; at 5% above $840,000; at 25% above $168,000. I never put a dollar on the dashboard breaking — I ask whether anyone believes the loss is smaller than that.

Numeric — Were you simply wrong the first time?

The first no was right on the information available: an unscoped ask with no interface list. What made it safe is that the no shipped with a funded study, so being wrong was cheap and fast to find out.

BEHAVIORAL PACK — FOUR QUESTIONS THIS STORY ANSWERS, EACH A DIFFERENT 90-SECOND MOMENT NEW

ARE RIGHT, A LOT **1. Tell me about a time you changed your mind based on data.**

Open (first two sentences): I refused a budget request, then reversed myself in public when the study came back.

A. I said no to an unscoped "potentially impacted" ask but funded a small joint study to turn it into a list. It showed real risk on two critical usages, so I funded the $42K option scoped to those files and told both programs I had changed my position and why. **R.** The usages crossed behind named gates and the public reversal bought credibility.

Land the number: $42K funded after a no | **Pivot to NACF:** At NACF: refuse the unscoped "this might affect inbound" and fund the study that names the interfaces.
Do NOT say: Claiming the dashboard would have broken — unknown and unpriced.

FRUGALITY **2. Tell me about a time you got more by spending less.**

Open (first two sentences): I refused to fund fear, but I did not refuse to fund evidence.

A. A small study replaced an open-ended budget; Option 2, $25K, was written as contingent and never pre-spent. I scoped tests to the files the critical usages consume. **R.** $42K spent against a risk, $25K held back.

Land the number: $42K, not $67K | **Pivot to NACF:** Inbound: spend on the lanes that matter, hold the rest.
Do NOT say: Calling it a savings — it is a refused open-ended ask.

EARN TRUST **3. Tell me about a time you admitted you were wrong.**

Open (first two sentences): The global program had a legitimate concern and I had said no.

A. I reversed on the record as soon as the evidence arrived, and said so to both programs rather than quietly funding. The first no had shipped with a study, so it never read as a stonewall. **R.** The reversal bought more credibility than the refusal had cost.

Land the number: public reversal | **Pivot to NACF:** Same posture with any planning partner: name the change of mind.
Do NOT say: Saying I was right all along.

DIVE DEEP **4. Tell me about a time you designed a test that caught what totals miss.**

Open (first two sentences): The tests had to find a regression that looks plausible, so I started from how migrations actually fail.

A. They drop rows, double rows through a bad join, or shift a mapping. Count catches drops, sum catches doubles, per-store catches shifts that cancel in the total — three cheap checks, each owning one failure mode. **R.** Delta extracts per interface passed or failed on arithmetic.

Land the number: 3 checks, 3 failure modes | **Pivot to NACF:** Inbound: the same per-door and per-lane matching, not the building total.
Do NOT say: Illustrative row counts as real.

SAME STORY, 3 ANGLES — DIFFERENT OPENER, HEADLINE, LP NEW

DATA / ANALYTICAL

“The totals matched and two stores were wrong by two hundred each.”

1,200 / 800 vs 1,000 / 1,000

LP: Dive Deep

Count 5,000 vs 4,950 catches 50 dropped rows; sum ratio 2.0 catches a double.

Per-store matching catches the shift: −200 and +200.

Three checks, each owning one failure mode.

CONFLICT / INFLUENCE / BACKBONE

“I told another program’s lead no.”

17 → 2

LP: Have Backbone · Frugality

“Potentially impacted” is not a scope.

The no came with a funded study.

I did not fund the unscoped ask.

FAILURE / LEARNING

“Was I wrong the first time?”

$42K

LP: Are Right, A Lot

The first no was right on the information available, and shipped with a study so being wrong was cheap.

The cost was calendar and some credibility; the public reversal repaid it.

Differently: ask for the interface list on day one of the program, not on the first request.

TRAP DRILLS — THE THREE NASTIEST FOLLOW-UPS, WITH THE EXACT SENTENCE NEW

Trap — “You said no to a legitimate risk. How is that not a failure?”

The no was to the unscoped version; the same answer funded a study. A wrong no costs a regression on a global-critical asset; a wrong yes costs $42K — so I ran the loop that made the no cheap: refuse, study, reverse on evidence. I would run it again.

Say exactly: “The no shipped with a funded study, so being wrong was cheap and fast.”

Trap — “What did you cut from the testing scope, and what could you never cut?”

I cut from the paid scope the interfaces the study cleared — they did not feed the critical usages — and held the mapping-adaptation option as contingent. I never cut the per-store match, because it is the check that catches the shift that cancels in the total.

Say exactly: “I cut the interfaces the study cleared; I never cut the per-store check.”

Trap — “How do you know the tests protected anything?”

I do not and I do not claim to — I deliberately do not price the avoided regression. What I can say is the arithmetic gates were in place and passed, and that I know exactly which three failure modes they would have caught.

Say exactly: “I can’t prove a regression that didn’t happen; I can show the three failures the gates would have caught.”

Where this lands at Amazon / NACF. Whenever a change touches a shared feed — at Amazon scale, an inbound system change under planners’ dashboards — the move is the same: refuse the unscoped ask, fund the list, test with counts, sums and per-node matches.

Peer operating drills — scope, dependencies, what I would cut, what I would never cut NEW

A peer tests cross-functional operating judgment: not whether the design was clever, but how I scope, who I depend on, what I would cut when capacity or time halves, and what I refuse to cut. Cut orders below are my judgment from the record's gates and exit bars; they are not claims of something that happened.

STORY	CUT FIRST	NEVER CUT	BINDING DEPENDENCY
S1	Long-tail node × channel cells to the second go-lives	Bucket owners; lower-of-two; signed FSD row for each no	Three signatory orgs; the council's three decisions
S3	Ship-to-home source — rides the overnight path longer	Matched-only clear; SNB never negative; the aging report with a named owner	S/4 Aggregation# on the posting; Inventory Accounting
S1-a	The 60-second budget on non-sale events	Idempotency key; reversal names its original; one bucket per event	POS, OMS, WMS and S/4 all honouring one contract
S3-a	Non-POS sources first (omni, ship-to-home)	Σ lines = posting to the cent; the duplicate-header guard	Order# join keys from the POS
S1-b	Internal Audit's read of every correction → a monthly sample (the audit row is still written on every one)	The audit row; T+1 clean; the drift report	Inventory Control's daily capacity (6.5 of 7.2 hours)
S1-c	Breadth of old-vs-new reconciliation (start with high-volume nodes)	Reserve-copy equality; one publisher; the no as an FSD row	Planning's reserve feed (nightly plus holds)
S1-d	Scope between go-lives (a node type or channel)	Blocker + Critical open = 0; armed rollback; owners named	The sync to S/4 (a rollback trigger)
S1-e	Interfaces the study cleared from the paid scope	The per-store match; the contingent option held	The global program's own owners deciding what is critical

EIGHT DRILLS, WITH THE EXACT SENTENCE NEW

Trap — "You have half the people and the date does not move. What do you cut?"

Scope before gates, sources before controls, breadth before depth. Cut: long-tail node × channel cells to the second go-live, the ship-to-home source, non-sale latency budgets, and the breadth of the old-versus-new reconciliation. Never cut: the bucket owners, matched-only clear, idempotency, the audit row on a correction, Blocker + Critical open = 0, or an armed rollback.

Say exactly: "I cut scope and sources, never the controls: owners, idempotency, the audit row, the severity gate and the rollback stay."

Trap — "Which dependency worried you most, and what did you do about it?"

The ISM-to-S/4 sync, because it was a rollback trigger — sync below 99% for four hours — and because the S/4 side owned MB52 and the posting. I made it visible instead of trusting it: a daily drift report at tolerance zero, an owner per open question with a date, and dated E2E days 28 and 41 days apart so each cycle could close before the next.

Say exactly: "The sync to S/4 — so I made it a rollback trigger and a daily drift report, not an assumption."

Trap — "Everyone wants their thing in. How do you scope?"

By consumers, not candidates. S1-e: seventeen candidate interfaces became the files behind two critical usages. S1: epics slotted by business criticality × effort, and no story enters QA without a matrix cell reference. Scope is a list with an owner for each line, and the council is asked for three decisions, not a review.

Say exactly: "Scope is a list of what is consumed, with an owner per line — not a list of what might be touched."

Trap — “What would you stop doing in your first ninety days at NACF?”

I would stop myself from building a new dashboard before one number has one owner. First: the book of record and the primary metric; then a stocks-and-flows map of inbound; then the invariant already broken — typically a unit in two stocks or a number with two owners. This is a plan, not a record.

Say exactly: “I would not add a dashboard before one inbound number has one owner and one daily drift check.”

Trap — “What do you want me to just fix, and what do you want me to bring to you?”

Fix what is inside the envelope and reversible today: a stale reservation corrected on the screen with a reason code, a vendor or door re-slot. Bring policy: a rule change, a bucket definition, a threshold — those ride the release train and are signed by the same signatories, never a side configuration.

Say exactly: “Fix reversible acts with an audit row; bring anything that changes a definition or a threshold.”

Trap — “A metric we both touch goes red. Who speaks first?”

The owner of the input speaks first, with the decomposition already written. In ISM that is the owning control role — Inventory Control for drift and negatives — and a red drift row arrives with a reason code and an owner before the review, not in it.

Say exactly: “The owner of the input speaks first, with the cause and the date already written.”

Trap — “Two teams’ priorities conflict and neither reports to you. How do you resolve it?”

Price both sides and write the number that would flip the answer. S1-c: promising held stock earns about $19 a unit if it releases in time and costs about $30 an incident if not; break-even 61%. S1-e: $42K against a regression I never price, read as a threshold at a chosen probability. Then the decision lives in a document.

Say exactly: “I put a price on both sides and write down the number that would change my mind.”

Trap — “How would you know a cut was a mistake?”

Each cut has a tripwire written in advance — the way the rollback triggers were: a threshold, a duration and an owner. For example, if the ship-to-home source stays on the overnight path and its unmatched rows exceed the aging-report capacity — about 10 a day at the 0.5% KPI on an illustrative 2,000 postings — I pull it back in.

Say exactly: “Every cut ships with a tripwire: a threshold, a duration and an owner.”

NACF bridge — AFS on the inbound side: what the promise can see NEW

Use this in the last minute of any S/4 answer, or when asked how the work transfers. Every inbound number below is from the inbound booklet's worked week (illustrative there: one vendor, three SKUs, 12,000 units on six trailers); the S/4 side is record design from S1.

S/4 / ISM CONCEPT	INBOUND ANALOG	INBOUND BOOKLET ANCHOR	LINE TO SAY
AFS = Basis − committed buckets	Inbound AFS = opening on-hand + units that clear receive and stow	F: planned 1,000 + 6,000 = 7,000; actual 1,000 + 5,880 = 6,880	The promise sees only what cleared stow.
Shipped-Not-Billed (sold, not yet posted)	Received-not-stowed (on the dock or in receiving)	Availability starts at stow, not at the dock	A unit on the dock is the SNB of a building.
Lower-of-two basis	Promise on the lower of ASN and dock count	ASN 6,000 vs counted 5,880: promise 5,880	Wrong on the safe side, visible.
Merch reserve: visible, not promisable	Stranded / held stock: present, not sellable	"Stranded — physically present, not sellable… inbound's silent inventory"	A stow without a sellable record is not availability.
Drift report, tolerance zero	Dock count vs system receipt, daily, per building	Accuracy 99.0% overall hides 98.0% on F	An aggregate hides the line.
Aggregate posting vs lines (S3-a)	ASN header vs ASN lines	0.02 × 6,000 = 120 units short on F alone	Lines must sum to the header or the row ages.
Negative AFS screen (S1-b)	Promised more than received-and-stowed	Cover 7.00 days → 6.76 days	Never clamp a negative promise; show it.

WORKED ON THE BOOKLET'S FAST MOVER — FOUR SMALL TABLES NEW

Table 1 — before: the problem in numbers.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Planned inbound AFS	`1,000 + 6,000 = 7,000`	opening on-hand 1,000; PO 6,000 (booklet, illustrative)	7,000 units
2	Dock count short 2%	`0.02 × 6,000 = 120`	2% short on F only	120 units
3	Received	`6,000 − 120 = 5,880`	step 2	5,880 units
4	Aggregate accuracy	`(5,880 + 4,000 + 2,000) ÷ 12,000 = 0.990`	M and S at plan (booklet)	99.0% overall
5	Line accuracy on F	`5,880 ÷ 6,000 = 0.980`	the aggregate hides this	98.0% on F

Table 2 — the fix: step by step, values substituted.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Promise on the lower of two	`min(6,000, 5,880) = 5,880`	ASN vs dock count — the lower-of-two asymmetry	promise 5,880
2	Inbound AFS after stow	`1,000 + 5,880 = 6,880`	only units that cleared receive and stow	6,880 units
3	Gap to plan	`7,000 − 6,880 = 120`	the shortage equals the whole miss	120 units
4	Cover at the close	`1,000 + 5,880 − 3,500 = 3,380; 3,380 ÷ 500 = 6.76`	week sales 3,500; demand 500 a day	6.76 days

Table 3 — after: the metric, before → after.

STEP	METRIC	BEFORE	AFTER	Δ / CHECK	BASIS
1	Inbound AFS vs plan	7,000	6,880	`7,000 − 6,880 = 120`	booklet, illustrative
2	Cover (days)	7.00	6.76	`7.00 − 6.76 = 0.24`	booklet, illustrative

STEP	METRIC	BEFORE	AFTER	Δ / CHECK	BASIS
3	Accuracy	99.0% overall	98.0% on F	`99.0 − 98.0 = 1.0 pt`	booklet, illustrative

Table 4 — the dollars, every multiplication shown.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Shortage as demand	`120 ÷ 500 = 0.24`	days of F demand	0.24 days
2	Shortage as a share of the week	`120 ÷ 3,500 = 0.0343`	F weekly demand 3,500	3.43%
3	Pass-through check	`0.24 × 500 = 120`	the days lost equal the units short	120 units

THE BRIDGE IN 20 SECONDS

“At Sephora the promise could see only what the ledger said was available: lower of two books, minus what was already committed. Inbound is the same discipline one step earlier — the promise sees opening on-hand plus the units that cleared receive and stow, on the lower of the ASN and the dock count, and a unit on the dock is just a shipped-not-billed unit of the building. One owner per state, and a daily check that the count and the system agree.”

Question bank — likely questions for this interviewer, with a one-line answer plan

UPGRADED

Twenty-four questions (the original twelve plus twelve new). Each plan points at a story in this packet or a bench story; stories allocated to other interviewers are pointers only.

#	LIKELY QUESTION	ANSWER PLAN (ONE LINE)	STORY
1	**Tell me about a time you aligned multiple organizations on one definition.**	S1 pitch — three truths, one signed FSD 6/5/25; pre-wiring; objections as rows; oversells recurring → zero.	S1
2	**Tell me about getting Finance to trust an operational number.**	S3 pitch — one FOH; failure direction; the aging report given to Finance to own.	S3
3	**How do you influence without authority?**	Pre-wire every row with its owner's LOE; working sessions to one table; three decisions not a review; a no that lives in a document. (S5 pointer.)	S1 (S5 pointer)
4	**Tell me about a time you were wrong.**	S4 (bench) — the in-transit rule; or S1-e — my first no, reversed on evidence. Keep S4 to the pointer unless the room asks for the full walk.	S4 / S1-e
5	**How would our two teams share a process / a metric here?**	One definition, one owner per bucket (S1-a); drift report tolerance zero with a named owner (S1-b); rule changes ride a release train.	S1-a / S1-b
6	**Tell me about designing for failure.**	S3 — failure direction before the happy path; unmatched ⇒ under-sell; aging report at T+1; late line −2, never +2.	S3
7	**How do you handle a stakeholder who wants a capability that sells more?**	S1-c — merch reserve promisable: priced (560 vs 710), visibility offered, no written as a row; break-even 61%.	S1-c
8	**How do you decompose a cross-functional process problem?**	Master §3.99 — book of record and metric first; every bucket a unit can sit in; flows and their documents; the broken invariant; minimal design; sequence.	Master §3.99
9	**What is the business value of the book-of-record migration?**	≈$2.5M/yr of program impact: $1M MCS run cost (user anchor) + ≈$1.5M of modeled avoided cost with its three rows; incidents to zero is the record movement.	S1
10	**How do you keep a design from being re-litigated after sign-off?**	Three signatories; the no as an FSD row (S1-c); changes ride the release train as FSD changes (S4 pointer).	S1 / S1-c
11	**Tell me about a release / change-management discipline you set up.**	Rule changes ride the monthly cadence; the matrix is the regression gate; no story enters QA without a cell reference; go/no-go from counted rows (S1-d).	S1 / S1-d
12	**What would a peer say is hard about working with you?**	I convert verbal agreements into rows people have to sign — slower in the week, faster over the program. S1 pre-wiring as the example.	S1
13	**Walk me through available-for-sale on one SKU — and how an event moves it.** NEW	S1-a — 980 − 420 = 560; commit → 555; pick, pack, ship move it between buckets; the posting moves both books together.	S1-a

#	LIKELY QUESTION	ANSWER PLAN (ONE LINE)	STORY
14	**How do you reconcile an aggregate to its lines? What about one cent?** NEW	S3-a — $100.00 over three units: last line absorbs → $33.33 + $33.33 + $33.34; missing line and duplicate-header guards.	S3-a
15	**Who owns the exceptions after you leave?** NEW	S1-b — Negative AFS to Inventory Control, Internal Audit and Finance; AFS Correction to the ISM app team; drift report to Inventory Control.	S1-b
16	**Tell me about a no you wrote down.** NEW	S1-c — the FSD row, the pricing, the number that would flip it.	S1-c
17	**How do you decide go or no-go?** NEW	S1-d — a checklist read from counted rows: Blocker + Critical open = 0; reconciliation clean N days; rollback armed; UDC log 836 rows, 441 severe.	S1-d
18	**A neighboring team asks for a test budget. How do you scope it?** NEW	S1-e — refuse the unscoped ask, fund a study, reverse on evidence; $42K; count, sum, per-store.	S1-e
19	**Half the capacity, same date: what do you cut?** NEW	Peer operating drills — cut scope and sources, never the controls; each cut has a tripwire.	Drills
20	**Which dependency worried you, and what did you do?** NEW	The sync to S/4 (rollback trigger); daily drift; dated E2E days; owner and date on every open question.	Drills / S1-d
21	**How does this transfer to NACF inbound?** NEW	NACF bridge — AFS = opening + units that cleared receive and stow; lower of ASN and dock count; one owner per state.	Bridge
22	**Tell me about a time you missed a deadline.** NEW	Honest: there is no missed S/4 date in the record — do not invent one; offer S4 (a rule I corrected) and say why **none on record**.	S4 pointer
23	**Tell me about developing someone.** NEW	Thin: S1-b — control owners handed the screens; measure by whether the drift report runs without me **no mentoring record**.	S1-b
24	**Tell me about a time you changed your mind.** NEW	S1-e — refuse, study, reverse in public; S4 — back-test before publishing.	S1-e / S4

Cheat-numbers box — canon values with basis labels UPGRADED

Canon numbers for the eight S/4 primaries and the bench stories, with basis labels, from the master packet's §4.1 ledger. The S1-a to S1-e rows are re-cuts of S1 and S3 facts; none adds a new dollar.

METRIC	VALUE TO SAY	BASIS	NEVER
S1 primary metric	customer-facing oversell incidents recurring across three node types → zero through two go-lives	record	an oversell count or rate as record
S1 headline $	≈$2.5M/yr of program impact = $1M MCS run cost retired + ≈$1.5M of modeled avoided cost (oversell ≈$0.9M · recon ≈$0.33M · change premium ≈$0.3M); range $1.7–5.2M	user anchor + modeled	presenting the $1.5M as record; adding upside into the headline
S1 signatures / dates	FSD v1.1 (3/5/25) → v1.2 (5/16/25) → signed 6/5/25 by Supply Chain Strategy, Inventory Management, FSC Inventory Accounting	record	one signatory
S1 gate chain	kickoff 3/29 → council 5/15 → LOE 7/2 → sign-off 7/31 → FSD 8/14; E2E 11/22 · 12/20 · 1/30; January planning + May execution go-lives; UDC log 836 rows / 85.4% closed	record (UDC program artifacts; dates shared with S2/S5)	—
S1 design	lower-of-two basis; AFS = Basis − PickRes − EcomPick − EcomPack − SNB; idempotency key per event; daily ISM-vs-MB52 drift tolerance 0; OMS reads, never computes	record design	"I configured S/4" / "I wrote ABAP"
S1 reconciliation FTE	≈4 FTE by hand → ≈1 FTE reading the drift report	modeled	stating it as record
S3 primary metric	Financial On-Hand truths 2 → 1; AFS decremented at sale/ship; over-sell from unmatched postings = 0 by construction	record (FSD 6/5/25)	—
S3 dollars	≈$0.6M/yr (cancels 0.45 · recon labor 0.14 · close 0.05); range $0.25–1.2M	modeled — no pre-state drift dollar in record	any record drift dollar; double-counting S1's cancel rate
S3 design KPI	match rate ≥ 99.5% within 24 h; aging report empty or explained at T+1; Σ per-line COGS = posting COGS to the cent	record design	a days-to-tie figure
S1-a contract NEW	ten buckets (incl. derived OOB and Basis), one owner each; one event → one bucket; 60 s publish; WMS feed ≤ 15 min; consumer lag < 5 min; drift tolerance 0	record design (walk units illustrative)	a measured defect rate; the 25% slice as fact
S3-a rounding NEW	$100.00 over 3 units: $33.33 + $33.33 + $33.34; $175.00 over 8 units: $65.63 + $65.63 + $43.74	record rule; amounts illustrative	claiming rounding mode / line order are in the record
S1-b controls NEW	Negative AFS → Inventory Control, Internal Audit, Finance; AFS Correction → ISM app team; drift → Inventory Control; Negative AFS clean at T+1	record	a maker-checker claim; workload numbers as fact

METRIC	VALUE TO SAY	BASIS	NEVER
S1-c reserve NEW	merch reserve outside S4_Avail; visible, not promisable; 560 vs 710 on the worked SKU; break-even 30 ÷ (19 + 30) = 61%	record design; units and benchmarks illustrative	a measured release rate; a hold count
S1-d cutover NEW	UDC log 836 rows; 85.4% closed; Blocker 164 + Critical 277 = 441 = 52.8%; triggers sync < 99% 4 h, p95 > 2× 1 h; E2E 28 then 41 days apart	record (derived arithmetic labeled)	a fired trigger; a cost of delay as record
S1-e interfaces NEW	17 interfaces; 2 critical usages; $42K funded; $25K contingent, not pre-spent; count, sum, per-store	record (loop packet S25 — not in master canon)	a price on the avoided regression
S4 (bench)	starved lanes → restored; oversell +0; kill switch never fired; ≈$0.5M/yr modeled margin	record / modeled	a dollar for the conversion lost while I was wrong
S5 (bench)	surprise escalations → 0; ~80 names across ~15 functions; 5/15 → 7/31 → 8/14; ≈$0.5M one-time modeled	record / modeled	—
SAP tenure	S/4HANA integration program lead 3+ years (Apr 2023 – present); client-SAP data integration at Elementum 2017–19; 15 years full-lifecycle SCM systems	record	“10 years of SAP configuration”; “SAP functional owner”

Never-say, one glance before the call.

- Never “I configured S/4,” “I deploy SAP,” or “I wrote ABAP” — integration program lead; SAP teams and configurators built to the signed FSD. Manhattan DOM is the only platform personally configured in canon.
- Never present the ≈$1.5M (S1) or the ≈$0.6M (S3) as record; the $1M MCS figure is the user anchor and the only non-modeled dollar here. Never add the variants’ slices to S1 or S3.
- Never give the S1 and S3 cancel rates as two results in one room — keep S1’s, drop S3’s row ($0.19M left).
- Never invent a days-to-tie figure, a direct-report headcount, a pre-state oversell dollar, or a price for the avoided regression (S1-e).
- Never inflate a multiple: customers 7 → 68 is 9.7×. Say “six sites,” never a regional qualifier on it. Finished-goods grain only — no component or bill-of-materials stories. Never turn a single dollar figure into a range you cannot source.
- S4 and S5 are bench stories; S14 is the bar raiser’s; S2/S18 are Goforth’s — “one umbrella program” if asked, no floor story here.

Backup bench — use only if a primary doesn't fit

Four fallbacks for a peer Senior Manager — S4 and S5 are the question bank's pointers, here in full bench form: S4 (Sephora — the in-transit promise rule I got wrong, fixed through the same three signatories; the failure / Are Right, A Lot question), S5 (Sephora — the architecture council without owning a system; Backbone and influence without authority — same program family as S1, so use only when the question is about governance rather than the definition) and S9 (Elementum — the 597-lane Pareto two levels up; Backbone / Think Big outside the Sephora program); plus a doc-derived fourth from the Flexport program record — the FFA Audit Logs PRD (controls and process integrity for a partner-facing platform; record-from-docs labels; same platform as Tripathy's S12, facts kept identical to her packet). Tell S1 and S3 first.

Bench discipline. The two (or four) primaries above are the stories to tell first. Reach for a bench story only when the question genuinely does not fit a primary, or when the interviewer asks for "another example" on an LP already spent. Each entry is a 60-second pitch, a compressed STAR, one results line and two probes — enough to answer and stop; no full probing bank or derivation. Where a bench story is another interviewer's primary, the red note says so: use it only if the question demands it, and keep the facts identical to that packet.

BACKUP — ONLY IF ASKED S4 — Sephora S/4HANA — the in-transit promise rule I got wrong, fixed through the signed policy, not a side configuration

LPs this backup can serve: Are Right, A Lot (the failure question) · Customer Obsession (backup) · Earn Trust (owning the miss) · Learn and Be Curious — the designated policy-failure story

Basis. The v1 rule was mine; the indictment came from my own program's reports. Lane σ ≈ 0.3 day and the ~0.49-day window are the record's promise math; the record carries no lane counts and no dollar for the conversion lost while I was wrong — I will not invent the loss figure.

60-SECOND PITCH — STOP ON THE PRIMARY METRIC

Situation. My first ISM ruleset barred every channel from promising units in transit; the Aging In-Transit report showed dotcom starved of fast movers one day from the building. **Task.** Correct my own policy without weakening oversell protection, through the same three signatories. **Action.** Discovery: conversion-loss evidence from the aging and replenishment reports — the entry bar to reopen the policy; I stated that the v1 rule was mine. Plan & Design: promise math — lane σ ≈ 0.3 day, a ~0.49-day window ≈ 95% coverage; Promisable = in-transit × 1[ETA ≤ window] × 1[on-time ≥ floor]; Aging In-Transit as the kill switch; both blanket options rejected in writing. Build & Test: back-test against receiving history; four regression cells added to the node × channel matrix. Cutover: released on the normal monthly cadence as a signed FSD change, lane by lane — no hotfix, no side configuration. Hypercare: kill switch armed; never fired. Lesson: both sides of a rule get a number before it ships. **Result.** Primary metric: sellable units on qualifying in-transit lanes — starved → restored inside a per-lane promise window, oversell increase = zero, kill switch never fired. ≈$0.5M/yr modeled margin recovered — the record has no dollar for the conversion lost while I was wrong, and I will not invent one.

COMPRESSED STAR

- **S —** v1 of my own ISM ruleset: no channel could promise in-transit units — safe-feeling, and starving dotcom of fast movers reliably one day from the building.
- **T —** Fix my own signed policy without weakening oversell protection; take the miss back through the same three signatories.
- **A —** Evidence pack from Aging In-Transit and Replenishment; reframed as promise math (σ ≈ 0.3 d → ~0.49-d window ≈ 95%); rule with on-time floor and kill switch; back-test against receiving history; four regression cells;

released lane by lane on the monthly cadence as an FSD change.

- **R —** Starved lanes restored inside the window; oversell increase = 0; kill switch never fired; re-signed by the same three signatories; lesson institutionalized — the Locus 98% gate (S2) and the 0.5% tolerance (S18) were set from data before the event.

RESULTS LINE

Sellable units on qualifying in-transit lanes starved → restored inside a back-tested ~0.49-day window (record) · oversell increase = 0; kill switch armed, never fired (record) · four regression cells added (record) · released as a signed FSD change on the normal cadence, no hotfix (record) · ≈$0.5M/yr modeled margin recovered (inputs illustrative).

THE MATH — STEP-BY-STEP DERIVATION

Table 1 — before: the problem in numbers.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	The v1 rule	`Promisable = 0 for ALL in-transit inventory`	equivalent to claiming σ → ∞ or on-time rate → 0 — infinite variability (record)	everything in transit hidden
2	One concrete lane	`Promisable = 120 × 0 = 0`	lane with InTransit = 120 units, ETA 0.4 day away (illustrative)	120 sellable units hidden
3	Across qualifying lanes	`≈300 fast-mover units/day × 0 promisable`	illustrative volume	≈300 units/day sat unsellable one day from the building
4	What starving costs (one SKU)	`100 × 60% × $20 = $1,200/day`	100 units one day out; 60% would convert; $20 margin (illustrative, not program data)	$1,200/day under-sell
5	The ignored other side	`100 × (1 − 0.97) = 3 late units; 3 × $15 = $45/day`	on-time rate 97%; $15 expedite (illustrative)	$45/day over-sell exposure — vs $1,200: the “safe” rule priced one side only

Table 2 — the fix: step by step, values substituted.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Measure lane variability	`σ = 0.3 day`	trailing receiving history (record)	σ = 0.3 day
2	Choose coverage	`z at 95% = 1.645`	normal approximation on lane transit (toolkit label)	z = 1.645
3	Promise window	`w = z × σ = 1.645 × 0.3 = 0.4935`	rows 1–2	w ≈ 0.49 day
4	Coverage check	`Φ(w / σ) = Φ(0.49 / 0.3) = Φ(1.63) ≈ 0.95`	share of arrivals inside the window	≈95% covered
5	Corrected rule	`Promisable = InTransit × 1[ETA ≤ w] × 1[OTR ≥ floor]`	OTR = on-time rate over trailing receipts; floor = per-lane admission threshold, 95% here (floor illustrative)	lane by lane, not all-or-nothing
6	Lane that qualifies	`ETA 0.4 ≤ 0.49 → 1; OTR 97% ≥ 95% → 1; Promisable = 120 × 1 × 1 = 120`	the Table 1 lane, re-evaluated (ETA and OTR illustrative)	its 120 units become promisable
7	Lane that fails	`ETA 1.2 > 0.49 → 0; Promisable = 80 × 0 = 0`	a slower lane with 80 units in transit (illustrative)	stays at 0 — correctly out
8	Sample-size gate	`n = z²p(1−p)/e² = 3.84 × 0.95 × 0.05 / 0.05² = 3.84 × 0.0475 / 0.0025 ≈ 73`	OTR estimated within ±5% at 95% CI (illustrative)	a lane qualifies after ~73 receipts

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
9	Kill switch, values in	`trigger = 2 × p90 = 2 × 2.0 = 4.0 days; any unit in transit > 4.0 days → lane out`	lane p90 = 2.0 days (illustrative); armed, reviewed weekly by Inventory Control (record)	never fired (record)
10	Back-test	`rule replayed against receiving history; 4 regression cells`	before release (record)	no oversell in the back-test

Table 3 — after: the metric, before → after.

STEP	METRIC	BEFORE	AFTER	BASIS
1	Sellable units on qualifying in-transit lanes	0	≈300 units/day restored inside the window	record mechanism; volume illustrative
2	Promise window	none — blanket exclusion	w ≈ 0.49 day at ~95% cover	record (promise math)
3	Oversell increase from the window	—	0	record
4	Kill switch	not armed	armed, never fired	record
5	Rule changes after	ad hoc	same 3 signatories, back-test + matrix regression	record — both sides priced from then on

Table 4 — the dollars, every multiplication shown.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Would-be sales	`300 × 60% = 180 sales/day`	units restored (illustrative) × conversion (illustrative)	180 sales/day
2	Truly lost without the fix	`180 × 40% = 72 sales/day`	40% lost-sale share — the rest would have waited (illustrative)	72 lost sales/day
3	Margin recovered	`72 × $20 = $1,440/day`	$20 unit margin (illustrative)	≈$1,440/day
4	Annualized	`$1,440 × 365 = $525,600`	modeled; the mechanism is record	≈$0.5M/yr
5	Cost while v1 stood	`no dollar in the record — not claimed`	record	not claimed

TWO PROBES

> **Why not just hotfix it — you knew it was wrong?**
>
> Because the error was in the policy, so the fix had to be too — in the governed artifact, with regression, through the same signatures. A side configuration would have fixed the symptom and left the policy lying about what it allowed.

> **What did the miss cost?**
>
> The record carries no dollar for the conversion lost while v1 stood, and I will not invent one. What I can say is the mechanism: fast movers reliably one day from the building were unsellable on dotcom until the window existed; the ≈$0.5M/yr recovered is my model on illustrative inputs, labeled.

Canon numbers (basis labels): Lane σ ≈ 0.3 d; ~0.49-d window ≈ 95% coverage (record promise math) · oversell increase = 0 (record) · 4 regression cells (record) · kill switch never fired (record) · same three signatories re-signed (record) · ≈$0.5M/yr (modeled). Never a lane count; never a dollar for the loss while wrong.

BACKUP — ONLY IF ASKED S5 — Sephora — the architecture council without owning any system: three decisions and a fence

LPs this backup can serve: Have Backbone; Disagree & Commit (third) · Earn Trust (influence without authority) · Think Big (architecture governance). Same program family as S1/S2 — the governance story

Basis. I owned none of the systems; the design decision was the deliverable. Zero escalations and the dated gate chain are the record; the dollars are my model of what governance avoided.

60-SECOND PITCH — STOP ON THE PRIMARY METRIC

Situation. The UDC location and inventory model had to clear architecture council — a challenge gate — with ~80 named participants across ~15 functions, and I owned none of those systems. **Task.** Take the design from discovery through council to sign-off and FSD without a surprise escalation. **Action.** Discovery: both inventory books profiled before any design argument; business sessions closed 4/30. Plan & Design: every impacted row pre-wired with a signed LOE by 7/2; the council asked for three separable decisions plus an out-of-scope fence (reserve math untouched; MCS decommission belongs to Global ERP). Build & Test: the "build" was the pack — architecture, staged timeline, peak-freeze boundary. Cutover: 5/15 challenge → 7/31 sign-off → 8/14 FSD. Hypercare: the fence held in writing for the whole program; a second live warehouse became a monitored shadow copy, not an argument. Lesson: a no that lives in a document survives. **Result.** Primary metric: surprise escalations on a 60+ stakeholder design — the norm on prior programs → zero; the dated gate chain held. ≈$0.5M one-time modeled (avoided rework on ~3 escalations that never happened + six weeks of benefit not deferred + a fenced second reserve engine not built).

COMPRESSED STAR

- **S —** About a dozen impacted systems (Apple C4), ~80 participants across ~15 functions, a council that re-litigates designs this size — and I owned none of the systems.
- **T —** Discovery (kickoff 3/29, sessions 4/30) → council 5/15 → sign-off 7/31 → FSD 8/14, without a surprise escalation.
- **A —** Profiled both books first — their numbers, not my opinion; systems-impacted table with an owner on every row; one LOE template, one deadline (7/2), one private conversation per owner before the room; three separable decisions plus the fence; the second-truth objection converted into a monitored shadow copy.
- **R —** Surprise escalations → 0 for the program's life; 100% pre-wire coverage by 7/2; gate chain held to the day with zero re-litigation; fence breaches 0; handed clean into the build (S2).

RESULTS LINE

Surprise escalations → 0 for the program's life (record) · gate chain 3/29 → 4/30 → 5/15 → 7/2 → 7/31 → 8/14 held to the day (record) · pre-wire coverage 100% — every impacted row with a signed LOE by 7/2 (record) · both scope refusals held in writing; fence breaches 0 (record) · ≈$0.5M one-time (modeled — what did not happen).

THE MATH — STEP-BY-STEP DERIVATION

Table 1 — before: the problem in numbers.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Design size	`~80 named participants × ~15 functions; about a dozen impacted systems`	record (Apple C1/C4)	the blast radius
2	Norm for this size	`~3 surprise escalations per design`	industry benchmark 2–4	~3 expected

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
3	Price of one escalation	`960 rework hours × $110/h = $105,600`	modeled	≈$0.1M each
4	Bundle decision risk	`P(all pass) = 0.9 × 0.9 × 0.9 = 0.73`	one up/down vote on the whole design; p = 0.9 per decision (illustrative)	27% chance one objection unwinds everything

Table 2 — the fix: step by step, values substituted.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Three separable decisions	`architecture · timeline · peak boundary → 2³ = 8 outcomes`	any partial approval preserves progress (record design)	a no on one cannot unwind all
2	Expected approvals	`E = 0.9 + 0.9 + 0.9 = 2.7 of 3`	vs the bundle's 0.73 chance of everything (illustrative p)	2.7 vs 0.73
3	Pre-wire coverage	`rows with signed LOE / impacted rows = 100% by 7/2`	one template, one deadline, one private conversation per owner (record)	100% before the room
4	Gate chain held	`3/29 → 4/30 → 5/15 → 7/2 → 7/31 → 8/14 = 77 days council→sign-off`	held to the day (record)	0 re-litigation

Table 3 — after: the metric, before → after.

STEP	METRIC	BEFORE	AFTER	BASIS
1	Surprise escalations	~3 (the norm)	0 for the program's life	record; norm is benchmark
2	Pre-wire coverage before council	partial	100% by 7/2	record
3	Fence breaches in the log	—	0	record
4	Re-litigation after sign-off	—	0 across 77 days	record

Table 4 — the dollars, every multiplication shown.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Avoided rework	`3 × 960 h = 2,880 h; 2,880 × $110 = $316,800`	~3 escalations (benchmark) × rework hours × rate (modeled)	≈$0.3M
2	Slip weeks avoided	`6 / 52 = 0.115 yr; 0.115 × $900,000 = $103,800`	six weeks of S2's ≈$0.9M/yr labor value not deferred (modeled)	≈$0.1M
3	Fenced engine not built	`2 eng × 12 wk × 40 h = 960 h; 960 × $110 = $105,600`	second reserve engine (modeled)	≈$0.1M
4	Total one-time	`0.3 + 0.1 + 0.1 = $0.5M`	all modeled — the dollars are what did not happen	≈$0.5M

TWO PROBES

> **How do you get sixty people to sign LOE you don't control?**
>
> A template, a deadline and a private conversation each — so nobody learned their cost in the room. The council was then asked for three separable decisions it could take one at a time; separability was the design.

> **What did you refuse, and did it hold?**
>
> Retuning the reserve math and folding in the MCS decommission — both written, signed and held; later scope requests were answered from the pack in writing. The legitimate concern (a second live warehouse becoming a second truth) was converted into the design as a monitored shadow copy, because converting a concern costs less than winning it.

Canon numbers (basis labels): ~80 participants / ~15 functions; about a dozen impacted systems (record, Apple C4; never “20+”) · 5/15 → 7/31 → 8/14; LOE 7/2 (record dates) · escalations 0; fence breaches 0 (record) · ≈$0.5M one-time (modeled). Never a dollar as record; never pair with S1 or S2 for the same interviewer.

BACKUP — ONLY IF ASKED S9 — Elementum — 1,000 supplier/carrier routes asked, 597 built first: X12 EDI onboarding behind a ≥95% valid-ASN cohort gate

LPs this backup can serve: Have Backbone; Disagree & Commit (primary — 597 vs 1,000, two levels up) · Deliver Results · Insist on the Highest Standards (the gate held twice) · Think Big (the staged 1,000-lane vision) · Broad Responsibility (small carriers kept inside the standard by the portal floor)

Basis. I owned the scope analysis and the onboarding standard, influencing an executive two levels above my sponsor; engineers built the maps and harness to the standard; Transportation owned the carrier relationships. Cohort sizes and the number of cohorts are not in the record — do not quote them.

60-SECOND PITCH — STOP ON THE PRIMARY METRIC

Situation. The Global Head of Supply Chain wanted inbound visibility on 1,000 lanes in two months; shipment-data completeness was 60% and high-value lanes were blind. **Task.** I owned the scope analysis and the onboarding standard, influencing an executive two levels above my sponsor. **Action.** Discovery: a spend Pareto — 597 lanes carried 80% of freight spend — pre-validated with Transportation and Procurement before I took it up. Plan & Design: the onboarding standard (X12 850/855/856/940/945/214/997 with an XML/portal floor — "the standard was the data, not the transport") and a cohort gate of ≥95% valid ASN. Build & Test: a mapping-and-test harness and a scorecard per lane. Cutover: 597 live in month one; one cohort failed the gate twice on a timestamp format and I held it both times. Hypercare: completeness scorecards to 95%; ~200 lanes the next wave behind the same gate. Lesson: automating dirty data manufactures confident wrong answers — the gate is the product. **Result.** Primary metric: inbound shipment-data completeness 60% → 95%. Freight spend −10% ≈ $2M in six months (user anchor); 597 lanes = 80% of spend live in month one (user anchor); adoption 8× (user anchor). Apple F1: “a thousand lanes onboarded across 25 European locations, goal hit.”

COMPRESSED STAR

- **S —** 1,000 global lanes asked in two months — scope, schedule and quality all fixed; completeness 60%; lanes differed by orders of magnitude in spend and carriers differed in EDI capability.
- **T —** Own the scope analysis and the onboarding standard; take the re-scope two levels up with the evidence attached.
- **A —** Spend Pareto (597 = 80% of spend) pre-validated so no lane owner could contradict it in the room; standard + ≥95% valid-ASN gate published to carriers before cohort one; harness so later cohorts were mostly configuration; held a failing cohort twice; ~400 lanes carrying ~20% of spend deferred — priced in spend share, not lane count.
- **R —** Completeness 60 → 95%; 100% visibility on high-value lanes; 597 lanes live month one; freight −10% ≈ $2M in six months (user anchor); method written up as the sourcing SOP.

RESULTS LINE

Completeness 60% → 95% (record) · 597 of 1,000 lanes = 80% of freight spend live in month one (user anchor — GFiber v8) · freight spend −10% ≈ $2M in six months (user anchor — never annualized) · adoption 8× (user anchor) · ≈$0.4M/yr document-handling labor (modeled add-on).

THE MATH — STEP-BY-STEP DERIVATION

Table 1 — before: the problem in numbers.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Completeness	`480 / 800 = 60%`	expected = 200 shipments × 4 milestones = 800; received 480 (illustrative volumes; 60% record)	60% — 2 of 5 events missing

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
2	What blindness costs	`off-guide spend unseen · premium modes by habit · chase calls per late shipment`	record problem	no instrument, no lever
3	Manual documents	`597 lanes × ~50 shipments × 3 docs ≈ 90,000 docs/yr keyed by hand`	illustrative volumes	~90K docs/yr

Table 2 — the fix: step by step, values substituted.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Spend Pareto	`S(k) = cumulative spend share; S(597) = 80%`	lanes sorted by annual freight spend (user anchor)	597 of 1,000 lanes carry 80% of spend — do them first
2	Cohort gate	`validASN rate ≥ 95% in test before activation`	published to carriers before cohort one (record)	one cohort failed twice on a timestamp format — held both times
3	Long-tail floor	`XML / portal when EDI is absent`	“the standard was the data, not the transport” (record)	small carriers kept inside the standard
4	Completeness after	`760 / 800 = 95%`	illustrative volumes; 95% record	95%; high-value lanes 100%

Table 3 — after: the metric, before → after.

STEP	METRIC	BEFORE	AFTER	BASIS
1	Shipment-data completeness	60%	95%; high-value lanes 100%	record
2	Lanes live month one	0	597 = 80% of spend	user anchor
3	Freight spend on onboarded lanes	base	−10% ≈ $2M in 6 months	record
4	Document handling	~90K docs/yr keyed	X12/XML validated	modeled

Table 4 — the dollars, every multiplication shown.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Implied addressable base	`$2,000,000 / 10% = $20,000,000/yr`	derived from the record pair, stated as such	≈$20M/yr
2	The record result	`$20M × 10% = $2M in 6 months`	record	≈$2M
3	Slice — routing-guide compliance	`~$0.8M of the $2M`	off-guide spend caught by 856/214 (modeled split)	~$0.8M
4	Slice — mode discipline	`~$0.6M of the $2M`	premium modes → ocean/ground when ETA is trusted (modeled split)	~$0.6M
5	Slice — expedites + detention	`~$0.6M of the $2M`	events replace chase calls (modeled split)	~$0.6M
6	Split check	`0.8 + 0.6 + 0.6 = $2.0M`	the slices decompose the record — they never add to it	= the record $2M
7	Document labor	`90,000 × 6 min = 540,000 min; / 60 = 9,000 h; × $45 = $405,000`	illustrative volumes (modeled add-on)	≈$0.4M/yr

TWO PROBES

> **You defied a Global Head of Supply Chain's number. How?**
>
> With the spend curve and the lane owners' validation attached — 597 lanes carried 80% of the spend, so the phased plan was argued in spend share, not lane count. He accepted the phasing off the curve; I then committed to the remainder behind the same gate and delivered it.

> **Is the $2M annualized?**
>
> No — it is the record as measured: minus ten percent on the onboarded lanes' spend in six months. My three-way split (routing-guide compliance ~$0.8M, mode discipline ~$0.6M, expedites and detention ~$0.6M) decomposes the record; it never adds to it.

Canon numbers (basis labels): 60% → 95% completeness (record) · 597 / 1,000 lanes = 80% of spend (user anchor) · −10% ≈ $2M in 6 months (user anchor) · ≥95% valid-ASN gate, held twice (user anchor) · 8× adoption (user anchor). Apple F1 safe line: “a thousand-lane program across 25 European locations; I phased by value” · ≈$0.4M/yr labor (modeled). Never: cohort counts or sizes; never double the $2M into a run-rate.

BACKUP — ONLY IF ASKED Doc bench — Flexport — the FFA Audit Logs PRD: an immutable who/what/when trail across ~100 partner orgs, nine categories at a 100%-captured bar (controls & process integrity)

LPs this backup can serve: Insist on the Highest Standards (primary — 100% capture as the written bar) · Ownership (compliance as a launch feature) · Dive Deep (the event taxonomy to the grain) · Earn Trust (sign-off over assumption)

Basis. Doc-derived micro-story from the program record (the FFA Audit Logs PRD, 2022) — every figure is RECORD-FROM-DOCS. Same platform as Tripathy's S12: use only for a controls / process-integrity question and keep S12's canon numbers identical to her packet. Program-lead voice: I owned the PRD, the event taxonomy and the sign-off chain; engineering built the pipeline and dashboards. Not in the master packet's S-numbering; supplements canon, never overrides it.

60-SECOND PITCH — STOP ON THE PRIMARY METRIC

Situation. The Freight Forwarding App serves external partner organizations — roughly a hundred of them, making thousands of changes a day to schedules, routes, documents and permissions. When a partner disputes "who changed this?", the answer has to come from an immutable trail — and the security organization's compliance bar ("important events") was deliberately vague. **Task.** Define the audit-log product: what gets captured, at what standard, phased how — so compliance is a launch feature, not an incident response. **Action.** I wrote the PRD: an immutable who/what/when record; nine event categories — logins/logouts, work-item changes, route updates, messages, document uploads/downloads, schedule updates, exception changes, assignment-rule changes, permission changes — each held to a written 100%-captured bar; priorities P0/P1/P2 per category; delivery phased Alpha (a login/logout dashboard on the log platform, Q4 2022) → Beta → MVP (Q1 2023); and the vague-bar risk mitigated structurally — PM buy-in plus CISO sign-off written into the plan, the "trusted date" for audits recorded in a governed document, and the eight dependent FFA teams mapped before build. **Result.** The capture standard was defined, prioritized and signed before build: 100% of nine categories, a phased path that did not pretend everything could be captured on day one, and zero ambiguity about who approved the scope. RECORD-FROM-DOCS

COMPRESSED STAR

- **S —** A partner-facing platform with ~100 partner orgs and thousands of changes a day; no immutable trail; a compliance bar stated vaguely enough to move.
- **T —** Own the audit-logs PRD: taxonomy, capture standard, phasing, and the sign-off chain that makes the bar stick.
- **A —** Nine event categories at a written 100%-captured bar with P0/P1/P2 priorities; phased Alpha → Beta → MVP so capture breadth grew honestly; CISO sign-off and PM buy-in as the mitigation for the vague "important events" definition; the "trusted date" for audits recorded in a governed document; eight dependent teams mapped.
- **R —** Capture standard defined and signed before build (record-from-docs). The record carries no post-launch compliance dollar, and none is claimed.

RESULTS LINE

Nine event categories at a 100%-captured bar (record-from-docs) · ~100 partner orgs · thousands of changes/day (record-from-docs — platform-wide; never merge with canon's 53-partner import migration set) · phased Alpha (Q4 2022) → Beta → MVP (Q1 2023) with P0/P1/P2 (record-from-docs) · CISO sign-off as the written mitigation (record-from-docs) · no post-launch dollar in the record — never invent one.

THE MATH — STEP-BY-STEP DERIVATION

Table 1 — before: the problem in numbers.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Audit coverage	`0 of 9 event categories captured — no dashboard, no tool, no store`	PRD problem statement (record-from-docs)	"who changed this data?" unanswerable
2	Organizations affected	`≈100 partner organizations on the app`	PRD hypotheses (record-from-docs)	every org's changes invisible
3	Change volume projected	`thousands of changes/day once all partners and internal operators are live`	PRD hypotheses (record-from-docs)	untracked writes at scale
4	Explicit non-goals	`reads/views, IP addresses, automated cascades — out of scope this iteration`	PRD scope section (record-from-docs)	scope held to writes by humans

Table 2 — the fix: step by step, values substituted.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Capture bar	`9 of 9 categories at 100% capture, each with a sample query`	PRD metrics section (record-from-docs)	login/logout · work-item edits · route updates · messages · document up/downloads · schedule updates · exception changes · assignment-rule changes · permission changes
2	Event schema	`8 parameters per event = user FID + date + time + action type + object + field + previous value + current value`	PRD Epic 1 (record-from-docs)	every change fully described
3	Immutability	`deletion of captured events = 0 — by construction`	PRD acceptance criteria (record-from-docs)	the log cannot be rewritten
4	Phase arithmetic	`Alpha (Q4 2022): 1 category live — login/logout on a Sumo Logic dashboard → Beta (Q1 2023): 9 of 9 tracked with full parameters → MVP (Q1 2023): every product area publishes links + sample queries`	PRD launch phases (record-from-docs)	proof → coverage → discoverability
5	Trusted-start mechanism	`each team registers a target date in Confluence; history is trusted only from the latest start date`	PRD risk table (record-from-docs)	honest about partial history
6	Volume sizing	`2,000 changes/day × 365 = 730,000 events/yr; × 8 parameters = 5,840,000 stored values/yr`	"thousands/day" from Table 1 row 3 (record-from-docs); 2,000 as the low-end reading and the arithmetic are illustrative	≈0.73M immutable events/yr

Table 3 — after: the metric, before → after.

STEP	METRIC	BEFORE	AFTER	BASIS
1	Categories captured at 100%	0 of 9	9 of 9 — the stated success metric	record-from-docs
2	Alpha exit	—	Sumo Logic dashboard with login/logout metrics	record-from-docs
3	MVP exit	—	all product areas linked, with sample queries	record-from-docs
4	Security sign-off	vague CISO requirements	CISO representatives sign the "important events" definition before development	record-from-docs

Table 4 — the dollars, every multiplication shown.

STEP	QUANTITY	FORMULA → SUBSTITUTION	INPUTS (BASIS)	VALUE
1	Investigation deflection	`2 h × $60/h = $120 per "who changed this" hunt; × 100 hunts/yr = $12,000/yr`	every input illustrative — no rate or count exists in the record	≈$12K/yr, illustrative only
2	Compliance exposure	`no dollar in the record — not claimed`	record-from-docs	not claimed
3	Post-launch dollar	`none in the record — never invent one`	record-from-docs	not claimed

TWO PROBES

Why 100% capture as the bar — isn't that over-engineering?

Because an audit trail with holes is worse than none — it lends false confidence to a dispute answer. The 100% bar is per category, phased: Alpha proved the pipeline on the highest-volume, simplest events (logins), then breadth grew by priority. The bar was written down precisely so "good enough" could not drift downward under delivery pressure.

What did the vague "important events" definition actually risk?

Scope that moves after build — the classic compliance trap where the auditor's definition surfaces at the audit, not at the PRD. The mitigation was structural, not heroic: PM buy-in and CISO sign-off on the taxonomy up front, so the definition of "important" was the signed document, not someone's memory. When the bar can move, pin it to a signature.

Canon numbers (basis labels): nine categories at 100% capture (record-from-docs) · ~100 partner orgs · thousands of changes/day (record-from-docs) · Alpha Q4 2022 → MVP Q1 2023 (record-from-docs) · CISO sign-off + "trusted date" in a governed document (record-from-docs). Never: merge ~100 partner orgs with S12's 53-partner migration set; never claim a prevented-breach dollar (none in the record).

VERIFY BEFORE YOU SAY IT — every tagged item NEW

Every VERIFY-WITH-RAJESH tag in the new stories, in one place. For each: confirm it from your own memory or the source, or drop it. Anything left unconfirmed should be said as “illustrative” or not said.

WHERE	ITEM TO CONFIRM	IF YOU CANNOT CONFIRM
S1-a	Ownership pairings for PickRes, MerchRes and LostFound — the master names the feeds (OMS, WMS, planning) but not every bucket-to-system pairing; “ten buckets” is my count of the named fields including derived OOB and Basis.	Confirm or say “ten fields.”
S1-a	The 25% event-level share of S1’s oversell row ($225,000) and 5,000 SKUs × 10 nodes = 50,000 rows.	Drop both unless you have counts.
S3-a	Rounding mode (half-up) and line-order key for “the last line absorbs”; 2,000 postings a day; the 30% share of bridge hours ($42,432); who proposed a rounding account (none claimed).	Say “the rule is last line absorbs”; keep the rest illustrative.
S1-b	Whether AFS Correction needs a second approver (maker-checker); the rehearsal list for control owners; 40 negative rows × 6 min and 10 drift rows × 15 min; 1,800 productive hours.	Say “screen, role, audit row” only.
S1-c	The 150-unit hold, 60% sell-through, 200 holds a year; the exact wording and tolerance of the four reserve-copy invariants; the 20,000 / 1,400 old-vs-new row counts; whether the FSD row carries a review trigger.	Keep as worked example, labeled illustrative.
S1-d	The 12,000 events an hour; the number of clean days N; “sync” = events applied ÷ received (my reading); the order-flow budget = the 60-second publish budget (my reading); the snapshot at which 85.4% was true; whether MCS run cost stops at cutover; whether any forum disagreed; whether a breach rehearsal was run.	Say the two trigger thresholds only; derive nothing else.
S1-e	All S25 figures (17, two usages, $42K, $25K, six core files plus the supply-chain family) come from the loop packet, not the master canon; the study’s cost; whether the $25K was ever spent; how many of the 17 were in scope. LOOP-CONSISTENCY: the loop packet lists S25 as a Bar Raiser backup.	Confirm figures; do not use if spent with the Bar Raiser.
Router	Hire and Develop the Best and Strive to be Earth’s Best Employer are thin; there is no manager-disagreement, missed-deadline or mentoring story in the S/4 record.	Keep answers short and honest; pivot to the nearest real story.
Drills	The cut orders and tripwires are my judgment from the record’s gates, not events.	Say “I would,” never “I did.”
Bridge	Inbound numbers (7,000, 6,880, 99.0%, 98.0%, 6.76) are the booklet’s illustrative week, not Amazon data.	Always say “in a worked example.”

NUMBERS I MUST NOT GET WRONG — with basis labels NEW

The canon and worked numbers used in this packet. Read the basis column aloud in your head before each one.

NUMBER	WHAT IT IS	BASIS	STORY
≈$2.5M/yr	$1.0M MCS run cost retired + ≈$1.5M modeled	user anchor + modeled	S1
range $1.7–5.2M	low / high of the same model; upside shrink ≈$0.3M + time-to-market ≈$0.2M excluded	modeled	S1
FSD 3/5/25 → 5/16/25 → signed 6/5/25	three signatory orgs	record	S1
E2E 11/22 · 12/20 · 1/30	28 then 41 days apart; January planning and May execution go-lives	record	S1, S1-d
836 rows / 85.4% / 52.8%	Blocker 164 + Critical 277 = 441 of 836 — the UDC issues log	record-from-docs (UDC log, recomputed with pandas)	S1-d
sync < 99% for 4 h · p95 > 2× for 1 h	armed both go-lives, never fired	record	S1, S1-d
FOH truths 2 → 1	AFS decrement at sale/ship; match ≥ 99.5% in 24 h	record / design KPI	S3
≈$0.6M/yr	cancels $0.45M + labor $0.14M + close $0.05M; range $0.25–1.2M	modeled	S3
560 = 980 − 420	worked SKU; raw S/4 over-promises 440, raw WMS 420	illustrative units	S1, S1-a
555 · 550 · 550 · 565	five-unit order and three bad events	illustrative units	S1-a
$100.00 = 33.33 + 33.33 + 33.34	last line absorbs the cent	record rule; amounts illustrative	S3-a
−20 vs +20	lower-of-two vs raw S/4 on S/4 500, WMS 460, committed 480	illustrative units	S1-b
61% = 30 ÷ 49	break-even release reliability for promising held stock	benchmark inputs	S1-c, S1 trap
$42K · $25K · 17 · 2	funded · contingent · interfaces · critical usages	record (loop packet)	S1-e
6.5 of 7.2 hours	one reader's daily workload on illustrative volumes	illustrative	S1-b